

Lecture

Operating System

5. Process API

5. The Process API

1. **fork(), wait() and exec() System Call**
2. **fork() and kill() going bad**
3. **pipe() and dup2()**



5. The Process API

1. **fork(), wait() and exec() System Call**
2. fork() and kill() going bad
3. pipe() and dup2()



Two ways to create a process

- Build a new empty process from scratch
- Copy an existing process and change it appropriately

Option 1: New process from scratch

■ Steps

- Load specified code and data into memory;
Create empty call stack
- Create and initialize PCB (make look like context-switch)
- Put process on ready list

■ **Advantages:** No wasted work

- **Disadvantages:** Difficult to setup process correctly and to express all possible options
 - Process permissions, where to write I/O, environment variables
 - Example: WindowsNT has call with 10 arguments

Option 2: Clone existing process and change

■ Example: Unix `fork()` and `exec()`

- `fork()`: Clones calling process
- `exec(char *file)`: Overlays file image on calling process

■ `fork()`

- Stop current process and save its state
- Make copy of code, data, stack, and PCB
- Add new PCB to ready list
- Any changes needed to child process?

■ `exec(char *file)`

- Replace current data and code segments with those in specified file

■ **Advantages:** Flexible, clean, simple. Fork is fast with CoW!

■ **Disadvantages:** Wasteful to perform copy and then overwrite of memory like in `exec()`.

The fork() System Call

- Create a new process
 - The newly-created process has its own copy of the **address space**, **registers**, and **PC**.
 - Syntax: `retval = fork()`
 - `fork()` creates identical copy of (parent-)process
 - Difference between child and parent: `retval` !
 - in parent-process: `PID` of child
 - in child-process: `0`

The fork() System Call: Example

fork.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (original process)
        printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());
    }
    return 0;
}
```


The fork() System Call: Executed

```
...  
} else if (rc == 0) {  
    // child (new process)  
    printf("hello, I am child (pid:%d)\n", (int) getpid());  
} else {  
    // parent goes down this path (original process)  
    printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());  
}  
return 0;  
}
```



Result (Not deterministic)

```
prompt> ./fork
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

```
prompt> ./fork
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

getpid()

- Why no error check of `getpid()` ?
 - From man page: “The `getpid()` and `getppid()` functions are always successful, and no return value is reserved to indicate”

The wait() System Call

This system call won't return until child has run and exited

wait.c

```
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        sleep(1);
    } else {
        // parent goes down this path (original process)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
            rc, wc, (int) getpid());
    }
    return 0;
}
```

Result (Deterministic)

```
prompt> ./wait  
hello world (pid:29266)  
hello, I am child (pid:29267)  
hello, I am parent of 29267 (wc:29267) (pid:29266)  
prompt>
```

The exec() System Call Family

Run a program that is different from the calling program

exec.c

```
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("exec.c"); // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // runs word count
        printf("this shouldn't print out");
    } else {
        // parent goes down this path (original process)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
            rc, wc, (int) getpid());
    }
    return 0;
}
```


All of the above in C with redirection

redirect.c

```
int main(int argc, char *argv[])
{
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child: redirect standard output to a file
        close(STDOUT_FILENO);
        open("./redirect.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU);
        // now exec "wc" ...
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("redirect.c"); // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // runs word count
    } else {
        // parent goes down this path (original process)
        int wc = wait(NULL);
    }
    return 0;
}
```

```
prompt> less redirect.output
      31      109      857 redirect.c
```

How are Unix shells implemented?

shell.c

```
While (1) {
    Char *cmd = getcmd();
    Int retval = fork();
    If (retval == 0) {
        // This is the child process
        // Setup the child's process environment here
        // E.g., where is standard I/O, how to handle signals?
        exec(cmd);
        // exec does not return if it succeeds
        printf("ERROR: Could not execute %s\n", cmd);
        exit(1);
    } else {
        // This is the parent process; Wait for child to finish
        int pid = retval;
        wait(pid);
    }
}
```

Terminate a process

- Programmer terminates:
 - in Unix by `exit()`
 - Windows: `ExitProcess()`
- OS terminates:
 - Bug: access violation, division by zero, ...
- Another process terminates in
 - UNIX by `kill()`, or from shell by **`kill -0 PID`**
 - Windows by `TerminateProcess()`

5. The Process API

1. `fork()`, `wait()` and `exec()` System Call
- 2. `fork()` and `kill()` going bad**
3. `pipe()` and `dup2()`



fork and kill: how badly things can go

bad.c

```
#include <signal.h>
#include <unistd.h>

int main(void) {
    pid_t child = fork();
    if (child) { // in parent
        sleep(5);
        kill(child, SIGKILL);
    } else { // in child
        for (;;) // loop until killed
        }

    return 0;
}
```

man page of kill:

If pid is -1, **sig shall be sent to all processes** (excluding an unspecified set of system processes) for which the process has permission to send that signal.

This program compiles with no errors or warnings, not even with `-Wall -Wextra -Werror`.


5. The Process API

1. `fork()`, `wait()` and `exec()` System Call
2. `fork()` and `kill()` going bad
3. **`pipe()` and `dup2()`**



The pipe() System Call

- `pipe(int fd[])`
 - Ordinary pipe (half duplex or duplex)
 - Typical: Producer -> Consumer



The pipe() System Call: Example

pipe.c

```
int main(void) {
    int fd[2], nbytes;
    pid_t childpid;
    char string[] = "Hello new Process!\n";
    char readbuffer[80];

    pipe(fd);

    if ((childpid = fork()) == -1) {
        perror("fork");
        exit(1);
    }
    ...
}
```

The pipe() System Call: Example

pipe.c

...

```
if (childpid == 0) /* Child */
{
    close(fd[1]); /* Close Output */
    nbytes = read(fd[0], readbuffer, sizeof(string) + 1);
    printf("Received string: %s", readbuffer);
    exit(0);
} else /* Elternprozess */
{
    close(fd[0]); /* Close Input */
    write(fd[1], string, strlen(string) + 1);
    wait(0); // wait for child, so process shutdowns
}
return (0);
}
```

```
> ./pipe
Received string: Hello new Process!
```

pipe() between processes

- Example: `ps ax | grep task1`
 - `ps: stdout` → to pipe
 - `grep: stdin` ← from pipe
- But how to connect `stdout` to `stdin` via pipe?
 - Systemcall: `dup2()`
 - `dup2(fd[1], STDOUT_FILENO)`
 - close `STDOUT ('1')` and reopen bound to write end of pipe
- Programming this Example means:
 - The shell has to wait for `grep task1` to finish
 - `grep task1` has to wait for `ps ax` to finish.

Thanks

Questions

