

Lecture

# Operating System

## 6. Mechanism: Limited Direct Execution



## 6. Mechanism: Limited Direct Execution

- 1. Goal**
- 2. Direct Execution**
- 3. Limited Direct Execution**
- 4. Context Switch**



# 6. Mechanism: Limited Direct Execution

- 1. Goal**
2. Direct Execution
3. Limited Direct Execution
4. Context Switch



# Virtualization of CPU by Time Share

- Goal:
  - Give each process impression it alone is actively using CPU
  - Resources can be shared in **time** and **space**
    - time: time-sharing the CPU
    - space: e.g. memory, disc sharing (later)
  - The OS needs to share the physical CPU by **time sharing**.
    - Issue:
      - *Performance*: How can we implement virtualization without adding excessive overhead to the system?
      - *Control*: How can we run processes efficiently while retaining control over the CPU?



# 6. Mechanism: Limited Direct Execution

1. Goal
- 2. Direct Execution**
3. Limited Direct Execution
4. Context Switch



# Direct Execution

- Just run the program directly on the CPU.

OS	Program
<ol style="list-style-type: none"><li>1. Create entry for process list</li><li>2. Allocate memory for program</li><li>3. Load program into memory</li><li>4. Set up stack with <code>argc / argv</code></li><li>5. Clear registers</li><li>6. Execute call <code>main()</code></li></ol>	<ol style="list-style-type: none"><li>7. Run <code>main()</code></li><li>8. Execute return from <code>main()</code></li></ol>
<ol style="list-style-type: none"><li>9. Free memory of process</li><li>10. Remove from process list</li></ol>	

**Without *limits* on running programs,  
the OS wouldn't be in control of anything and  
thus would be “just a library”**

# Problems of Direct Execution?

1. What if a process wishes to perform some kind of **restricted operation** such as ...
  - Issuing an I/O request to a disk
  - Gaining access to more system resources such as CPU or memory
2. How can the OS **regain control** of the CPU so that it can switch between processes?
  - A cooperative Approach: Wait for system calls
  - A Non-Cooperative Approach: The OS takes control

# Problem 1: Restricted Operation

- Solution: Using **protected control transfer**
  - **User mode:** Applications do not have full access to hardware resources.
  - **Kernel mode:** The OS has access to the full resources of the machine
- Switch from User mode to Kernel mode: **System Call**
  - Allow the kernel to carefully expose certain key pieces of functionality to user program, such as ...
    - Accessing the file system
    - Creating and destroying processes
    - Communicating with other processes
    - Allocating more memory



# System Call

- **Trap** instruction
  - Jump into the kernel
  - Raise the privilege level to kernel mode
- **Return-from-trap** instruction
  - Return into the calling user program
  - Reduce the privilege level back to user mode

## 6. Mechanism: Limited Direct Execution

1. Goal
2. Direct Execution
- 3. Limited Direct Execution**
4. Context Switch



# Limited Direction Execution Protocol

## OS@boot (Kernel Mode)

## Hardware

**Initialize trap table**

remember address of syscall handler

...

...

...

## OS@run (Kernel Mode)

## Hardware

## Program (User Mode)

Create entry for process list  
Allocate memory for program  
Load program into memory  
Setup user stack with `argv`  
Fill kernel stack with reg/PC

**return-from -trap**

restore regs from kernel stack  
move to user mode  
jump to main

Run `main()`

...

Call **system trap** into OS

--- (Cont.) ---



# Limited Direction Execution Protocol

OS@run (Kernel Mode)	Hardware	Program (User Mode)
	--- --- --- (Cont.) --- --- ---	
	save regs to kernel stack move to kernel mode jump to trap handler	
Handle trap Do work of syscall <b>return-from-trap</b>		
	restore regs from kernel stack move to user mode jump to PC after trap	
		... return from main trap (via <code>exit()</code> )
Free memory of process Remove from process list		

## Problem 2: Switching Between Processes

- How can the OS **regain control** of the CPU so that it can switch between processes?
  - A cooperative Approach: **Wait for system calls**
  - A Non-Cooperative Approach: **The OS takes control**

# Wait for system calls

- Processes **periodically give up the CPU** by making system calls such as `yield`.
  - The OS decides to run some other task.
  - Application also transfer control to the OS when they do something illegal.
    - Divide by zero
    - Try to access memory that it shouldn't be able to access
- Example:
  - Early versions of the Macintosh OS, The old Xerox Alto system

**A process gets stuck in an infinite loop.**  
→ **Reboot the machine**



# OS Takes Control

### ■ A timer interrupt

- During the boot sequence, the OS start the timer.
- The timer **raise an interrupt** every so many milliseconds.
- When the interrupt is raised :
  - The currently running process is halted.
  - Save enough of the state of the program
  - A pre-configured interrupt handler in the OS runs.

**A timer interrupt gives OS the ability to run again on a CPU.**

# Saving and Restoring Context

- **Scheduler** makes a decision:
  - Whether to continue running the **current process**, or switch to a **different one**.
  - If the decision is made to switch, the OS executes **context switch**.

## 6. Mechanism: Limited Direct Execution

1. Goal
2. Direct Execution
3. Limited Direct Execution
- 4. Context Switch**





# Context Switch

- A low-level piece of assembly code
  - **Save a few register values** for the current process onto its kernel stack
    - General purpose registers
    - PC
    - kernel stack pointer
  - **Restore a few register values** for the *soon-to-be-executing* process from its kernel stack
  - **Switch to the kernel stack** for the *soon-to-be-executing* process

# Limited Direction Execution Protocol

## Timer Interrupt

OS@boot (Kernel Mode)	Hardware
<b>Initialize trap table</b>	remember address of syscall handler and timer handler
<b>start interrupt timer</b>	start timer interrupt CPU in X ms
...	...                      ...

# Limited Direction Execution Protocol

## Timer Interrupt

OS@run (Kernel Mode)	Hardware	Program (User Mode)
		Process A ....
	<b>timer interrupt</b> save regs(A) to k-stack(A) move to kernel mode jump to trap handler	
Handle the trap Call switch() routine save regs(A) to proc-struct(A) restore regs(B) from proc-struct(B) switch to k-stack(B) <b>return-from-trap (into B)</b>		
	restore regs(B) from k-stack(B) move to user mode jump to B's PC	
		Process B ...



## 2 Stacks per Process


- **Every** process has 2 different stacks
  - **User stack**
    - User code is executed using the normal 'stack' called user stack
  - **Kernel stack**
    - If a trap to a syscall is done, the code of the OS runs.
    - This OS code is executed using the kernel stack

# Example Context Switch

- Running in user mode, SP points to user stack

**Representation of Kernel  
Stack (Memory)**

**SP -> UserStack**



## Example Context Switch

- Take an exception, syscall, or interrupt, and we switch to the kernel stack



# Example Context Switch

- We push a trapframe on the stack
  - Also called exception frame, user-level context....
  - Includes the user-level PC and SP





# Example Context Switch

- Call 'C' code to process syscall, exception, or interrupt
- Results in a 'C' activation stack building up, since code of the OS is executed



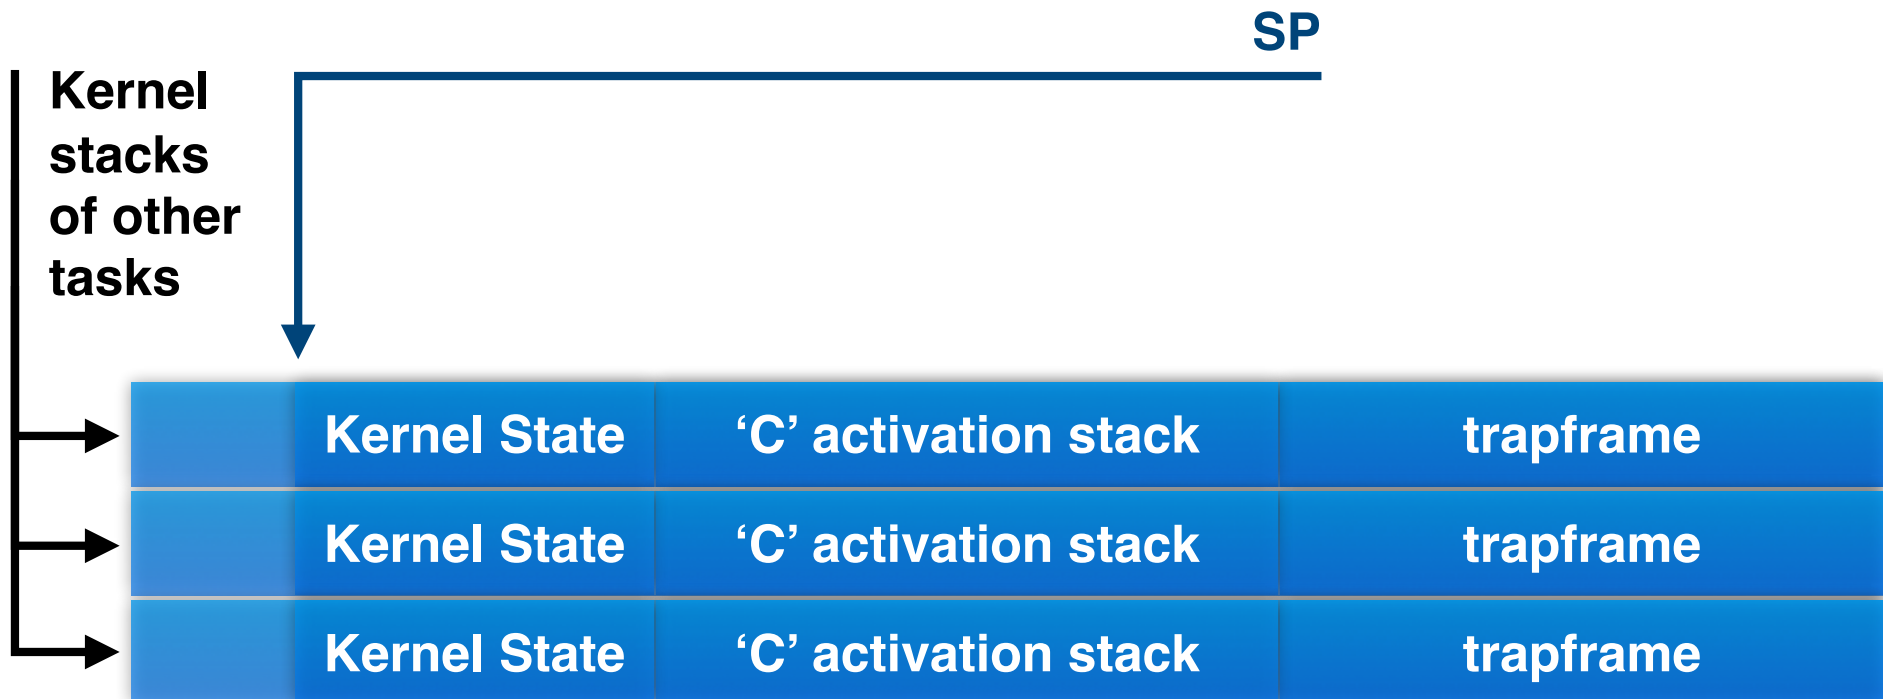
# Example Context Switch

- Not the kernel decides to perform a context switch
- It chooses a target task (e.g. process)
- It pushes remaining kernel context onto the stack



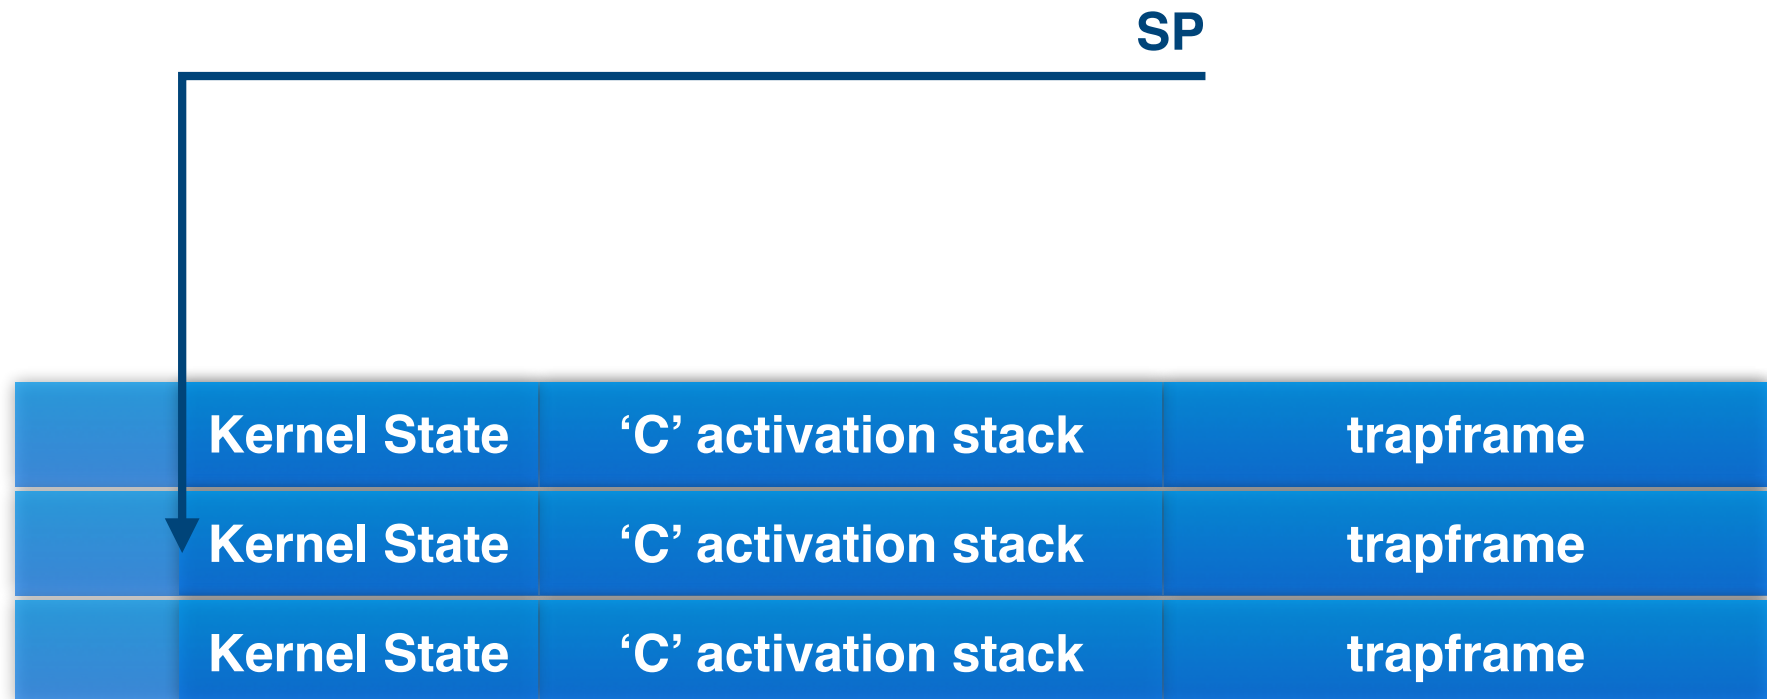
# Example Context Switch

- Any other existing task must
  - be in kernel mode (on a uni processor) and
  - have a similar stack layout to the stack we are currently using



# Example Context Switch

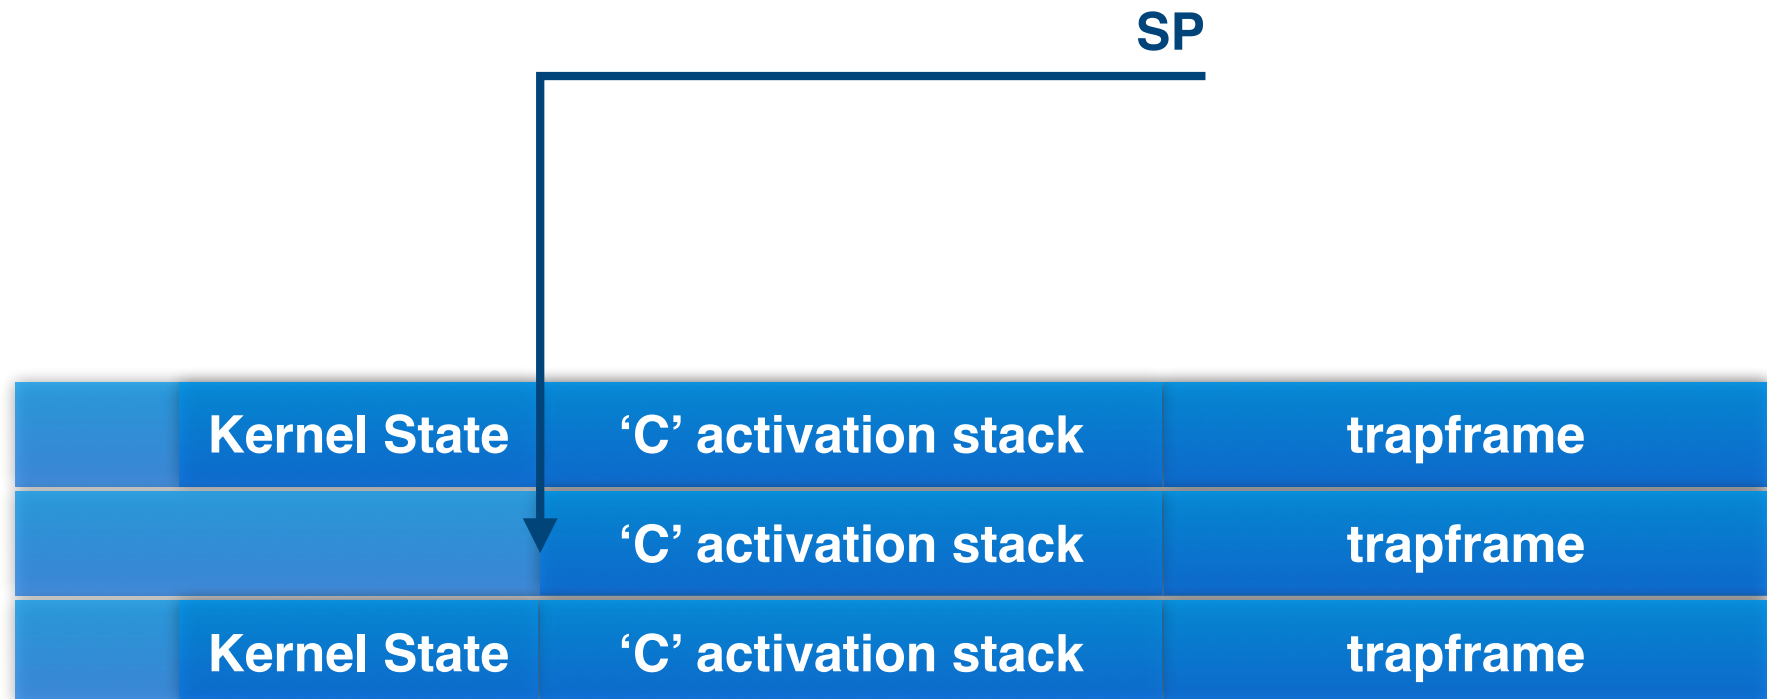
- We save the current SP in the PCB (or TCB) and
- load the SP of the target thread.
- Thus we have switched contexts





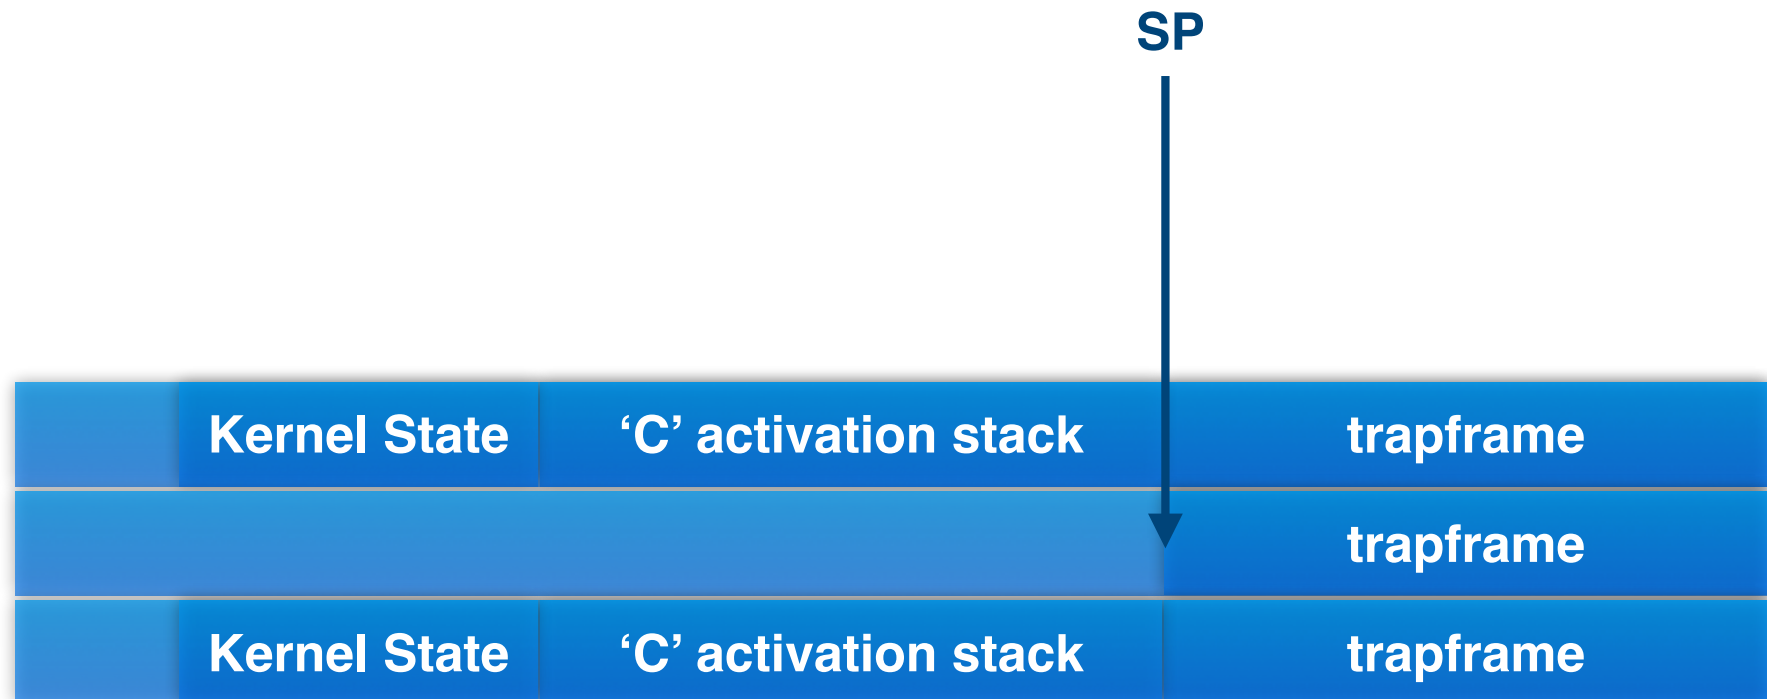
# Example Context Switch

- Load the target thread's previous context, and return to C



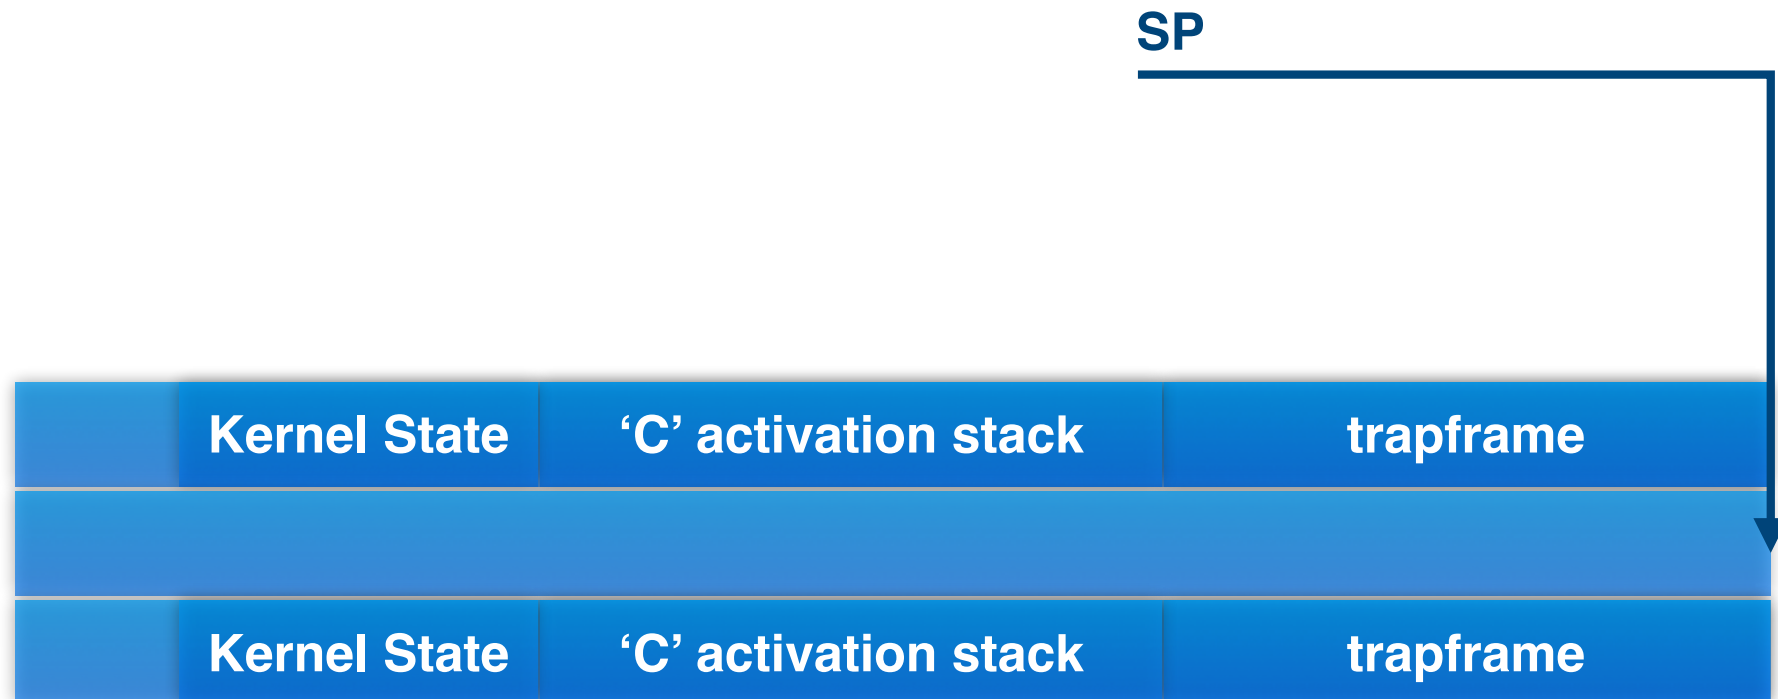
## Example Context Switch

- Load the target thread's previous context, and return to C
  - C-Code of the operating system will be finished, where we left it because of an different context switch (long time before ...)



# Example Context Switch

- The user-level context is restored
  - Trapframe is saved back, preparing to run the user task



# Example Context Switch

- SP points to user stack
  - The process continues as if nothing had happened

**SP -> UserStack**





# Worried About Concurrency?

- What happens if, during interrupt or trap handling, another interrupt occurs?
- OS handles these situations:
  - **Disable interrupts** during interrupt processing
  - Use a number of sophisticated **locking** schemes to protect concurrent access to internal data structures.