

Operating System

7. Scheduling: Introduction

7. Scheduling: Introduction

1. **Scheduler: Policy to determine which process gets CPU when**
2. **Policies Examples**

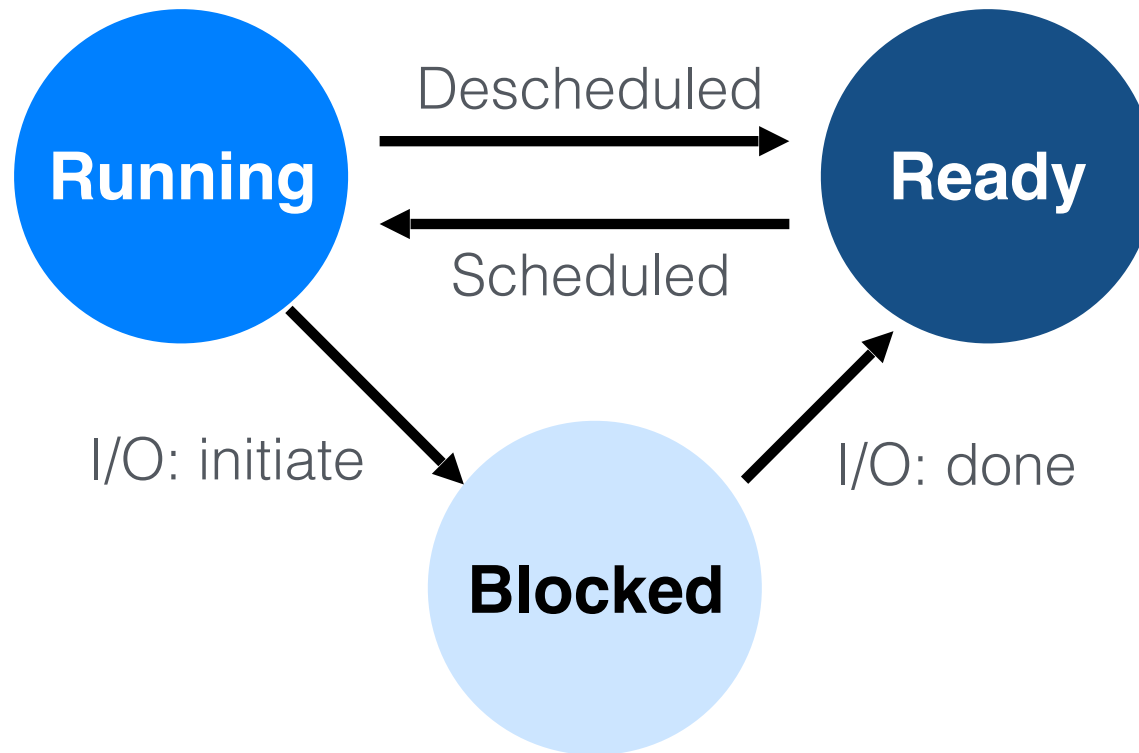


7. Scheduling: Introduction

- 1. Scheduler: Policy to determine which process gets CPU when**
2. Policies Examples



State Transitions



How to transition? ("mechanism")

When to transition? ("policy")

Vocabulary

- **Workload**: set of job descriptions (arrival time, run_time)
 - Job: View as current CPU burst of a process
 - Process alternates between CPU and I/O
process moves between ready and blocked queues
- **Scheduler**: logic that decides which ready job to run
- **Metric**: measurement of scheduling quality

Workload Assumptions

1. Each job runs for the **same amount of time**.
2. All jobs **arrive** at the same time.
3. All jobs only use the **CPU** (i.e., they perform no I/O).
4. The **run-time** of each job is known.

Scheduling Metrics

- Performance metric: **Turnaround time**
 - The time at which the **job completes** minus the time at which the **job arrived** in the system.

$$T_{turnaround} = T_{completion} - T_{arrival}$$

- Another metric is **fairness**.
 - Performance and fairness are often at odds in scheduling.

More Metrics

■ **Minimize Response Time**

- Schedule interactive jobs promptly so users see output quickly

$$T_{response} = T_{firstrun} - T_{arrival}$$

■ **Minimize Waiting Time**

- Do not want to spend much time in Ready queue

■ **Maximize Throughput**

- Want many jobs to complete per unit of time

■ **Maximize Resource Utilization**

- Keep expensive devices busy

■ **Minimize Overhead**

- Reduce number of context switches

7. Scheduling: Introduction

1. Scheduler: Policy to determine which process gets CPU when

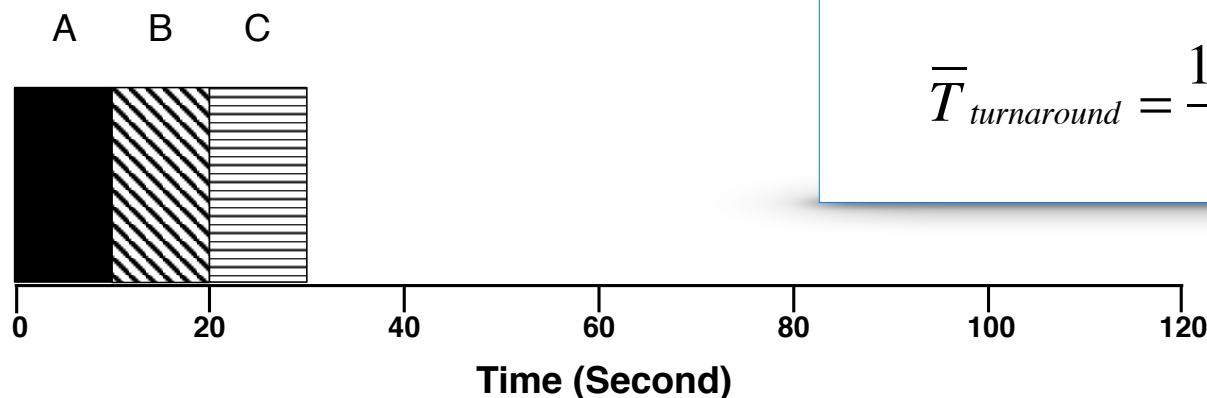
2. Policies Examples



First In, First Out (FIFO)

- First Come, First Served (FCFS)
 - Very simple and easy to implement
 - **Non-preemptive** scheduler
- Example:
 - A arrived just before B which arrived just before C.
 - Each job runs for 10 seconds.

Job	T _{arrival}	T _{runtime}
A	0	10
B	0	10
C	0	10



$$\overline{T}_{turnaround} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

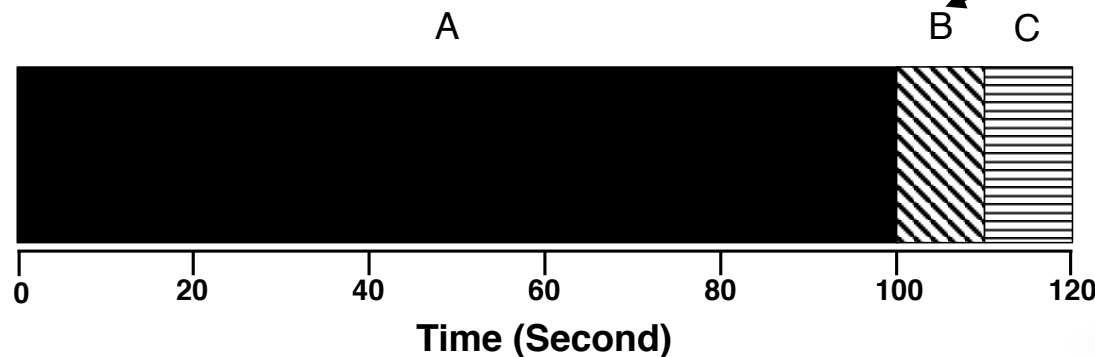
Workload Assumptions

1. ~~Each job runs for the **same amount of time**.~~
2. All jobs **arrive** at the same time.
3. All jobs only use the **CPU** (i.e., they perform no I/O).
4. The **run-time** of each job is known.

Why FIFO is not that great?

- Each job no longer runs for the same amount of time.
- Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.

Job	T_{arrival}	T_{runtime}
A	0	100
B	0	10
C	0	10



Convoy effect

$$\begin{aligned} \overline{T}_{\text{turnaround}} &= \frac{100 + 110 + 120}{3} = 110 \text{ sec} \end{aligned}$$

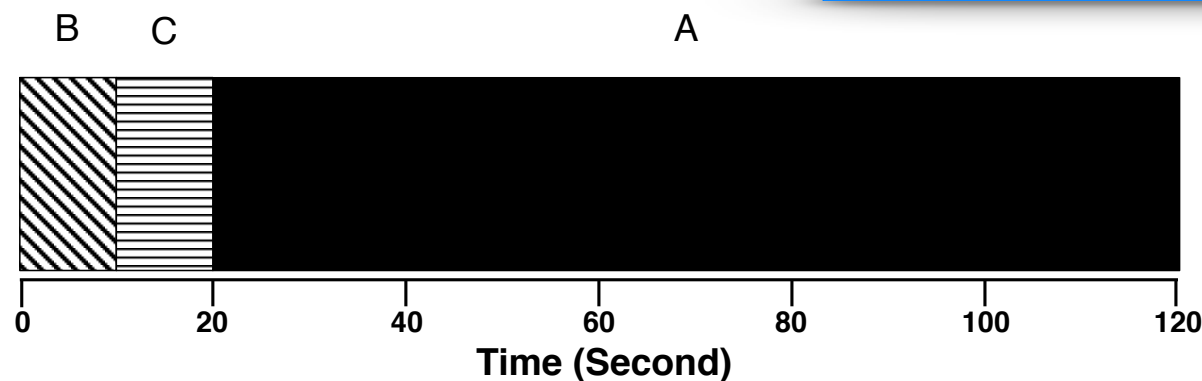
Convoy Effect: How to solve?

- Problem with Previous Scheduler:
 - FIFO: Turnaround time can suffer when short jobs must wait for long jobs
- New scheduler:
 - SJF (Shortest Job First)
 - Choose job with smallest run_time

Shortest Job First (SJF)

- Run the shortest job first, then the next shortest, and so on
 - **Non-preemptive** scheduler
- Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.

$$\overline{T}_{turnaround} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$



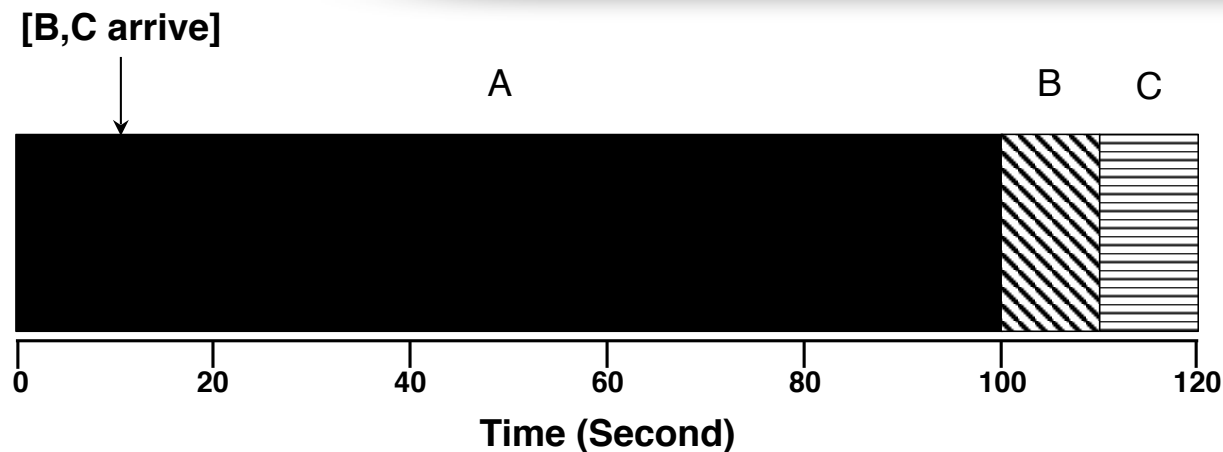
Workload Assumptions

1. ~~Each job runs for the **same amount of time**.~~
2. ~~All jobs **arrive** at the same time.~~
3. All jobs only use the **CPU** (i.e., they perform no I/O).
4. The **run-time** of each job is known.

SJF with Late Arrivals from B and C

- Jobs can arrive at any time.
- Example:
 - A arrives at $t=0$ and needs to run for 100 seconds.
 - B and C arrive at $t=10$ and each need to run for 10 seconds

$$\overline{T}_{turnaround} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103,33\text{sec}$$



Job	$T_{arrival}$	$T_{runtime}$
A	0	100
B	10	10
C	10	10

Shortest Time-to-Completion First (STCF)

- Add **preemption** to SJF
 - Also known as Preemptive Shortest Job First (PSJF)
 - or as Shortest Remaining Time (SRT)
- A new job enters the system:
 - Determine of the remaining jobs and new job
 - Schedule the job which has the less time left

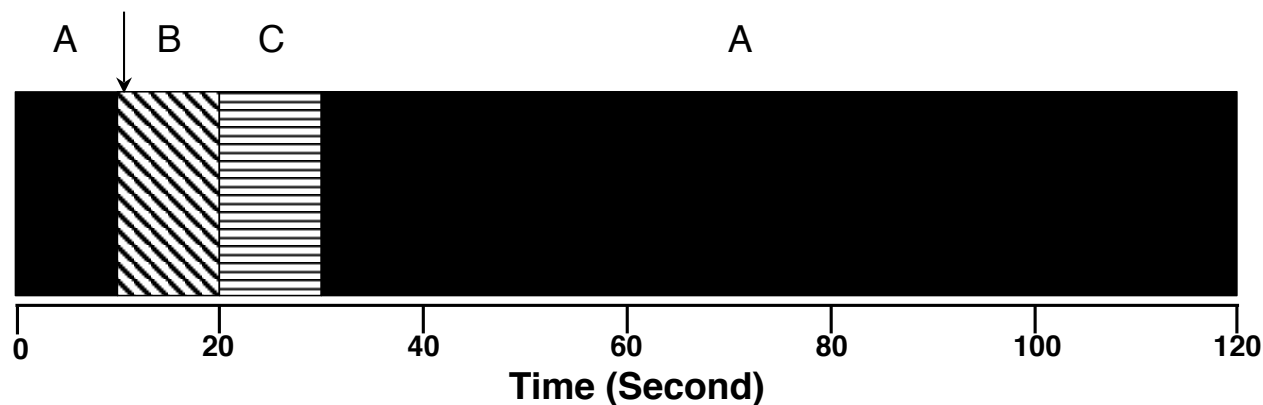
Shortest Time-to-Completion First (STCF)

■ Example:

- A arrives at $t=0$ and needs to run for 100 seconds.
- B and C arrive at $t=10$ and each need to run for 10 seconds

$$\bar{T}_{turnaround} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

[B,C arrive]



Job	$T_{arrival}$	$T_{runtime}$
A	0	100
B	10	10
C	10	10

New scheduling metric: Response time

- The time from when the job arrives to the first time it is scheduled.
- STCF and related disciplines are not particularly good for response time.

$$T_{response} = T_{firstrun} - T_{arrival}$$

How can we build a scheduler that is
sensitive to response time?

Round Robin (RR) Scheduling

- Time slicing Scheduling
 - Run a job for a **time slice** and then switch to the next job in the **run queue** until the jobs are finished.
 - Time slice is sometimes called a [scheduling quantum](#).
 - It repeatedly does so until the jobs are finished.
 - The length of a time slice must be a [multiple](#) of the **timer-interrupt period**.

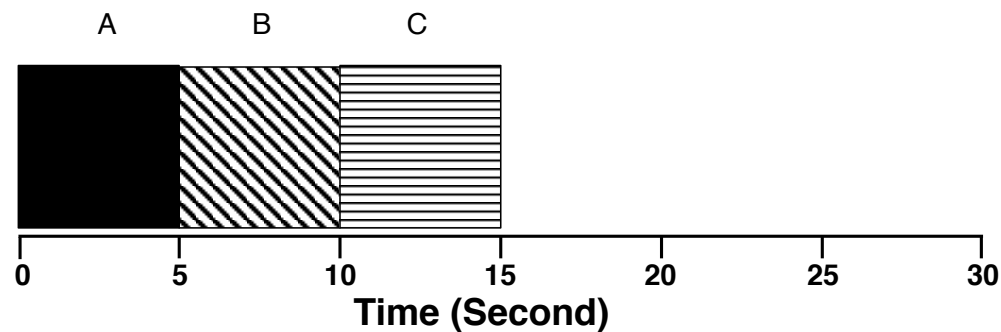
RR is fair, but performs poorly on metrics such as turnaround time

RR Scheduling Example

- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.

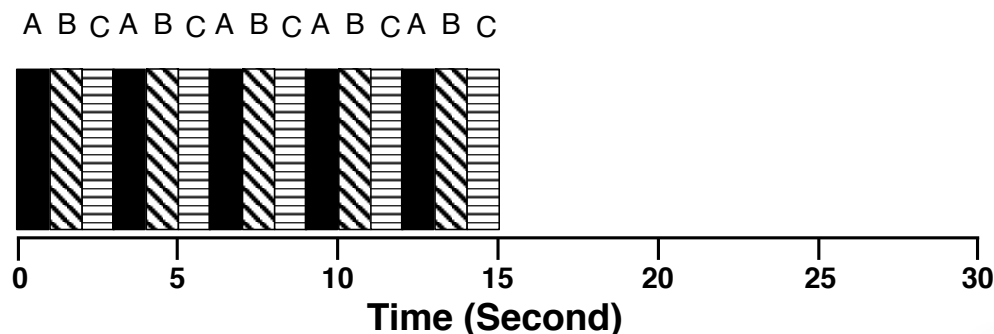
Job	T _{arrival}	T _{runtime}
A	0	5
B	0	5
C	0	5

SJF (Bad for Response Time)



$$\bar{T}_{response} = \frac{0 + 5 + 10}{3} = 5 \text{ sec}$$

RR with a time-slice of 1sec (Good for Response Time)



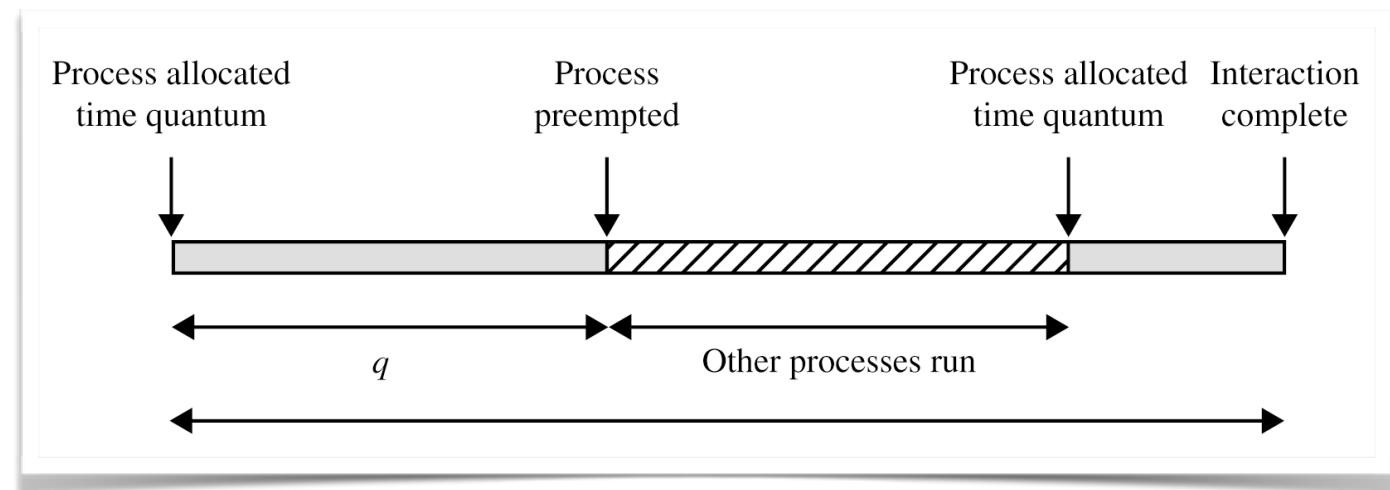
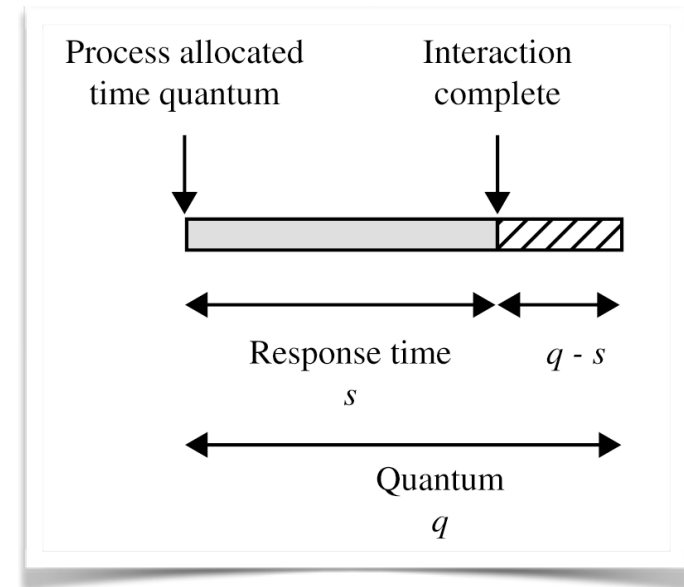
$$\bar{T}_{response} = \frac{0 + 1 + 2}{3} = 1 \text{ sec}$$

The length of time slice is critical

- The **shorter** time slice
 - Better response time
 - The cost of context switching will dominate overall performance.
- The **longer** time slice
 - Amortize the cost of switching
 - Worse response time

The length of time slice is critical

- Deciding on the **length** of the **time slice**
 - presents a trade-off to a system designer.
 - should be compared to typical 'interaction'.



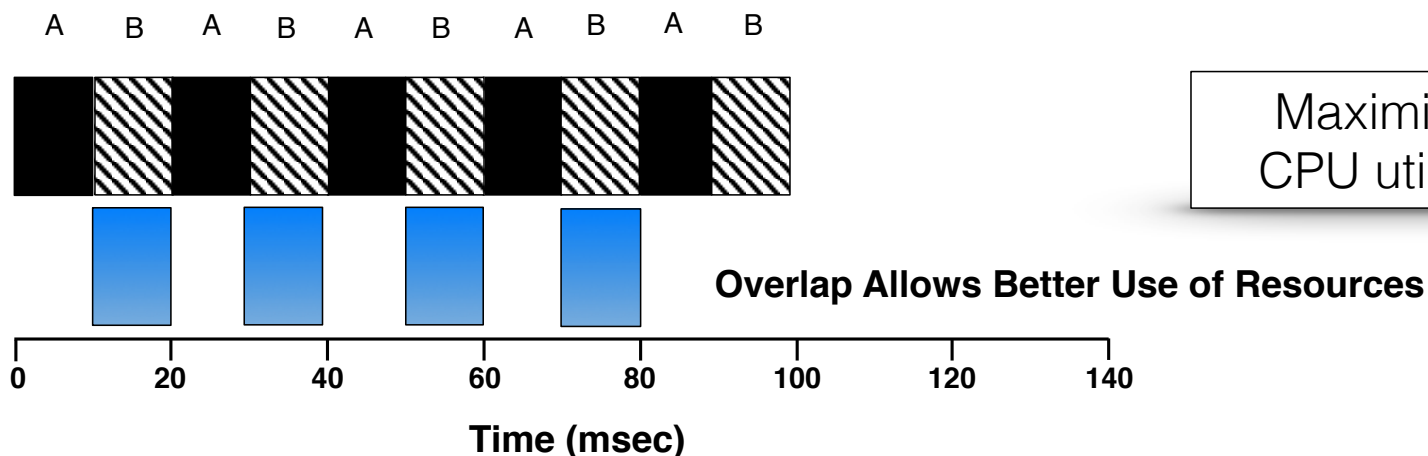
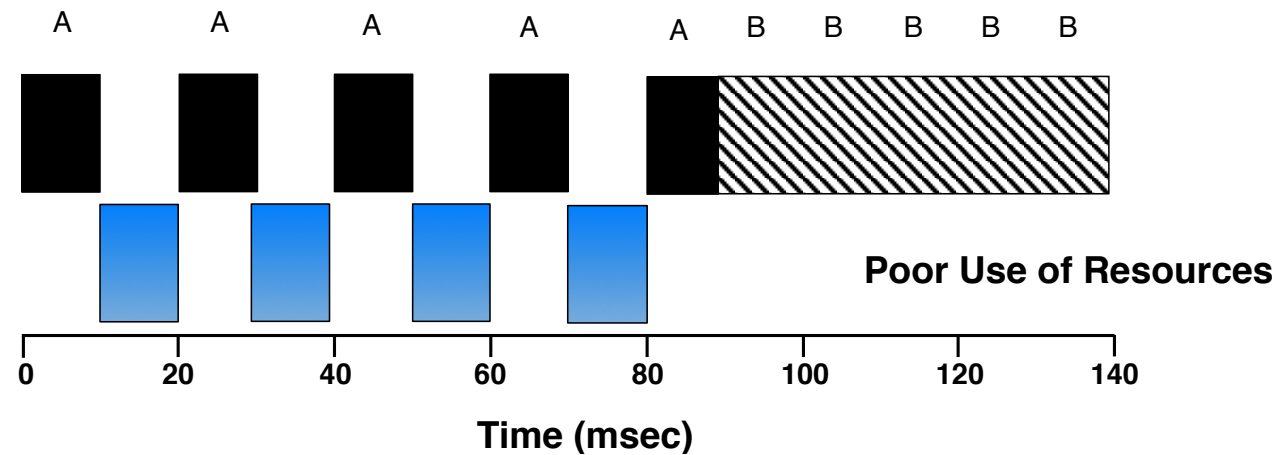
Workload Assumptions

1. ~~Each job runs for the **same amount of time**.~~
2. ~~All jobs **arrive** at the same time.~~
3. ~~All jobs only use the **CPU** (i.e., they perform no I/O).~~
4. The **run-time** of each job is known.

Incorporating I/O

- All programs perform I/O
- Example:
 - A and B need 50ms of CPU time each.
 - **A** runs for 10ms and then issues an **I/O** request
 - I/Os each take 10ms
 - **B** simply uses the CPU for 50ms and performs **no I/O**
 - The scheduler runs A first, then B after

Incorporating I/O (Cont.)



Maximize the
CPU utilization

Incorporating I/O (Cont.)

- When a job **initiates** an I/O request.
 - The job is **blocked** waiting for I/O completion.
 - The scheduler should schedule another job on the CPU.
- When the I/O **completes**
 - An interrupt is raised.
 - The OS moves the process from blocked back to the **ready** state.

Preemptive / Non-Preemptive

- **Non-Preemptive** schedulers:
 - Only schedule new job when previous job voluntarily relinquishes CPU (performs I/O or exits)
 - FIFO and SJF are non-preemptive
- **Preemptive** schedulers
 - Potentially schedule different job at any point by taking CPU away from running job
 - STCF and RR are preemptive
 - STCF: Always run job that will complete the quickest
 - RR: Schedule after Time Slice

The background of the slide features a close-up, slightly blurred image of a clock face. The clock has a white face with black numbers and hands. The numbers 4, 10, 11, 12, 17, 18, 19, 25, 26, and 27 are visible. The clock hands are black, and the center has a small black dot. In the top left corner, there is a small, partially visible calendar grid showing dates and numbers.

Thanks

Questions?