

Operating System

8. Scheduling: The Multi-Level Feedback Queue

8. Scheduling: The Multi-Level Feedback Queue

Goal: general-purpose scheduling

Must support two job types with distinct goals:

- **interactive** programs care about response time
- **batch** programs care about turnaround time



Workload Assumptions

1. ~~Each job runs for the **same amount of time**.~~
2. ~~All jobs **arrive** at the same time.~~
3. ~~All jobs only use the **CPU** (i.e., they perform no I/O).~~
4. ~~The **run-time** of each job is known.~~

History

- Use past behavior of process to predict future behavior
 - Common technique in systems
- Processes alternate between **I/O** and **CPU** work
- Guess how CPU burst (job) will behave based on past CPU bursts (jobs) of this process

Multi-Level Feedback Queue (MLFQ)

- A Scheduler that learns from the past to predict the future.
- Objective:
 - Optimize **turnaround time** → Run shorter jobs first
 - Minimize **response time** without *a priori knowledge of job length*.

MLFQ: Basic Rules

- MLFQ has a number of distinct **queues**.
 - Each queue is assigned a *different priority level*.
- A job that is ready to run is on a single queue.
 - A job **on a higher queue** is chosen to run.
 - Use round-robin scheduling among jobs in the same queue

Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).

Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.

MLFQ: How to Change Priority

- MLFQ varies the priority of a job based on **its observed behavior**.

Rule 3: When a job enters the system, it is placed at the highest priority


Rule 4a: If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down on queue).

Rule 4b: If a job gives up the CPU before the time slice is up, it stays at the same priority level

In this manner, MLFQ approximates SJF

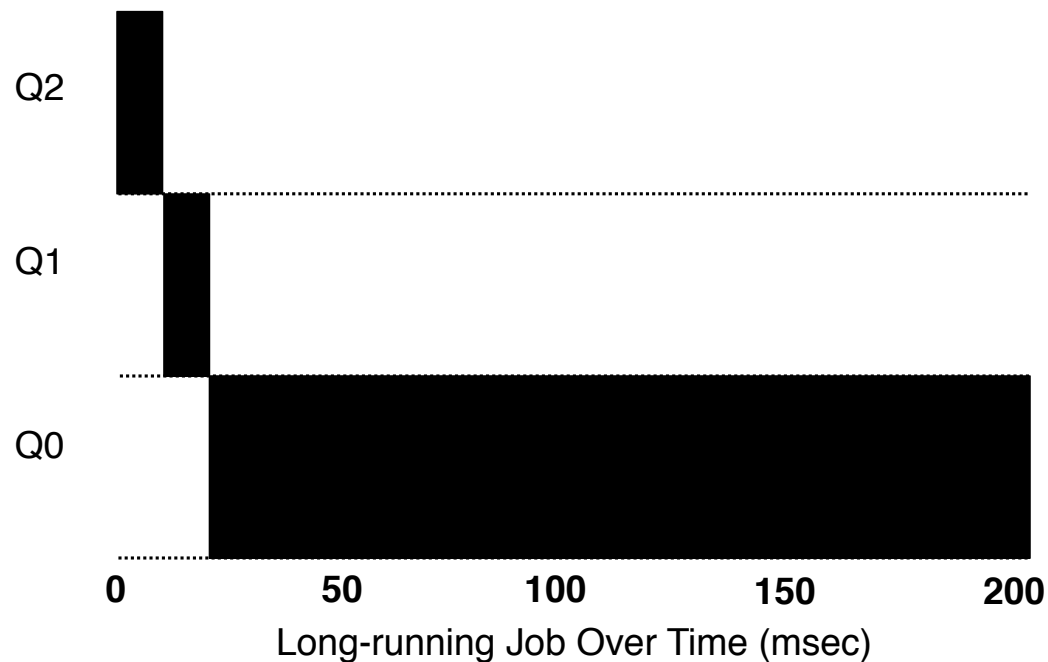
MLFQ: Overview

- A job repeatedly relinquishes the **CPU** while waiting **IOs**
→ Keep its priority *high*
- A job uses the **CPU** intensively for long periods of time
→ *Reduce* its priority.



Example 1: A Single Long-Running Job

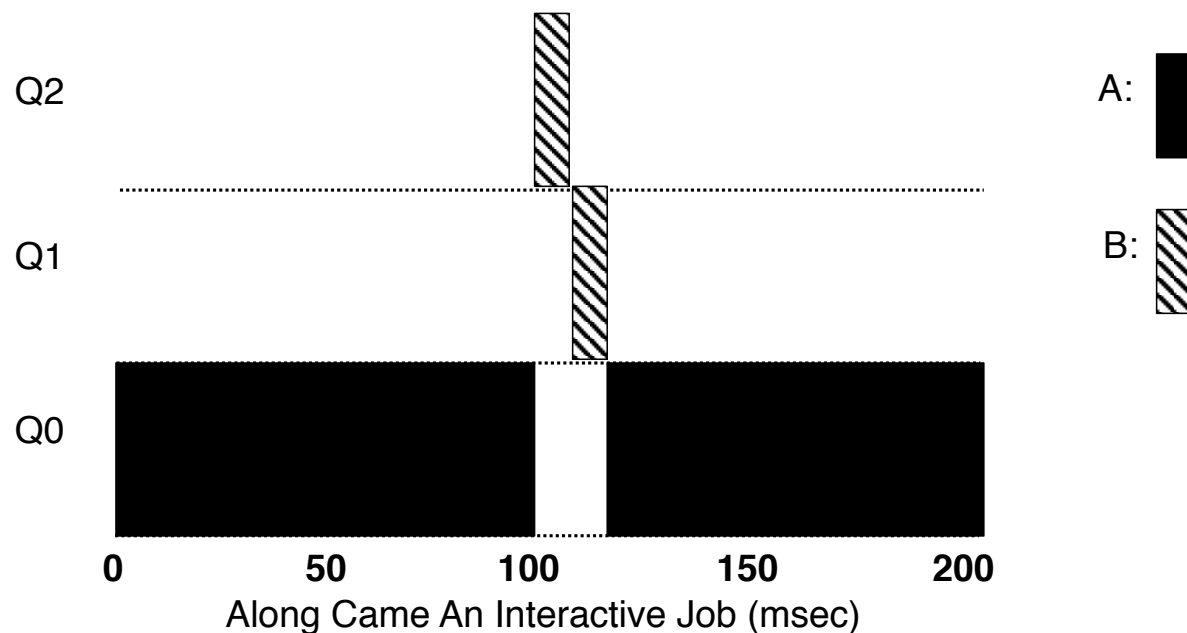
- A three-queue scheduler with time slice 10ms



Example 2: Along Came a Short Job

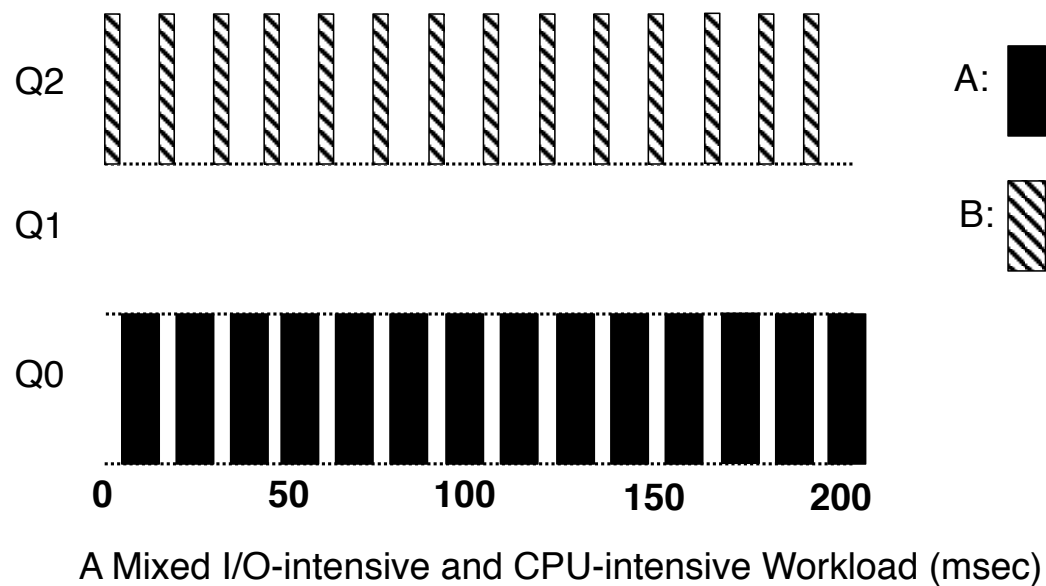
- Assumption:

- **Job A:** A long-running CPU-intensive job
- **Job B:** A short-running interactive job (20ms runtime)
- A has been running for some time, and then B arrives at time $T=100$.



Example 3: What About I/O?


- Assumption:
 - **Job A:** A long-running **CPU**-intensive job
 - **Job B:** An interactive job that need the CPU only for 1ms before performing an **I/O**




The MLFQ approach keeps an interactive job at the highest priority


Taskset Analysis

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2



MLFQ Analysis (q=1)



MLFQ Analysis ($q=1$)

	A	B	C	D	E	Mean
Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time (T_s)	3	6	4	5	2	
Finish Time	4	20	16	19	11	
Turnaround Time (T_r)	4	18	12	13	3	10.00
T_r/T_s	1.33	3.00	3.00	2.60	1.5	2.29

Problems with the Basic MLFQ

■ **Starvation**

- If there are “too many” interactive jobs in the system.
- Lon-running jobs will never receive any CPU time.



■ **Game** the scheduler



- After running 99% of a time slice, issue an I/O operation.
- The job gain a higher percentage of CPU time.
- A program may **change its behavior** over time.
 - **CPU** bound process → **I/O** bound process

The Priority Boost

■ Example:

- A long-running job(A) with two short-running interactive job(B, C)

A:  B:  C: 




Rule 5: After some time period S , move all the jobs in the system to the topmost queue.

Better Accounting


- How to prevent gaming of our scheduler?

Rule 4 (Rewrite 4a and 4b): Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down on queue).



Tuning MLFQ And Other Issues

- The high-priority queues → Short time slices
 - E.g., 10 or fewer milliseconds
- The Low-priority queue → Longer time slices
 - E.g., 100 milliseconds



Lower Priority,
Longer Quanta

MLFQ: Summary

Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).

Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.

Rule 3: When a job enters the system, it is placed at the highest priority

Rule 4: Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down on queue).

Rule 5: After some time period S, move all the jobs in the system to the topmost queue.

The background of the slide features a close-up, slightly blurred image of a clock face. The clock has a white dial with black numbers and hands. In the upper left corner, a portion of a calendar grid is visible, showing dates and days of the week. Overlaid on the left side of the clock face are two white rectangular boxes containing text.

Thanks

Questions?