

# Operating System

## 9. Scheduling: Proportional Share

## 9. Scheduling: Proportional Share

**How to share CPU in a promotional manner?**

**What are the key mechanism for doing so?**

**How effective are they?**



# Proportional Share Scheduler

- **Fair-share** scheduler
  - Guarantee that **each job** obtain a certain **percentage** of CPU time.
  - **Not optimized** for turnaround or response time

# Basic Concept

### ■ Tickets

- **Represent the share** of a resource that a process should receive
- The percent of tickets represents its share **of the system resource** in question.

### ■ Example

- There are two processes, A and B.
  - Process A has 75 tickets → receive 75% of the CPU
  - Process B has 25 tickets → receive 25% of the CPU

## 9. Scheduling: Proportional Share

**1. Lottery Scheduling**

**2. Stride Scheduling**



# 9. Scheduling: Proportional Share

## 1. Lottery Scheduling

## 2. Stride Scheduling



# Lottery scheduling

- The scheduler **randomly** picks a winning ticket.
  - Load the state of that winning process and runs it.
- Example
  - There are 100 tickets
    - Process A has 75 tickets: 0 ~ 74
    - Process B has 25 tickets: 75 ~ 99

Scheduler's winning tickets: 63 85 70 39 76 17 29 41 36 39 10 99 68 83 63

Resulting scheduler: A B A A B A A A A A A B A B A

The longer these two jobs compete,  
The more likely they are to achieve the desired percentages.



# Ticket Mechanisms

### ■ Ticket currency

- A user allocates tickets among their own jobs in whatever currency they would like.
- The system converts the currency into the correct global value.
- Example
  - There are 200 tickets (Global currency): Process A has 100 tickets and Process B has 100 tickets

**User A**       $\rightarrow 500$  (A's currency) to A1  $\rightarrow 50$  (global currency)  
                  $\rightarrow 500$  (A's currency) to A2  $\rightarrow 50$  (global currency)

**User B**       $\rightarrow 10$  (B's currency) to B1  $\rightarrow 100$  (global currency)



## Ticket Mechanisms (Cont.)

### ■ **Ticket transfer**


- A process can temporarily hand off its tickets to another process.

### ■ **Ticket inflation**

- A process can temporarily raise or lower the number of tickets it owns.
- If any one process needs more CPU time, it can boost its tickets.

# Implementation

- Example: There are three processes, A, B, and C.
  - Keep the processes in a list:



- Algorithm
  - Pick winning ticket
  - walk through list and update a counter by sum up **Tix**
    - until sum is  $>$  winning ticket  $\rightarrow$  Job found

# Implementation (Cont.)

```
1  // counter: used to track if we've found the winner yet
2  int counter = 0;
3
4  // winner: use some call to a random number generator to
5  // get a value, between 0 and the total # of tickets1
6  int winner = getrandom(0, totaltickets);
7
8  // current: use this to walk through the list of jobs
9  node_t *current = head;
10
11 // loop until the sum of ticket values is > the winner
12 while (current) {
13     counter = counter + current->tickets;
14     if (counter > winner)
15         break; // found the winner
16     current = current->next;
17 }
18 // 'current' is the winner: schedule it...
```

1: <http://stackoverflow.com/questions/2509679/how-to-generate-a-random-number-from-within-a-range>

# Lottery Fairness Study

- **$U$** : unfairness metric

- The time the first job completes divided by the time that the second job completes.

- Example:


- There are two jobs, each jobs has runtime 10.
    - First job finishes at time 10
    - Second job finishes at time 20

$$U = \frac{10}{20} = 0.5$$

- **$U$**  will be close to 1 when both jobs finish at nearly the same time.

# Lottery Fairness Study (Cont.)

- There are two jobs.
- Each jobs has the same number of tickets (100).



When the job length is not very long, average unfairness can be quite severe.

## 9. Scheduling: Proportional Share

1. Lottery Scheduling

**2. Stride Scheduling**



# Stride Scheduling

- Problem: How to assign tickets to jobs
  - Randomness (Lottery Scheduling)
    - but no fair if jobs are short
  - By the user?
- Another Solution: Stride of each process
  - $(\text{A large number}) / (\text{the number of tickets of the process})$
  - Example: A large number = 10,000
    - Process A has 100 tickets  $\rightarrow$  stride of A is 100
    - Process B has 50 tickets  $\rightarrow$  stride of B is 200
- A process runs, **increment** a counter(=pass value) for it **by its stride**.
  - Pick the process to run that has the lowest pass value



# A Pseudo Code implementation

```

current = remove_min(queue);           // pick client with minimum pass
schedule(current);                     // use resource for quantum
current→pass += current→stride;        // compute next pass using stride
insert(queue, current);                // put back into the queue

```

Pass(A) (stride=100)	Pass(B) (stride=200)	Pass(C) (stride=40)	Who runs?
0	0	0	A
100	0	0	B
100	200	0	C
100	200	40	C
100	200	80	C
100	200	120	A
200	200	120	C
200	200	160	C
200	200	200	...

If new job enters with pass value 0,  
It will monopolize the CPU!

# Summary

- Proportional-Share Scheduling
  - Lottery (using randomness)
  - Stride (deterministic)
- Conceptually interesting but not wide-spread adaption
  - does not work well with I/O
  - Problem: How to assign tickets?
- MLFQ and similar schedulers do so more gracefully



Thanks

Questions?