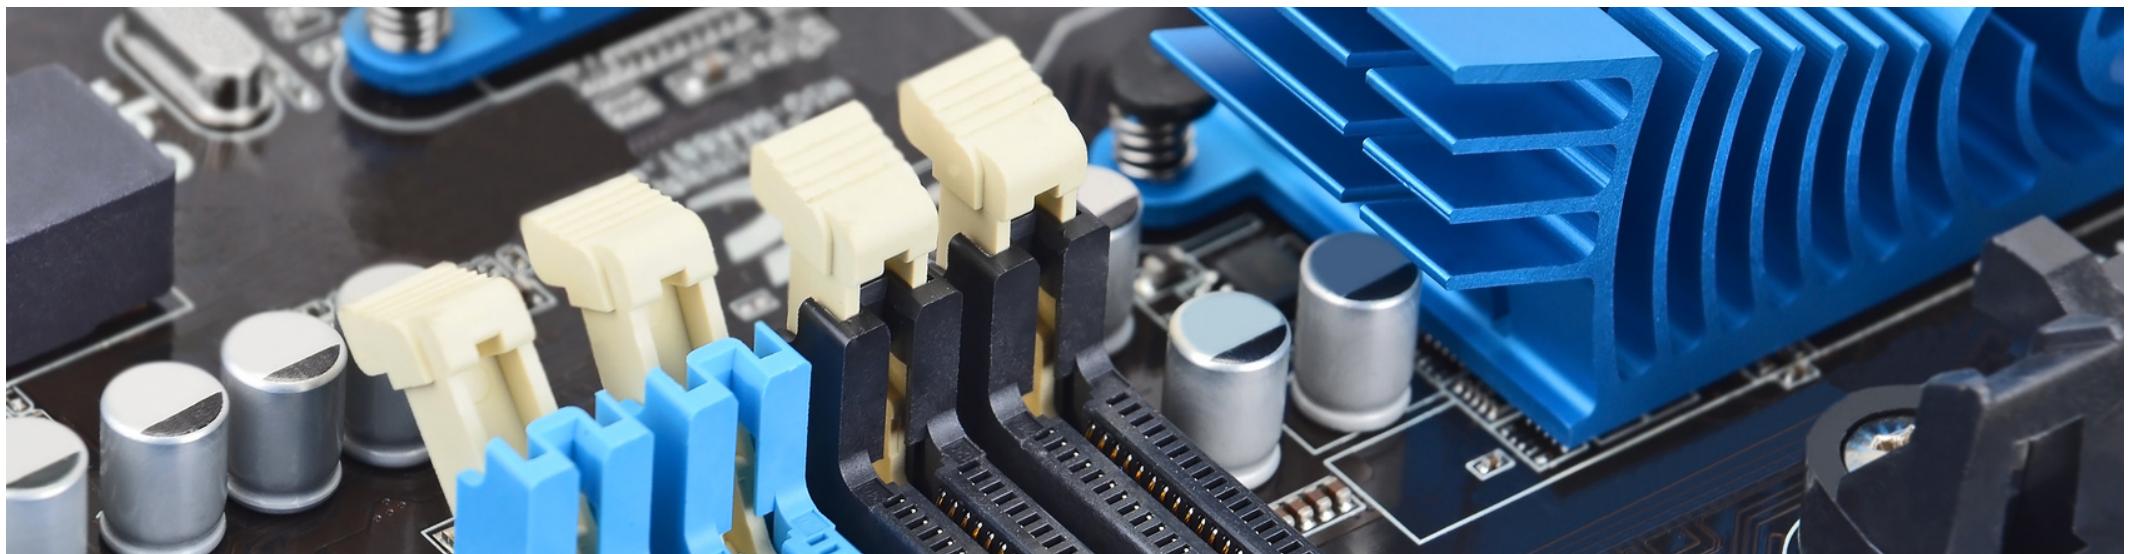


# Lecture Operating System

## 14. Memory API

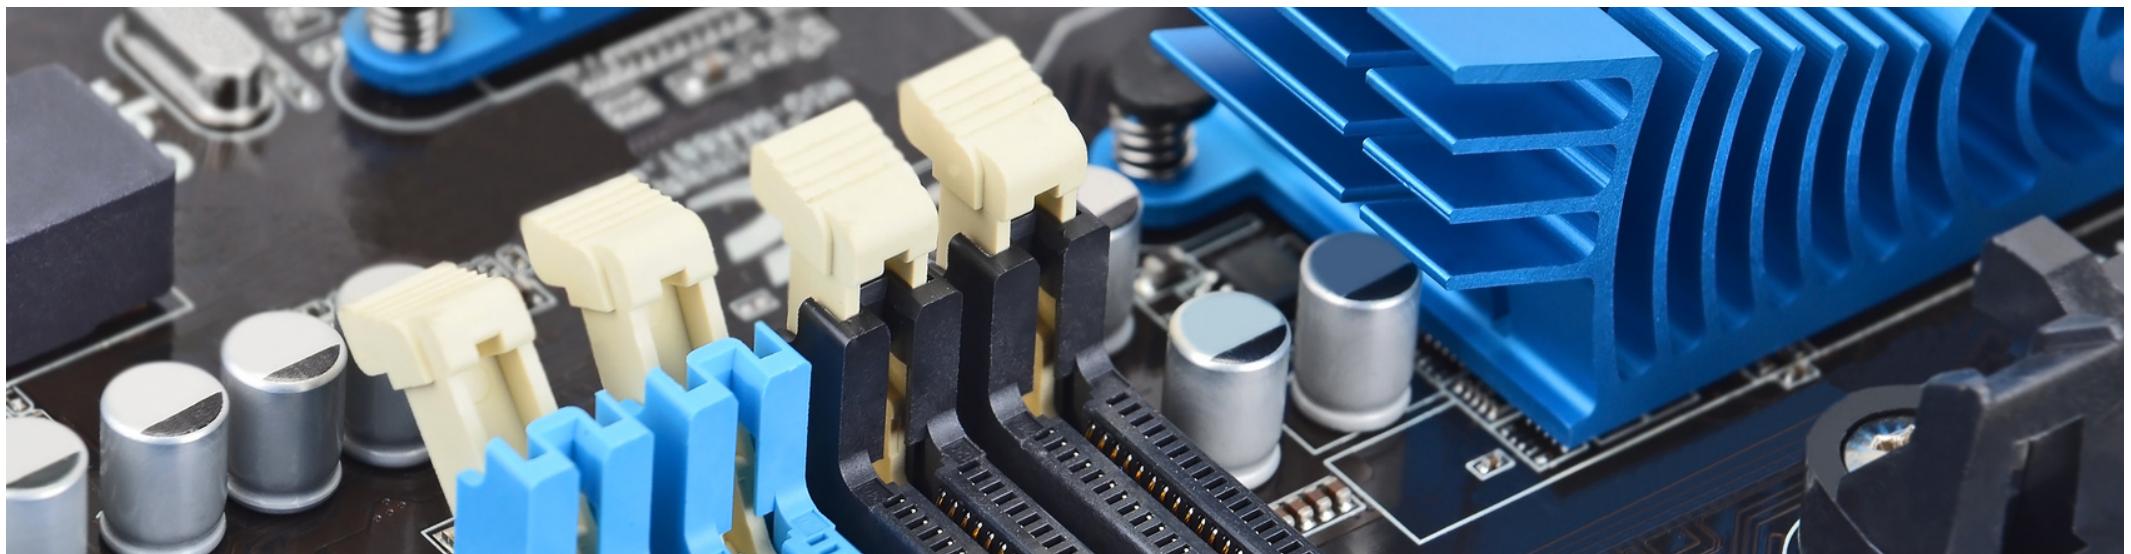
# 14. Memory API

- 1. malloc, sizeof and free**
- 2. Common Errors**
- 3. Underlying OS Support**
- 4. Garbage Collector**



# 14. Memory API

- 1. `malloc`, `sizeof` and `free`**
2. Common Errors
3. Underlying OS Support
4. Garbage Collector



# Memory Region on the stack

- Declaring memory on the stack in C is easy.
- Integer in function
  - Declare int
  - The compiler does the rest
    - After return from the function, the compiler deallocates the memory

```
void func(int argc, char *argv[]){  
    int x = 3; // declares an integer on the stack  
    ...  
}
```

# Memory Region on the heap

- `malloc()`: Allocate a memory region on the heap.
- Argument
  - `size_t size` : size of the memory block(in bytes)
  - `size_t` is an unsigned integer type.
- Return
  - Success : a `void` type pointer to the memory block allocated by `malloc`
  - Fail : a `null` pointer

```
#include <stdlib.h>

void free(void* ptr)
```

# Size of Memory on the heap

- Routines and macros are utilized for `size` in `malloc` instead typing in a number directly.
- Two types of results of `sizeof` with variables
  - The actual size of `x` is known at run-time.
  - The actual size of `x` is known at compile-time.

heap1.c

```
int *x = malloc(10 * sizeof(int));
printf("%d\n", sizeof(x));
```

```
prompt> ./heap1
4
```

heap2.c

```
int x[10];
printf("%d\n", sizeof(x));
```

```
prompt> ./heap2
40
```

# Free memory region on heap

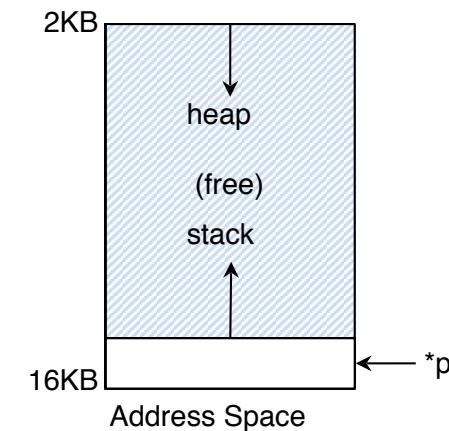
- **free( )**: Free a memory region allocated by a call to **malloc**.
  - Argument
    - **void \*ptr** : a pointer to a memory block allocated with **malloc**
  - Return
    - none

```
#include <stdlib.h>

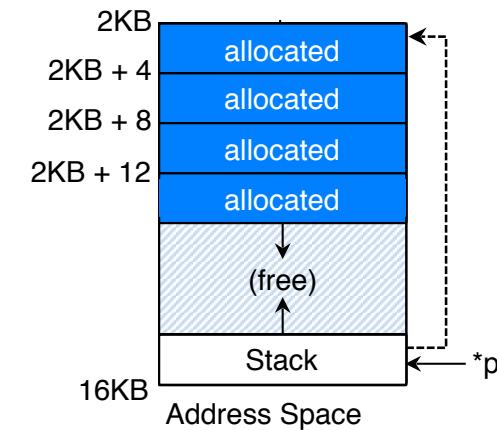
void free(void* ptr)
```

# Memory Allocation

```
int *x; // local variable
```



```
pi = (int *) malloc(sizeof(int)*4);
```



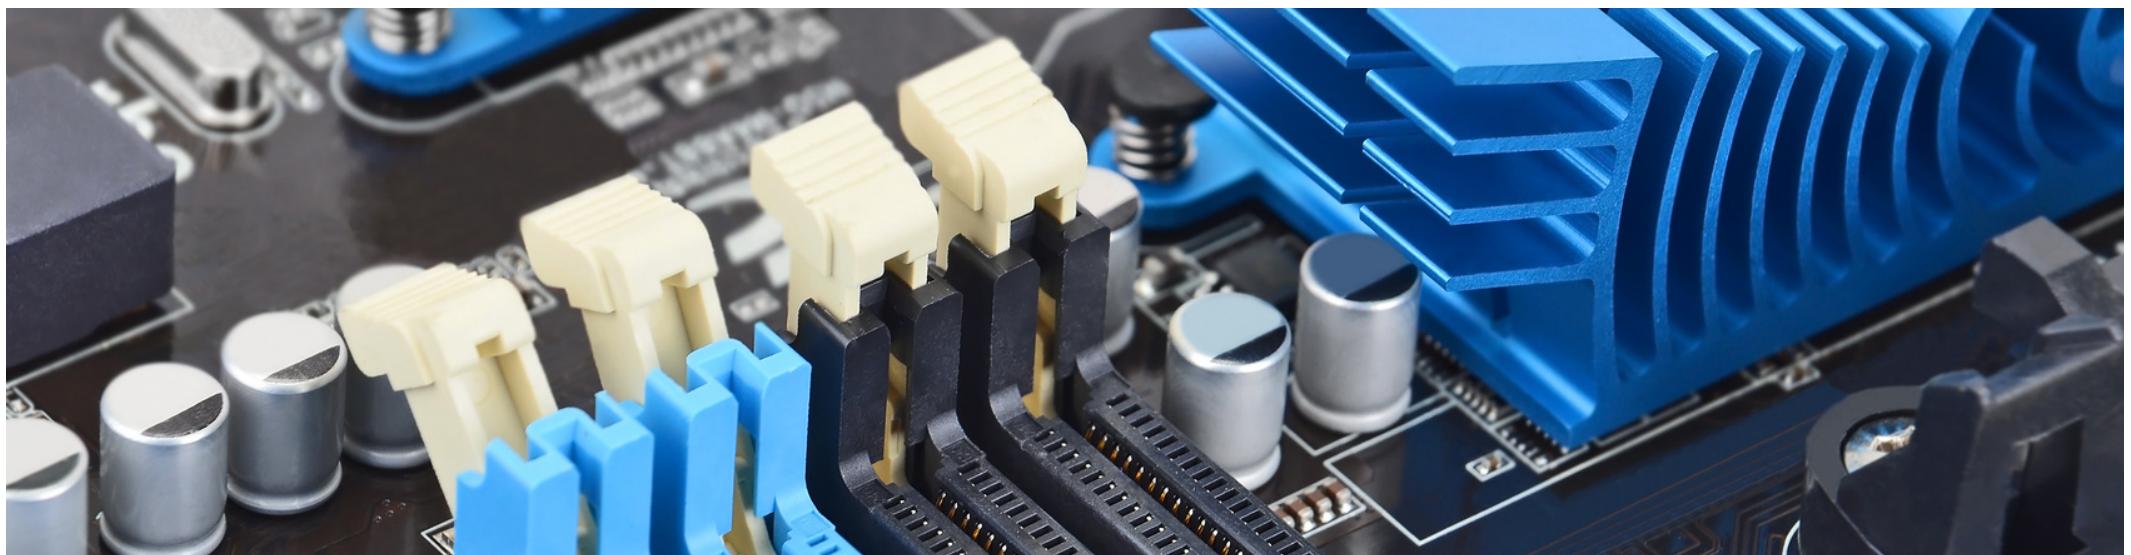
# 14. Memory API

1. `malloc`, `sizeof` and `free`

2. Common Errors

3. Underlying OS Support

4. Garbage Collector

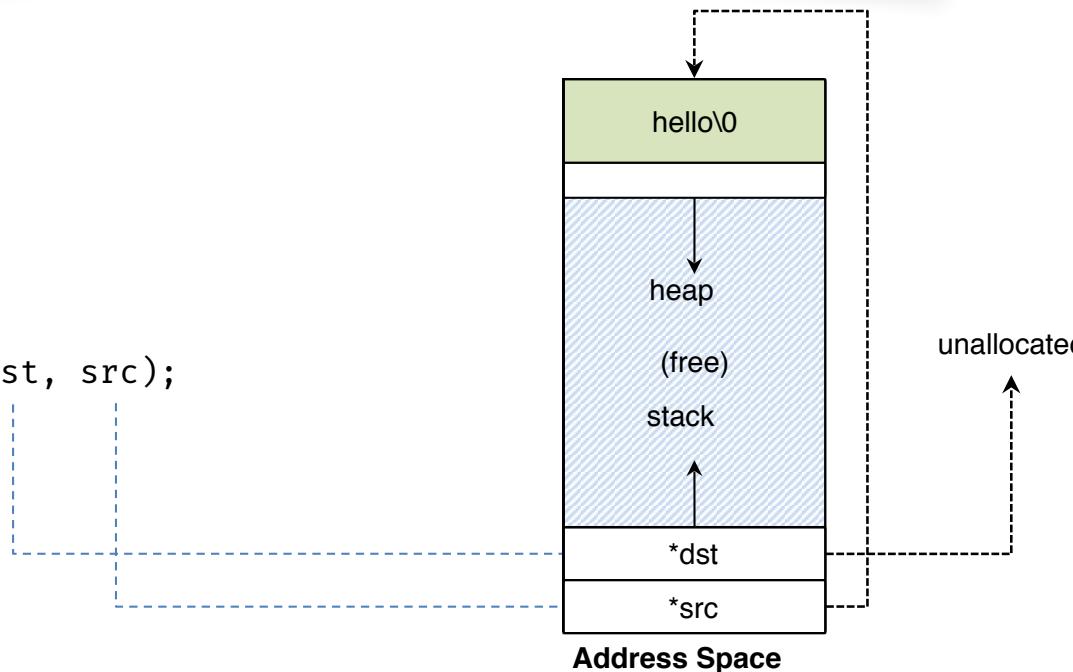


# Forgetting To Allocate Memory



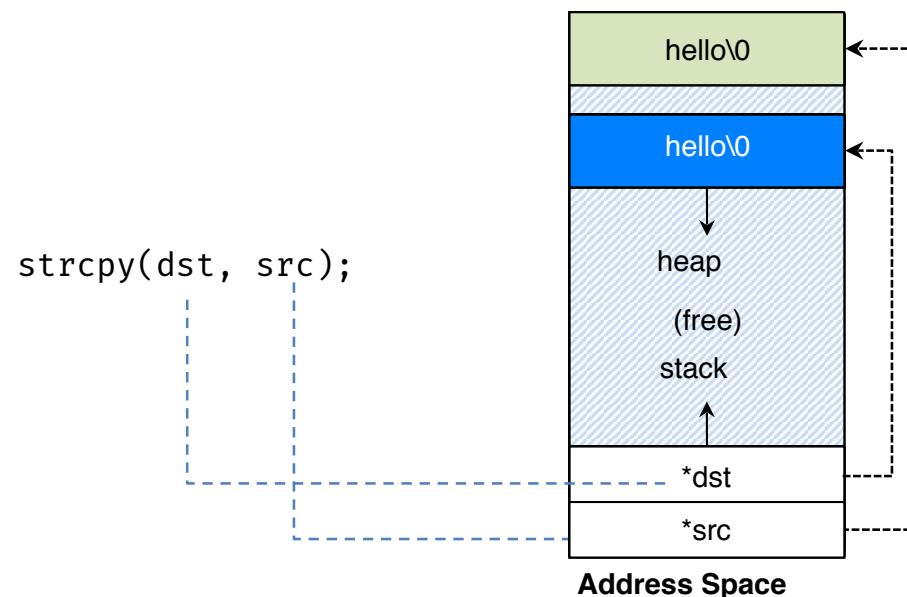
```
char *src = "hello"; //character string constant
char *dst;           //unallocated
strcpy(dst, src);   //segfault and die
```

```
strcpy(dst, src);
```



# NOT Forgetting To Allocate Memory

```
char *src = "hello"; //character string constant  
char *dst = (char *)malloc(strlen(src) + 1); // allocated  
strcpy(dst, src); //work properly
```

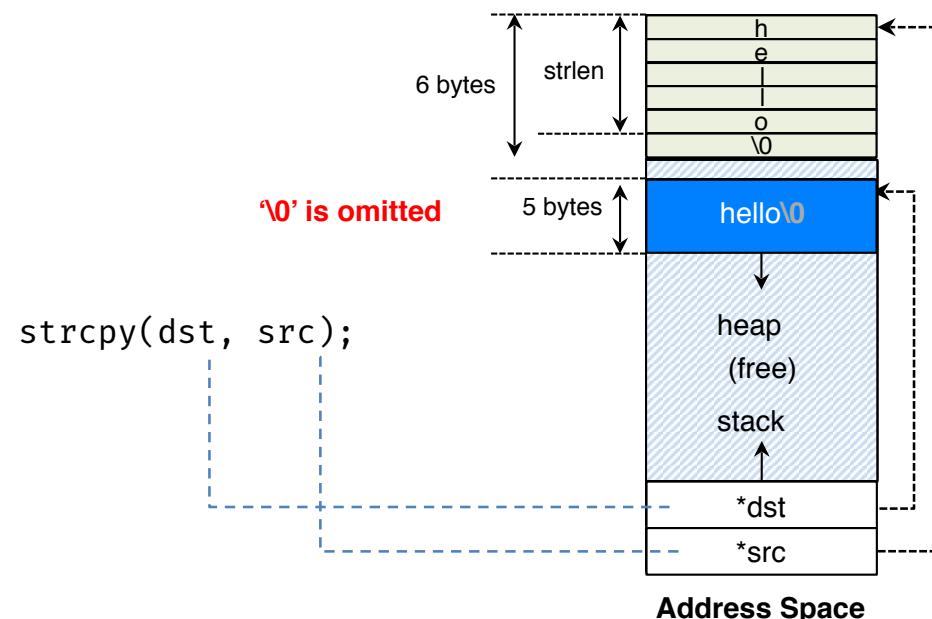


# Not Allocating Enough Memory

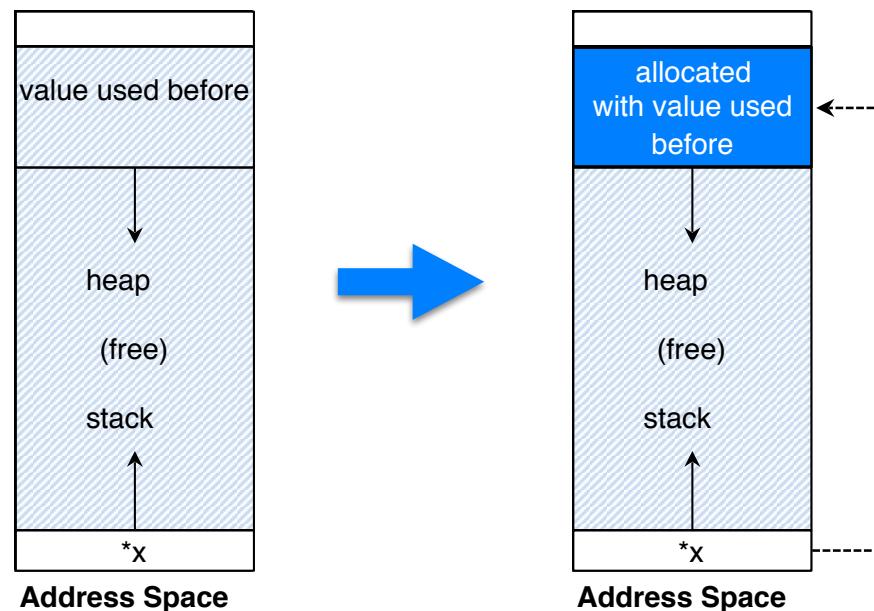
BUG

- Incorrect code, but strcpy works properly

```
char *src = "hello"; //character string constant
char *dst = (char *)malloc(strlen(src)); // too small
strcpy(dst, src); //work properly
```

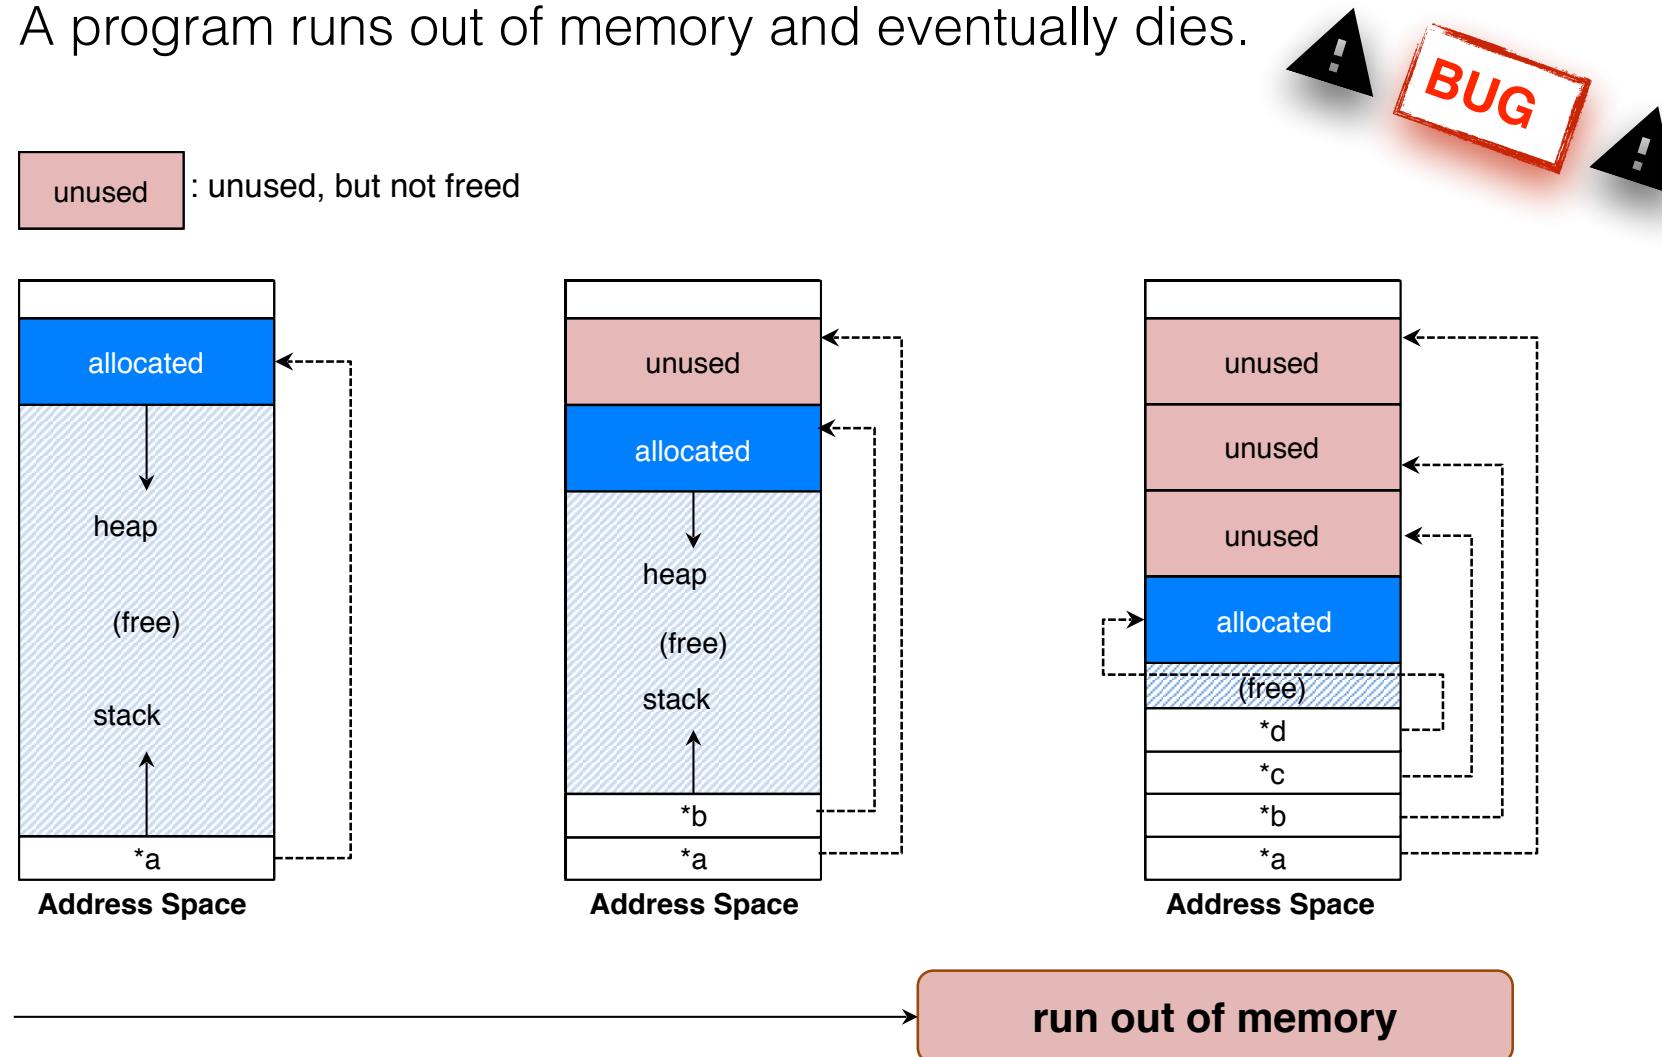


# Forgetting To Initialize



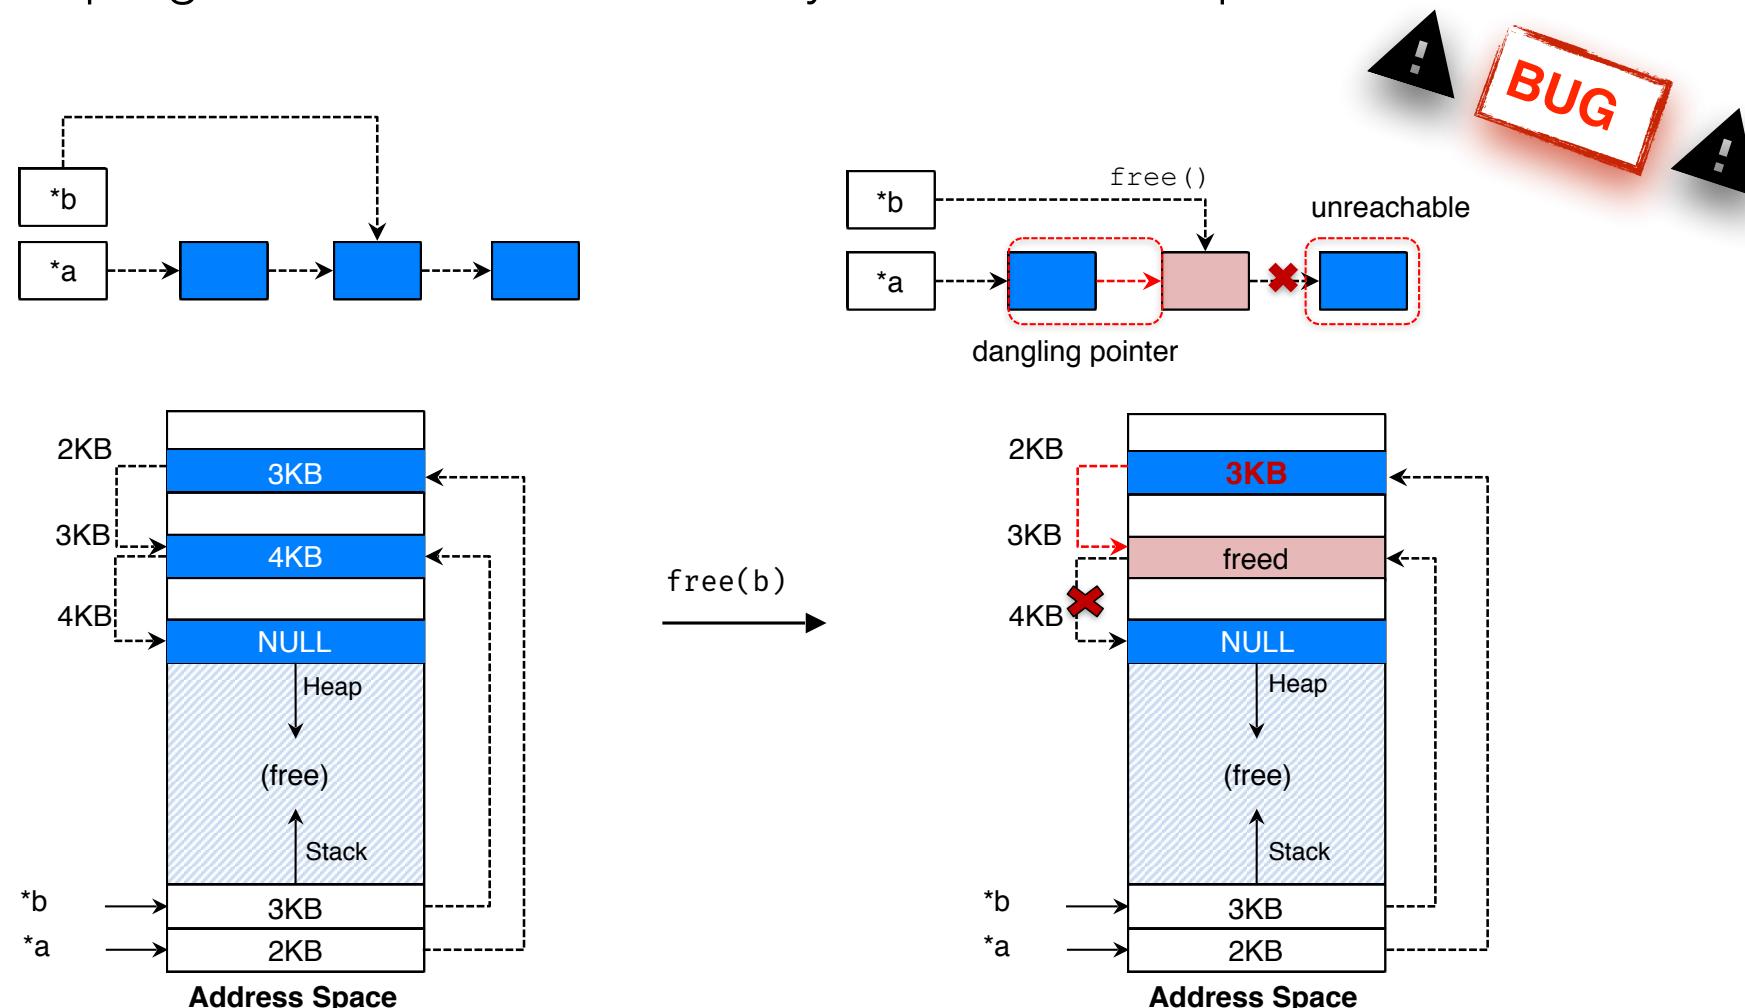
# Memory Leak

A program runs out of memory and eventually dies.



# Dangling Pointer

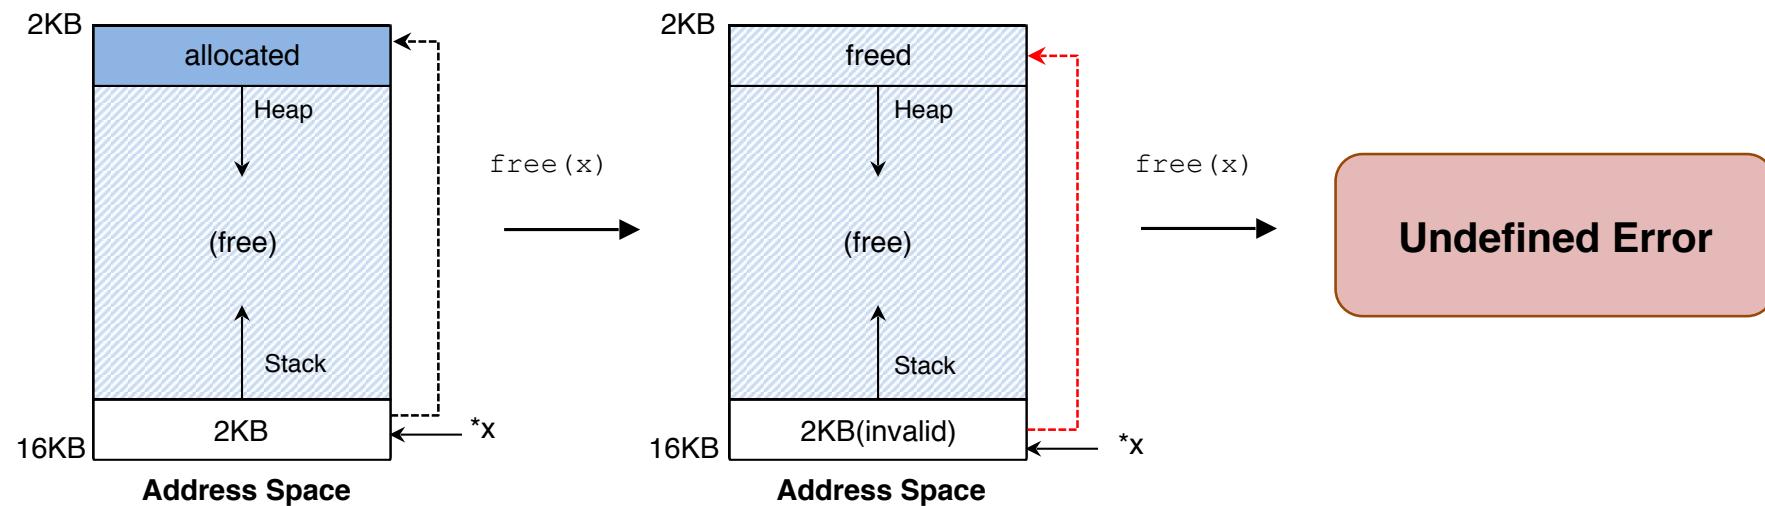
A program accesses to memory with an invalid pointer



# Double Free

Free memory that was freed already.

```
int *x = (int *)malloc(sizeof(int)); // allocated  
free(x); // free memory  
free(x); // free repeatedly
```

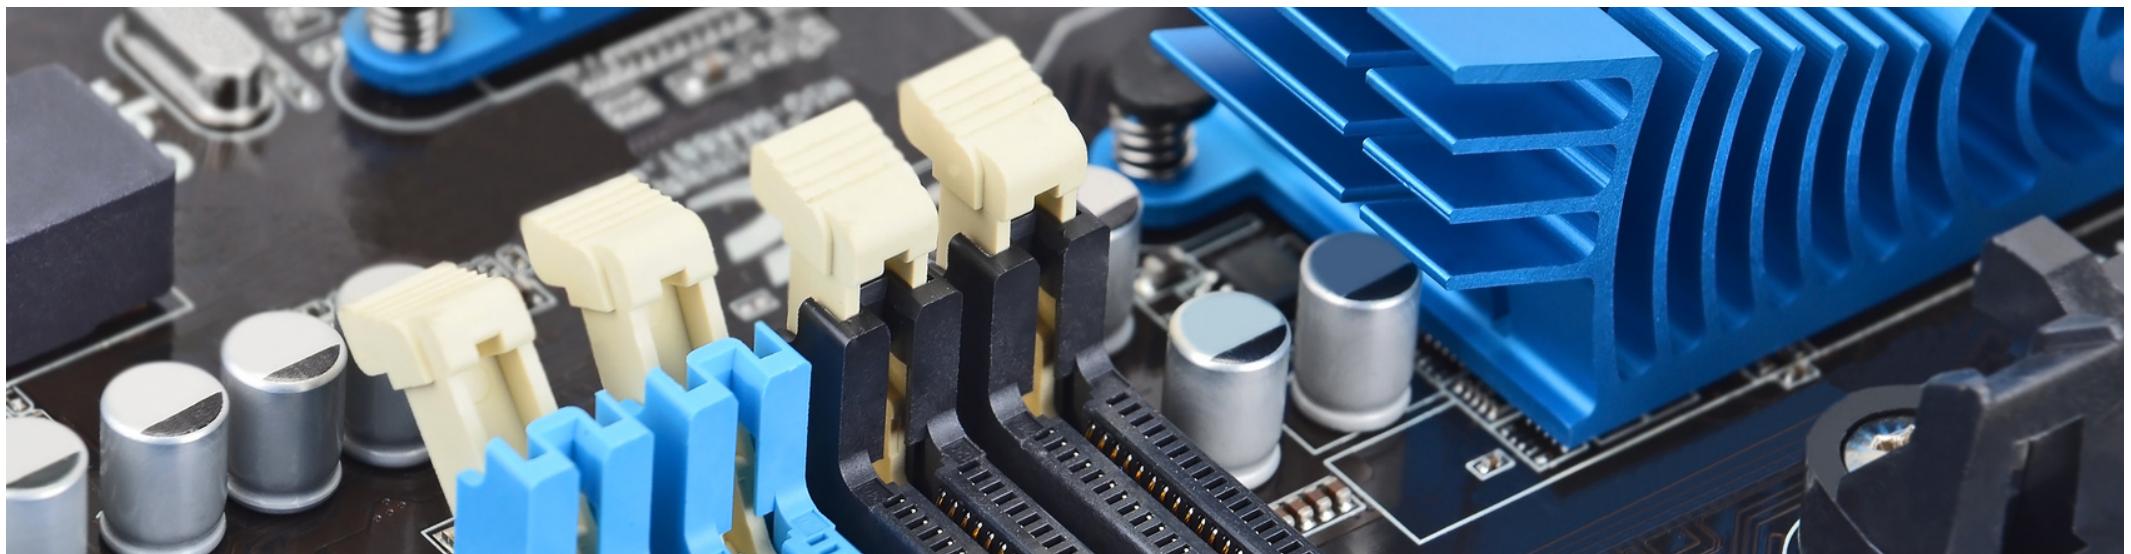


# Other Memory APIs: realloc()

- Change the size of memory block.
  - A pointer returned by **realloc** may be either the same as **ptr** or a new.
  - Argument
    - **void \*ptr**: Pointer to memory block allocated with **malloc**, **calloc** or **realloc**
    - **size\_t size**: New size for the memory block(in bytes)
- Return
  - Success: Void type pointer to the memory block
  - Fail : Null pointer

# 14. Memory API

1. `malloc`, `sizeof` and `free`
2. Common Errors
3. **Underlying OS Support**
4. Garbage Collector



# Underlying OS Support

- `malloc()` and `free()`
  - are library calls, **not system calls**
- `malloc` library call use `brk` system call.
  - `brk` is called to expand the program's *break*.
    - *break*: The location of **the end of the heap** in address space
  - `sbrk` is an additional call similar with `brk`.
- Programmers **should never directly call** either `brk` or `sbrk`.

# mmap()

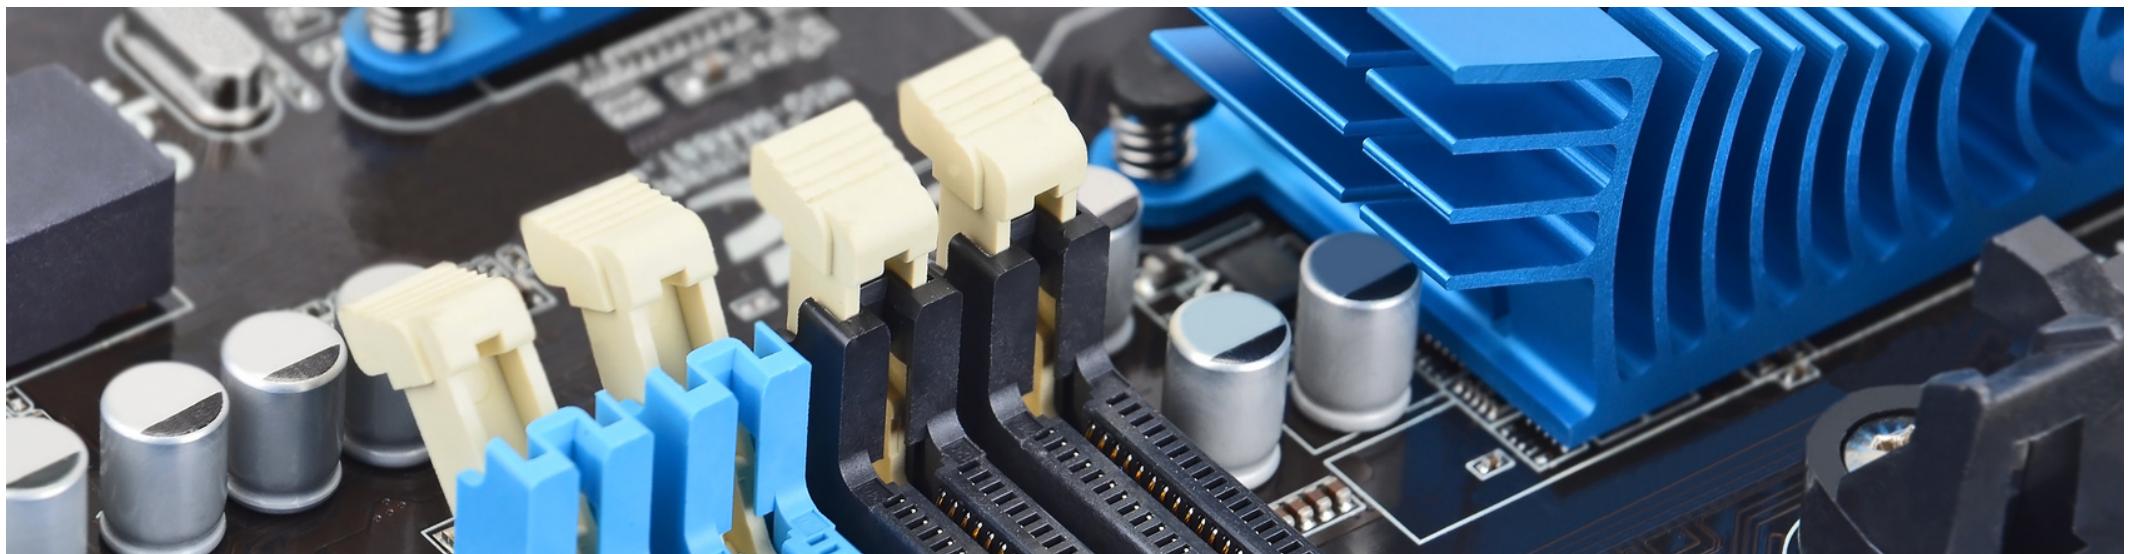
- Finally, you can also obtain memory from the operating system via the `mmap()` call
  - `mmap` system call can create **an anonymous** memory region.
    - a region which is not associated with any particular file but rather with swap space.
    - This memory can then also be treated like a heap and managed as such.
    - is used to create ‘shared’ Memory

```
#include <sys/mman.h>

void *mmap(void *ptr, size_t length, int port, int flags, int fd, off_t offset)
```

# 14. Memory API

1. `malloc`, `sizeof` and `free`
2. Common Errors
3. Underlying OS Support
4. **Garbage Collector**



# Garbage Collector?

Reference: „[Why I hate garbage collectors](#)“



or why manual resource management still matters

# Spot the Leak: C/C++

old C++

```
class Enemy {  
private:  
    EventFilter *filter; // "I own this"  
  
public:  
    Enemy(World *w, EnemyType t) {  
        // we only want to receive specific events  
        this->filter = new EventFilter(t);  
        w.register_event_handler(filter, this); // who owns this?  
    }  
};
```

where's the delete?

shouldn't delete

# Spot the Leak: Java

Java

```
class Enemy {  
  
    private EventFilter filter;  
  
    public Enemy(World w, EnemyType t) {  
        // we only want to receive specific events  
        this.filter = new EventFilter(t);  
        w.register_event_handler(filter, this);  
    }  
};
```

```
List<Enemy> enemies;  
World w;  
  
if (random())  
    enemies.add(new Enemy(weak));  
  
// remove all dead enemies  
enemies.removeIf(e → e.isDead());
```

GC deletes filter

this still in world

# Rules!

- Important Rules:
  - For every `malloc` there must be a `free`
  - For every `new` there must be a `delete`
  - Don't you dare to cross those and don't `free` twice
  - Be clear about ownership
  - C++: **RAII** (we'll come to that later)
- Rules?



# ~~GC to the rescue~~



Garbage Collector



well .... something ....



# What are resources?

- examples:
  - file
  - socket
  - database handle
  - ...
- **All with „open()/close() semantics“ is a resource!**



# Resource Acquisition Is Initialization

- The constructor acquires resources
- The destructor releases resources
- The destructor is guaranteed to be called
- Stack mechanics to the rescue!
- RAII in detail:
  - RAII relies on an automatic call to a common function at a defined time.
  - C++ uses a function that is already known to everyone = the destructor.
  - Destructor (in Rust drop()) is guaranteed to be called
- **All three types of leaks can be avoided!**

# GC and RAII

- Idea: GC for memory leaks, RAII for everything else
- Java: finalize()
  - not reliable!
  - is only called when GC is in the mood ...
- Fix the leak: Java

```
class Foo implements Closable {  
    private EventFilter f;  
  
    public Foo(int a) {  
        f = new EventFilter(a);  
        GlobalEvents.Get().Register(f, this);  
    }  
  
    public void Close() {  
        GlobalEvents.Get().Unregister(this);  
    }  
};
```

```
void UseFoo() {  
    Foo f;  
    try {  
        f = new Foo(10);  
    } finally {  
        f.Close();  
    }  
}
```

```
void UseFoo() {  
    try (Foo f = new Foo(10)) {  
        // code goes here  
    }  
}
```

# GC with RAII

## Memory Leak

forgot to **delete**/  
**free**

## Lingering Object

referenced but  
unused

## Resource Leak

forgot to close

- ~~Finalizer?~~ **unreliable!**
- Defer (do, D)?
- Disposable (C#) / Closable (Java)?

# Manual resource management

- Traditional Languages
  - Static code analysis to the rescue
  - Some things can only be seen at runtime ...
    - ... unless you extend the language
    - ... and let the **compiler detect things.**
- Hello: **Rust**
  - If you are not used to think about ownership **you will forget things**
  - If you have no enforced rules **you will miss things**

# With GC

- Developers get a false sense of safety
- Things tend to get even more complicated
- The problem is delayed, not solved
- Technical implications of GCs are often very bad for realtime or embedded applications