

A detailed close-up photograph of a computer motherboard. The image shows various components including blue plastic connectors, yellow RAM modules, and a large blue heat sink. The circuitry of the motherboard is visible, with numerous capacitors and integrated circuits. The lighting is bright, highlighting the textures and colors of the hardware.


# Lecture Operating System

## **19. Paging: Faster Translations**



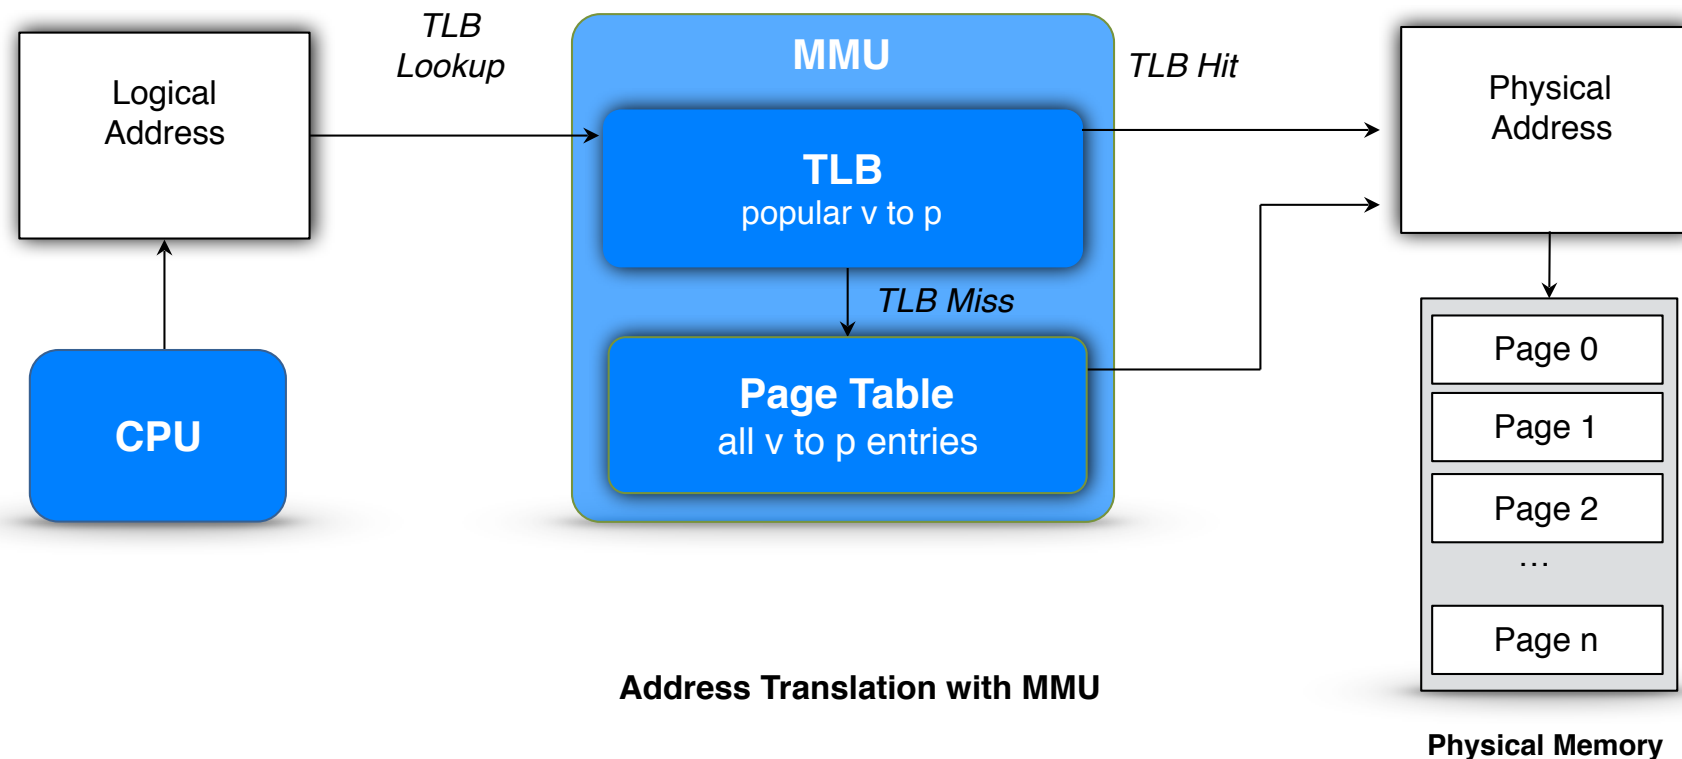
# 19. Paging: Faster Translations

1. How can we **speed up address translation**, and generally avoid the extra memory reference that paging seems to require?
2. What hardware support is required?
3. What OS involvement is needed?



# TLB: Translation Lookaside Buffer

- Part of the chip's memory-management unit (MMU).
- A hardware cache of **popular** virtual-to-physical address translation.



# TLB Basic Algorithms

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else
11     PTEAddr = PTBR + (VPN * sizeof(PTE))
12     PTE = AccessMemory(PTEAddr)
13     if (PTE.Valid == False)
14         RaiseException(SEGMENTATION_FAULT)
15     else if (CanAccess(PTE.ProtectBits) == False)
16         RaiseException(PROTECTION_FAULT)
17     else
18         TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
19     RetryInstruction()
```

extract the virtual page number(VPN).

check if the TLB holds the translation for this VPN.

extract the page frame number from the relevant TLB entry, and form the desired physical address and access memory.

The hardware accesses the page table to find the translation.

updates the TLB with the translation.

# Example: Accessing An Array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```

0: int sum = 0 ;
1: for( i=0; i<10; i++){
2:     sum+=a[i];
3: }

```


The TLB improves performance  
due to **spatial locality**

3 misses and 7 hits.  
Thus **TLB hit rate** is 70%.

# Locality


### ■ Temporal Locality

- An instruction or data item that has been recently accessed will likely be re-accessed soon in the future.




### ■ Spatial Locality

- If a program accesses memory at address x, it will likely soon access memory near x.



# TLB entry

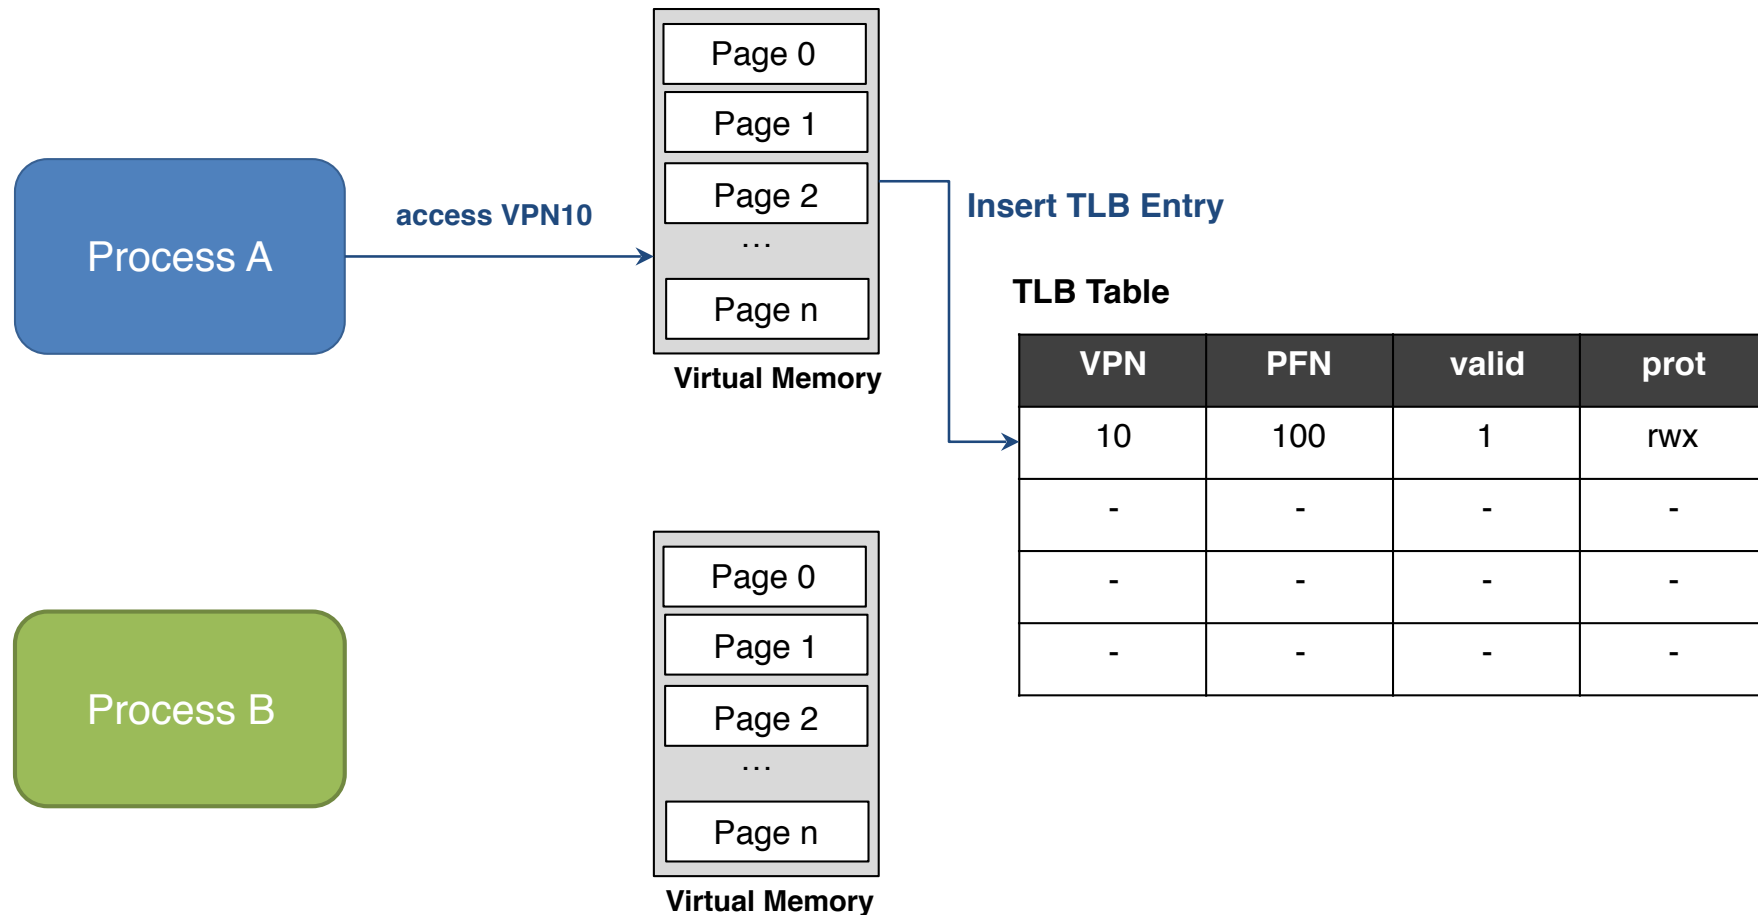
- TLB is managed by **Full Associative** method.
  - A typical TLB might have 32, 64, or 128 entries.
  - Hardware search the entire TLB in **parallel** to find the desired translation.
  - other bits: valid bits<sup>1</sup>, protection bits, address-space identifier, dirty bit



Typical TLB entry look like this

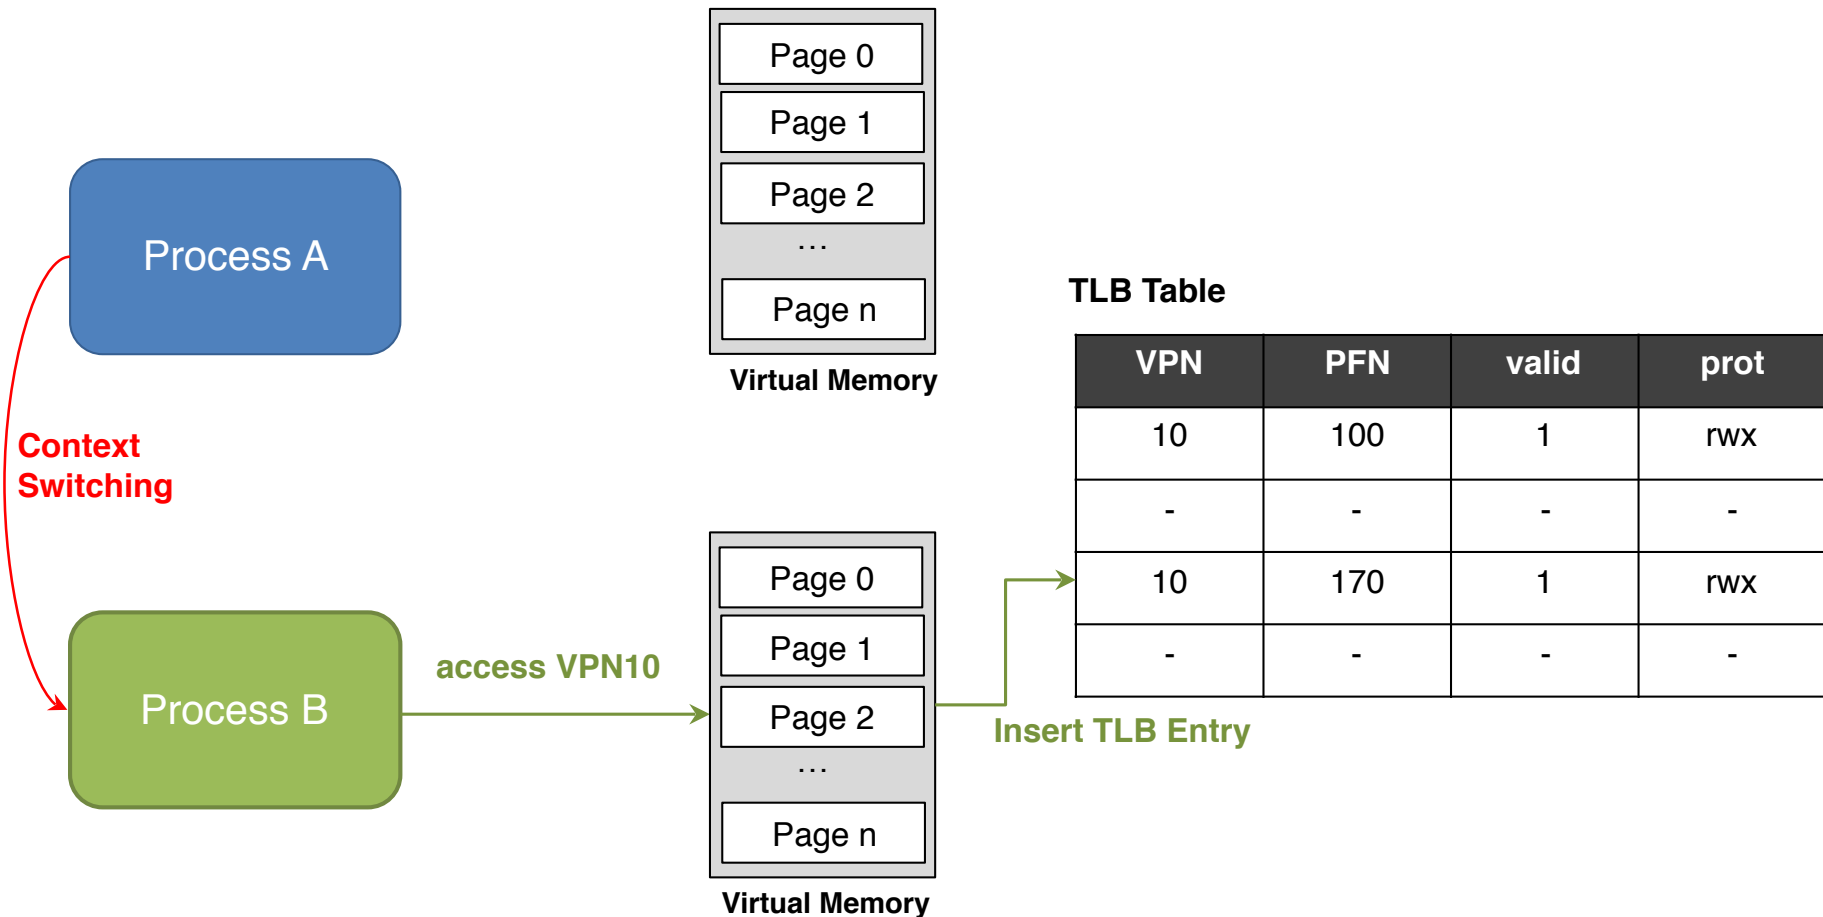
1: valid has in TLB different meaning than in Page Table!

# TLB Issue: Context Switching

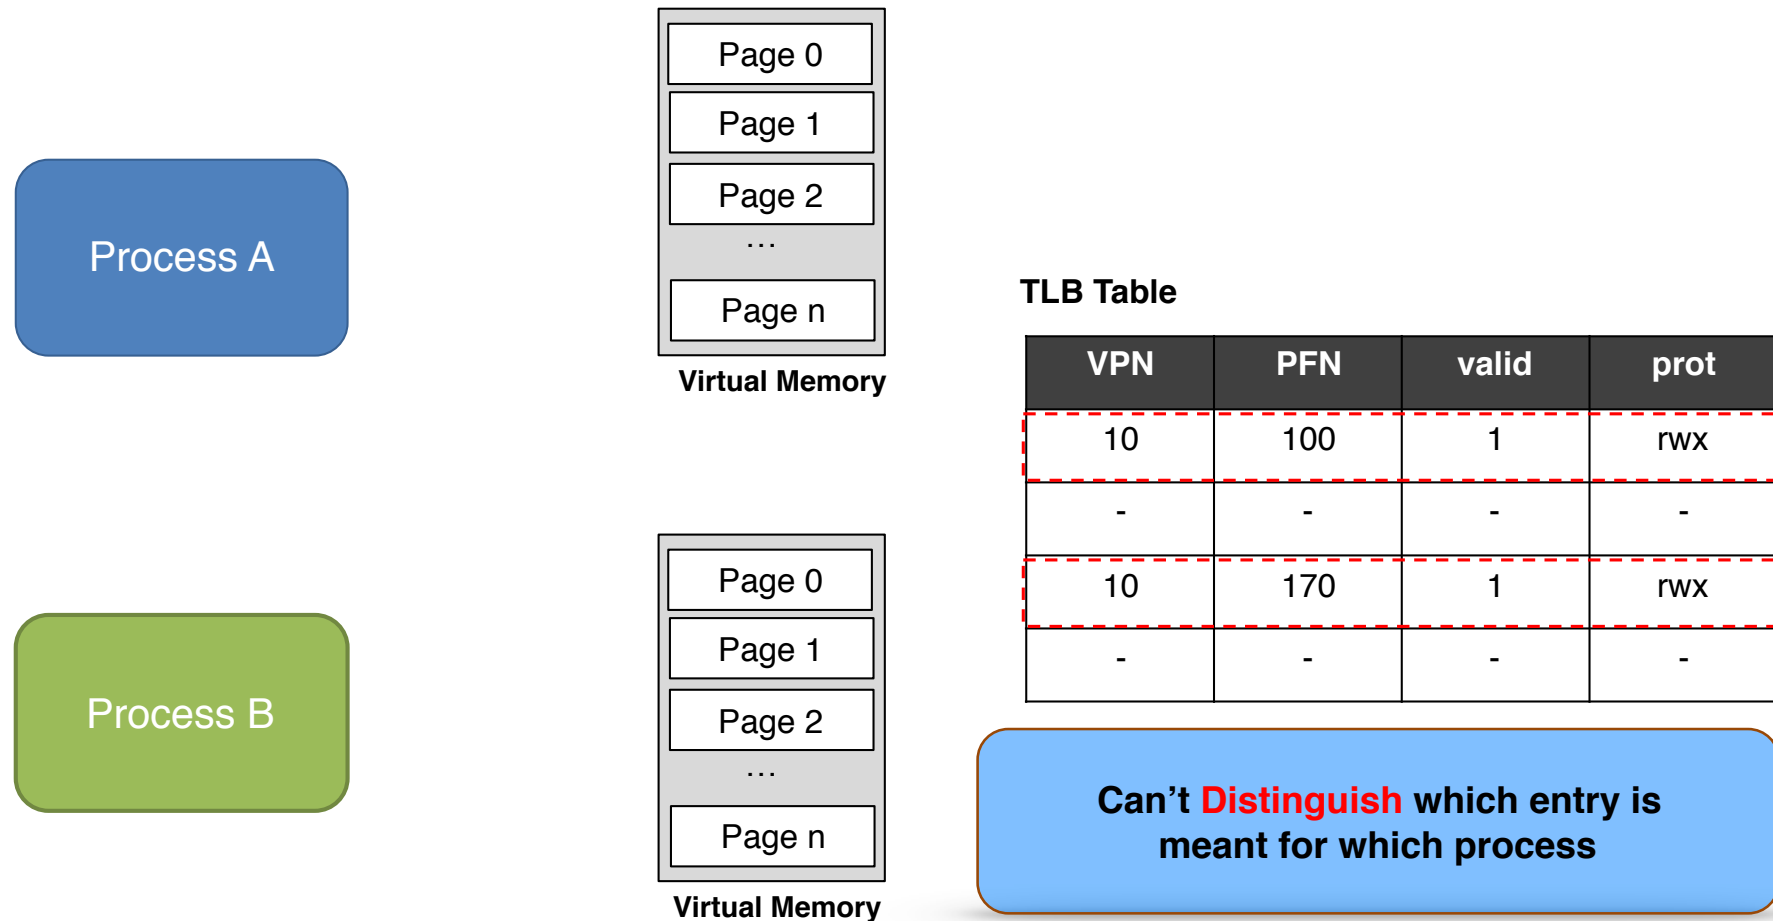




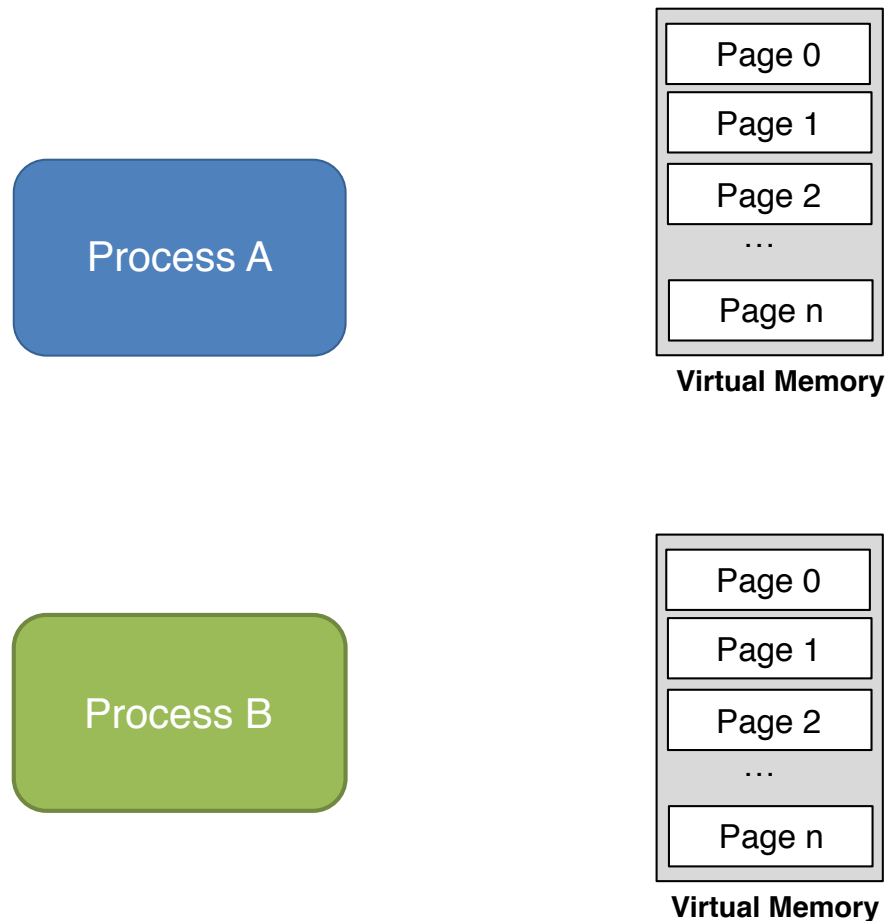
# TLB Issue: Context Switching



# TLB Issue: Context Switching



# To Solve Problem

**TLB Table**

VPN	PFN	valid	prot	ASID
10	100	1	rwX	1
-	-	-	-	-
10	170	1	rwX	2
-	-	-	-	-

Provide an address space identifier (ASID) field in the TLB.

## Another Case

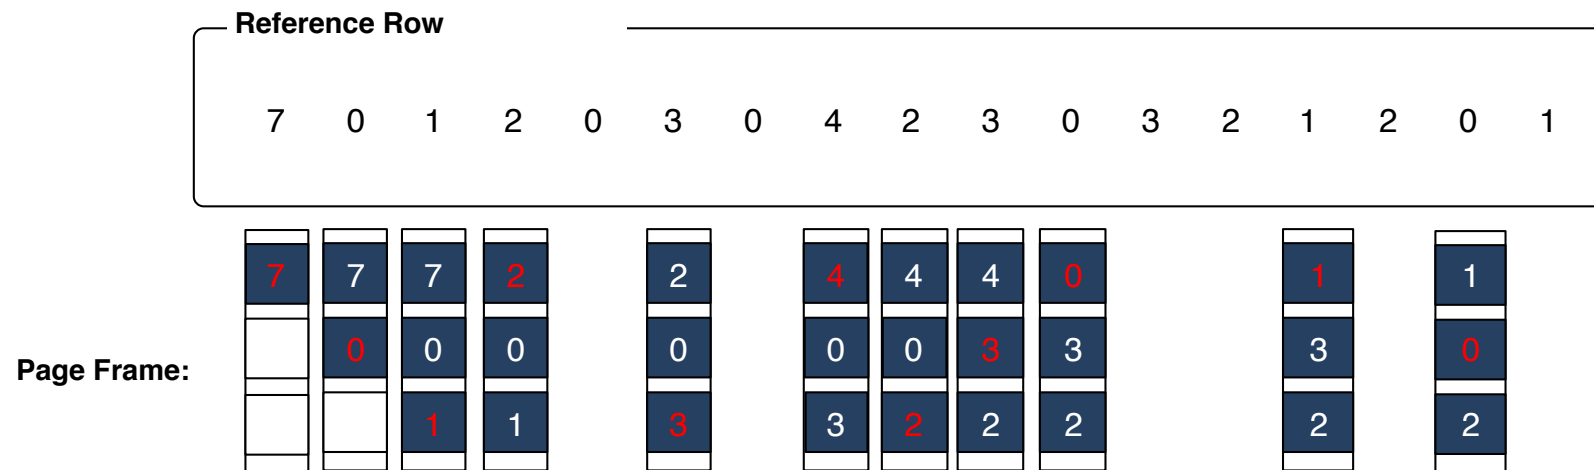
- Two processes **share a page**.
  - Process 1 is sharing physical page 101 with Process2.
  - P1 maps this page into the 10th page of its address space.
  - P2 maps this page to the 50th page of its address space.

VPN	PFN	valid	prot	ASID
10	101	1	rwX	1
-	-	-	-	-
50	101	1	rwX	2
-	-	-	-	-

Sharing of pages is **useful** as it reduces the number of physical pages in use.

# TLB Replacement Policy


- LRU(Least Recently Used)
  - Evict an entry that has not recently been used.
  - Take advantage of *locality* in the memory-reference stream.



Total **11** TLB miss



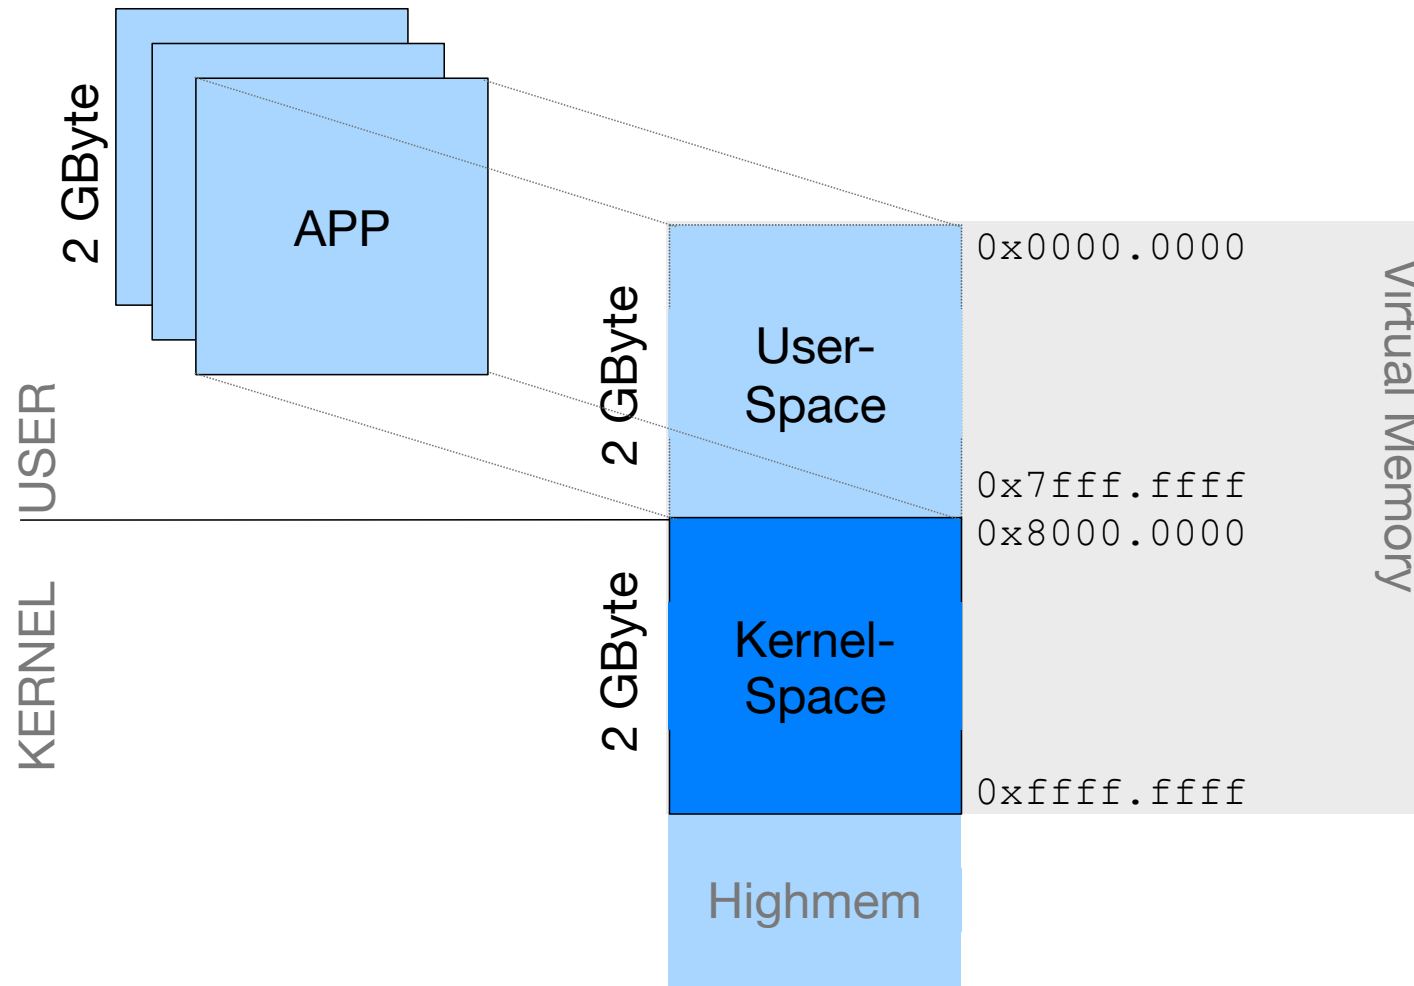
# A Real TLB Entry (64 bits MIPS)



Flag	Content
19-bit VPN	The rest reserved for the kernel.
24-bit PFN	Systems can support with up to 64GB of main memory ( $2^{24} * 4\text{KB}$ pages)
Global bit(G)	Used for pages that are globally-shared among processes.
ASID	OS can use to distinguish between address spaces.
Coherence bit(C)	determine how a page is cached by the hardware.
Dirty bit(D)	marking when the page has been written.
Valid bit(V)	tells the hardware if there is a valid translation present in the entry.

All 64 bits of this TLB entry (example of MIPS R4000)

# VPNs: Shared between User and Kernel Address



# Example: 64 Bit Linux System

**cat /proc/8080/maps**

# address	Perms	Offset	Dev	Inode	Pathname
00400000-00406000	r-xp	00000000	03:05	1437882	/usr/sbin/metalog
00505000-00506000	rw-p	00005000	03:05	1437882	/usr/sbin/metalog
00506000-00527000	rw-p	00506000	00:00	0	<b>[heap]</b>
2b76d29df000-2b76d29f5000	r-xp	00000000	03:02	61500	/lib64/ld-2.3.6.so
2b76d29f5000-2b76d29fc000	rw-p	2b76d29f5000	00:00	0	
2b76d2af5000-2b76d2af6000	r--p	00016000	03:02	61500	/lib64/ld-2.3.6.so
2b76d2af6000-2b76d2af7000	rw-p	00017000	03:02	61500	/lib64/ld-2.3.6.so
2b76d2af7000-2b76d2b0c000	r-xp	00000000	03:05	311804	/usr/lib64/libpcre.so.0.0.1
2b76d2b0c000-2b76d2c0c000	---p	00015000	03:05	311804	/usr/lib64/libpcre.so.0.0.1
2b76d2c0c000-2b76d2c23000	rw-p	00015000	03:05	311804	/usr/lib64/libpcre.so.0.0.1
2b76d2c23000-2b76d2d41000	r-xp	00000000	03:02	62921	/lib64/tls/libc-2.3.6.so
2b76d2d41000-2b76d2e41000	---p	0011e000	03:02	62921	/lib64/tls/libc-2.3.6.so
2b76d2e41000-2b76d2e44000	r--p	0011e000	03:02	62921	/lib64/tls/libc-2.3.6.so
2b76d2e44000-2b76d2e47000	rw-p	00121000	03:02	62921	/lib64/tls/libc-2.3.6.so
2b76d2e47000-2b76d2e4e000	rw-p	2b76d2e47000	00:00	0	
7ffffd80b5000-7ffffd80cb000	rw-p	7ffffd80b5000	00:00	0	<b>[stack]</b>
fffffffffff600000-fffffffffff600000	---p	00000000	00:00	0	<b>[vdso]</b>

**cat /proc/790/maps**

00400000-004af000	r-xp	00000000	03:02	62290	/bin/bash
005ae000-005b9000	rw-p	000ae000	03:02	62290	/bin/bash
005b9000-00728000	rw-p	005b9000	00:00	0	<b>[heap]</b>
.....					

# TLB Summary

- **Hardware** can help us **make address translation faster**.
  - By TLB memory references often will be handled without having to access the page table in main memory.
  - the performance of the program will be almost as if memory isn't being virtualized at all.
- Locality depends on program:
  - **exceeding** the **TLB coverage**
  - TLB access can easily become a **bottleneck** in the **CPU pipeline**

A close-up photograph of a computer motherboard. The image shows various components including blue heat sinks, yellow RAM modules, and blue SATA connectors. The motherboard is black with visible circuitry and components like capacitors.

# Thanks

## Questions?