

Lecture

# Operating System

## 30. Condition Variables

# Concurrency Objectives

- **Mutual exclusion** (e.g., A and B don't run at same time)
  - solved with locks
- **Ordering** (e.g., B runs after A does something)
  - solved with **condition variables** and semaphores



# 30. Condition Variables

- 1. Join**
- 2. Bounded Buffer**
- 3. Broadcast**



# 30. Condition Variables

- 1. Join**
2. Bounded Buffer
3. Broadcast



# Condition Variables

- There are many cases where a thread wishes to **check** whether a **condition** is true before continuing its execution.
- Example:
  - A parent thread might wish to check whether a child thread has **completed**.
  - This is often called a `join()`.
- The crux: How to wait for a condition?

# Parent waiting for Child

```
void *child(void *arg) {  
    printf("child\n");  
    // XXX how to indicate we are done?  
    return NULL;  
}  
  
int main(int argc, char *argv[]) {  
    printf("parent: begin\n");  
    pthread_t c;  
    Pthread_create(&c, NULL, child, NULL); // create child  
    // XXX how to wait for child?  
    printf("parent: end\n");  
    return 0;  
}
```

What we would like to see here is:

```
parent: begin  
child  
parent: end
```

# join: Spin-based Approach

```
volatile int done = 0;

void *child(void *arg) {
    printf("child\n");
    done = 1;
    return NULL;
}

int main(int argc, char *argv[]) {
    printf("parent: begin\n");
    pthread_t c;
    Pthread_create(&c, NULL, child, NULL); // create child
    while (done == 0)
        ; // spin
    printf("parent: end\n");
    return 0;
}
```

This is hugely inefficient as the parent spins and wastes CPU time.



# How to wait for a condition

- Condition variable
  - **Waiting** on the condition
    - *An explicit queue* that threads can put themselves on when some state of execution is not as desired.
  - **Signaling** on the condition
    - Some other thread, *when it changes said state*, can wake one of those waiting threads and allow them to continue.



# Definition and Routines

- **Declare** condition variable

- **pthread\_cond\_t c;**

- Proper initialization is required.

- **Operation** (the POSIX calls)

- **pthread\_cond\_wait(pthread\_cond\_t \*c, pthread\_mutex\_t \*m);**
  - **pthread\_cond\_signal(pthread\_cond\_t \*c);**

- The wait() call takes a mutex as a parameter.

- The wait() call release the lock and put the calling thread to sleep.
  - When the thread wakes up, it must re-acquire the lock.

# Parent waiting for Child:

## Use a condition variable

```
int done = 0;
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t c = PTHREAD_COND_INITIALIZER;

void thr_exit() {
    Pthread_mutex_lock(&m);
    done = 1;
    Pthread_cond_signal(&c);
    Pthread_mutex_unlock(&m);
}

void *child(void *arg) {
    printf("child\n");
    thr_exit();
    return NULL;
}

void thr_join() {
    Pthread_mutex_lock(&m);
    while (done == 0)
        Pthread_cond_wait(&c, &m);
    Pthread_mutex_unlock(&m);
}
```

```
int main(int argc, char *argv[]) {
    printf("parent: begin\n");
    pthread_t p;
    Pthread_create(&p, NULL, child, NULL);
    thr_join();
    printf("parent: end\n");
    return 0;
}
```

# Parent waiting for Child:

## Use a condition variable

- Parent:
  - Create the child thread and continues running itself.
  - Call into `thr_join()` to wait for the child thread to complete.
    - Acquire the lock
    - Check *if* the child is *done*
    - Put itself to sleep by calling `wait()`
    - Release the lock
- Child:
  - Print the message “child”
  - Call `thr_exit()` to wake the parent thread
    - Grab the lock
    - Set the state variable *done*
    - `Signal()` the parent thus waking it.



# Importance of state variable done

```
void thr_exit() {  
    Pthread_mutex_lock(&m);  
    Pthread_cond_signal(&c);  
    Pthread_mutex_unlock(&m);  
}  
  
void thr_join() {  
    Pthread_mutex_lock(&m);  
    Pthread_cond_wait(&c, &m);  
    Pthread_mutex_unlock(&m);  
}
```

- Imagine the case where the child runs immediately.
  - The child will `signal()`, but there is no **thread asleep** on the condition.
  - When the parent runs, it will call `wait()` and be stuck.
  - **No thread will ever wake it.**

# Another poor implementation

- The issue here is a subtle race condition.
- The parent calls `thr_join()`.
  - The parent checks the value of `done`.
  - It will see that it is 0 and try to go to sleep.
  - **Just before** it calls wait to go to sleep, the parent is **interrupted** and the child runs.
- The child changes the state variable `done` to 1 and signals.
  - But no thread is waiting and thus no thread is woken.
  - When the parent runs again, it sleeps forever.

```
void thr_exit() {  
    done = 1;  
    Pthread_cond_signal(&c);  
}  
  
void thr_join() {  
    if (done == 0)  
        Pthread_cond_wait(&c);  
}
```

# 30. Condition Variables

1. Join
- 2. Bounded Buffer**
3. Broadcast





# The Producer / Consumer (Bounded Buffer) Problem


- Producer
  - **Produce** data items
  - Wish to place data items in a buffer
- Consumer
  - Grab data items out of the buffer **consume** them in some way
- Example: Multi-threaded web server
  - A **producer** puts HTTP requests in to a work queue
  - **Consumer** threads take requests out of this queue and process them

## Bounded buffer

- A bounded buffer is used when you *pipe the output* of one program into another.
  - Example: **grep foo file.txt | wc -l**
    - The **grep** process is the producer.
    - The **wc** process is the consumer.
    - Between them is an in-kernel **bounded buffer**.
- Bounded buffer is Shared resource
  - **Synchronized access** is required.

# UNIX pipe

- A pipe may have many writers and readers
- Internally, there is a finite-sized buffer
- Writers add data to the buffer
  - Writers have to wait if buffer is full
- Readers remove data from the buffer
  - Readers have to wait if buffer is empty





# The Put and Get Routines

- Only put data into the buffer when **count** is zero.
  - i.e., when the buffer is **empty**.
- Only get data from the buffer when **count** is one.
  - i.e., when the buffer is **full**.

```
int buffer;  
int count = 0; // initially,  
empty  
  
void put(int value) {  
    assert(count == 0);  
    count = 1;  
    buffer = value;  
}  
  
int get() {  
    assert(count == 1);  
    count = 0;  
    return buffer;  
}
```

# Producer/Consumer Threads

## ■ Producer

- puts an integer into the shared buffer loops number of times.

## ■ Consumer

- gets the data out of that shared buffer.

```
void *producer(void *arg) {  
    int i;  
    int loops = (int) arg;  
    for (i = 0; i < loops; i++) {  
        put(i);  
    }  
}  
  
void *consumer(void *arg) {  
    int i;  
    while (1) {  
        int tmp = get();  
        printf("%d\n", tmp);  
    }  
}
```

# Producer/Consumer:

## Single CV and If Statement

```
cond_t cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // p1
        if (count == 1)                       // p2
            Pthread_cond_wait(&cond, &mutex); // p3
        put(i);                               // p4
        Pthread_cond_signal(&cond);           // p5
        Pthread_mutex_unlock(&mutex);         // p6
    }
}
```

p1-p3: A producer waits for the buffer to be empty



# Producer/Consumer:

## Single CV and If Statement

```
cond_t cond;
mutex_t mutex;

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // c1
        if (count == 0)                       // c2
            Pthread_cond_wait(&cond, &mutex); // c3
        int tmp = get();                      // c4
        Pthread_cond_signal(&cond);           // c5
        Pthread_mutex_unlock(&mutex);         // c6
        printf("%d\n", tmp);
    }
}
```

c1-c3: A consumer waits for the buffer to be full.

# Producer/Consumer:

## Single CV and If Statement

- With just a single producer and a single consumer, the code works.
- If we have more than one of producer and consumer?

```
cond_t cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // p1
        if (count == 1)                       // p2
            Pthread_cond_wait(&cond, &mutex); // p3
        put(i);                               // p4
        Pthread_cond_signal(&cond);           // p5
        Pthread_mutex_unlock(&mutex);         // p6
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // c1
        if (count == 0)                       // c2
            Pthread_cond_wait(&cond, &mutex); // c3
        int tmp = get();                      // c4
        Pthread_cond_signal(&cond);           // c5
        Pthread_mutex_unlock(&mutex);         // c6
        printf("%d\n", tmp);
    }
}
```

# Thread Trace

## Single CV and If Statement

T <sub>c1</sub>	State	T <sub>c2</sub>	State	T <sub>p</sub>	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep		Ready	p1	Running	0	
	Sleep		Ready	p2	Running	0	
	Sleep		Ready	p4	Running	1	Buffer now full
	Ready		Ready	p5	Running	1	T <sub>c1</sub> awoken
	Ready		Ready	p6	Running	1	
	Ready		Ready	p1	Running	1	
	Ready		Ready	p2	Running	1	
	Ready		Ready	p3	Sleep	1	Buffer full; sleep
	Ready	c1	Running		Sleep	1	T <sub>c2</sub> sneaks in ...
	Ready	c2	Running		Sleep	1	
	Ready	c4	Running		Sleep	0	... and grabs data
	Ready	c5	Running		Ready	0	T <sub>p</sub> awoken
	Ready	c6	Running		Ready	0	
c4	Running		Ready		Ready	0	Oh oh! No data

# Producer/Consumer:

## Single CV and While Statement

- Try to fix with `while` instead of `if`
- Better, but still broken
  - but it solves the problem before. Why?
  - Still a problem, because of only one CV ...

```
cond_t cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // p1
        while (count == 1)                    // p2
            Pthread_cond_wait(&cond, &mutex); // p3
        put(i);                               // p4
        Pthread_cond_signal(&cond);           // p5
        Pthread_mutex_unlock(&mutex);         // p6
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // c1
        while (count == 0)                    // c2
            Pthread_cond_wait(&cond, &mutex); // c3
        int tmp = get();                      // c4
        Pthread_cond_signal(&cond);           // c5
        Pthread_mutex_unlock(&mutex);         // c6
        printf("%d\n", tmp);
    }
}
```

# Thread Trace

## Single CV and While Statement

T <sub>c1</sub>	State	T <sub>c2</sub>	State	T <sub>p</sub>	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep	c1	Running		Ready	0	
	Sleep	c2	Running		Ready	0	
	Sleep	c3	Sleep		Ready	0	Nothing to get
	Sleep		Sleep	p1	Running	0	
	Sleep		Sleep	p2	Running	0	
	Sleep		Sleep	p4	Running	1	Buffer now full
	Ready		Sleep	p5	Running	1	T <sub>c1</sub> awoken
	Ready		Sleep	p6	Running	1	
	Ready		Sleep	p1	Running	1	
	Ready		Sleep	p2	Running	1	
	Ready		Sleep	p3	Sleep	1	Must sleep (full)
c2	Running		Sleep		Sleep	1	Recheck condition
c4	Running		Sleep		Sleep	0	T <sub>c1</sub> grabs data
c5	Running		Ready		Sleep	0	<b>Oops! Woke T<sub>c2</sub></b>



# Thread Trace (Cont.)

## Single CV and While Statement

T <sub>c1</sub>	State	T <sub>c2</sub>	State	T <sub>p</sub>	State	Count	Comment
...	...	...	...	...	...	...	(cont.)
c6	Running		Ready		Sleep	0	
c1	Running		Ready		Sleep	0	
c2	Running		Ready		Sleep	0	
c3	Sleep		Ready		Sleep	0	Nothing to get
	Sleep	c2	Running		Sleep	0	
	Sleep	c3	Sleep		Sleep	0	<b>Everyone asleep</b>

### ■ Solution:

#### ■ Use **two** condition variables and **while**

- Producer threads wait on the condition **empty**, and signals **fill**.
- Consumer threads wait on **fill** and signal **empty**.

# The single Buffer Producer/Consumer Solution

```
cond_t empty, fill;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 1)
            Pthread_cond_wait(&empty, &mutex);
        put(i);
        Pthread_cond_signal(&fill);
        Pthread_mutex_unlock(&mutex);
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 0)
            Pthread_cond_wait(&fill, &mutex);
        int tmp = get();
        Pthread_cond_signal(&empty);
        Pthread_mutex_unlock(&mutex);
        printf("%d\n", tmp);
    }
}
```

# The Final Producer/Consumer Solution

- More **concurrency** and **efficiency**
  - Add more buffer slots.
    - Allow concurrent production or consuming to take place.
    - Reduce context switches.

```
int buffer[MAX];
int fill = 0;
int use = 0;
int count = 0;

void put(int value) {
    buffer[fill] = value;
    fill = (fill + 1) % MAX;
    count++;
}

int get() {
    int tmp = buffer[use];
    use = (use + 1) % MAX;
    count--;
    return tmp;
}
```

# The Final Producer/Consumer Solution

```
cond_t empty, fill;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // p1
        while (count == MAX)                  // p2
            Pthread_cond_wait(&empty, &mutex); // p3
        put(i);                               // p4
        Pthread_cond_signal(&fill);           // p5
        Pthread_mutex_unlock(&mutex);         // p6
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // c1
        while (count == 0)                    // c2
            Pthread_cond_wait(&fill, &mutex); // c3
        int tmp = get();                      // c4
    }
}
```

- **p2**: A **producer** only sleeps if all buffers are currently filled.
- **c2**: A **consumer** only sleeps if all buffers are currently empty.

# 30. Condition Variables

1. Join
2. Bounded Buffer
- 3. Broadcast**





# Covering Conditions

- Multithreaded Memory Allocation Library
- Assume there are zero bytes press
  - Thread  $T_a$  calls `allocate(100)`.
  - Thread  $T_b$  calls `allocate(10)`.
  - Both,  $T_a$  and  $T_b$  wait on the condition and go to sleep.
  - Thread  $T_c$  calls `free(50)`.
- **Which waiting thread should be woken up?**

# Example: Code

```
// how many bytes of the heap are free?
int bytesLeft = MAX_HEAP_SIZE;

// need lock and condition too
cond_t c;
mutex_t m;

void *
allocate(int size) {
    Pthread_mutex_lock(&m);
    while (bytesLeft < size)
        Pthread_cond_wait(&c, &m);
    void *ptr = ... ;           // get mem from heap
    bytesLeft -= size;
    Pthread_mutex_unlock(&m);
    return ptr;
}

void free(void *ptr, int size) {
    Pthread_mutex_lock(&m);
    bytesLeft += size;
    Pthread_cond_signal(&c);    // whom to signal??
    Pthread_mutex_unlock(&m);
}
```

# Solution for Covering Conditions

- Replace `pthread_cond_signal()` with `pthread_cond_broadcast()`
- `pthread_cond_broadcast()`
  - Wake up **all waiting threads**.
  - **Cost**: too many threads might be woken.
    - Threads that shouldn't be awake will simply wake up, re-check the condition, and then go back to sleep.

# Solution: Code

```
// how many bytes of the heap are free?
int bytesLeft = MAX_HEAP_SIZE;

// need lock and condition too
cond_t c;
mutex_t m;

void *
allocate(int size) {
    Pthread_mutex_lock(&m);
    while (bytesLeft < size)
        Pthread_cond_wait(&c, &m);
    void *ptr = ... ;           // get mem from heap
    bytesLeft -= size;
    Pthread_mutex_unlock(&m);
    return ptr;
}

void free(void *ptr, int size) {
    Pthread_mutex_lock(&m);
    bytesLeft += size;
    Pthread_cond_broadcast(&c); // send all!
    Pthread_mutex_unlock(&m);
}
```

# Summary

- Keep state in addition to CV's
- Always do wait/signal with lock held
- Whenever thread wakes from waiting, recheck state
  - Possible for another thread to grab lock in between signal and wakeup from wait



# Thanks

## Questions?

