

Lecture

# Operating System

## 33. Event-based Concurrency



# 33. Event-based Concurrency

- 1. Basic Idea**
- 2. Problems**
- 3. Signals**



# 33. Event-based Concurrency

## 1. Basic Idea

## 2. Problems

## 3. Signals



# Event-based Concurrency

- A different style of **concurrent programming**
  - Used in GUI-based applications, some types of internet servers.
- **The problem** that event-based concurrency addresses is two-fold.
  - **Managing concurrency correctly** in multi-threaded applications.
    - Missing locks, deadlock, and other nasty problems can arise.
  - The developer has little or no control over **what is scheduled** at a given moment in time.

# The Crux

- How can we build a **concurrent** server **without** using **threads**
  - but **retain control** over concurrency
  - and **avoid** some of the **problems** that seem to plague multi-threaded applications?

# The Basic Idea: An Event Loop

- The approach:
  - **Wait** for something (i.e., an “*event*”) to occur.
  - When it does, **check** what type of event it is.
  - **Do** the small amount of work it requires.
- Example:

```
while(1){  
    events = getEvents();  
    for( e in events )  
        processEvent(e); // event handler  
}
```

A canonical event-based  
server (Pseudo code)

How exactly does an event-based server  
determine which events are **taking place**.



# An Important API: select() (or poll())

- Check whether there is any **incoming I/O** that should be attended to.

- `select()`

```
int select(int nfd,  
           fd_set * restrict readfds,  
           fd_set * restrict writefds,  
           fd_set * restrict errorfds,  
           struct timeval * restrict timeout);
```

- Lets a server determine that a **new packet has arrived** and is in need of processing.
- Let the service know when **it is OK to reply**.
- `timeout`
  - `NULL`: Cause `select()` to block indefinitely until some descriptor is ready.
  - `0`: Use the call to `select()` to return immediately.

# Using `select()`

```
while (1) {  
    // initialize the fd_set to all zero  
    fd_set readFDs;  
    FD_ZERO(&readFDs);  
  
    // now set the bits for the descriptors  
    // this server is interested in  
    // (for simplicity, all of them from min to max)  
    int fd;  
    for (fd = minFD; fd < maxFD; fd++)  
        FD_SET(fd, &readFDs);  
  
    // do the select  
    int rc = select(maxFD+1, &readFDs, NULL, NULL, NULL);  
  
    // check which actually have data using FD_ISSET()  
    for (fd = minFD; fd < maxFD; fd++)  
        if (FD_ISSET(fd, &readFDs))  
            processFD(fd);  
}
```



# Why Simpler? No Locks Needed

- The event-based server **cannot be interrupted** by another thread:
  - With a **single CPU** and an **event-based application**.
  - It is decidedly **single threaded**.
  - Thus, **concurrency bugs** common in threaded programs **do not manifest** in the basic event-based approach.

# 33. Event-based Concurrency

1. Basic Idea
- 2. Problems**
3. Signals



# A Problem: Blocking System Calls

- What if an event requires that you issue **a system call** that might block?
  - There are no other threads to run: *just the main event loop*
  - The entire server will do just that: **block until the call completes**.
  - **Huge potential waste of resources**

In event-based systems: **no blocking calls** are allowed.



# A Solution: Asynchronous I/O

- Enable an application to issue an I/O request and **return control immediately** to the caller, before the I/O has completed.
- Example:

```
struct aiocb {  
    int aio_fildes;           /* File descriptor */  
    off_t aio_offset;         /* File offset */  
    volatile void *aio_buf;   /* Location of buffer */  
    size_t aio_nbytes;        /* Length of transfer */  
};
```

- An Interface provided on *Mac OS X*
- The APIs revolve around a basic structure, the `struct aiocb` or **AIO control block** in common terminology.

## Example: Async I/O

- To issue an asynchronous read to a file

```
int aio_read(struct aiocb *aiocbp);
```

- If successful, it returns right away and the application can continue with its work.
- Checks whether the request referred to by `aiocbp` has completed.

```
int aio_error(const struct aiocb *aiocbp);
```

- An application can **periodically poll** the system via `aio_error()`.
  - If it has completed, returns success.
  - If not, `EINPROGRESS` is returned.

# Example: Async I/O (Signaling)

- Interrupt
  - Remedy **the overhead to check** whether an I/O has completed
  - Using **UNIX signals** to inform applications when an asynchronous I/O completes.
  - Removing the need to *repeatedly ask the system*.



# Another Problem: State Management

- The code of event-based approach is generally **more complicated** to write than *traditional thread-based* code.
- It must package up some program state for the next event handler to use when the I/O completes.
- The state the program needs is on the stack of the thread.
  - **manual stack management** !

# Example: Problem with 'state'

- Example (an thread-based system):

```
int rc = read(fd, buffer, size);  
rc = write(sd, buffer, size);
```

- How to do in asynchronous systems?
  - First **issue** the read asynchronously.
  - Then, **periodically check** for completion of the read.
  - That call informs us that the **read is complete**.
  - How does the event-based server know **what to do**?

# Solution to 'state' Problem

- Solution: **continuation**
  - **Record** the needed information to finish processing this event in *some data structure*.
  - When the event happens (i.e., when the disk I/O completes), **look up** the needed information and process the event.



# What is still difficult with Events.

- Systems moved from a single CPU to **multiple CPUs**.
  - Some of the simplicity of the event-based approach disappeared.
- It **does not integrate well** with certain kinds of systems activity.
  - **Ex. Paging:** A server will not make progress until page fault completes (implicit blocking).
- Hard to manage overtime: The exact semantics of various routines changes.
- Asynchronous disk I/O **never quite integrates with asynchronous network I/O** in as simple and uniform a manner as you might think.

# 33. Event-based Concurrency

1. Basic Idea
2. Problems
- 3. Signals**



# UNIX signals

- Provide a way to communicate with a process.
  - HUP (hang up), INT(interrupt), SEGV(segmentation violation), and etc.
  - Example: When your program encounters a segmentation violation, the OS sends it a SIGSEGV.



# Example: Standard Signals

```
static void SigHandler( int SignalNumber )
{
    printf("got signal %d\n", SignalNumber );
}
```

```
int main( int argc, char **argv )
{
    int i;

    signal( SIGINT, SigHandler );
    printf("my PID is %d\n", getpid() );
    for( i=1; i<100; i++ ) {
        printf("%d seconds\n", i);
        sleep( 1 );
    }
    return( 0 );
}
```

# Unreliable Signals

- Standard Signals
  - poor performance
  - not reliable
  - possible race conditions
  - not all system calls are 'signal' safe

Value	Name	Description
01	SIGHUP	Hang up; sent to process when kernel assumes that the user of that process is doing no useful work
02	SIGINT	Interrupt
03	SIGQUIT	Quit; sent by user to induce halting of process and production of core dump
04	SIGILL	Illegal instruction
05	SIGTRAP	Trace trap; triggers the execution of code for process tracing
06	SIGIOT	IOT instruction
07	SIGEMT	EMT instruction
08	SIGFPT	Floating-point exception
09	SIGKILL	Kill; terminate process
10	SIGBUS	Bus error
11	SIGSEGV	Segmentation violation; process attempts to access location outside its virtual address space
12	SIGSYS	Bad argument to system call
13	SIGPIPE	Write on a pipe that has no readers attached to it
14	SIGALARM	Alarm clock; issued when a process wishes to receive a signal after a period of time
15	SIGTERM	Software termination
16	SIGUSR1	User-defined signal 1
17	SIGUSR2	User-defined signal 2
18	SIGCLD	Death of a child
19	SIGPWR	Power failure

# Reliable POSIX Signals

- `sigaction()`: Install a signal handler
- `sigprocmask()`: Mask signals
- `sigsendset()`: Send signal
- `sigsuspend()`: Block until signal received

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>
#include <string.h>

#define message "SIGINT caught\n"

void signal_handler(int value)
{
    // printf is not signalsafe
    write( 1, message, strlen(message) );
}
```

# Reliable POSIX Signals (Cont.)

```
int main(int argc, char **argv, char **envp )
{
    struct sigaction new_action;
    int time_to_sleep;

    new_action.sa_handler = signal_handler;
    sigemptyset( &new_action.sa_mask );
    new_action.sa_flags = 0;

    sigaction( SIGINT, &new_action, NULL );

    printf("pid: %d\n", getpid() );
    time_to_sleep = 10;
    while( time_to_sleep )
        time_to_sleep = sleep( time_to_sleep );

    return 0;
}
```

# Signale Multithreading

- Which thread of a multithreaded app receives the signal?
  - Synchronous Signal (e.g. `SIGFPE`) ?
  - from `kill()` ?
- Another example:
  - `mutex1` protects critical section, also used in signal handler

```
void signal_handler( int sig )
{
    ...
    pthread_mutex_lock( &mutex1 );
    ...
    pthread_mutex_unlock( &mutex1 );
    ...
}
```



# Thanks

## Questions?

