

The background of the slide is a close-up, blue-tinted photograph of a hard drive's internal components. A shiny, circular metal platter is visible in the center-right, with a complex actuator arm and its read/write heads positioned above it. The image is out of focus, emphasizing the mechanical precision of the device.

Lecture Operating System

36. I/O Devices

36. I/O Devices

- 1. System Architecture**
- 2. Devices**
- 3. Case Study**



36. I/O Devices

1. System Architecture

2. Devices

3. Case Study




Input/Output (I/O) Devices

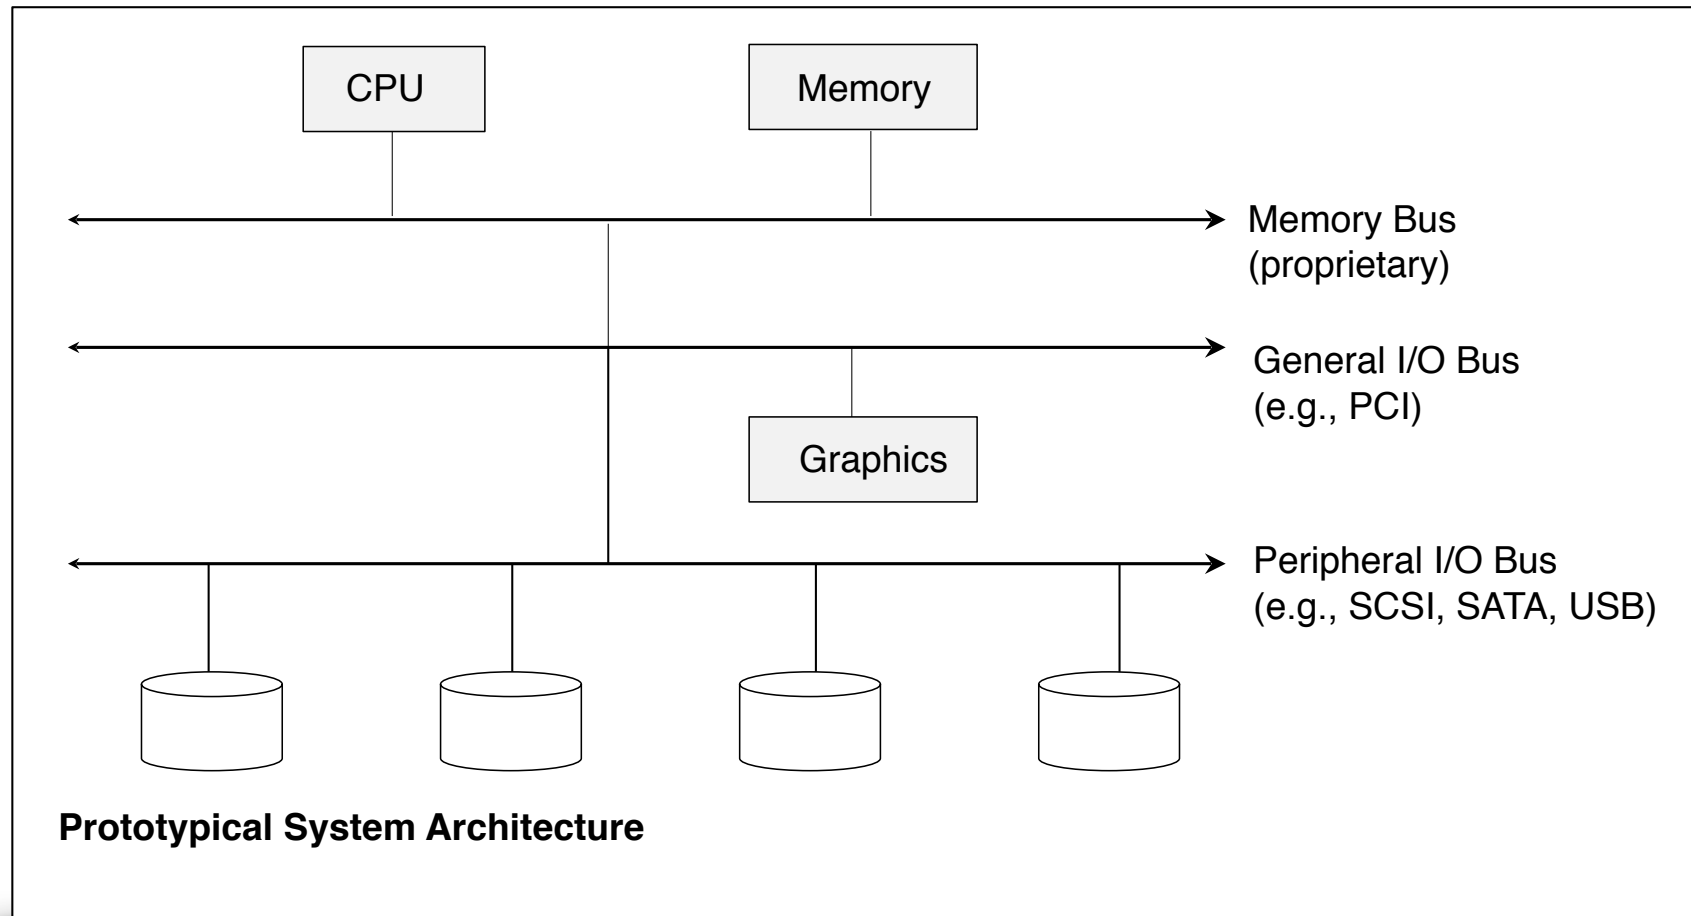
- What good is a computer without any I/O devices?
 - keyboard, display, disks
 - I/O is **critical** to computer system to **interact with systems**.
- We want:
 - H/W that will let us plug in different devices
 - OS that can interact with different combinations
- Issue :
 - How should I/O be integrated into systems?
 - What are the general mechanisms?
 - How can we make it efficiently?

I/O Devices are different!

■ e.g. data rate:



System Architecture



CPU is attached to the main memory of the system via some kind of memory **bus**.
Some devices are connected to the system via a general **I/O bus**.

Bus Architecture

■ Buses

- **Data paths** that provided to enable information between CPU(s), RAM, and I/O devices.

■ I/O bus

- Data path that **connects** a **CPU** to an **I/O** device.
- I/O bus is connected to I/O device by three hardware components:
 - I/O ports,
 - interfaces and
 - device controllers.

36. I/O Devices

1. System Architecture

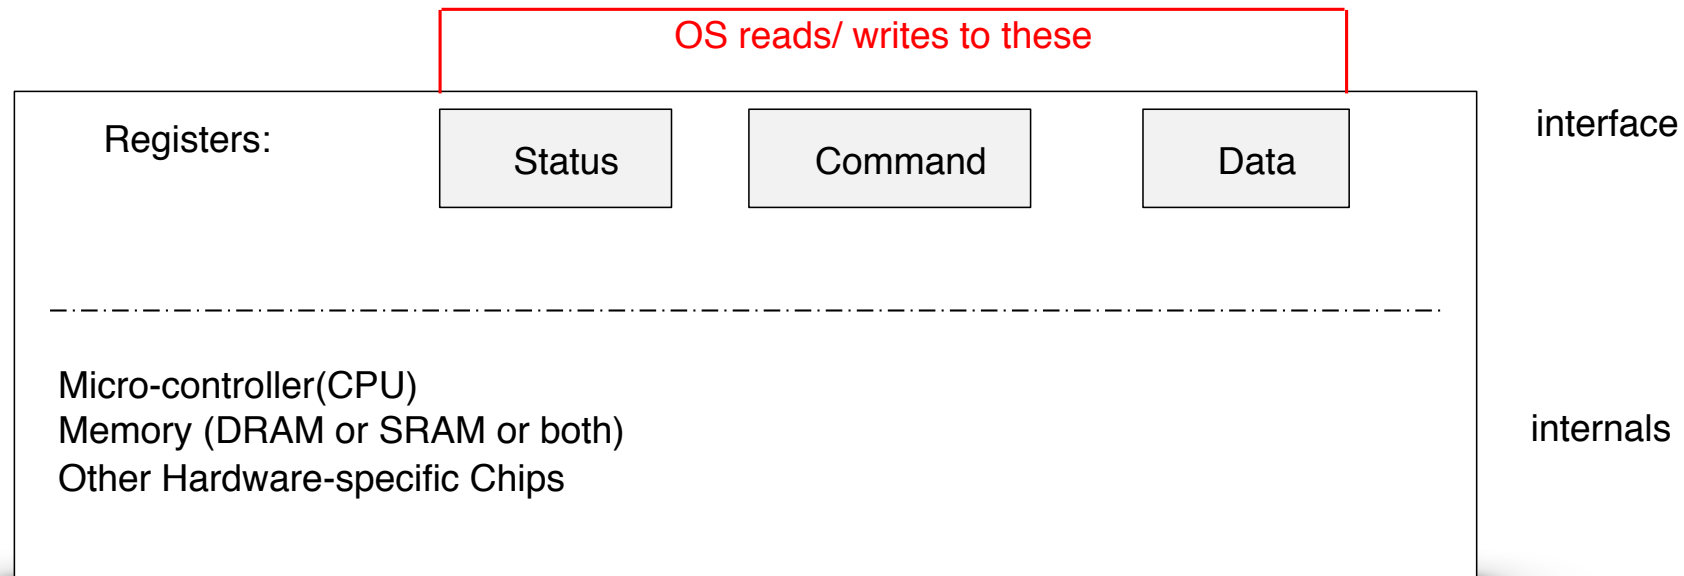
2. Devices

3. Case Study



Canonical Device

- Canonical Devices has two important components.
 - **Hardware interface** allows the system software to control its operation.
 - **Internals** which is implementation specific.



Canonical Device

Hardware interface of Canonical Device

■ **status register**

- See the current status of the device

■ **command register**

- Tell the device to perform a certain task

■ **data register**

- Pass data to the device, or get data from the device

By reading and writing above **three registers**,
the operating system can **control device behavior**.

Hardware interface of Canonical Device

Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    Doing so starts the device and executes the command
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```

A wants to do I/O



CPU:




Disk:



Hardware interface of Canonical Device

Typical interaction example

```
while ( STATUS = BUSY)                                     (1)
    ; //wait until device is not busy
write data to data register
write command to command register
    Doing so starts the device and executes the command
while ( STATUS = BUSY)
    ; //wait until device is done with your request
```




Hardware interface of Canonical Device

Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    Doing so starts the device and executes the command
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```

(2)



Hardware interface of Canonical Device


Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
```

(3)

Doing so starts the device and executes the command

```
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```




Hardware interface of Canonical Device

Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    Doing so starts the device and executes the command
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```


(4)



Hardware interface of Canonical Device

Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    Doing so starts the device and executes the command
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```



Polling

- Operating system waits **until** the device is **ready** by **repeatedly reading** the status register.
- Positive aspect is simple and working.
- However, it wastes CPU time just waiting for the device.
 - Switching to another ready process is better utilizing the CPU.

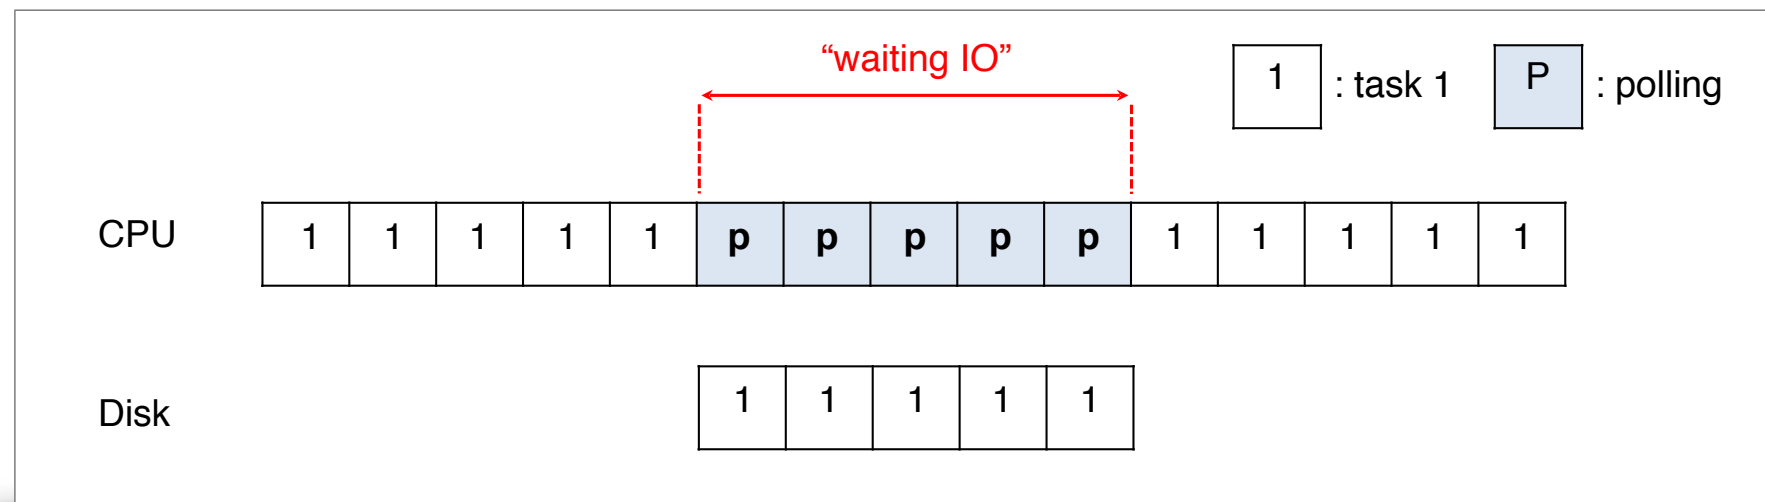


Diagram of CPU utilization by polling

Interrupts

- Put the I/O request process to **sleep** and context switch to another.
- When the device is finished, **wake** the process waiting for the I/O by interrupt.
 - Positive aspect is allow to CPU and the disk are properly utilized.




Diagram of CPU utilization by interrupt

Polling vs interrupts

- However, “**interrupts is not always the best solution**”
 - If, device performs very quickly, interrupt will “slow down” the system.
 - Because **context switch is expensive** (switching to another process)

If a device is fast → **poll** is best.
If it is slow → **interrupts** is better.

CPU is once again over-burdened

- CPU **wastes a lot of time** to copy the *a large chunk* of data from memory to the device.

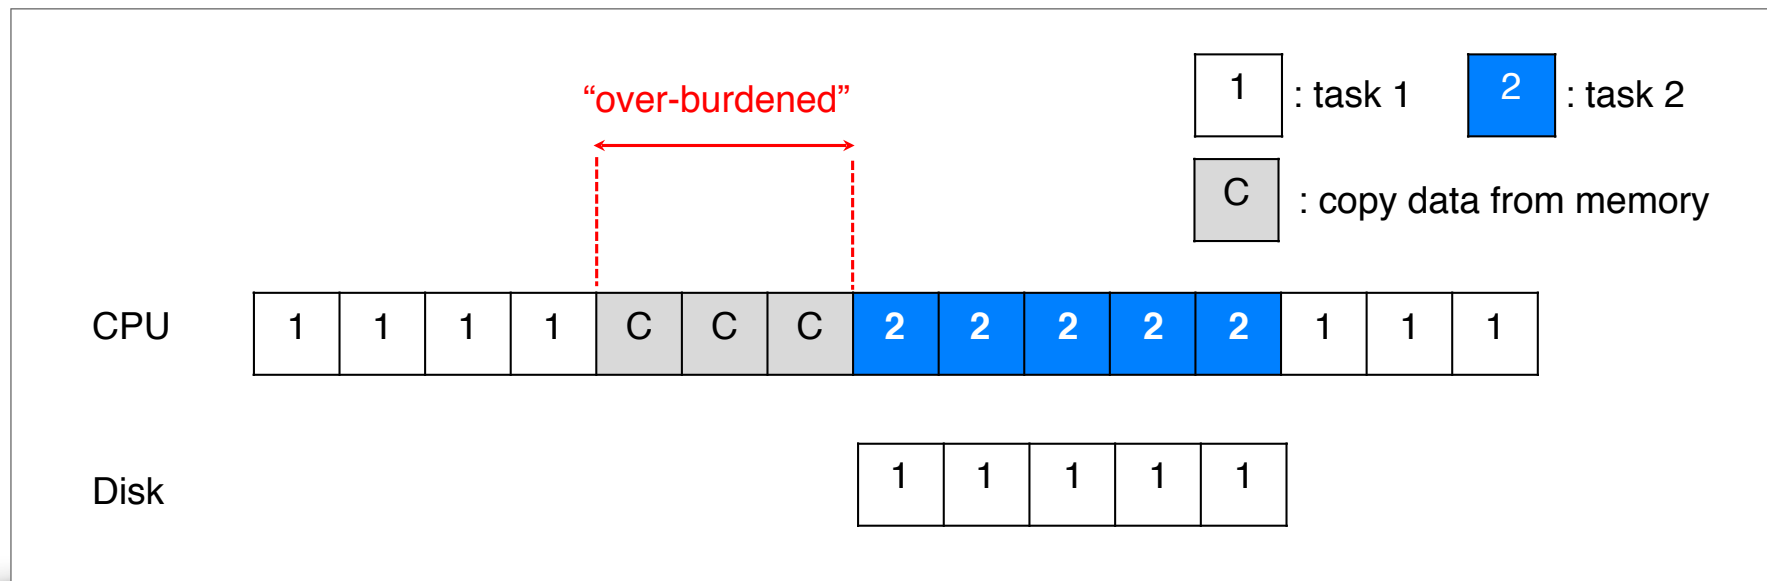


Diagram of CPU utilization

DMA (Direct Memory Access)

- **Copy data** in memory by knowing “where the data lives in memory, how much data to copy”
- When completed, DMA raises an interrupt, I/O begins on Disk.

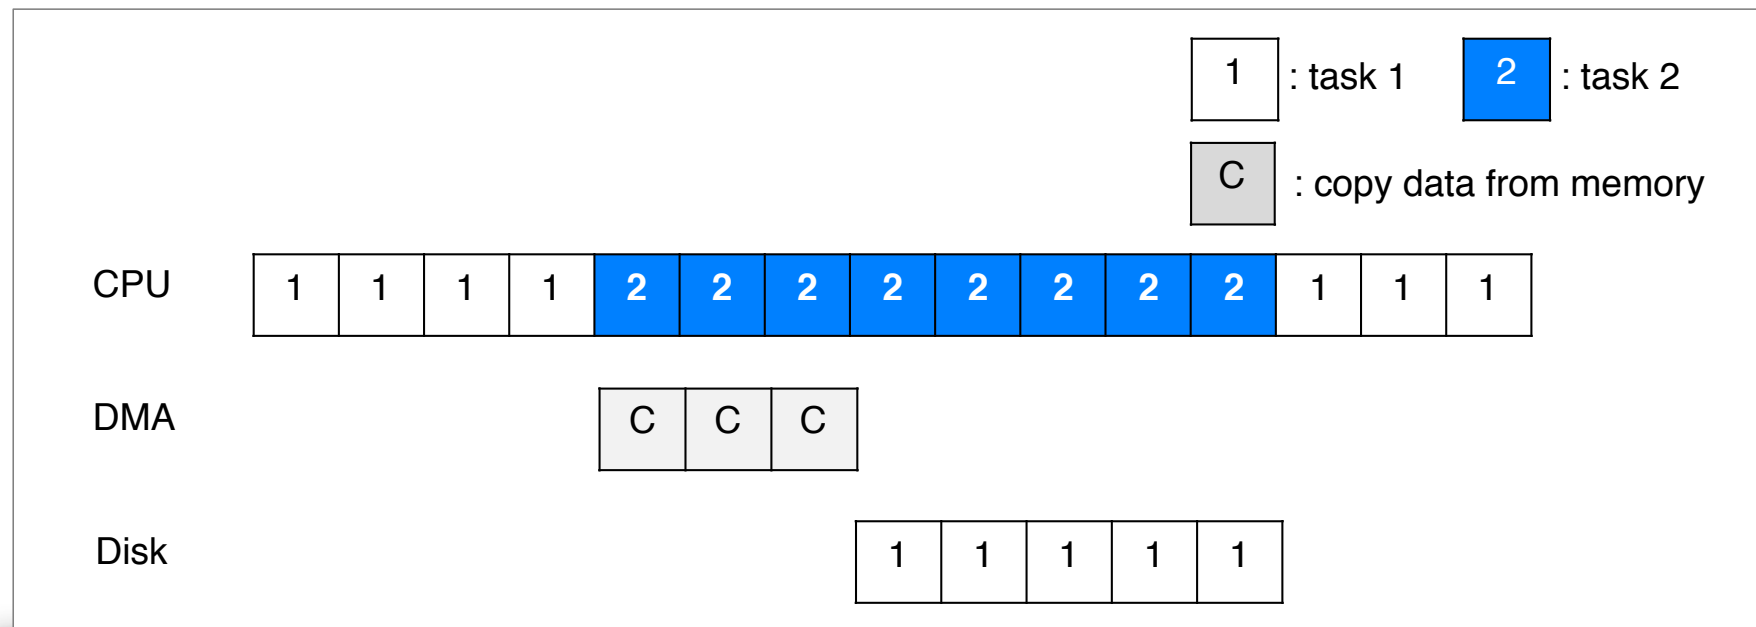
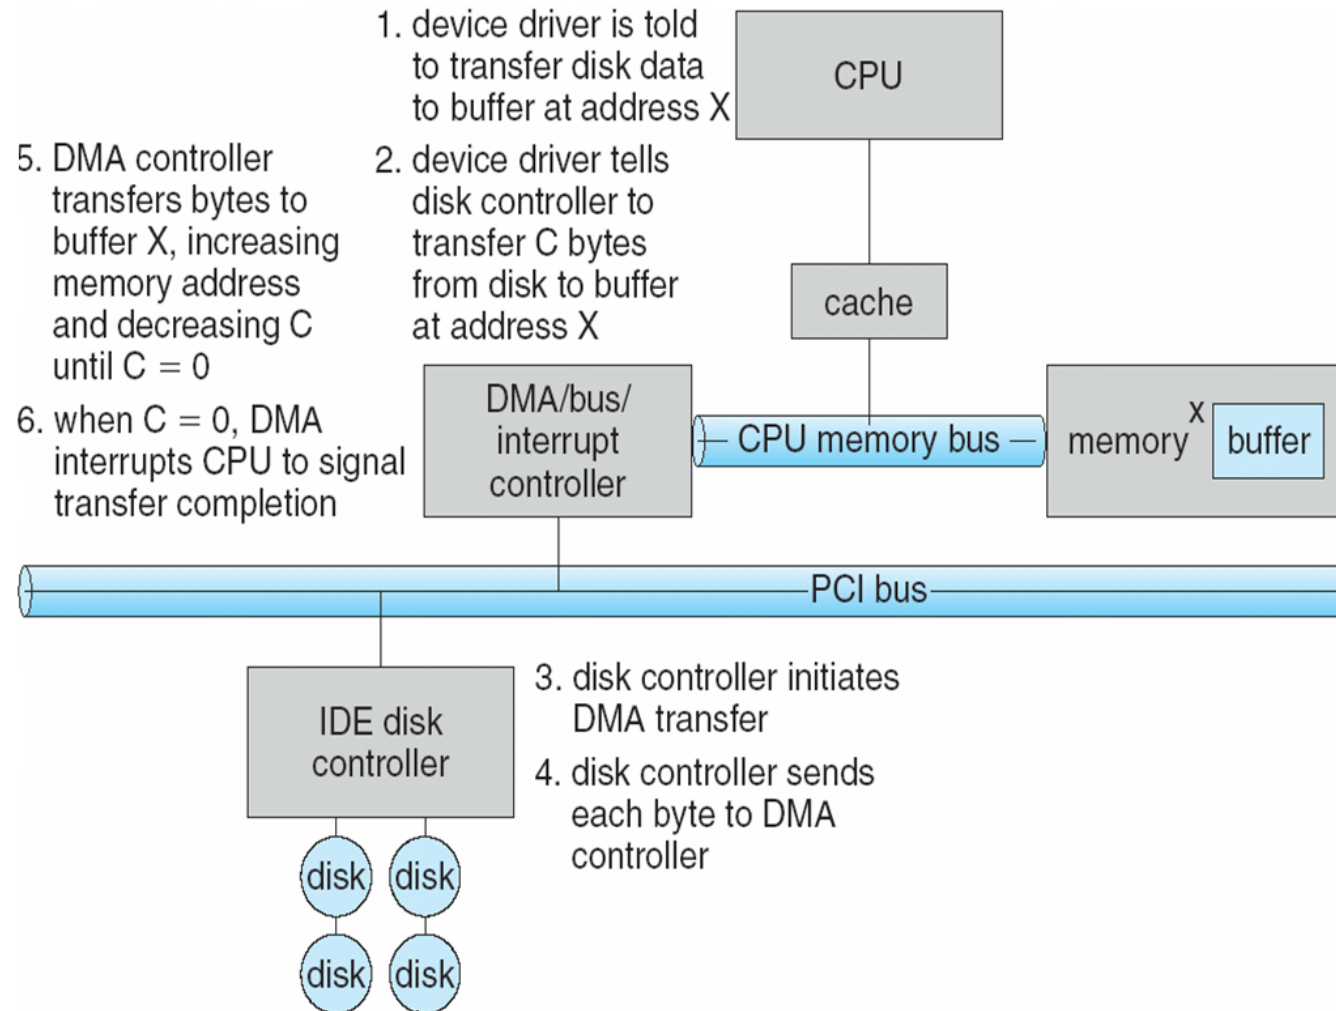


Diagram of CPU utilization by DMA

DMA Handshake



Example: OS Code

Polling

```
copy_from_user(buffer, p, count);          /* p is the kernel bufer */
for (i = 0; i < count; i++) {              /* loop on every character */
    while (*printer_status_reg != READY) ; /* loop until ready */
    *printer_data_register = p[i];         /* output one character */
}
return_to_user();
```

Interrupt

```
copy_from_user(buffer, p, count);
enable_interrupts();
while (*printer_status_reg != READY)
    *printer_data_register = p[0];
scheduler();
```

ISR

```
if (count == 0) {
    unblock_user();
} else {
    *printer_data_register = p[i];
    count = count - 1;
    i = i + 1;
}
acknowledge_interrupt();
return_from_interrupt();
```

DMA

```
copy_from_user(buffer, p, count)
set_up_DMA_controller();
scheduler();
```

ISR

```
acknowledge_interrupt();
unblock_user();
return_from_interrupt();
```

Device interaction

- How the OS communicates with the **device**?
- Solutions
 - **I/O instructions**: a way for the OS to send data to specific device registers.
 - Ex) `in` and `out` instructions on x86
 - **memory-mapped I/O**
 - Device registers available as if they were memory locations.
 - The OS `load` (to read) or `store` (to write) to the device instead of main memory.
- Doesn't matter much (both are used)

Protocol Variants

- **Status** checks:
 - polling vs. interrupts
- **Data**:
 - PIO vs. DMA
- **Control**:
 - special instructions vs. memory-mapped I/O

Device interaction (Cont.)

- How the OS interact with **different specific interfaces**?
 - Ex) We'd like to build a file system that worked on top of SCSI disks, IDE disks, USB keychain drivers, and so on.
- Solutions: **Abstraction**
 - Abstraction encapsulate **any specifics of device interaction**.

Device Drivers

- **Block** device drivers

- Manages devices that read and write data in blocks
 - Magnetic disc, CD-ROMS, etc.

- **Character** device drivers

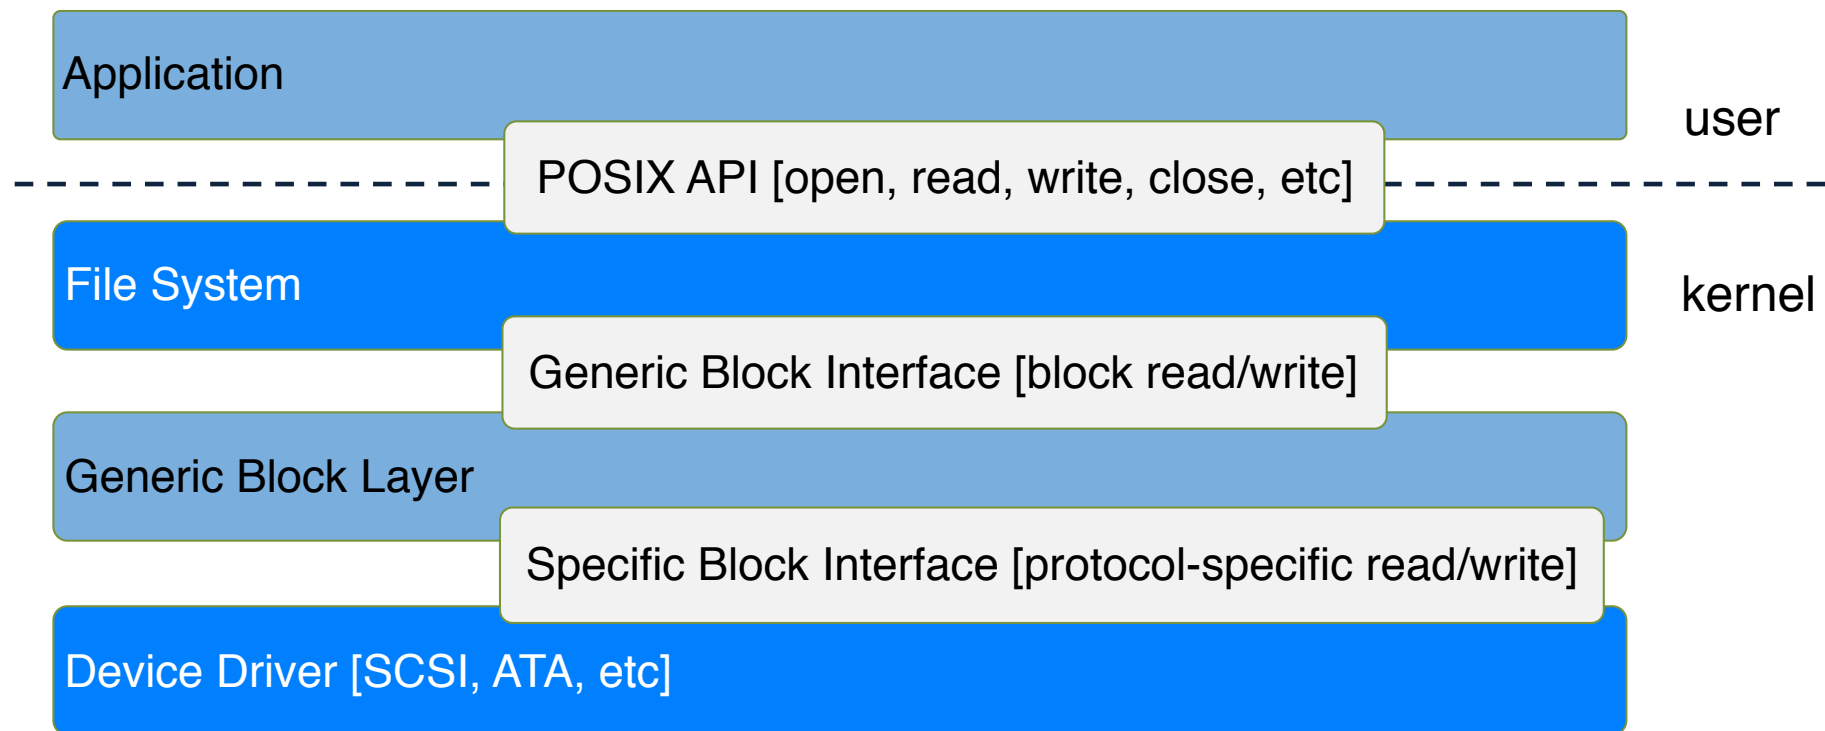
- Manages devices that require the kernel to send or receive data one byte at a time
 - terminals, modems, and printers.

- **Network** device drivers

- Manages network interface cards and communication ports that connect to network devices such as bridges and routers

File system Abstraction

- File system **specifics** of which disk class it is using.
 - Ex) It issues **block read** and **write** request to the generic block layer.

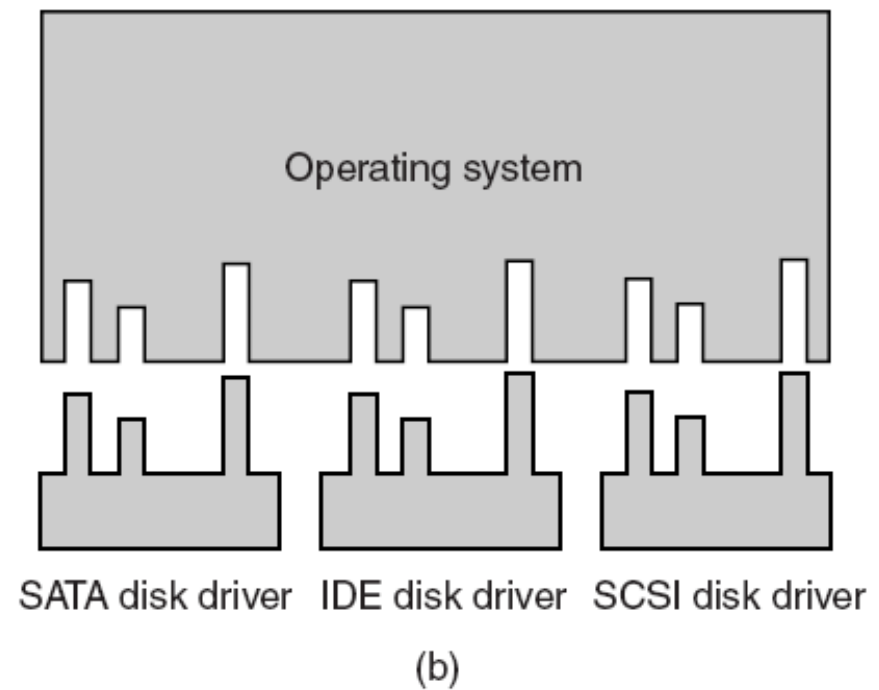



The File System Stack

Problem of File system Abstraction

- If there is a device having many special capabilities, these capabilities **will go unused** in the generic interface layer.
- **Over 70% of OS** code is found in device drivers.
 - Any device drivers are needed because you might plug it to your system.
 - They are primary contributor to **kernel crashes**, making **more bugs**.

Importance of Device Driver Interface



(a) Without a Standard Device Driver Interface

(b) With Standard Device Driver Interface

36. I/O Devices

1. System Architecture
2. Devices
- 3. Case Study**



A Simple IDE Disk Driver

- Four types of register
 - Control, command block, status and error
 - Memory mapped IO
 - `in` and `out` I/O instruction

Control / Command Register

■ **Control Register:**

- Address $0 \times 3F6 = 0 \times 80$ ($0000\ 1RE0$):
R=reset, E=0 means "enable interrupt"

■ **Command Block Registers:**

- Address $0 \times 1F0$ = Data Port
- Address $0 \times 1F1$ = Error
- Address $0 \times 1F2$ = Sector Count
- Address $0 \times 1F3$ = LBA low byte
- Address $0 \times 1F4$ = LBA mid byte
- Address $0 \times 1F5$ = LBA hi byte
- Address $0 \times 1F6$ = 1B1D TOP4LBA: B=LBA, D=drive
- Address $0 \times 1F7$ = Command/status

Status and Error Register

■ Status Register (Address 0x1F7):

7	6	5	4	3	2	1	0
BUSY	READY	FAULT	SEEK	DRQ	CORR	IDDEX	ERROR

■ Error Register (Address 0x1F1): (check when Status ERROR==1)

7	6	5	4	3	2	1	0
BBK	UNC	MC	IDNF	MCR	ABRT	T0NF	AMNF

- BBK = Bad Block
- UNC = Uncorrectable data error
- MC = Media Changed
- IDNF = ID mark Not Found
- MCR = Media Change Requested
- ABRT = Command aborted
- T0NF = Track 0 Not Found
- AMNF = Address Mark Not Found

Basic Protocol

- **Wait for drive to be ready.** Read Status Register (0x1F7) until drive is not busy and READY.
- **Write parameters to command registers.** Write the sector count, logical block address (LBA) of the sectors to be accessed, and drive number (master=0x00 or slave=0x10, as IDE permits just two drives) to command registers (0x1F2-0x1F6).
- **Start the I/O.** by issuing read/write to command register. Write READ—WRITE command to command register (0x1F7).
- **Data transfer** (for writes): Wait until drive status is READY and DRQ (drive request for data); write data to data port.
- **Handle interrupts.** In the simplest case, handle an interrupt for each sector transferred; more complex approaches allow batching and thus one final interrupt when the entire transfer is complete.
- **Error handling.** After each operation, read the status register. If the ERROR bit is on, read the error register for details.

A close-up, blue-tinted photograph of a precision industrial machine, likely a lathe or mill. A polished, cylindrical metal part is being machined, with a sharp tool bit visible in the background. The machine's structure is composed of various metal components, including a perforated blue plate in the foreground.

Thanks

Questions?