

# Programmierkurs 3

## Systemnahe Programmierung

Marcel Waldvogel

University of Konstanz

Winter 2016/2017

These slides are based on previous lectures, held by Alexander Holupirek, Roman Byshko, and especially Stefan Klinger.

0

Organisational stuff

## 0.1 Personnel

Marcel Waldvogel Distributed Systems Laboratory (di.sy)  
mail/xmpp [marcel.waldvogel@uni-konstanz.de](mailto:marcel.waldvogel@uni-konstanz.de)

Klaus Herberth Distributed Systems Laboratory (di.sy)  
mail/xmpp [klaus.herberth@uni-konstanz.de](mailto:klaus.herberth@uni-konstanz.de)

Help Channel XMPP  
xmpp [info3@conference.uni-konstanz.de](mailto:info3@conference.uni-konstanz.de)

## 0.2 Informatik 3

- ▶ Together with **“Betriebssysteme”** (OS, *aka.* “Operating Systems”), this lecture forms the module “Informatik 3”
- ▶ You need to gain **50% of the exercise points** in **each** of the exercises be admitted for the final exam (“Three strikes and you’re out”).
- ▶ There will be  $\approx 30\%$  PK3 questions in the joint exam.
- ▶ **Registration**
  - For the exercises, form groups of two.
  - Sign up for OS **and** PK3 via StudIS<sup>1</sup>, as usual<sup>2</sup>.

---

<sup>1</sup><https://studis.uni.kn>

<sup>2</sup><http://www.informatik.uni.kn/studieren/studium/pruefungen/pruefungsanmeldung/>

## 0.3 Coordinates

### Weekly schedule

Lectures Monday

OS tutorials Thursday

PK3 tutorials Friday

Handin Thursday, 04:00 via **git**

Handout Thursday, ≈10:00 via **git**

**Teaching materials** are available via git<sup>3</sup> on GitLab<sup>4</sup>. One member of the group forks the repository<sup>5</sup> into his personal space and

- ▶ **grants access** to the other member (at least 'Developer'),
- ▶ **shares** it with the group 'info3' (at 'Master' level).

You can then browse all files in the repository, which includes instructions on how to clone it into a local repository.

<sup>3</sup><https://git-scm.com/>

<sup>4</sup><https://git.uni.kn>

<sup>5</sup><https://git.uni.kn/info3/bspk2016>

## The main command is

```
1 $ git clone git@git.uni.kn:user-name/bspk2016.git
2 Cloning into 'bspk2016'...
3 remote: Zähle Objekte: 20, Fertig.
4 remote: Komprimiere Objekte: 100% (18/18), Fertig.
5 remote: Total 20 (delta 3), reused 0 (delta 0)
6 Receiving objects: 100% (20/20), 1.09 MiB | 0 bytes/s, done.
7 Resolving deltas: 100% (3/3), done.
8 Checking connectivity... done.
```

This works only after having set up [ssh](#) (with keys) and [git](#). You will learn how to do this in the tutorials.

## 0.4 PK3 Assignments

- ▶ You need to form groups of two to work on the exercises. This is **organised during the BS lecture**, see there.
- ▶ One **PK3**-assignment every week
  - Every new assignment is released on **Thursday at 10:00**,
  - due the following Thursday, before 04:00.
  - Discussed on **Thursday/Friday** in the tutorials.
- ▶ Submit your exercises via [git](#):
  - Each group has r/w access only to their fork.
  - Commit your solutions to: [/ass\\_pk](#), where ass is the 2-digit assignment number.<sup>6</sup>

Especially in systems programming, just because your code works does not mean it is correct!

---

<sup>6</sup>Put your solution in the directory mentioned before — we will not look in other places.

► We are **quite strict** about compiler errors and warnings:

- Compile (we will) your code with

```
1 $ gcc -std=c99 -g -Wall -Wextra -Wpedantic -Wbad-function-cast \  
2 > -Wconversion -Wwrite-strings -Wstrict-prototypes source.c
```

- You will gain **no points at all** for a programming exercise if the compiler stops with an **error**.
- We will subtract **3 points** for every **compiler warning**.

► **Bonus points** for nice code:

- Use [checkpatch.pl](#) from the Linux kernel project<sup>7</sup> against your source:

```
1 $ checkpatch.pl --no-tree --no-signoff -f --ignore NEW_TYPEDEFS,\  
2 > AVOID_EXTERNS,GLOBAL_INITIALISERS,BLOCK_COMMENT_STYLE source.c  
3 total: 0 errors, 0 warnings, 44 lines checked
```

- Your point score for a program will **increase** by 5% if [checkpatch](#) generates warnings but **no errors**, and by **10%** if there are neither errors **nor warnings**!

► The Tutors will show you on Friday how to use the compiler and the [checkpatch](#) script.

<sup>7</sup>Included in your repository under [/pk\\_code](#). Source: <https://github.com/torvalds/linux/blob/c5595fa/scripts/checkpatch.pl>



- It makes sense to set up some aliases in you `~/ .bashrc`:

```
1 # a Very Picky C Compiler
2 alias pk-cc='gcc -std=c99 -g -Wall -Wextra -Wpedantic -Wbad-function-cast \
3 -Wconversion -Wwrite-strings -Wstrict-prototypes';
4
5 # using checkpatch from the Linux kernel
6 alias pk-chk='checkpatch.pl --no-tree --no-signoff -f [...]';
```

- and then use

```
1 $ pk-cc source.c                                     # compiles
2 $ pk-chk source.c
3 total: 0 errors, 0 warnings, 44 lines checked           # and looks nice
4
5 source.c has no obvious style problems and is ready for submission.
```

## 0.5 Literature

- ▶ Brian W. Kernighan, Dennis M. Ritchie. *The C Programming Language*. 1978, Prentice Hall Software Series. Uni-KN kid 248 k27.



- ▶ Peter van der Linden. *Expert C Programming — Deep C Secrets*. 1994, Sunsoft Press, Prentice Hall. ISBN 0-13-177429-8, Uni-KN kid 248 v16.
- ▶ Randal E. Bryant, David O'Hallaron. *Computer Systems — A Programmer's Perspective*. 2003, Pearson Education International, Prentice Hall. ISBN 0-13-178456-0, Uni-KN kid 100n b79.

## 0.6 What is this course about?

### System Programming

- ▶ With **system** we mean *operating system*.
- ▶ With **programming** we mean *using the interface* an operating system (OS) provides.
- ▶ With OS we mean UNIX-like OSs, *i.e.*, Linux.

### Operating System

- ▶ Layer of software on top of bare hardware
- ▶ Shields programmers from the complexity of the hardware
- ▶ Presents an interface (of a virtual machine) that is easier to understand and program

## Systems vs. Kernel programming

- ▶ Black Box Modell is suitable for systems programming.
- ▶ Knowledge about the system's internals, however, is beneficial to use the system properly and to not work against it.
- ▶ Providing the system services is (mostly) kernel programming.

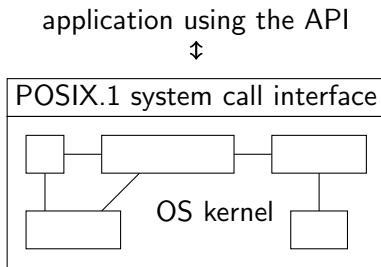
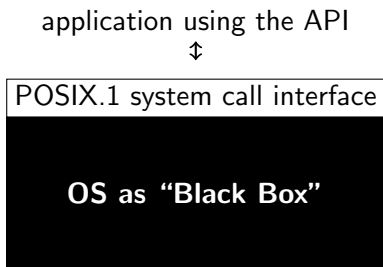


Figure: Black Box vs. White Box View of a UNIX System

## 0.7 Our working environment

- ▶ Get GNU/Linux up and running on your machine.
  - For PK3 you will work *under* Linux.
  - For BS you will work *on* Minix.
- ▶ Use the virtual machine set up as part of the tutorials.

### You will need

- ▶ An editor, e.g., [vim](#), [emacs](#), [nano](#), [geany](#), ... Eclipse is not recommended!
- ▶ [gcc](#) — the GNU project C (and C++) compiler
- ▶ ...

1

# Gentle introduction to C

## 1.1 C standardization

- ▶ ISO/IEC 9899:1990 Programming Language C, (C89 or C90)
- ▶ ISO/IEC 9899:1999 Programming Language C, (C99)
- ▶ ISO/IEC 9899:2011 Programming Language C, (C11)

**Note** We will focus on C99, *i.e.*, use `-std=c99` as compiler flag.

## C popularity

### ► Requirements that make C mandatory:

- embedded systems (close to hardware, scarce resources)
- extreme performance (better usage of resources)
- the world is built on C and C++ (with C++ being a superset of C)  
— Herb Sutter. C++ and Beyond.<sup>8</sup>
- C is simple & powerful  
— Damien Katz (CouchDB). The Unreasonable Effectiveness of C.<sup>9</sup>

### ► Programming Languages Rankings

- 2nd place in TIOBE<sup>10</sup> (October 2015)
- 9th place in RedMonk<sup>11</sup>, with C++ ranking 5th (June 2015)

---

<sup>8</sup><https://www.youtube.com/watch?v=xcwxGzbTyms>

<sup>9</sup>[http://damienkatz.net/2013/01/the\\_unreasonable\\_effectiveness\\_of\\_c.html](http://damienkatz.net/2013/01/the_unreasonable_effectiveness_of_c.html)

<sup>10</sup><http://www.tiobe.com/index.php/content/paperinfo/tpci/index.html>

<sup>11</sup><http://redmonk.com/sograzy/2015/07/01/language-rankings-6-15/>



## 1.2 First C Program

### Print the sentence: “Hello world!”

- ▶ Create the program text
- ▶ Compile it successfully
- ▶ Run it
- ▶ Get the output

```
1 #include <stdio.h>
2
3 main()
4 {
5     printf("Hello world!\n");
6 }
```

## Compilation on a UNIX-like OS

```
1 $ gcc hello.c
2 hello.c:3:1: warning: return type defaults to 'int' [-Wimplicit-int]
3     main()
4     ~
5 $ ls
6 a.out  hello.c
7 $ ./a.out
8 Hello world!
```

engine	filename	description
	hello.c	source code
preprocessor	hello.i	source w/ preproc. directives expanded
compiler	hello.s	assembler code
assembler	hello.o	object code ready to be linked
linker	a.out	executable

(Use `-save-temps` to preserve these files)

## 1.3 Program structure

### Basic building blocks

- ▶ *functions* contain *statements*
- ▶ *statements* specify computing operations to be done
- ▶ *variables* store values used during computation
- ▶ *arguments* (one way to) communicate data between functions

## Building blocks of our example

```
1 #include <stdio.h>
2
3 int main(void)
4 {
5     printf("Hello world!\n");
6     return 0;
7 }
```

- ▶ line 1: include the standard input/output library
- ▶ line 3–7: define a function `main`
  - `main` is special, a program begins execution here
  - `main` will usually call other functions to help perform its job
  - You can define your own functions, but don't name them `main`
  - `main` returns an `int`. (It has to!)
- ▶ line 3: Parentheses after the function name surround the argument list, `void` means no arguments.
- ▶ line 5: `main` calls library function `printf`, which prints this sequences of characters; `\n` represents the newline character.

# Arithmetics

Fahrenheit-Celsius:  $^{\circ}\text{C} = (5/9)(^{\circ}\text{F}-32)$

```
1 /* print fahrenheit-celsius table for fahrenheit = 0, 20, ..., 300 */
2
3 #include <stdio.h>
4
5 int main(void)
6 {
7     int fahr, celsius;
8     int lower, upper, step;
9
10    lower = 0;    /* lower limit */
11    upper = 300;  /* upper limit */
12    step = 20;    /* step size */
13
14    fahr = lower;
15    while (fahr <= upper) {
16        celsius = 5 * (fahr - 32) / 9;
17        printf("%d\t%d\n", fahr, celsius);
18        fahr = fahr + step;
19    }
20    return 0;
21 }
```

## Running:

```
1 $ ./a.out
2 0      -17
3 20     -6
4 40      4
5 60     15
6 80     26
7 100    37
8 120    48
9 140    60
10 160    71
11 180    82
12 200    93
13 220   104
14 240   115
15 260   126
16 280   137
17 300   148
```

## Declarations and assignment statements

- ▶ A *declaration* announces the properties of variables. Consists of *type name* and a *list of variables*, such as:

```
7     int fahr, celsius;  
8     int lower, upper, step;
```

Range/ size of data types depends on machine

- ▶ *Assignment statements* set the variables to their initial values.

```
10    lower = 0;    /* lower limit */  
11    upper = 300;  /* upper limit */  
12    step = 20;    /* step size */
```

## Basic data types

`char` a single byte. By definition, this is the unit of measurement for memory size.

`int` an integer, typically reflecting the natural size of integers on the host machine

`float` single-precision floating point

`double` double-precision floating point

`short` and `long` are *qualifiers* that can be applied to integers:

```
short int i;  
long int f;  
unsigned long d;
```

The qualifiers `signed` and `unsigned` can be applied to `char` and any integer.

## The while loop

Each line in the result table is computed the same way:

```
15 while (fahr <= upper) {  
16     celsius = 5 * (fahr - 32) / 9;  
17     printf("%d\t%d\n", fahr, celsius);  
18     fahr = fahr + step;  
19 }
```

**Note** that  $^{\circ}\text{C} = (5/9)(^{\circ}\text{F} - 32)$  is computed as

```
16 celsius = 5 * (fahr - 32) / 9;
```

- **Integer division truncates**, i.e., any fractional part is discarded. Since 5 and 9 are integers,  $5/9$  would be truncated to zero and so all the Celsius temperature would be reported as zero.

⇒ Be careful with integer divisions.



## printf revisited

```
#include <stdio.h>
int printf(const char *format, ...);
```

`printf`(3) is a general-purpose output formatting function.<sup>12</sup>

- ▶ 1<sup>st</sup> argument is the string of characters to be printed.
  - Each **%** indicates **where** one of the other arguments
  - and **in what form** it is to be printed.
- ▶ Each **%** in the 1st arg is paired with the 2nd, 3rd arg etc.

17

```
printf("%d\t%d\n", fahr, celsius);
```

- ▶ **%d**, for instance, specifies an integer argument, so `fahr` and `celsius` are printed with a tab (`\t`) between them.

---

<sup>12</sup>Not part of the C language, but defined in ANSI X3.159-1989 ("ANSI C")