

## Example

```
1 size_t strlen(const char *s);    /* available with #include <string.h> */
2
3 size_t (*fp)(const char *);
4 fp = &strlen;
5 printf("result = %zu\n", (*fp)("hello world"));
```

- ▶ This can be abbreviated:

```
5 printf("result = %zu\n", fp("hello world"));
```

- ▶ Nicer with `typedef`

```
2 typedef size_t (*func)(const char *);
3 func fp = &strlen;
```

**Note** Function pointers are heavily used in the real world! *E.g.*,

- ▶ pass a comparing function to a queue datastructure;
- ▶ installing signal handlers (*cf.* OS lecture, and later in this course); and
- ▶ abstractions (syscall interface, subclasses, VFS, ...).

**Question** Can you read `int ((*f)(void))[2]` ?

## 4.7 Unscrambling C declarations

**Precedence rules** for reading C declarations.

1. Parentheses group parts of the declaration.
2. Read **type specifiers** as atomic tokens, e.g.,
  - `double`,
  - `struct foo`, or
  - `unsigned short int`.
3. The keyword `const`:
  - If *next to* a type specifier, it belongs to that, making the **value** constant.
  - Otherwise, it belongs to the asterisk to its *left*, making the **pointer** constant.
4. The **postfix** operators, being one of
  - **parentheses** `(...)` indicating a function, or
  - **brackets** `[...]` indicating an array.
5. The **prefix** operator **asterisk** `*` indicating a pointer.

**Note** Inside parenthesis, a declaration may contain *further* declarations of function arguments! These do *not necessarily* have a name.

## An algorithm for reading declarations

1. Start at the leftmost identifier that is *not* a type specifier. That is being declared.
2. Do not leave parenthesis while:
  - 2.1 Handle the **postfix** operators, *i.e.*, optional (...) or [...] to the **right**, do so from left to right.
    - ▶ For a function, apply the whole algorithm to each parameter.
    - ▶ For an array, optionally note the size.
  - 2.2 Handle the **prefix** operators \* to the **left**, do so from right to left.
3. If inside parenthesis, leave them, and restart with 2.
4. Read the **type specifier** on the left.

*tl;dr* — look right, look left.

### Example `int *(*list[42])(void)`

- ▶ `int *(*list[42])(void)`    `list` is...
- ▶ `int *(*list[42])(void)`    ...an array of 42...
- ▶ `int *(*(list)[42])(void)`    ...pointers to

Leaving parenthesis, we're done with them. Goto step 2 of algorithm:

- ▶ `int *(*(list)[42])(void)`    ...function of ...
- ▶ `int *(*(list)[42])(void)`    ...no arguments...
- ▶ `int *(list)[42])(void)`    ...returning a pointer to...
- ▶ `int *(*(list)[42])(void)`    ...an integer.

## Example `int (*f)(const char *s)`

- ▶ `int (*f)(const char *s)`    `f` is...
- ▶ `int (*f)(const char *s)`    ...a pointer to...
- ▶ `int (*f)(const char *s)`    ...a function of (...)
- ▶ `int (*f)(const char *s)`    ...`s`, which is...
- ▶ `int (*f)(const char *s)`    ...a pointer to...
- ▶ `int (*f)(const char *s)`    ...a `constant character` )...
- ▶ `int (*f)(const char *s)`    ...returning an integer.

**Example** `void f(char *x[])`

- ▶ `void f(char *x[])`    `f` is a...
- ▶ `void f(char *x[])`    ...function of (...)
- ▶ `void f(char *x[])`    ...`x`, which is...
- ▶ `void f(char *x[])`    ...an array of unspecified size of...
- ▶ `void f(char *x[])`    ...pointers to...
- ▶ `void f(char *x[])`    ...character )...
- ▶ `void f(char *x[])`    ...not returning anything.

The declaration `void f(char **x)` is equivalent, specifying array dimensions does not make any sense in this case (*cf.* page 79).

## Example `void *f(char *(*p)[5])`

- ▶ `void *f(char *(*p)[5])`    `f` is a...
- ▶ `void *f(char *(*p)[5])`    ...function of (...)
- ▶ `void *f(char *(*p)[5])`    ...`p`, which is...
- ▶ `void *f(char *(*p)[5])`    ...a pointer to...
- ▶ `void *f(char *(*p)[5])`    ...an array of five...
- ▶ `void *f(char *(*p)[5])`    ...pointers to...
- ▶ `void *f(char *(*p)[5])`    ...character )...
- ▶ `void *f(char *(*p)[5])`    ...returning a pointer to...
- ▶ `void *f(char *(*p)[5])`    ...data of unspecified type.

In this case, specifying the array dimensions makes sense: In the body of `f`, `sizeof(*p)` will return 40 if the size of a pointer is 8. This also effects pointer arithmetics on `p`.

**Note** Function parameters need not be named in a **declaration**!

```
double (*f)(double x)  ≡  double (*f)(double)
```

This makes it occasionally hard to find out what is being declared.

**Example** `int f(char *[])` (Example from page 136)

- ▶ `int f(char *[])`    `f` is a...
- ▶ `int f(char *[])`    ...function of (...)

No identifier: So “it” is to the right of all `*`, and to the left of all `(...)` and `[...]`.

- ▶ `int f(char *[])`    ...an array of...
- ▶ `int f(char *[])`    ...pointers to...
- ▶ `int f(char *[])`    ...character )...
- ▶ `int f(char *[])`    ...returning an integer.

This is actually equivalent to `int f(char **)`.

**Question** What is this: `int f(char (*)[23])` ?



## Easily spotted mistakes

Some observations about **parentheses** in declarations (Note: ... is a meta-placeholder!):

- **Invalid types**, *i.e.*, in a type declaration, you will **never** see

`foo(...)(...)` Functions cannot return functions.

`foo(...) [...]` Functions cannot return arrays.

`foo[...](...)` Arrays cannot contain functions.

- **Valid types**

`int bar[...] [...]`; `bar` is an array of arrays.

`int (*fun(...))(...)`; Function `fun` returns a *pointer to* a function.

`int (*foo(...)) [...]`; Function `foo` returns a *pointer to* an array.

`int (*arr[...])(...)`; `arr` is an array of *pointers to* functions.

# 5

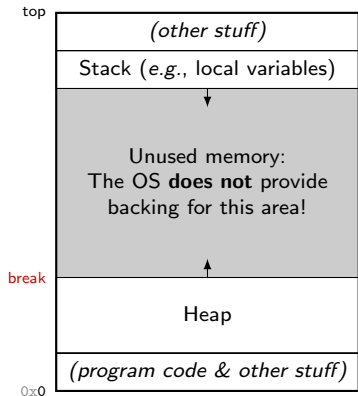
## Excursus: Inside Malloc

This excursus demonstrates:

- ▶ How to impose a meaning on a region of memory.
- ▶ Heavy use of pointer arithmetics.
- ▶ Glimpse under the hood of memory allocation.

## 5.1 Process memory layout

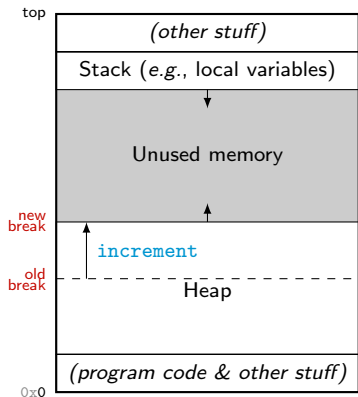
The program's view of its memory: **virtual RAM**.



(simplified picture)

- ▶ The stack contains the variables local to a function call. Grows downwards.
- ▶ The **program break** is the first location after the program's *data segment* (*cf.* later).
- ▶ Incrementing the break is **allocating heap memory**: Ask the OS to provide backing for the increased consumption of memory.

## Allocating heap memory



- ▶ `sbrk(2)` can **move the break**<sup>23</sup>.

```
1 void *sbrk(intptr_t increment);
```

- ▶ Returns address of old break, or `(void *)-1` on error.
- ▶ If the break was *increased*, then the returned value is a pointer to **newly allocated** memory, backed by the OS!
- ▶ (There are other system calls to get memory from the OS, *cf.* later)

**Note** Avoid using `brk()` and `sbrk()`: the `malloc(3)` memory allocation package is the portable and comfortable way of allocating memory.

<sup>23</sup>there's also `brk(2)`, for the same purpose — we use `sbrk(2)` only.

## 5.2 Implementing a memory allocator

How to write your own `malloc`<sup>25</sup>

- ▶ We know that we can get **fresh memory** from the OS via `sbrk(2)`.
- ▶ “The real” `malloc(3)` uses this, and other techniques. There are many different, *very sophisticated* implementations of memory allocators.
- ▶ We implement a **very simple** allocator.<sup>24</sup> Most prominently, we ignore data **alignment**.

### The Interface

```
1 void *kr_malloc(size_t b);  
2 void kr_free(void *ap);
```

- ▶ `kr_malloc` allocates `b` bytes and returns a pointer to the allocated memory, or `NULL` on error.
- ▶ `kr_free` frees memory pointed to by `ap`, which must have been returned by a previous call to `kr_malloc`.

<sup>24</sup>adapted from: Kernighan, Ritchie. The C Programming Language. Prentice Hall Software Series. Section 8.7, *A Storage Allocator*.

<sup>25</sup>Just because you can — in general, this is not a smart idea.

## The rough plan


- ▶ System calls are **expensive**: Avoid using them often.
- ▶ So `kr_malloc` tries to get a **big chunk** of memory from the OS, and hands **smaller pieces** of that to the calling program.
- ▶ So we need to maintain a **list of free memory chunks**, that
  - have been allocated **from the OS** via `sbrk(2)`, but
  - have **not yet** been handed to the **program**,
  - or have been **returned from** the program.
- ▶ If the program frees memory, `kr_free` adds that to the list of free memory, but does not return it to the OS.
  - Obvious weakness: Memory consumption of the process never shrinks.

**Note** The functions `sbrk(2)`, `kr_malloc` and `kr_free` implement a concept of **transferring ownership** of memory between the OS, the allocator, and the program.

## Chunks of free memory

Our allocator maintains a list of **free memory chunks**. These are not currently used by the program, *i.e.* they

- ▶ lie in memory allocated from the OS via `sbrk`,
- ▶ have not been given to the program by returning from `kr_malloc`, or have been given back to the allocator via `kr_free`.




- ▶ **Circular list:** Every chunk points to the next chunk.
- ▶ List is **ordered by memory address**, with the obvious exception.
- ▶ A pointer `freep` is the **entry point** into the list. It may point to any chunk, and we will move it around quite a bit.

To maintain the list of free chunks, we install a header **at the start** of each chunk:

```
1 typedef struct header Header;  
2 struct header {  
3     Header *next;  
4     size_t size;  
5 };  
6  
7 #define BLOCKSIZE (sizeof(Header))
```

- ▶ **next** points to the next chunk in the circular list.
- ▶ **size** is the size of the **entire** chunk, given in the unit **BLOCKSIZE** bytes.



**Questions** If **Header \*p** points to the header, then

- ▶ where does **p + 1** point to?
- ▶ where does **p + p->size** point to?



## 5.3 The kr\_malloc() function

```

1 Header *freep = NULL, /* a global pointer to the free list */
2     base; /* and a dummy for the empty list */
3
4 void *kr_malloc(size_t bytes) {
5
6     /* number of blocks required, including one more for the header */
7     size_t reqd = 1 + (bytes + BLOCKSIZE - 1) / BLOCKSIZE;
8     Header *prevp = freep; /* ptr to previous chunk */
9
10    if (!freep) { /* make empty list if called for the first time */
11        base.next = freep = prevp = &base;
12        base.size = 0;
13    }
14
15    for (Header *p = prevp->next; ; prevp = p, p = p->next) {

```

Check the chunk `*p`. `return` a pointer if it is big enough. See the following slides.

```

27         if (p == freep) { /* if we have unsuccessfully traversed the whole list... */
28             p = morecore(reqd); /* ...get more from the OS (cf. page 154)... */
29             if (p == NULL) return NULL; /* ...or fail. */
30         }
31     }
32 }

```

## `kr_malloc()` — using a chunk that fits exactly

- ▶ Header `*p` points to current chunk, `*prevp` to the previous one.
- ▶ We need `reqd` blocks of free memory.

```

16 if (p->size >= reqd) { /* this chunk is large enough */
17     if (p->size == reqd) /* it fits exactly */
18         prevp->next = p->next; /* remove chunk from free list */
19     else {


```

Split the chunk `p` points to. See the following slides.

```


23     }
24     freep = prevp; /* next search continues from here */
25     return p + 1; /* memory address the program may write to */
26 }

```



## kr\_malloc() — split a chunk that is too large


Assume we need `reqd = 23` blocks, but the chunk has 42...



20  
21  
22

```
p->size -= reqd;
p += p->size;
p->size = reqd;
```

/\* Enjoy: We just make up a header here! \*/




## 5.4 The `kr_free()` function

Where in the list should the freed chunk be linked?

```

1 void kr_free(void *ap) {
2     Header *p,
3     *headp = (Header *)ap - 1; /* determine header of chunk *ap */
4
5     /* Find p so, that headp belongs between p and p->next. */
6     for (p = freep; !(p < headp && headp < p->next); p = p->next)
7         if (p >= p->next && (p < headp || headp < p->next)) break;

```




**Question** Now it would be easy to hook the freed chunk into the list:

```
headp->next = p->next;    p->next = headp;
```

Why may this not be the smartest thing to do?


## `kr_free()` — fuse/link with the following chunk




```

10  /* If the chunk is adjacent to the following chunk, fuse the two into one... */
11  if (headp + headp->size == p->next) {
12      headp->size += p->next->size;
13      headp->next = p->next->next;
14  } else
15      headp->next = p->next; /* ...otherwise just link without fusing. */

```



## `kr_free()` — fuse/link with the previous chunk




```

16  /* If the chunk is adjacent to the previous chunk, fuse the two into one... */
17  if (p + p->size == headp) {
18      p->size += headp->size;
19      p->next = headp->next;
20  } else
21      p->next = headp; /* ...otherwise just link without fusing. */

```

/\* Exercise: Why is this required? \*/



## `kr_free()` — after linking into the list

```
25     /* We set 'freep' to point just before, or at the freed chunk.  Used by morecore. */  
26     freep = p;  
27 }
```