

Asynchronously handling signals

- ▶ We have used `pause(2)` to **wait for** a signal.
- ▶ Signals can occur **at any time**, and interrupt processing *at any point*. Even, e.g., during a `malloc`.
 - What should happen if a **system call** is interrupted...
 - ...or a function that modifies **global variables**...
 - ...or the signal **handler** itself?

Solutions *signal handling is messy*

- ▶ A function is called **reentrant** or **async-signal-safe**, if it can safely be called from an interrupt handler *while* itself is being interrupted by a signal.
 - *Only* use reentrant functions in signal handlers!
 - A list is in `signal(7)`: No `printf`, no `malloc`, ...
- ▶ You might even need to maintain & restore a copy of `errno`!
- ▶ When interrupted by a signal, **system calls** and **library functions** shall fail with `EINTR`, or restart automatically.
- ▶ A program may **block signals** during critical sections, e.g., signal handlers.

Restart system calls — and other flags

```
1 struct sigaction {  
2     void (*sa_handler)(int);  
3     int sa_flags;  
4  
5     /* still incomplete */  
6 };
```

- ▶ The `sa_flags` member further tunes the behavior of signal handling.
- ▶ It is the bitwise OR of zero or more of the following:

SA_NOCLDWAIT If used with the **SIGCHLD** signal, do not transform children into **zombies** when they terminate. See `wait(2)`, *Notes*!

SA_RESETHAND* Trigger the handler **only once**, then reset handler.

SA_RESTART* Some system calls or library functions may be **restarted** instead of failing with **EINTR**. See `signal(7)` for further details.

- ▶ `wait(2)` is restarted if **SA_RESTART** is set.
- ▶ `pause(2)` is never restarted, otherwise it would never return.

SA_SIGINFO cf. page 288.

... more flags are described in `sigaction(2)`.

Note Marked* flags are not available with `gcc -std=c99`.

Blocking signals

- ▶ By default, a signal is blocked while a signal of **the same type** (*i.e.*, the same signal number) is handled.
- ▶ Signals occurring while they are blocked, become **pending**.
 - Pending signals are delivered **when unblocked**.
 - **Only one** signal of each type may be pending, so it's rather a **flag**.

(There's an exception for the signals above 32)

```
1 $ ./signal0
2 pid is 13789
3 ^Ccaught signal 2      /* Repeat C-c... */
4 ^C^C^C^Ccaught signal 2 /* ...fast */
```

- ▶ The signal handler has a delay of 1s built in.
- ▶ **Only two** signals are caught!

- ▶ What if **other signals** arrive while inside the handler?

```
1 $ kill -USR1 13789 & kill -INT 13789    #try this
```

- ▶ We need means to block more signals inside a handler.

The signal mask

```
1 #include <signal.h>
2
3 int sigemptyset(sigset_t *set);          int sigaddset(sigset_t *set, int sig);
4 int sigfillset(sigset_t *set);          int sigdelset(sigset_t *set, int sig);
5
6 int sigprocmask(int how, const sigset_t *set, sigset_t *old);
```

- ▶ `sigemptyset(3)` **initializes** an empty set of signals, `sigfillset(3)` a set containing all signals.
- ▶ `sigaddset(3)` adds a `signal` to a set, `sigdelset(3)` removes it.
- ▶ `sigprocmask(2)` sets/queries the set of currently blocked signals:
 - If not `NULL`, the `set` pointed to is used.
 - If not `NULL`, the old set is passed back via `old`.
 - `how` specifies how `set` is used:
 - `SIG_BLOCK` Add the `set` to currently blocked signals.
 - `SIG_UNBLOCK` Remove the `set` from currently blocked signals.
 - `SIG_SETMASK` Block exactly the signals in the `set`.

Example

```
1 #define _POSIX_C_SOURCE 199309L
2
3 #include <signal.h>
4 #include <unistd.h>
5
6 int main(void)
7 {
8     sigset_t ss;
9
10    sigfillset(&ss);
11    sigprocmask(SIG_BLOCK, &ss, NULL);
12
13    pause();
14    return 0;
15 }
```

Questions

- ▶ What does this program (not) do?
- ▶ To block additional signals in a signal handler, why can we not use `sigprocmask`?

The signal mask

- ▶ A signal we want to block might be delivered to the handler **just before masking** is completed!

Solution

- ▶ The `struct sigaction` has a member `sa_mask` to specify signals that should be blocked:

```
1 struct sigaction {  
2     void (*sa_handler)(int);  
3     int sa_flags;  
4     sigset_t sa_mask; /* Additional set of signals to block inside handler */  
5  
6     /* still incomplete */  
7 };  
8
```

Example: Implement sleep, 1st try

```
1 void wakeup(int s)
2 {
3     (void)s; /* nothing to do, just mark variable as used */
4 }
5
6 unsigned int sleep1(unsigned int n)
7 {
8     struct sigaction sa;
9
10    memset(&sa, 0, sizeof(sa));
11    sa.sa_handler = wakeup;
12
13    if (sigaction(SIGALRM, &sa, NULL) != 0)
14        return n; /* not slept at all */
15
16    alarm(n); /* start the timer */
17    pause(); /* wait for any signal */
18    return alarm(0); /* turn off timer, return remaining time */
19 }
```

Question What flaws does this code have?

```
1 unsigned int sleep1(unsigned int n)                                /* same code again */
2 {
3     struct sigaction sa;
4
5     memset(&sa, 0, sizeof(sa));
6     sa.sa_handler = wakeup;
7
8     if (sigaction(SIGALRM, &sa, NULL) != 0)
9         return n;          /* not slept at all */
10
11     alarm(n);              /* start the timer */
12     pause();               /* wait for any signal */
13     return alarm(0);       /* turn off timer, return remaining time */
14 }
```

- ▶ What if there already was a handler for **SIGALRM**, ...
- ▶ ...if **SIGALRM** was blocked, or...
- ▶ ...if an alarm goes off between the calls to **sigaction** and **pause**?

pause with modified mask

```
1 #include <signal.h>
2 int sigsuspend(const sigset_t *set);
```

- ▶ `sigsuspend(2)` waits for any signal in the `set`.
 - This is like `pause(2)` with a **temporarily modified** signal mask.

Motivation Why do we need this? Blocking is not enough.

```
1 sigset_t newmask, oldmask;
2
3 sigemptyset(&newmask);                                /* block SIGALRM and save current signal mask */
4 sigaddset(&newmask, SIGALRM);
5 if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0) err(1, "sigprocmask");
6
7 /* critical region of code: sigaction set up here */
8
9 /* reset signal mask, which should unblock SIGALRM */
10 if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0) err(1, "sigprocmask");
11 /* window is open */
12 pause(); /* wait for signal to occur */
13 /* continue processing */
```

A more robust implementation of sleep

```
1 unsigned int sleep2(unsigned int n)
2 {
3     struct sigaction sa, old_sa;
4     memset(&sa, 0, sizeof(sa));
5     sa.sa_handler = wakeup;
6
7     sigset_t newmask, oldmask;          /* save current mask, block SIGALRM */
8     sigemptyset(&newmask);
9     sigaddset(&newmask, SIGALRM);
10    sigprocmask(SIG_BLOCK, &newmask, &oldmask);
11
12    sigaction(SIGALRM, &sa, &old_sa);    /* missing: handle error, cleanup */
13    alarm(n);
14
15    sigset_t suspmask = oldmask;          /* wait for any signal... */
16    sigdelset(&suspmask, SIGALRM);        /* ...SIGALRM is not blocked... */
17    sigsuspend(&suspmask);                 /* ...only while suspending... */
18                                          /* ...then it is blocked again */
19
20    unsigned int remain = alarm(0);
21    sigaction(SIGALRM, &old_sa, NULL);    /* reset previous action */
22    sigprocmask(SIG_SETMASK, &oldmask, NULL); /* reset signal mask */
23    return remain;
24 }
```

Getting more information about the signal

```
1 #define _POSIX_C_SOURCE 199309L
2 struct sigaction {
3     void (*sa_handler)(int);
4     int sa_flags;
5     sigset_t sa_mask; /* Additional set of signals to block inside handler */
6     void (*sa_sigaction)(int, siginfo_t *, void *);
7 };
```

- ▶ An **alternative** signal handler can be installed via the `sa_sigaction` member.
- ▶ Only assign to **either** `sa_sigaction`, **or** `sa_handler`. Never to both!
- ▶ Set `SA_SIGINFO` in `sa_flags`, to indicate the use of `sa_sigaction`.
- ▶ **Information** will be passed to `sa_sigaction` in a struct pointed to by `siginfo_t`.
- ▶ The `void` pointer can pass information about the context where the signal occurred. Mostly unused.

Signal information

```
1 struct siginfo_t {  
2     int      si_signo;    /* Signal number */  
3     int      si_code;     /* Signal code */  
4     pid_t    si_pid;     /* Sending process ID */  
5     uid_t    si_uid;     /* Real user ID of sending process */  
6     int      si_status;   /* Exit value or signal */  
7     int      si_int;     /* value passed by sender, cf. page 291 */  
8     void     *si_ptr;     /* value passed by sender, cf. page 291 */  
9     /* many more, cf. sigaction(2) */  
10 }
```

`si_signo` is the signal number.

`si_code` indicates, together with `si_signo`, how the signal was sent, and how the other fields need to be interpreted.

Example `si_code == SI_USER` — sent by `kill(2)`.

`si_pid` Sending process ID. Only with `SI_USER`.

`si_uid` Real user ID of sending process. Only with `SI_USER`.

► Some of the members may share storage \Rightarrow Obey `si_code`.

```
1 void action(int s, siginfo_t *si, void *v)
2 {
3     (void)v;    /* not used */
4
5     if (si->si_code == SI_USER)
6         printf("got signal %d via kill: pid %d, uid %d\n",
7             s, si->si_pid, si->si_uid);
8     else
9         printf("got signal %d\n", s);
10 }
11
12 int main(void)
13 {
14     printf("pid is %d\n", getpid());
15
16     struct sigaction sa;
17     sigemptyset(&sa.sa_mask);
18     sa.sa_sigaction = action;
19     sa.sa_flags = SA_SIGINFO;
20
21     sigaction(SIGINT, &sa, NULL); /* install handler for SIGINT */
22
23     while (1) pause(); /* sleep until signal arrives */
24     return 0;
25 }
```

Passing information to the handler

```
1 #define _POSIX_C_SOURCE 199309L
2 #include <signal.h>
3
4 int sigqueue(pid_t pid, int s, const union sigval value);
5
6 union sigval {
7     int    sival_int;
8     void *sival_ptr;
9 };
```

- ▶ `sigqueue(2)` sends signal `s` to process `pid`.
- ▶ It sets `si_code` to `SI_QUEUE`, and fills `si_int` and `si_ptr` with `sival_int` and `sival_ptr` respectively.
⇒ You can pass either a pointer, or an integer.

From the shell

`/bin/kill [-s sig] [-q val] pid...` Allows to send a value in the `si_int` field. This is **not** the shell builtin, cf. `kill(1)`

```
1 void action(int s, siginfo_t *si, void *v)
2 {
3     (void)v;
4
5     if (si->si_code == SI_USER)
6         printf("got signal %d via kill: pid %d, uid %d\n",
7             s, si->si_pid, si->si_uid);
8
9     else if (si->si_code == SI_QUEUE)
10         printf("queued signal %d: pid %d, uid %d, val %d\n",
11             s, si->si_pid, si->si_uid, si->si_int);
12
13     else
14         printf("got signal %d\n", s);
15 }
```

► Sending a signal with a value:

```
1 $ ./signal1
2 pid is 10799
3 queued signal 2: pid 10806, uid 1000, val 666
```

```
1 $ /bin/kill -INT -q 666 10799
```

12

Linking

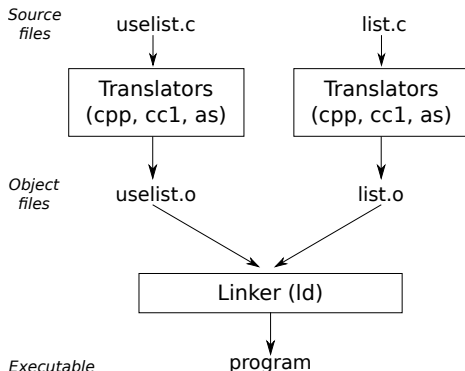
Some of the examples in this chapter are taken from the excellent book

- ▶ Randal E. Bryant, David O'Hallaron. *Computer Systems, A Programmer's Perspective*. 2003, Pearson Education, Prentice Hall. ISBN 0-13-178456-0.

12.1 Introduction

- ▶ We have already seen how separate compilation works (*cf.* page 183).
- ▶ The compiler driver `gcc(1)` employs a bunch of different tools for this task:

- preprocessor `cpp(1)` — removes comments, applies macros.
- compiler `cc1` — compiles into assembler code.
- assembler `as(1)` — translates into binary object file.
- linker `ld(1)` — **links together the compiled object files.**



- ▶ We'll have a closer look at linking now...

Object files

Object files contain chunks of data, (almost) ready to be copied to memory for execution.

- ▶ program code, *i.e.*, CPU instructions compiled from your program, and
- ▶ constant data (*e.g.*, string literals),

There are three kinds of object files:

- ▶ **Executable** object files can be executed directly, *cf.* page 305.
 - Generated by the linker, not by the compiler!
- ▶ **Relocatable** object files can be **linked** with other relocatable object files, to form an executable.
 - Symbols may change their position (*cf.* page 313), hence the name.
- ▶ **Shared** object files are relocatable object files that can be loaded into memory at runtime, and be shared amongst processes (*cf.* page 329).

The functions, global variables, and **static** variables defined in an object file, can be referred to by name: The **symbols**.

Linker Symbols

Relocatable object files come with a **symbol table**, that lists all the symbols an object file exposes.

- ▶ **Global** symbols are defined in the object file, and may be referenced from other object files.
- ▶ **External** symbols are referenced by the object file, but not defined. *I.e.*, the definition must be provided in another object file.
- ▶ **Local** symbols are defined and referenced only from within the object file.

Note *Local symbols* have nothing to do with function-local variables in a C-program. Unless **static**, they are never visible in the symbol table. (Compare debugger symbols.)

Example

```

1 extern int buf[];
2 int *bufp0 = &buf[0];
3 int *bufp1;
4
5 void swap(void)
6 {
7     int temp;
8     static int count = 42;
9
10    bufp1 = &buf[1];
11    temp = *bufp0;
12    *bufp0 = *bufp1;
13    *bufp1 = temp;
14
15    count++;
16 }

```

```

1 $ pk-cc -c swap.c
2 $ readelf -s swap.o                                # cf. readelf(1)
3 Symbol table '.symtab' contains 18 entries:
4   Num: Size Type      Bind      Ndx Name
5 # ...
6       5:    4 OBJECT    LOCAL      3 count.1597
7 # ...
8      14:    8 OBJECT    GLOBAL     3 bufp0
9      15:    0 NOTYPE    GLOBAL    UND buf
10     16:    8 OBJECT    GLOBAL    COM bufp1
11     17:   74 FUNC      GLOBAL     1 swap

```

(some lines and columns have been removed)

- ▶ The local symbol `count` (has its name extended to avoid name clashes) uses 4 bytes, and will be stored in section 3 (Section? cf. page 309)
- ▶ Object `bufp0` uses 8 bytes in section 3, function `swap` uses 74B in section 1.
- ▶ `buf` is **UN**Defined, *i.e.*, referenced by this module, but we have no idea where it will be in the compiled program.
- ▶ **COM**mon objects, like `bufp1`, are uninitialized, and not yet allocated