

Systems 3

Options and Dynamic Memory

Marcel Waldvogel

(Handout)

Department of Computer and Information Science
University of Konstanz

Winter 2019/2020

These slides are based on previous lectures held by Alexander Holupirek, Roman Byshko, and especially Stefan Klinger.

Chapter Goals

- How do processes get their information?
- What is the difference between arguments and environment variables?
- Why is one mechanism alone not sufficient?
- Why are processes not happy with static memory allocations?
- How do processes deal with dynamic memory requests?
- How is dynamic memory managed?
- Why are C's pointers and lack of checks necessary to implement dynamic memory management? How is this achieved?
- How to handle strings well?
- How to avoid (security) bugs when doing so?

Arguments and Environment

Command-line arguments

- The function called at program startup is named `main`.
- It shall be defined with a **return** type of `int`, and either zero, or two parameters:

```
1 int main(void);  
2 int main(int argc, char *argv[]);
```

Terminology (although other names may be used)


- `argc` stands for *argument count*
- `argv` stands for *argument vector*

If `argc` and `argv` are declared:

- The value of `argc` shall be **nonnegative**
- `argv[0]` represents the **program name** or `argv[0][0]` shall be the null character if the program name is not available.
- `argv[1]` to `argv[argc-1]` represent *program parameters*.
- `argv[argc]` shall be `NULL`, i.e., it may be accessed.

- When a program is executed, the process that starts the new program can pass command-line arguments to it.
- That is the normal operation for UNIX system shells.

```
1 $ echo hello, world
2 hello, world
```



Echo command-line arguments

```
1 #include <stdio.h>
2
3 int main(int argc, char *argv[])
4 {
5     for (int i = 0; i < argc; i++)
6         printf("argv[%d]: \"%s\"\n", i, argv[i]);
7
8     return 0;
9 }
```


```
1 $ ./a.out dsf ' dfsdf\' ' ' t
2 argv[0]: "./a.out"
3 argv[1]: "dsf"
4 argv[2]: " dfsdf\ "
5 argv[3]: "t"
```

Environment variables

- Each program is also passed an **environment list**.
- Like the argument list, it is an array of character pointers.
- Each pointing to a null-terminated C string, of the form
name=value
- The address of the array is contained in a global variable, the **environment pointer**:

```
1 extern char **environ;
```

environ:



History There once was an optional third argument to `main`, containing the environment.

Print the environment

```
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 extern char **environ;
5
6 int main(void)
7 {
8     for (char **env = environ; *env; ++env)
9         printf("%s\n", *env);
10
11     char *p = getenv("PATH"); /* see getenv\(3\) */
12     if (p)
13         printf("Current path is: %s\n", p);
14
15     return 0;
16 }
```


Dynamic memory management

Current situation Until now, we cannot change the amount of space available to store data:

- The **number of variables** in a C program is fixed in the source code.
- **Arrays** cannot grow, nor shrink.

⇒ Use **excessively large** arrays that are guaranteed to be big enough.
That's not nice!

Dynamic memory Get more memory **on demand**, and only if required.

- First figure out how much memory is needed, then request that from the OS (*aka. allocating*).
- Or guess how much is needed and allocate that. Adapt as necessary.
- **Return** unused memory to the OS.

Allocating memory

`malloc(3)` and `calloc(3)` **allocate** blocks of memory.

```
1 #include <stdlib.h>
2 void *malloc(size_t size);
3 void *calloc(size_t num, size_t size);
```

- `size_t`, defined in `stddef.h` is an unsigned integral type.
- `malloc` allocates a block of `size` bytes of memory.
 - The memory is **not initialised**.
 - Initialisation can be done using `memset(3)`.
- `calloc` allocates memory for an array of `num` elements of `size` bytes each.
 - The storage is **initialised to zero** (good practice).
- Both functions return a pointer to the (start of) the allocated memory, or `NULL` if the request cannot be satisfied (or the requested size is 0).
- `void *` is the proper type for a **generic pointer**. No casting needed.

```
1 int *ip;
2 ip = calloc(42, sizeof(int)); /* space for 42 ints */
```

Extend or reduce allocated memory

`realloc(3)` “modifies” the size of a block of memory previously allocated with `malloc(3)`.





```
1 #include <stdlib.h>
2 void *realloc(void *ptr, size_t size);
```

- Changes the size of the object pointed to by `ptr` to `size` bytes.
 - Note that it may be necessary to **move** (i.e. copy) all data to a new location!
- `realloc` returns a **new pointer** to the (possibly moved) object.
 - **Do not use the old pointer**, it is invalid!
- The contents will be **unchanged** in the range from the start of the region up to the minimum of the old and new sizes.
 - Freshly allocated memory is **not initialised**.
- **Note:** `ptr` must point to memory previously allocated with `malloc`, i.e., this will not work:

```
1 int arr[23];
2 int *p = arr;
3 p = realloc(p, 42 * sizeof(int)); /* wrong */
```

realloc() scenarios

```
1 void *inp, *outp;
2 outp = realloc(inp, newSize);
```



Allocated chunks are white, free chunks are grey. Addresses increase to the right.

Freeing allocated memory

`free(3)` frees memory previously allocated with `malloc(3)`.

```
1 void free(void *ptr);
```

- If `ptr` is a `NULL` pointer, no action occurs.
- It is an error to dereference something **after it has been freed**.
- It is important to free memory you do not need anymore.
 - In general, this is not an easy task.
 - There is **no garbage collector**.
 - If you do not `free`, you may **run out of memory**.
- **Note:** `ptr` must point to memory previously allocated with `malloc` etc, *i.e.*, this will not work:

```
1 int arr[23];  
2 free(arr); /* wrong */
```

Example

```
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 int main(void)
5 {
6     int *p = malloc(8 * sizeof(int)); /* allocate mem for 8 int */
7
8     for (int i = 0; i < 8; i++)          /* write some data */
9         p[i] = i*i;
10
11     p = realloc(p, 16 * sizeof(int)); /* get more space */
12     p[15] = 100;
13
14     p = realloc(p, 12 * sizeof(int)); /* free some memory */
15     /* p[15] = 7; */ /* invalid */
16
17     for (int i = 0; i < 12; i++)          /* print whole memory block */
18         printf("%2d\t%d\n", i, p[i]);    /* slots 8-11 contain garbage */
19
20     free(p); /* free all memory used by p */
21
22     return 0;
23 }
```

Caution



- **Always free allocated memory** when it's no longer used.
- It is a bug not to **check the return values** of `malloc(3)`, `calloc(3)`, or `realloc(3)` for error conditions.
Review the example on slide 15!
- One **must not access unallocated memory**, or memory after calling `free` on it.

Ignoring any of these rules **is a bug** that may, or may not, show up during testing. Even if the program behaves as expected, it is still buggy!

Handling strings

- With `#include <string.h>` you'll get access to a plethora of string handling functions, documented in `string(3)`.
- Example: Copy string pointed to by src, to buffer pointed to by dest.

```
1 char *strcpy(char *dest, const char *src);           /* cf. strcpy(3) */
```

Question How can we make a copy of a string?

```
1 const char *msg = "hello world\n";  
2  
3 char *copy;  
4 strcpy(copy, msg);
```

- What do you think about this approach?

```
1 const char *msg = "hello world\n";  
2  
3 char *copy; /* not initialized, points nowhere */  
4 strcpy(copy, msg);
```

Bad idea: The target pointer does not point to any allocated memory!

⇒ **Undefined behavior**¹

¹<https://blog.regehr.org/archives/213>, <https://blog.regehr.org/archives/970>

Question String copy: What about this one?

```
1 char *strcpy(char *dest, const char *src);  
2 const char *msg = "hello world\n";  
3  
4 char *copy = malloc(strlen(msg));  
5 strcpy(copy, msg);
```

```
1 char *strcpy(char *dest, const char *src);
2 const char *msg = "hello world\n";
3
4 char *copy = malloc(strlen(msg)); /* not enough */
5 /* return value unchecked */
6 strcpy(copy, msg);
```

- Unchecked if we got any memory at all.
- Even then, not enough memory is allocated: `strlen` returns length *excluding* NUL, but `strcpy` copies that as well!

⇒ **Undefined behavior**

Easy to fix:

```
1 #include <err.h>
2
3 char *copy = malloc(strlen(msg) + 1);
4 if (!copy)
5     err(1, "copy"); /* cf. err(3). Terminates with a message like */
6 /* a.out: copy: Cannot allocate memory */
```

Question String copy: Not correct. Why?

```
1  const char *msg = "Old MacDonald Had a Farm";
2
3  size_t len = strlen(msg) + 1;
4  char *cp1 = malloc(len),
5  *cp2 = malloc(len);
6
7  if (!cp1 || !cp2)
8      err(1, "cp1 or cp2");
9
10 for (size_t i = 0; i < 13; i++) /* copy only first two words */
11     cp1[i] = msg[i];
12
13 strcpy(cp2, cp1); /* copy that to cp2 */
```

```
13 strcpy(cp2, cp1); /* copy cp1 to cp2 */
```

- `strcpy` will copy bytes from `cp1` until the string ends, *i.e.*, until it sees a `'\0'` character.

- The source `cp1` may not be terminated by a `NUL` character!

⇒ `strcpy` may "fall over the edge", and overwrite adjacent memory!

⇒ **Undefined behavior**

Solution to all these cases:

- Use `strncpy(3)` instead, which will not write more than `n` bytes!

```
1 char *strncpy(char *dest, const char *src, size_t n);
```

Always be aware of the amount of data to be written!

- *Overflowing fixed-length string buffers leads to security vulnerabilities*
`strcpy(3)`



Note that `strncpy` may **not write** the terminating `NUL`!