

# Systems 3

## OS Structure

Marcel Waldvogel

(Handout)

Department of Computer and Information Science  
University of Konstanz

Winter 2019/2020



Photo by Kelly Sikkema on Unsplash

# Chapter Goals

- How does the OS (kernel) differ from user space?
- How to switch between privileged/unprivileged mode?
- How to interface with the kernel?
- What to put into the kernel?

# OS as Black Box

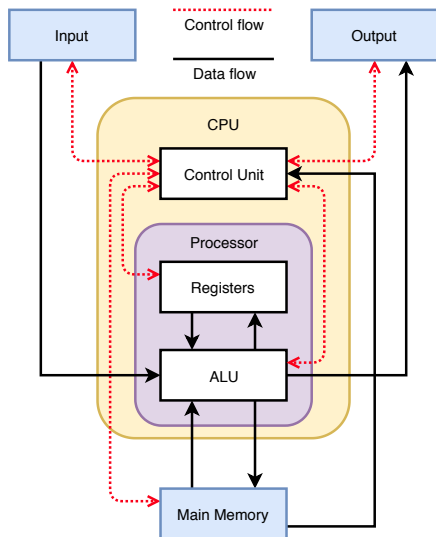
- System calls as API
- Scheduling
- Memory Management
- Resource Management

# Central Processing Unit (simplified)

What is a CPU and how does it work?

- Arithmetic Logic Unit (ALU)
- Processor registers
- Control Unit

*Figure: Lambtron (CC BY-SA 4.0)  
modified by Klaus Herberth*



# Instruction Cycle (simplified)

The classic RISC pipeline consists of the following operations:

- Fetch
- Decode
- Execute

**Question:** How does the CPU find the next instruction?

**Solution:** Program counter

## Heads up

In a few weeks, we will look into abusing the instruction cycle.

# Instructions

**r3=r2+r1**    ADD r3, r2, r1

**add(a, b)**    PUSH a  
              PUSH b  
              CALL add  
              ADD sp, sizeof(a)+sizeof(b)

**while (c) x();**    loop: CALL x  
                  TEST c  
                  BNE label

# Permissions

## Problem

If every program can execute every instruction, how can we prevent processes interfering with each other?

## Solution

Have some privileged instructions, which only the kernel can execute.

## Which ones?

- 1 Access to/manipulating hardware
- 2 Manipulating CPU state (outside of 'user-visible' registers), e.g., modify memory mapping, privilege level
- 3 Access memory not assigned to current process (kernel data structures, memory of other processes)



# Privileges

## Definitions

**User level** Access rights of a user process (no access to privileged instructions)

**Kernel level** Access rights of the kernel (access to privileged instructions)

**User space** Memory visible to user process

**Kernel space** Memory visible to kernel

Space/level distinction not always strict.

# Switching between User and Kernel

## Switching steps

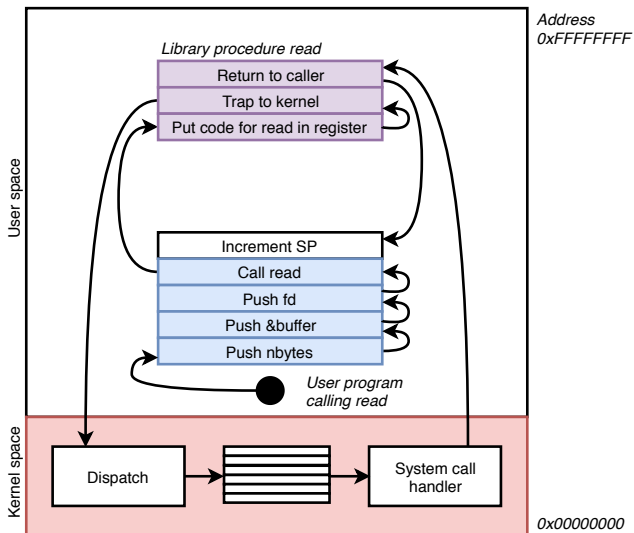
- 1 Address space
- 2 Privilege mode (bit in processor status register)
- 3 Program counter (to a kernel-defined address only)

## Requesting a switch

- Interrupt
- Trap (illegal instruction, illegal access, syscall, ...)

Return under kernel control

# Making a System Call



# What should be part of the Kernel?


## Criteria for inclusion

- 1 Direct hardware manipulation (I/O, CPU configuration)
- 2 Direct triggering by interrupt
- 3 Communication between processes

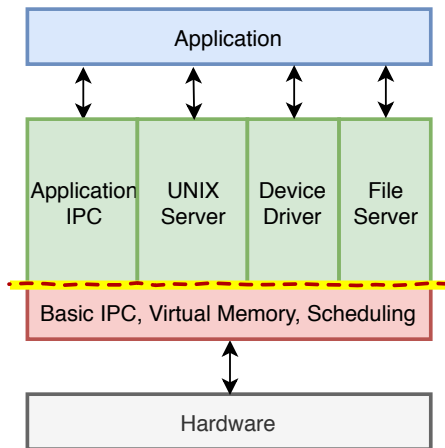
## Criteria for exclusion

- 1 Complex operation (error-prone)
- 2 Debuggable
- 3 Updateable
- 4 Restartable


# Monolithic Kernel



# Microkernel

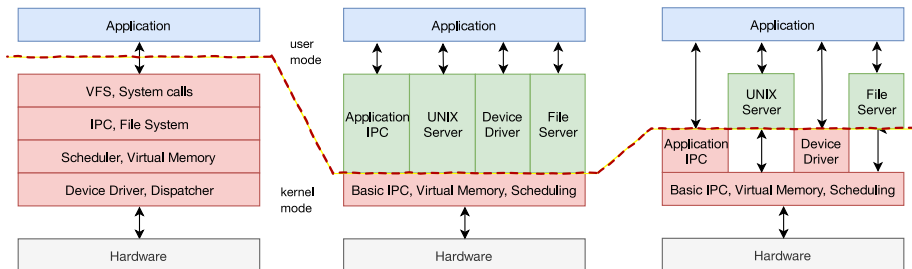


# Hybrid kernels



# Comparison

Criteria	Microkernel	Monolithic
Size	small	large
Speed	slower	fast
Extensible	yes	no
Reliability	high	lower <sup>1</sup>
Interface	messages	functions



Adapted from <https://www.geeksforgeeks.org/monolithic-kernel-and-key-differences-from-microkernel/> and <https://techdifferences.com/difference-between-microkernel-and-monolithic-kernel.html>.

<sup>1</sup>Crash of a service leads to system crash