

# Systems 3

## Libraries

Marcel Waldvogel

(Handout)

Department of Computer and Information Science  
University of Konstanz

Winter 2019/2020

These slides are based on previous lectures held by Alexander Holupirek, Roman Byshko, and especially Stefan Klinger.

# Chapter Goals

- How to manage big programs?
- How to split/structure them into modules?
- How modules can be separated/made to interact?
- How to compile big programs (efficiently)?
- What happens behind the scenes?
- How are programs loaded and run?
- The use of header files.
- The use (and dangers) of macros.
- Portable code and conditional compilation.

# Linking

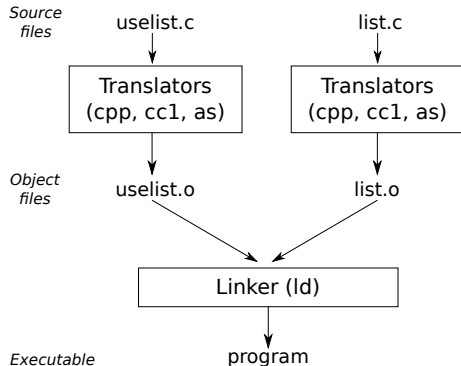
Some of the examples in this chapter are taken from the excellent book

- Randal E. Bryant, David O'Hallaron. *Computer Systems, A Programmer's Perspective*. 2003, Pearson Education, Prentice Hall. ISBN 0-13-178456-0.

# Introduction

- We have already seen how separate compilation works (Lecture on 'Big Programs').
- The compiler driver `gcc(1)` employs a bunch of different tools for this task:

- preprocessor `cpp(1)` — removes comments, applies macros.
- compiler `cc1` — compiles into assembler code.
- assembler `as(1)` — translates into binary object file.
- linker `ld(1)` — **links together the compiled object files.**



- We will have a closer look at linking now...

# Object files

Object files contain chunks of data, (almost) ready to be copied to memory for execution.

- program code, *i.e.*, CPU instructions compiled from your program, and
- constant data (*e.g.*, string literals),

There are three kinds of object files:

- **Executable** object files can be executed directly, *cf.* page 15.
  - Generated by the linker, not by the compiler!
- **Relocatable** object files can be **linked** with other relocatable object files, to form an executable.
  - Symbols may change their position (*cf.* page 23), hence the name.
- **Shared** object files are relocatable object files that can be loaded into memory at runtime, and be shared amongst processes (*cf.* page 39).

The functions, global variables, and `static` variables defined in an object file, can be referred to by name: The **symbols**.

# Linker Symbols

Relocatable object files come with a **symbol table**, that lists all the symbols an object file exposes.

- **Global** symbols are defined in the object file, and may be referenced from other object files (**no modifier**).
- **External** symbols are referenced by the object file, but not defined. *I.e.*, the definition must be provided in another object file (**extern**).
- **Local** symbols are defined and referenced only from within the object file (**static or compiler-generated**).

**Note** *Local symbols* have nothing to do with function-local variables in a C-program. Unless **static**, they are never visible in the symbol table. (Compare debugger symbols.)

# Example

```

1 extern int buf[];
2 int *bufp0 = &buf[0];
3 int *bufp1;
4
5 void swap(void)
6 {
7     int temp;
8     static int count = 42;
9
10    bufp1 = &buf[1];
11    temp = *bufp0;
12    *bufp0 = *bufp1;
13    *bufp1 = temp;
14
15    count++;
16 }

```

```

1 $ gcc -c swap.c
2 $ readelf -s swap.o                                # cf. readelf(1)
3 Symbol table '.symtab' contains 18 entries:
4   Num: Size Type      Bind   Ndx Name
5   #...
6       5:    4 OBJECT    LOCAL    3 count.1597
7   #...
8      14:    8 OBJECT    GLOBAL    3 bufp0
9      15:    0 NOTYPE    GLOBAL   UND buf
10     16:    8 OBJECT    GLOBAL   COM bufp1
11     17:   74 FUNC      GLOBAL    1 swap

```

(some lines and columns have been removed)

- The local symbol `count` (has its name extended to avoid name clashes) uses 4 bytes, and will be stored in section 3
- (Sections are explained on slide 19)
- Object `bufp0` uses 8 bytes in section 3, function `swap` uses 74B in section 1.
- `buf` is **UN**Defined, i.e., referenced by this module, but we have no idea where it will be in the compiled program.
- **COM**mon objects, like `bufp1` are uninitialized, and not yet allocated.

# Symbol resolution

- For each **local symbol**, the compiler guarantees exactly one definition. The name is modified to be unique (e.g. `count` above).
- If *the compiler finds no definition*, it expects it to come from another module, and leaves it to the linker, (e.g. `buf` above).
- When **the linker** resolves *global* symbols, several conditions can occur:
  - **No definition** is found in the symbol table of any input object file.
  - **Multiple definitions** are found in different object files, choose one.

**Example** No `main` function, and `buf` undefined.

```
1 $ gcc swap.o #without -c, try to build an executable
2 .../lib/crt1.o: In function '_start':
3 (.text+0x20): undefined reference to 'main'
4 swap.o: In function 'swap':
5 .../swap.c:12: undefined reference to 'buf'
6 swap.o(.data+0x0): undefined reference to 'buf'
7 collect2: error: ld returned 1 exit status
```

- The linker tries to link with `crt1.o`, which refers to the `main` function.



# Choosing one among multiple definitions

- The linker distinguishes weak and strong symbols:
  - Functions and initialized global variables are **strong symbols**.
  - Uninitialized global variables are **weak symbols**.
- If a conflict arises, the strategy is as follows:
  - Multiple strong symbols → raise error.
  - One strong, and multiple weak symbols → choose the strong one.
  - Multiple weak symbols → choose an arbitrary one.

## Example Two strong symbols

```
1 $ cc -c foo1.c bar1.c           #compilation is fine
2 $ cc foo1.o bar1.o
3 bar1.o: In function 'main':
4 .../bar1.c:2: multiple definition of 'main'
5 foo1.o:.../foo1.c:2: first defined here
6 collect2: error: ld returned 1 exit status
```

### foo1.c

```
1 int main(void)
2 {
3     return 0;
4 }
```

### bar1.c

```
1 int main(void)
2 {
3     return 0;
4 }
```

## Question What will happen here?

foo2.c

```
1 #include <stdio.h>
2
3 void f(void);
4 int x = 12345;
5
6 int main(void)
7 {
8     f();
9     printf("x = %d\n", x);
10
11     return 0;
12 }
```

bar2.c

```
1 int x = 54321;
2
3 void f(void)
4 {
5     x++;
6 }
```

## Example One strong, and one weak symbol.

### foo3.c

```
1 #include <stdio.h>
2
3 void f(void);
4
5 int x = 12345; /* strong */
6
7 int main(void)
8 {
9     f();
10    printf("x = %d\n", x);
11
12    return 0;
13 }
```

### bar3.c

```
1 int x; /* weak */
2
3 void f(void)
4 {
5     x = 54321;
6 }
```

The result is probably expected:

```
1 $ gcc -c foo3.c bar3.c
2 $ gcc foo3.o bar3.o
3 $ ./a.out
4 x = 54321
```

■ Maybe check out the symbol tables? `readelf -s {foo,bar}3.o`

## Example Two weak symbols.

### foo4.c

```
1 #include <stdio.h>
2
3 void f(void);
4
5 int x; /* weak */
6
7 int main(void)
8 {
9     x = 12345;
10    f();
11    printf("x = %d\n", x);
12
13    return 0;
14 }
```

### bar4.c

```
1 int x; /* weak */
2
3 void f(void)
4 {
5     x = 54321;
6 }
```

Again, no surprise:

```
1 $ gcc -c foo4.c bar4.c
2 $ gcc foo4.o bar4.o
3 $ ./a.out
4 x = 54321
```

**Note** The linker has unified both variables `x`, and makes references to both symbols address the same space in memory.

## Question What about this one?

### foo5.c

```
1 #include <stdio.h>
2
3 void f(void);
4
5 int x = 12345; /* strong */
6 int y = 54321; /* strong */
7
8 int main(void)
9 {
10     f();
11     printf("x = %d, y = %d\n", x, y);
12
13     return 0;
14 }
```

### bar5.c

```
1 double x; /* weak */
2
3 void f(void)
4 {
5     x = 1;
6 }
```

## Explain this:

```
1 $ gcc -c foo5.c bar5.c
2 $ gcc foo5.o bar5.o
3 /usr/bin/ld: Warning: alignment 4
4 of symbol 'x' in foo5.o is smaller
5 than 8 in bar5.o
6 $ ./a.out
7 x = 0, y = 1072693248
```

## Notes

- Not long ago, the linker would give **no warning** at all.
- It is extremely difficult to **debug** such code.

What else?

- After resolving symbols, the linker knows which definition belongs to each symbol.
- The linker does not know about the type, only about the size.

## Recall

- Machine code does not use variable names any more.
  - The compiler produced code that accesses variables and functions only by their **memory addresses**.
- ⇒ How does this go together with separate compilation and symbol resolution?

# The program in memory

How does a program start?

- When a program is run, it is **copied into memory** by the **loader**.
  - Copy **text segment**, *i.e.*, the actual machine code,
  - copy **initialized data**,
  - **initialize** uninitialized data,
  - *etc.*
- We want to minimize the amount of data to be copied!
  - Only load parts that are **actually required**,
  - and only load them **when** they are needed.
- We want to save memory!
  - Do not load the same code into memory multiple times.
  - **Share** already loaded code between processes.
- Avoid expensive **transformations**
  - Store program on disk in a **format** that allows fast setup of the process image.

# Virtual memory

- VM is a mapping from the process' **virtual address space** into the machine's physical address space (organized in **pages**).
- The VM system may flag pages as, e.g., **read only**, **executable**, or **private**, cf. `mmap(2)`.
- A physical page may reside on disk, until **loaded on demand**.
  - So we compile the memory layout into the executable file,
  - the loader just **maps the file** into the process' virtual address space, and
  - the VM system gets the pages into memory when actually referenced.
- Multiple running instances of a program share their text (machine code) through a *read only* mapping to **the same physical** address space.

**Note** To achieve all this, the structure of the program file depends on the process' memory layout!



# Example

- The mapping of virtual memory can be observed in `/proc/$PID/maps`.
- I have several `xterms` running, one with PID 23172.

```

1 $ cat /proc/23172/maps                                     # this is on a 64bit machine
2 # address          perms offset  dev   inode  pathname
3 00400000-00472000    r-xp 00000000 00:10 987773  /usr/bin/xterm
4 00672000-00673000    r--p 00072000 00:10 987773  /usr/bin/xterm
5 00673000-0067c000    rw-p 00073000 00:10 987773  /usr/bin/xterm
6 0067c000-0067e000    rw-p 00000000 00:00 0
7 006f7000-00718000    rw-p 00000000 00:00 0      [heap]
8 #...
9 7f05d802c000-7f05d81cc000 r-xp 00000000 00:10 953240  /usr/lib/libc-2.18.so
10 #...
11 ffffffff600000-ffffffff601000 r-xp 00000000 00:00 0 [vsyscall]


```

- The executable pages are readonly,
- the writable area is private (*i.e.*, copy-on-write).
- Other instances of `xterm` have the same mapping for the binary `xterm`,
- but may have others for the shared libraries.

# Process memory layout

When running, a process has the following **virtual memory layout**.

(This is for 32bit Linux)



- Kernel memory (1GiB) is not accessible by the process.
  - **Shared** among all processes.
- The stack maintains local variables and function calls.
- Shared libraries may even be added at runtime.
- The heap contains allocated memory.
- **.data** and **.bss** store global variables.
- **.text** and **.rodata** are marked **ro**, **so** can be shared with other processes.

# Object file layout

- There are various formats to store binary programs.
- Linux uses ELF, the **Executable and Linking Format**.
- COFF and `a.out` are others, the latter coined the name used by `gcc` for default binaries (in ELF on Linux!).
- All formats have the concept of **sections** in common.
- A section is the unit of organization in a binary.

**Some section names** (but there are many more)

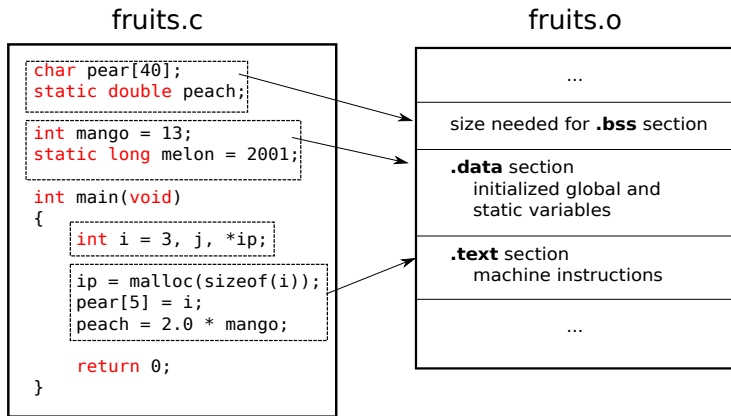
`.text` The program code, *i.e.*, processor instructions.

`.rodata` Read-only data, *e.g.*, string literals.

`.data` *Initialized global* variables.

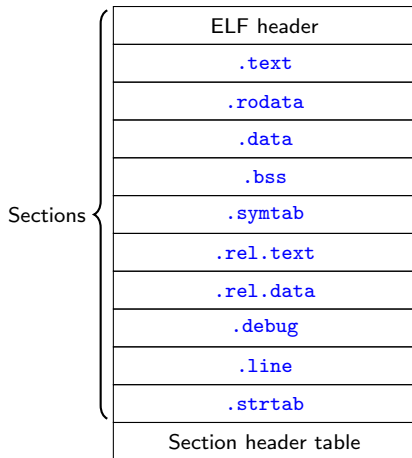
`.bss` *Uninitialized global* variables.

`.symtab` The symbol table, displayed with `readelf -s`.



# A typical relocatable object file

This is what the compiler produces out of the **individual C files**.

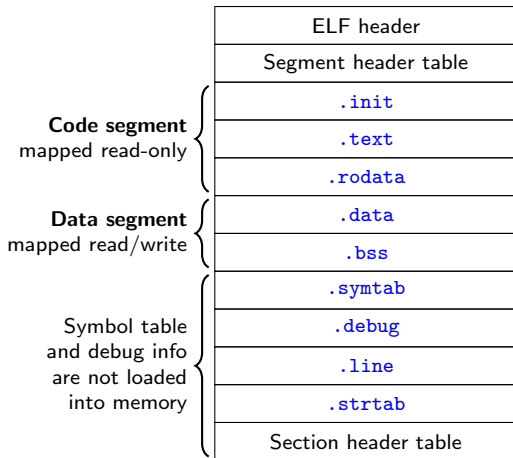


- The **ELF header** describes word size, endian, object file type, machine type, offset and format of the section header table, and other information.
- The **section header table** describes the locations of the various sections.
- Try

```
1 $ gcc -c -m32 swap.c
2 $ readelf -S swap.o
```

# A typical executable object file

That is what we want to have in the **final binary program**.




- The **segment header table** describes the mapping of contiguous file sections into memory.

- Try this

```
1 $ gcc -m32 -static swap.c \
2 > main.c
3 $ readelf -l a.out
```

# Relocation

- So after resolving the symbols, the linker needs to put all the code from the individual object files' sections into the final program's sections:



- This process is called **Relocation**.

Relocation involves **two tasks**:

**1** Relocating **sections** and symbol **definitions**.

- Merge sections of the same type.
  - Assign **run-time memory addresses** to the new aggregate sections, the input sections, and each symbol defined in the input.
- ⇒ After this step, every global symbol has a known run-time memory address.

**2** Relocating symbol **references** everywhere in the code.

- Modification of each reference in `.text` and `.data`, so that they point to correct location.

## Relocation entries

- The assembler does not know where data and code will be stored ultimately,
  - nor does it know addresses of the external objects.
- ⇒ In such situations a *relocation entry* is generated by the assembler.



# Relocation entries

```
1 typedef struct {  
2     int offset;      /* Offset of the reference to relocate. */  
3     int symbol: 24,  /* Symbol the reference should point to. */  
4     type:8;         /* Relocation type. */  
5 } Elf32_Rel;
```

- ELF defines 11 different relocation types
- We will look at [R\\_386\\_32](#) only.  
This is used to relocate 32bit absolute addresses.

# Example: Relocation at work

- Function `f` simply assigns `0xbeef` to the global variable `x`.
- The final memory location of `x` is not yet known.
- Relocation will fix the runtime address of `x`.
- First we **compile** `foo.c` into a *relocatable object file*:

`foo6.c`

```
1 int x = 0xdead;
2
3 void f(void)
4 {
5     x = 0xbeef;
6 }
```

```
1 $ gcc -m32 -c foo6.c
```

- Have a look at the `.data` section:

```
1 $ objdump -d -j.data foo6.o
2 Disassembly of section .data:
3 00000000 <x>:
4      0:    ad de 00 00                                ....
```

- Variable `x` appears at address `0x0` in the `.data` section.
- The value `0xdead` is stored there.

## ■ Have a look at the `.text` section:

```

1 $ objdump -d -j.text foo6.o
2 Disassembly of section .text:
3 00000000 <f>:
4   0:   55                push    %ebp
5   1:   89 e5             mov     %esp,%ebp
6   3:   c7 05 00 00 00 ef movl    $0xbeef,0x0
7   a:   be 00 00
8   d:   5d                pop     %ebp
9   e:   c3                ret

```

- In line 6, the value `0xbeef` is copied to address `0x0`.
- This address `0x0` appears at **offset 5** in the `.text` section.

## ■ These are the relocation entries:

```

1 $ objdump -r -j.text foo6.o
2 RELOCATION RECORDS FOR [.text]:
3 OFFSET      TYPE          VALUE
4 00000005 R_386_32                x

```

- So on relocation of symbol `x`, the absolute 32bit address at **offset 5** in the `.text` section must be updated.

- Then we **link** `foo.o` with something that uses `f`.

```
1 $ gcc -m32 foo6.o bar6.c /* main in bar6.c simply calls f in foo6.c */
```

- Have a look at the `.data` section **after relocation**:

```
1 $ objdump -d -j.data a.out
2 Disassembly of section .data:
3 08049698 <x>:
4 8049698:      ad de 00 00      ....
```

- Variable `x` has been moved to address `0x8049698` in the `.data` section.
- The value `0xdead` is stored there.

- Have a look at the `.text` section **after relocation**:

```
1 $ objdump -d -j.text a.out | grep -C3 beef
2 080483cd <f>:
3 80483cd:      55                push    %ebp
4 80483ce:      89 e5             mov     %esp,%ebp
5 80483d0:      c7 05 98 96 04 08 ef  movl    $0xbeef,0x8049698
6 80483d7:      be 00 00
7 80483da:      5d                pop     %ebp
8 80483db:      c3                ret
```

- The reference to variable `x` has been **updated to address** `0x8049698`.

## Example: Relocation in the `.data` section

- Sometimes, updating references in the `.text` section is not enough.
- Here we have a global variable `xp` initialized with an address!
- The compiler **cannot even fix a value** for `xp`!

`foo7.c`

```
1 int x = 0xdead;
2 int *xp = &x;
3
4 void f(void)
5 {
6     *xp = 0xbeef;
7 }
```

- Again, we compile and link our object files:

```
1 $ gcc -m32 -c foo7.c
2 $ gcc -m32 foo7.o bar6.c
```

## ■ Relocation in the `.data` section.

```

1 $ objdump -d -j.data foo7.o
2 Disassembly of section .data:
3 00000000 <x>:
4     0:    ad de 00 00                                ....
5 00000004 <xp>:
6     4:    00 00 00 00    # This value has to be updated on relocation of x!
7 $ objdump -r -j.data foo7.o    # Note: in .data this time!
8 RELOCATION RECORDS FOR [.data]:
9 OFFSET      TYPE              VALUE
10 00000004 R_386_32                    x
11 $ objdump -d -j.data a.out
12 Disassembly of section .data:
13 08049698 <x>:
14    8049698:    ad de 00 00                                ....
15 0804969c <xp>:
16    804969c:    98 96 04 08                                ....

```

## ■ Relocation in the `.text` section:

```

1  $ objdump -d -j.text foo7.o
2  Disassembly of section .text:
3  00000000 <f>:
4      0:  55                push    %ebp
5      1:  89 e5             mov     %esp,%ebp
6      3:  a1 00 00 00 00    mov     0x0,%eax
7      8:  c7 00 ef be 00 00    movl    $0xbeef, (%eax)
8      e:  5d                pop     %ebp
9      f:  c3                ret
10 $ objdump -r -j.text foo7.o
11 RELOCATION RECORDS FOR [.text]:
12 OFFSET      TYPE          VALUE
13 00000004 R_386_32          xp
14 $ objdump -d -j.text a.out | grep -C3 beef
15 80483cd:  55                push    %ebp
16 80483ce:  89 e5             mov     %esp,%ebp
17 80483d0:  a1 9c 96 04 08    mov     0x804969c,%eax
18 80483d5:  c7 00 ef be 00 00    movl    $0xbeef, (%eax)
19 80483db:  5d                pop     %ebp
20 80483dc:  c3                ret

```

■ Recall: `0x804969c` is the address of the variable `xp`!

# Static libraries

- A static library is a **collection of relocatable object files**.
  - Since the term “object file” is not correct in this context, the members of a library are referred to as **object modules** instead.
- Linking with a library means to link with all the **required** object files.

## Why use libraries at all?

- Why not put all library functions into one relocatable object file?
  - An object module is added to a program in its **entirety, or not at all**.
  - The potentially **large object file** would be added to every binary using one of the functions.  $\Rightarrow$  Waste of space.
- Why not copy only **required functions** from an object file?
  - Sections like `.text` and `.data` are merely binary blocks to the linker.
- Why not **explicitly link** all the required object files with the binary?
  - Tedious with object modules at the granularity of individual functions!

```
1 $ gcc -o main main.c printf.o atoi.o read.o write.o ...
```



# Making a static library

- A static library is an **archive** of relocatable object modules

`ar rcs archive [member...]` Create archive, containing the members.  
`nm archive` List symbols from object files or archives.

- The `ar(1)` command provides various means to modify an archive.
  - `r` Add the members to archive, **replace** existing with the same name.
  - `c` **Create** the archive if it does not exist.
  - `s` Write an object-file **index** into the archive.
  - `...` many others
- `nm(1)` displays the symbols defined in a module, or a library.
- By the way: See `info binutils` for an overview of tools for manipulating ELF binaries.

# Example

## addvec.c

```

1 #include "addvec.h"
2
3 void addvec(int n, int *x, int *y,
4             int *z)
5 {
6     for (int i = 0; i < n; i++)
7         z[i] = x[i] + y[i];
8 }

```

## dotproduct.c

```

1 #include "dotproduct.h"
2
3 int dotproduct(int n, int *x, int *y)
4 {
5     int r = 0;
6     for (int i = 0; i < n; i++)
7         r += x[i] * y[i];
8     return r;
9 }

```

- The header files just contain the respective function prototype.

```

1 $ gcc -c addvec.c                                     # usually, all this is arranged for in a Makefile
2 $ gcc -c dotproduct.c
3 $ ar rcs libvector.a addvec.o dotproduct.o
4 $ nm libvector.a   # a shiny new library
5 addvec.o:
6 0000000000000000 T addvec
7 dotproduct.o:
8 0000000000000000 T dotproduct

```

## main.c

```
1 #include "dotproduct.h"
2 #include <stdio.h>
3
4 int main(void)
5 {
6
7     int x[3] = { 1, 0, 0 }, y[3] = { 0, 1, 1 }, z[3] = { 1, 1, 0 };
8
9     printf("<x,y> = %d\n", dotproduct(3, x, y));
10    printf("<x,z> = %d\n", dotproduct(3, x, z));
11
12    return 0;
13 }
```

- The `-static` flag tells `gcc` to build a **statically linked** binary.

```
1 $ gcc -c main.c
2 $ gcc -static -omain main.o libvector.a
3 $ ./main
4 <x,y> = 0
5 <x,z> = 1
```

# Linker flags for static libraries

- Typically, libraries do not reside in the directory where they are used.
  - With `-lname` the library `libname.a` is searched for in the library search path.
  - With `-Ldir`, a directory is added to the **library search path**.
  - The library search path is searched in the order of the `-L` options.

```
1 $ gcc -static -omain main.o -L. -lvector      # link with the libvector.a library
```

- The **order** in which libraries are given on the **command line** is significant, and counter-intuitive:
  - The library providing a symbol must appear **after** the object using it.

```
1 $ gcc -static -omain libvector.a main.o
2 main.o: In function 'main':
3 /home/marcel/Repos//bspk/pk/lect/src/lib/main.c:10: undefined reference
4 to 'dotproduct'
5 /home/marcel/Repos//bspk/pk/lect/src/lib/main.c:12: undefined reference
6 to 'dotproduct'
7 collect2: error: ld returned 1 exit status
8 $ gcc -static -omain -L. -lvector main.o
9 # the same error message
```

**Input** : Files passed to the linker on the command line

**Output:** A statically linked binary

**Data** : The set of object modules  $O$  **to be linked** to the binary, the set of referenced but yet **unresolved symbols**  $U$ , and the set of already **defined symbols**  $D$

$O \leftarrow \emptyset$ ;  $U \leftarrow \emptyset$ ;  $D \leftarrow \emptyset$

**foreach** input file  $f$  given on the command line **do**

**if**  $f$  is an object file **then**

$O \leftarrow O \cup \{f\}$ ;  $D \leftarrow D \cup \text{global } f$ ;  $U \leftarrow (U \setminus D) \cup \text{external } f$

**else if**  $f$  is an archive **then**

**repeat**

**foreach** object module  $m$  which is a member of  $f$  **do**

**if**  $U \cap \text{global } m \neq \emptyset$  **then**

$O \leftarrow O \cup \{m\}$ ;  $D \leftarrow D \cup \text{global } m$ ;  $U \leftarrow (U \setminus D) \cup \text{external } m$

**until**  $U$  and  $D$  do not change anymore

**if**  $U \neq \emptyset$  **then**

    Fail with error message: Undefined references to all symbols in  $U$

**else**

    Relocate object modules in  $O$  and build executable.

(global  $m$  = symbols defined in  $m$ ; external  $m$  = symbols not defined in, but referenced by  $m$ )

## ■ Review the previous example: Why does this fail?

```
1 $ gcc -static -omain libvector.a main.o # Wrong!
```

- $U$  is empty when `libvector.a` is visited,
- so no object modules are added to  $O$ .
- When `main.o` is checked, `dotproduct` is added to  $U$ , but `libvector.a` is not visited again.

## ■ Sometimes, it is necessary to specify a library multiple times on the command line. Example (pseudocode!):

```
1 main.o
2     main() { foo(); }
```

```
3 libfoobar.a
4     foo.o
5         foo() { ding(); }
6     bar.o
7         bar() { dong(); }
```

```
8 libdingdong.a
9     ding.o
10         ding() { bar(); }
11     dong.o
12         dong() { foo(); }
```

```
1 $ gcc -static -omain main.o -L. -lfoobar -ldingdong -lfoobar -ldingdong
```

# Shared libraries

- Shared libraries are linked to the program **not until runtime**.
- **Different programs** can use the same shared library.
- The tool `ldd(1)` lists the **dynamic dependencies** of a thus linked binary.<sup>1</sup>

## Example

- The binary created in the previous section is quite big:

```
1 $ gcc -static -o main main.o -L. -lvectors
2 $ ls -l main
3 -rw----- 1 marcel users 826k Feb  3 18:37 main
4 $ ldd main
5          not a dynamic executable
```

- This is, because the `-static` flag enforces static linking, including way more than only `libvector`.

---

<sup>1</sup>Security risk: Runs program!

- Without `-static`, **dynamic linking** is used where possible:

```

1 $ gcc -o main main.o -L. -lvector
2 $ ls -l main
3 -rwx----- 1 marcel users 8.7k Feb  3 18:44 main
4 $ ldd main
5         linux-vdso.so.1 (0x00007fffc013a000)
6         libc.so.6 => /usr/lib/libc.so.6 (0x00007ff560b8c000)
7         /lib64/ld-linux-x86-64.so.2 (0x00007ff560f36000)

```

- `linux-vdso.so.1` (Virtual Dynamic Shared Objects) is a part of the kernel, providing **fast system calls**. It is not a shared library in the usual sense.
- `libc.so.6` is the **standard C library** on Linux systems.
- `ld-linux-x86-64.so.2` contains the ELF **dynamic linker and loader**.

Most Programs (unless compiled `-static`) will depend on these.

- Obviously, `libvector.a` is **not shared**, but still statically linked.  
⇒ How can we change this?



# Making a shared library

- The individual object modules have to be compiled with `-fPIC`.
  - This generates **position-independent code**, which allows relocation later on, *cf.* page 46.
- Instead of using `ar(1)`, the object files are linked together into a **single shared object file** with `.so` suffix.

**Example** To build a shared `libvector` instead of a static one:

```
1 $ gcc -c -fPIC addvec.c
2 $ gcc -c -fPIC dotproduct.c
3 $ gcc -shared -o libvector.so addvec.o dotproduct.o
4 $ nm libvector.so      # my first shared library
5 00000000000000665 T addvec
6 #...
7 000000000000006d5 T dotproduct
8 #...
```

# Using a shared library

- The shared library is used just like a static library:

```
1 $ gcc -c main.c
2 $ gcc -o main main.o -L. -lvector
```

- The generated binary now **depends** on `libvector.so`:

```
1 $ ldd main
2     linux-vdso.so.1 (0x00007fffe8eaa000)
3     libvector.so => not found
4     libc.so.6 => /usr/lib/libc.so.6 (0x00007f2cbc343000)
5     /lib64/ld-linux-x86-64.so.2 (0x00007f2cbc6ed000)
6 $ ./main
7 ./main: error while loading shared libraries: libvector.so: cannot open
8 shared object file: No such file or directory
```

- The dynamic linker `ld-linux(8)` only searches a **default path**:

- Falling back to `/lib`, and `/usr/lib`, but see the manual!
- The search path can be extended (prefixed) with `$LD_LIBRARY_PATH`.

```
1 $ LD_LIBRARY_PATH=. ./main
2 <x,y> = 0
3 <x,z> = 1
```

# Choosing a shared library at runtime

- Applications can **decide on which shared libraries to load at runtime**.

```
1 #include <dlfcn.h>
2
3 void *dlopen(const char *filename, int flag);
4 void *dlsym(void *handle, const char *symbol);
5 int dlclose(void *handle);
6 char *dlerror(void);
```

- **dlopen(3)** **loads** a shared library, and returns a handle to it.
  - See the manual for how the library is **searched**.
  - The flag indicates when/how to **resolve symbols**:
    - RTLD\_NOW** Before **dlopen** returns, or
    - RTLD\_LAZY** when the called function is needed.
    - ... further flags are available
- **dlsym(3)** returns a **pointer to the symbol** named.
- **dlclose(3)** **unloads** a shared library if it is not used anymore.
- **dlopen(3)** and **dlsym(3)** return **NULL** on failure. **dlerror(3)** returns a string describing the most recent error.

# Example

```
1 typedef int (*dotproduct_t)(int, int *, int *);
2
3 int main(int argc, char *argv[])
4 {
5     dotproduct_t dotproduct;
6     void *handle;
7
8     if (argc < 2) errx(1, "use: sick <lib>");
9
10    handle = dlopen(argv[1], RTLD_NOW);
11    if (handle == NULL) errx(1, "dlopen: %s", dlerror());
12
13    /* dotproduct = (dotproduct_t)dlsym(handle, "dotproduct"); */
14    *(void **)&dotproduct = dlsym(handle, "dotproduct");
15    if (dotproduct == NULL) errx(1, "dlsym: %s", dlerror());
16
17    int x[3] = { 1, 0, 0 }, y[3] = { 0, 1, 1 }, z[3] = { 1, 1, 0 };
18    printf("<x,y> = %d\n", dotproduct(3, x, y));
19    printf("<x,z> = %d\n", dotproduct(3, x, z));
20
21    dlclose(handle);
22    return 0;
23 }
```


```
1 $ ./sick ./libvector.so
2 <x,y> = 0
3 <x,z> = 1
4 $ ./sick ./libfake.so
5 <x,y> = 42
6 <x,z> = 42
```

- Some projects use this mechanism to provide a **plugin interface**.
- If the program is compiled with `-rdynamic`, then a loaded library can use the program's global symbols.

# Position-independent code


- **Different programs** should use a shared library simultaneously.
- Mapping is likely to happen to **different virtual memory regions**.
- Simple **relocation breaks sharing**, since it modifies `.text`.

## How can we solve this?




- Recall process memory layout:
  - Shared libraries' `.text` and `.data` is **not merged** with the main program's.
  - Instead, each shared library is loaded **somewhere above 0x40000000**.
- **Position-independent code** is compiled in a way that allows it to be executed at any address, without prior relocation.

# Memory layout of a shared library



- The data segment of a shared library is mapped **directly after** the code segment.
  - For each access to a local symbol, the **distance from instruction to variable** is **fix!**
  - This  $\Delta$  is known at **compile time**.
  - Variable access is implemented by **offset from the program counter**, instead of absolute addresses.
- How can we access variables in **other modules?**

# Memory layout of a shared library



- PIC adds **one level of indirection** to access external symbols.
  - A **global offset table (GOT)** is added at the **start of the data segment** of every module.
  - At compile time, references to variables are replaced by **indirect references** via the GOT.
  - The **dynamic loader fills the GOT** with the correct addresses (relocation) **at runtime**.
- Thus, relocation **happens in the private data segment!**
  - The code segment can be shared.



# Final Remarks

- Relocation for **function calls** also uses the GOT.
  - A more sophisticated algorithm, called **lazy binding**, reduces the overhead after the first function call.
- Shared libraries come with a **runtime overhead** for accessing any external symbols.
- Using shared libraries requires **expensive setup of all GOTs** when loading a program.
- Shared libraries **increase code sharing** (and page sharing) more than static libraries.
  - Static library code **cannot be shared between different programs**, only between different instances of the same program.
- An in-depth discussion about shared libraries can be found here:
  - Ulrich Drepper. *How To Write Shared Libraries*. December 2011, <http://www.akkadia.org/drepper/dsohowto.pdf>.