

A detailed close-up photograph of a computer motherboard. The image shows various components including blue plastic connectors, yellow RAM modules, and a large blue heat sink. The circuit board is black with intricate gold-colored traces and several silver capacitors. The lighting is bright, highlighting the textures and colors of the hardware.


Lecture

Operating System

17. Free Space Management

17. Free Space Management

- 1. Splitting and Coalescing**
- 2. malloc() and free()**
- 3. Free List Management**
- 4. Managing Free Space**




17. Free Space Management

1. Splitting and Coalescing

2. `malloc()` and `free()`

3. Free List Management


4. Managing Free Space



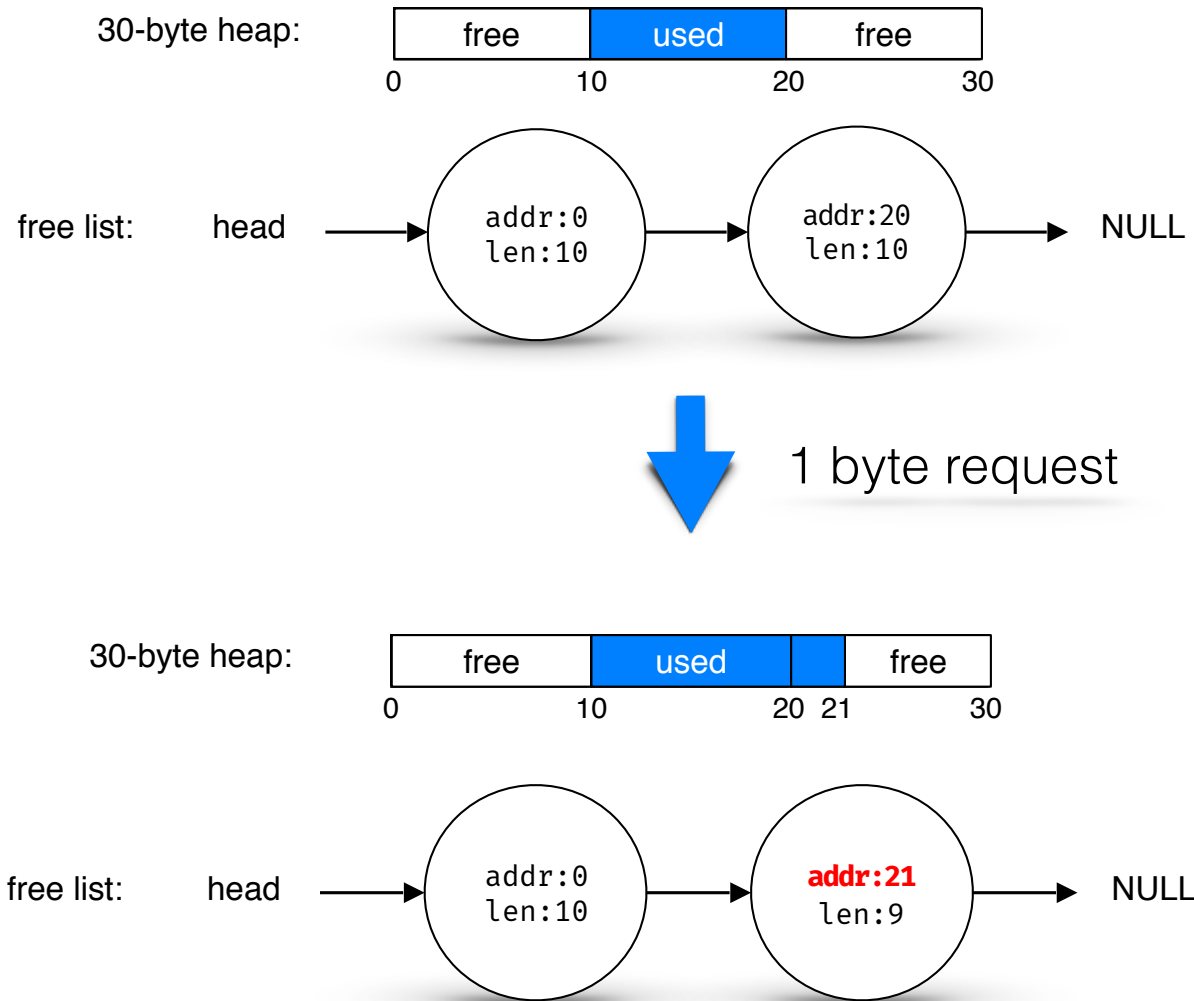
Splitting

- Finding a free chunk of memory that can satisfy the request and splitting it into two.
- When request for memory allocation is **smaller** than the size of free chunks.

30-byte heap:




Splitting (1-byte request)



Coalescing

- If a user requests memory that is **bigger than free chunk size**, the list will **not find** such a free chunk.
- **Coalescing**: **Merge** returning a free chunk with existing chunks into a large single free chunk if **addresses** of them are **nearby**.




Coalescing free chunks



17. Free Space Management

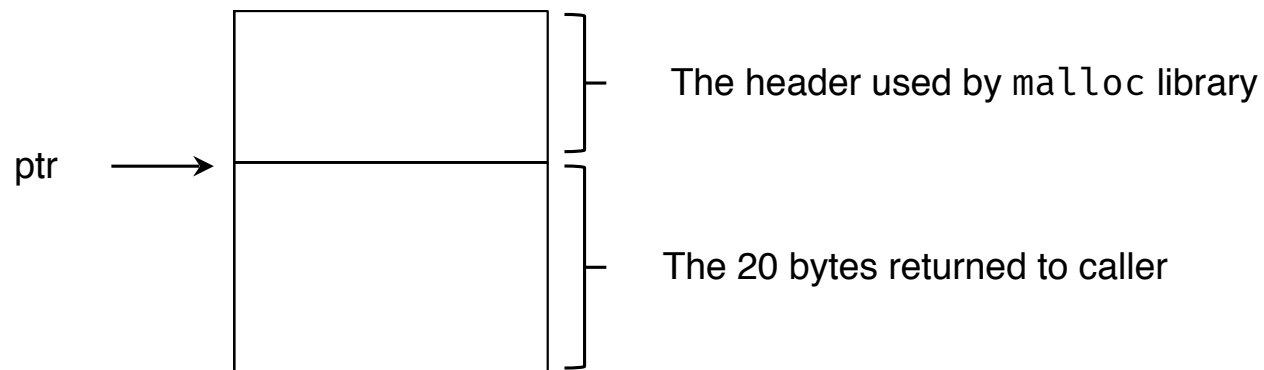
1. Splitting and Coalescing
- 2. malloc() and free()**
3. Free List Management
4. Managing Free Space



Tracking Size of Allocated Regions

- The interface to `free(void *ptr)` does **not take a size parameter**.
- How does the library **know the size** of memory region that will be back into free list?
- Most allocators store **extra information** in a **header block**.

```
ptr = malloc(20);
```




An Allocated Region Plus Header

Header of Allocated Memory Chunk

- The header minimally **contains the size** of the allocated memory region.
- The header may also contain
 - Additional pointers to speed up deallocation
 - A magic number for integrity checking
- The **size** for free region is the **size of the header plus the size of the space** allocated to the user.
 - If a user **request N bytes**, the library searches for a free chunk of **size N plus the size of the header**
- Simple pointer arithmetic to find the header pointer (`hptr`).

Header: Example


```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```



```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
}
```

17. Free Space Management

1. Splitting and Coalescing
2. malloc() and free()
- 3. Free List Management**
4. Managing Free Space



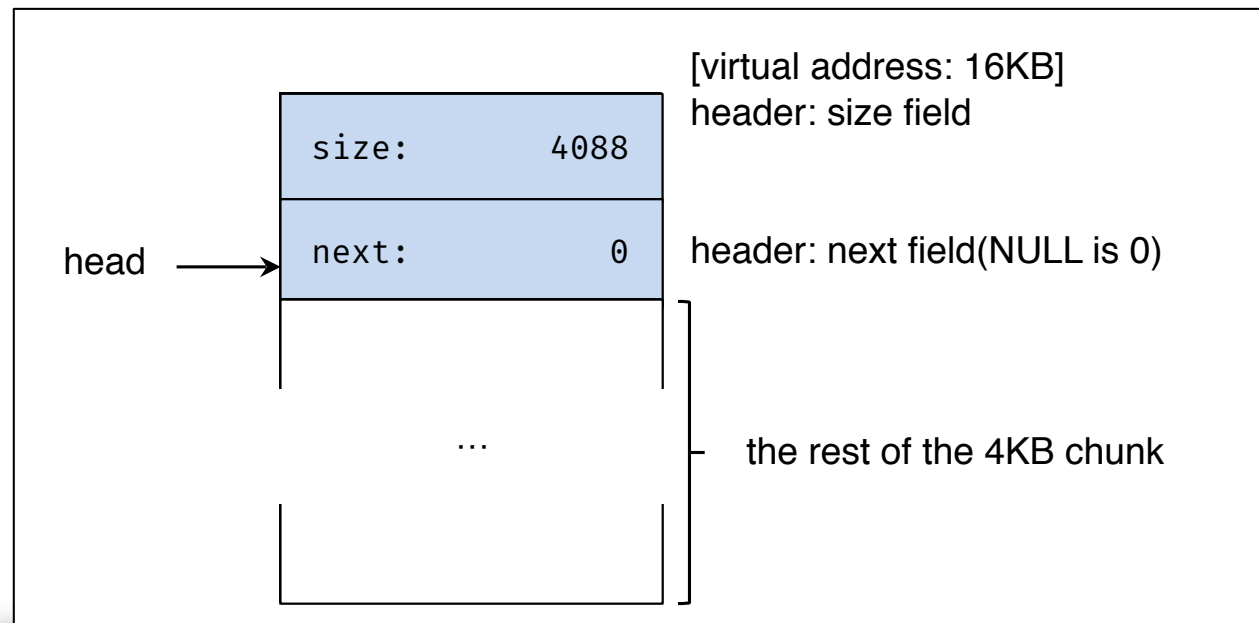
Embedding A Free List

- The memory-allocation library **initializes** the heap and puts the first element of the **free list** in the **free space**.
 - The library **can't use** `malloc()` to build a list **within itself**.
- Build a heap and put a free list into it.
 - Assume that the heap is built via `mmap()` system call.

A Heap With One Free Chunk

```
// mmap() returns a pointer to a chunk of free space
node_t *head = mmap(NULL, 4096, PROT_READ|PROT_WRITE,
                    MAP_ANON|MAP_PRIVATE, -1, 0);
head->size = 4096 - sizeof(node_t);
head->next = NULL;
```

```
typedef struct __node_t {
    int size;
    struct __node_t *next;
} node_t;
```




Embedding A Free List: Allocation

- If a chunk of memory is requested, the library **will first find** a chunk that is **large enough** to accommodate the request.
- The library will
 - **Split** the large free chunk into two.
 - **One** for the **request** and the **remaining** free chunk
 - **Shrink** the size of free chunk in the list.

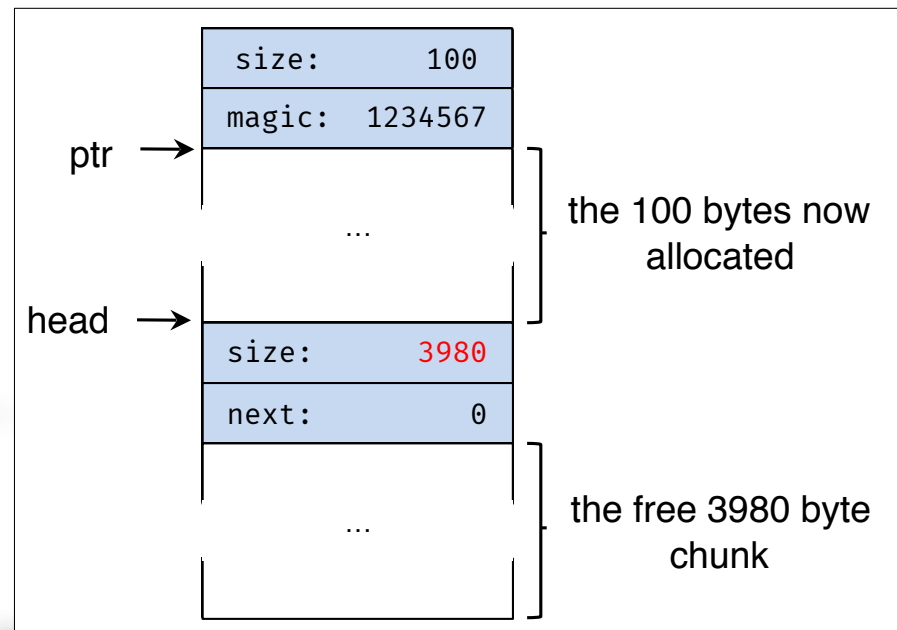
Allocation Example

- A request for 100 bytes by `ptr = malloc(100)`
 - Allocating 108 bytes out of the existing one free chunk.
 - shrinking the one free chunk to 3980(4088 minus 108).

A 4KB Heap With One Free Chunk

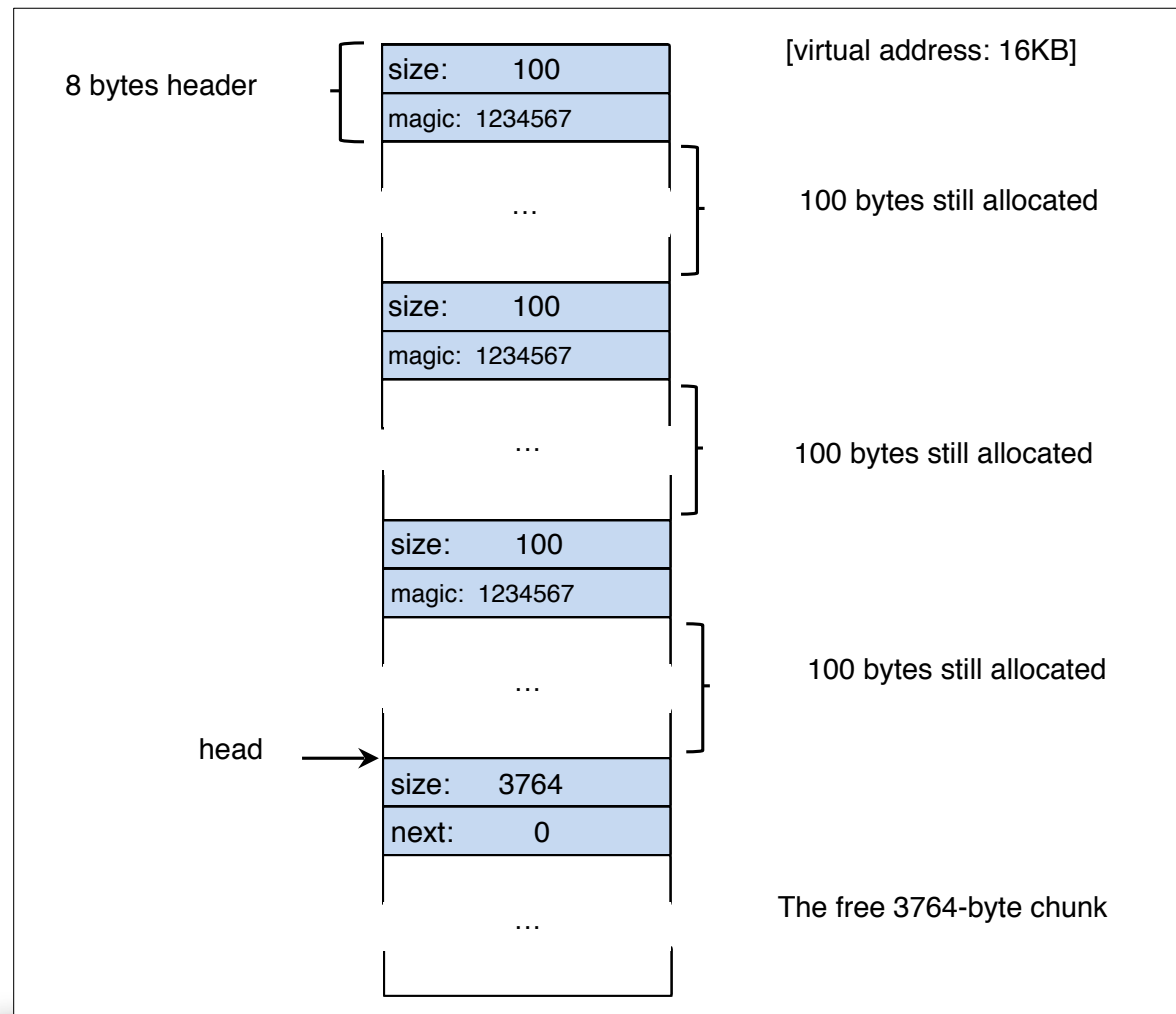


A Heap : After One Allocation



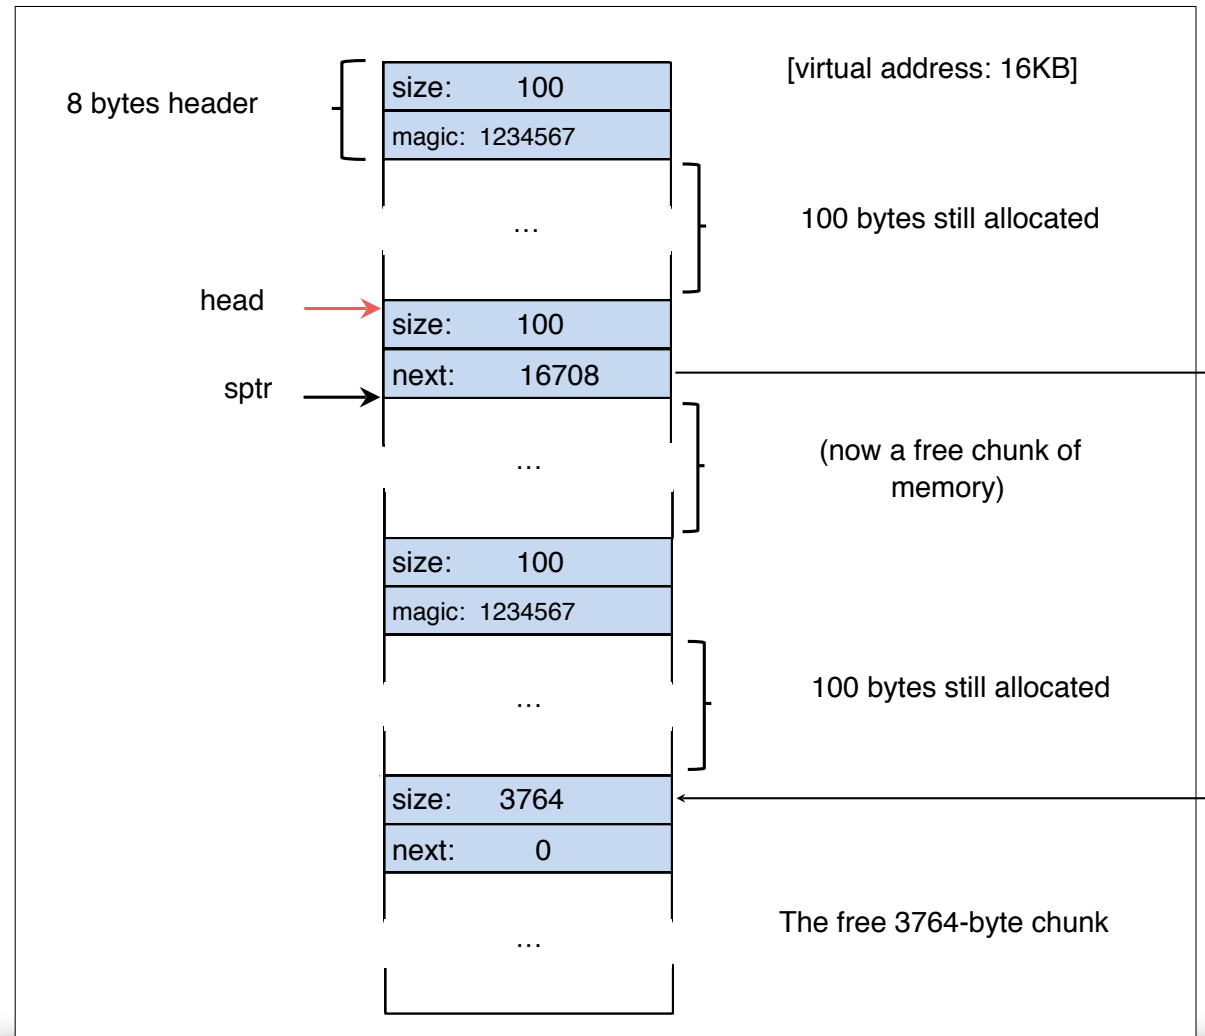
Free Space With Chunks Allocated

Free Space With Three Chunks Allocated



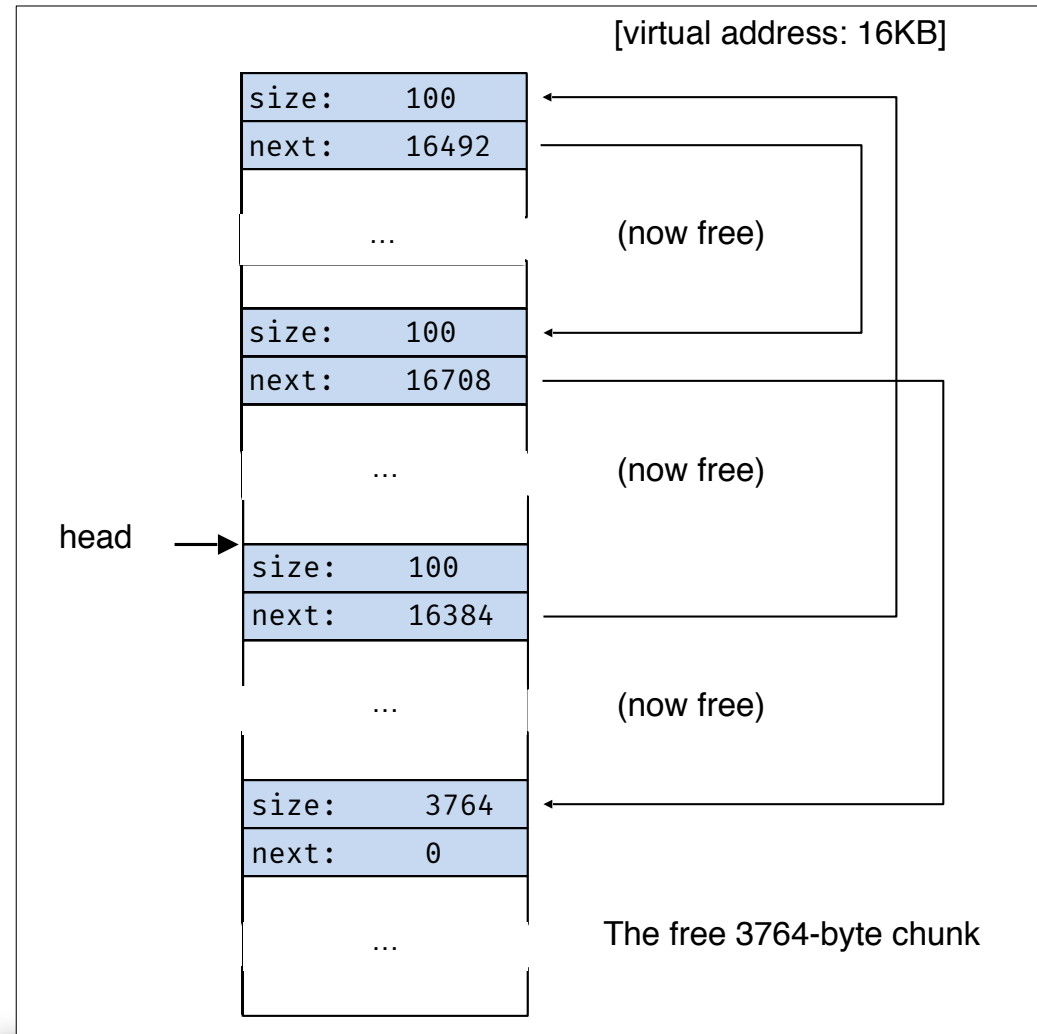
Free Space With `free()`

- The 100 bytes chunks is **back into** the free list.
- The free list will **start** with a **small chunk**.
 - The list header will point the small chunk



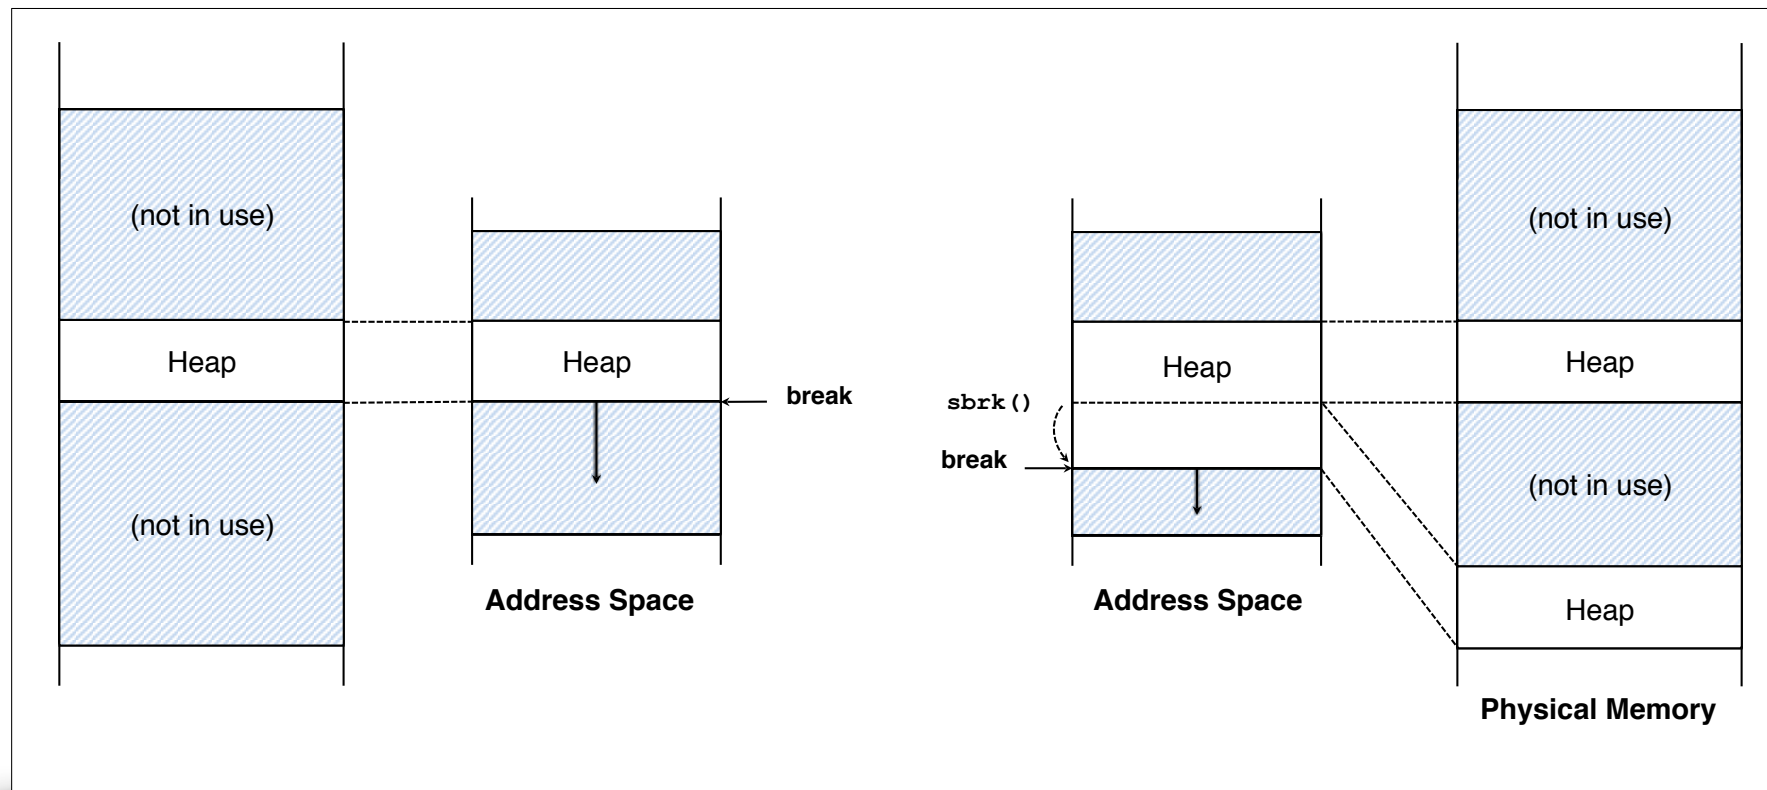
Free Space With Freed Chunks

- Let's assume that the last two in-use chunks are freed.
- **External Fragmentation** occurs.
 - **Coalescing** is needed in the list.




Growing The Heap

- Most allocators **start** with a **small-sized heap** and then **request more** memory from the OS when they run out.
- e.g., `sbrk()`, `brk()` in most UNIX systems.



17. Free Space Management

1. Splitting and Coalescing
2. malloc() and free()
3. Free List Management
- 4. Managing Free Space**



Managing Free Space: Basic Strategies

■ **Best Fit:**

- Finding free chunks that are **big or bigger than the request**
- Returning the **one of smallest** in the chunks **in the group** of candidates

■ **Worst Fit:**

- Finding the **largest free chunks** and allocation the amount of the request
- **Keeping the remaining chunk** on the free list.

Managing Free Space: Basic Strategies


■ **First Fit:**

- Finding the **first chunk** that is **big enough** for the request
- Returning the requested amount and remaining the rest of the chunk.

■ **Next Fit:**

- Finding the first chunk that is big enough for the request.
- Searching at **where one was looking** at instead of the beginning of the list.

Example: Best, First and Next Fit



Other Approaches: Segregated List

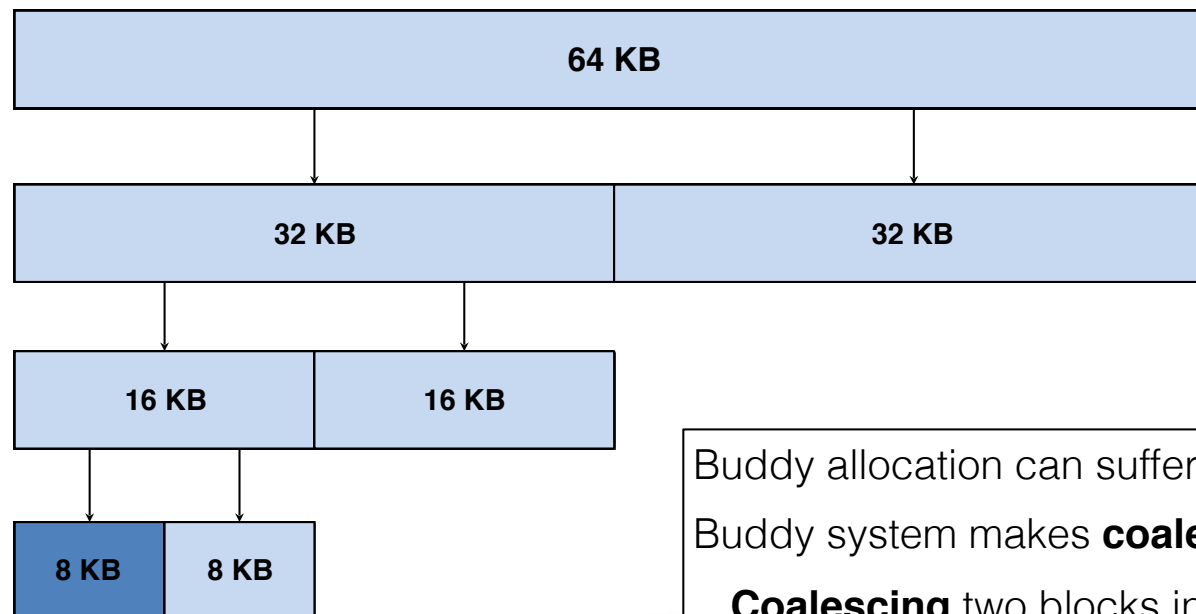
■ Segregated List:

- Keeping free chunks in different size in a separate list for the size of popular request.
- New Complication:
 - **How much** memory should dedicate to **the pool of memory** that serves **specialized requests** of a given size?
- Slab Allocator will solve this!

Other Appr.: Buddy Allocation

- Binary Buddy Allocation
 - The allocator **divides free space** by two **until a block** that is big enough to accommodate the request is **found**.

64KB free space for 7KB request



Buddy allocation can suffer from **internal fragmentation**.
Buddy system makes **coalescing** simple:
Coalescing two blocks in to the next level of block.

Example: Buddy Allocation

free

used

1 Mbyte block

1 M

Request 100 K

A = 128 K	128 K	256 K	512 K
-----------	-------	-------	-------

Request 240 K

A = 128 K	128 K	B = 256 K	512 K
-----------	-------	-----------	-------

Request 64 K

A = 128 K	C = 64 K	64 K	B = 256 K	512 K
-----------	----------	------	-----------	-------

Request 256 K

A = 128 K	C = 64 K	64 K	B = 256 K	D = 256 K	256 K
-----------	----------	------	-----------	-----------	-------

Release B

A = 128 K	C = 64 K	64 K	256 K	D = 256 K	256 K
-----------	----------	------	-------	-----------	-------

Release A

128 K	C = 64 K	64 K	256 K	D = 256 K	256 K
-------	----------	------	-------	-----------	-------

Request 75 K

E = 128 K	C = 64 K	64 K	256 K	D = 256 K	256 K
-----------	----------	------	-------	-----------	-------

Release C

E = 128 K	128 K	256 K	D = 256 K	256 K
-----------	-------	-------	-----------	-------

Release E

512 K	D = 256 K	256 K
-------	-----------	-------

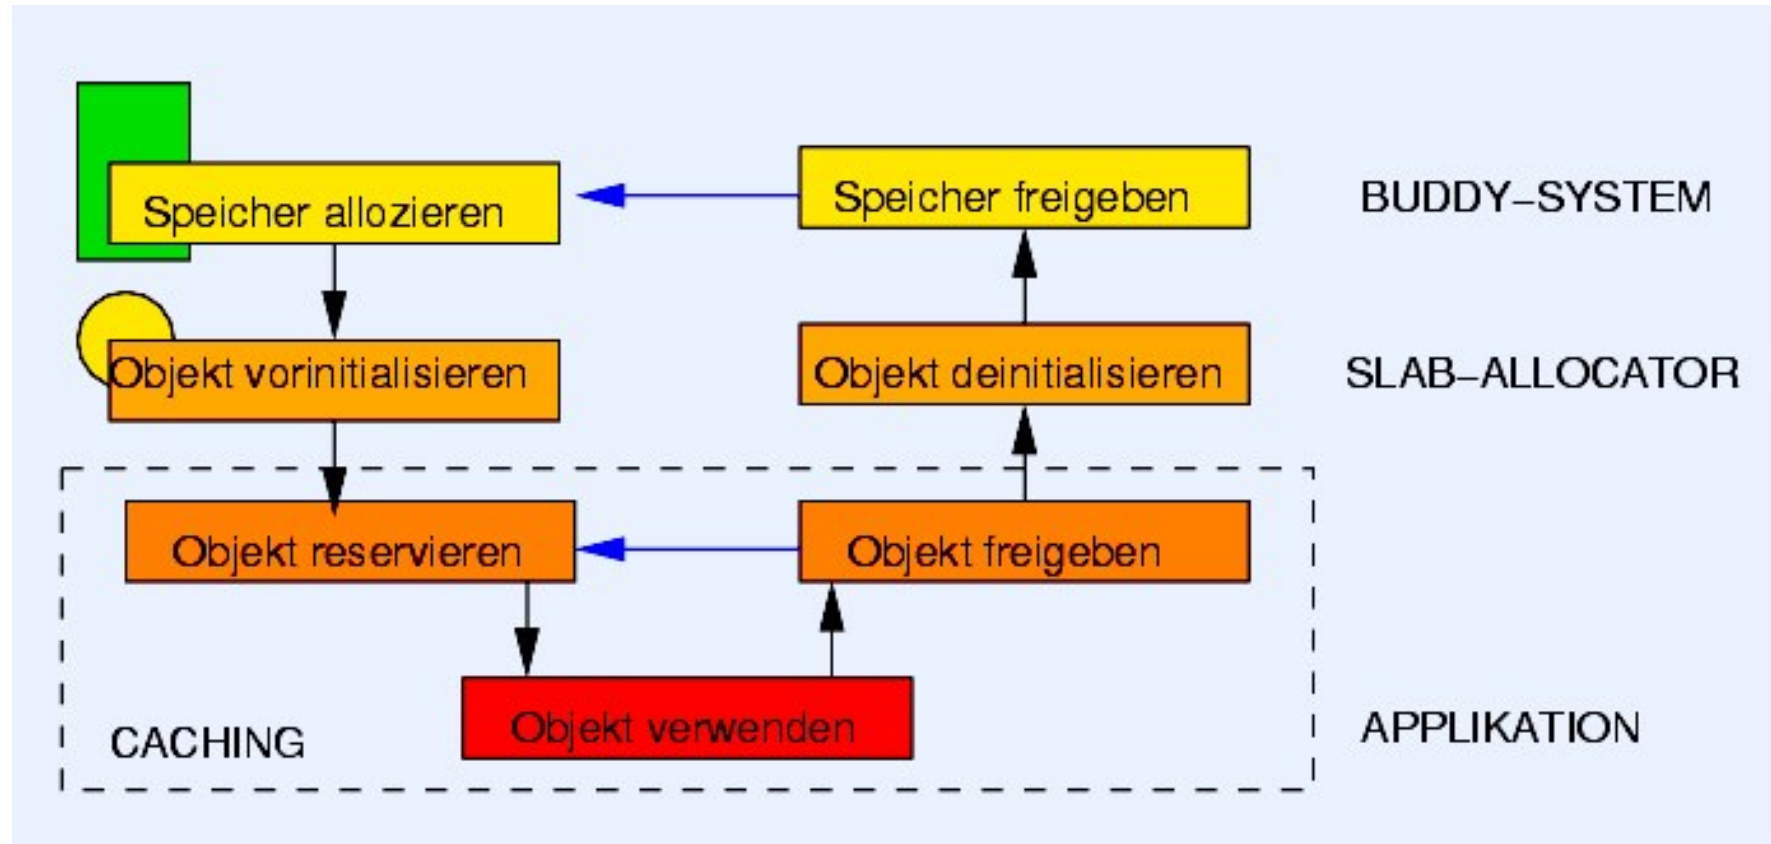
Release D

1 M

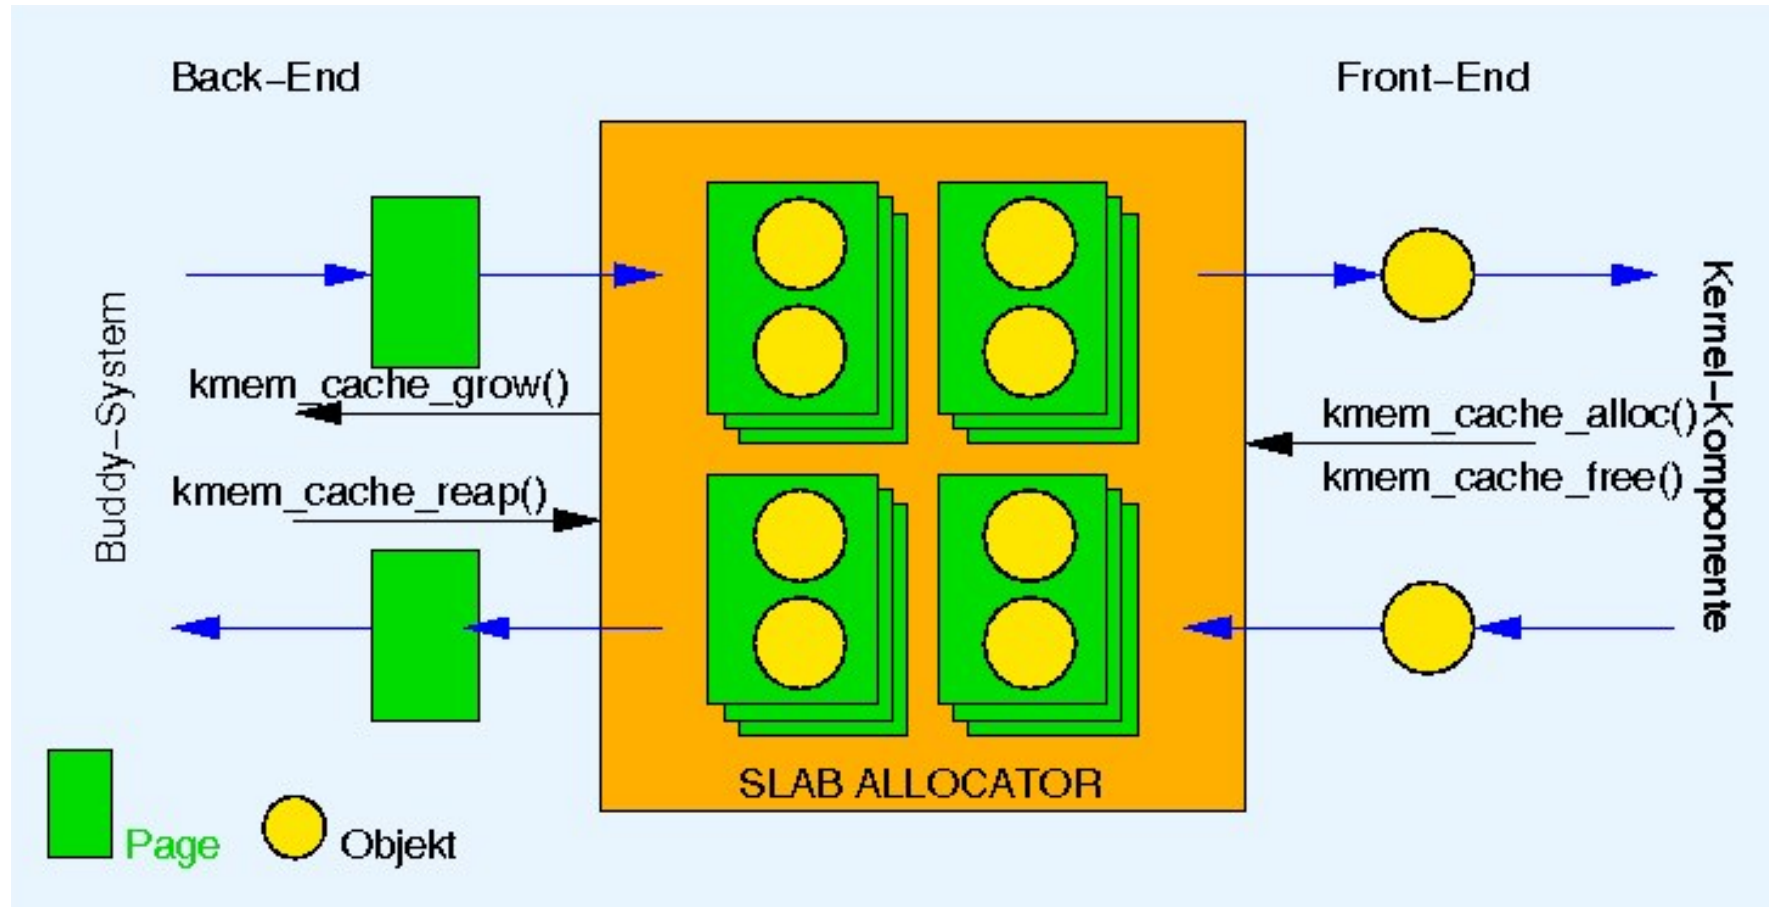
Other Approaches: SLAB Allocator

- Allocate a number of object caches.
 - The objects are likely to be requested frequently.
 - e.g., locks, file-system inodes, etc.
- **Request some memory** from a more general memory allocator when **a given cache is running low** on free space.

SLAB: Overview



SLAB: Cache Structure



SLAB: Fragmentation

Slab bestehend aus 1 Page (4KByte)



Page

$$\text{Fragmentierung} = \frac{4096\text{Bytes} - 2052\text{Bytes}}{1 \text{ Objekt}} = 2044 \frac{\text{Bytes}}{\text{Objekt}}$$

Slab bestehend aus 2 Pages (2*4KByte)

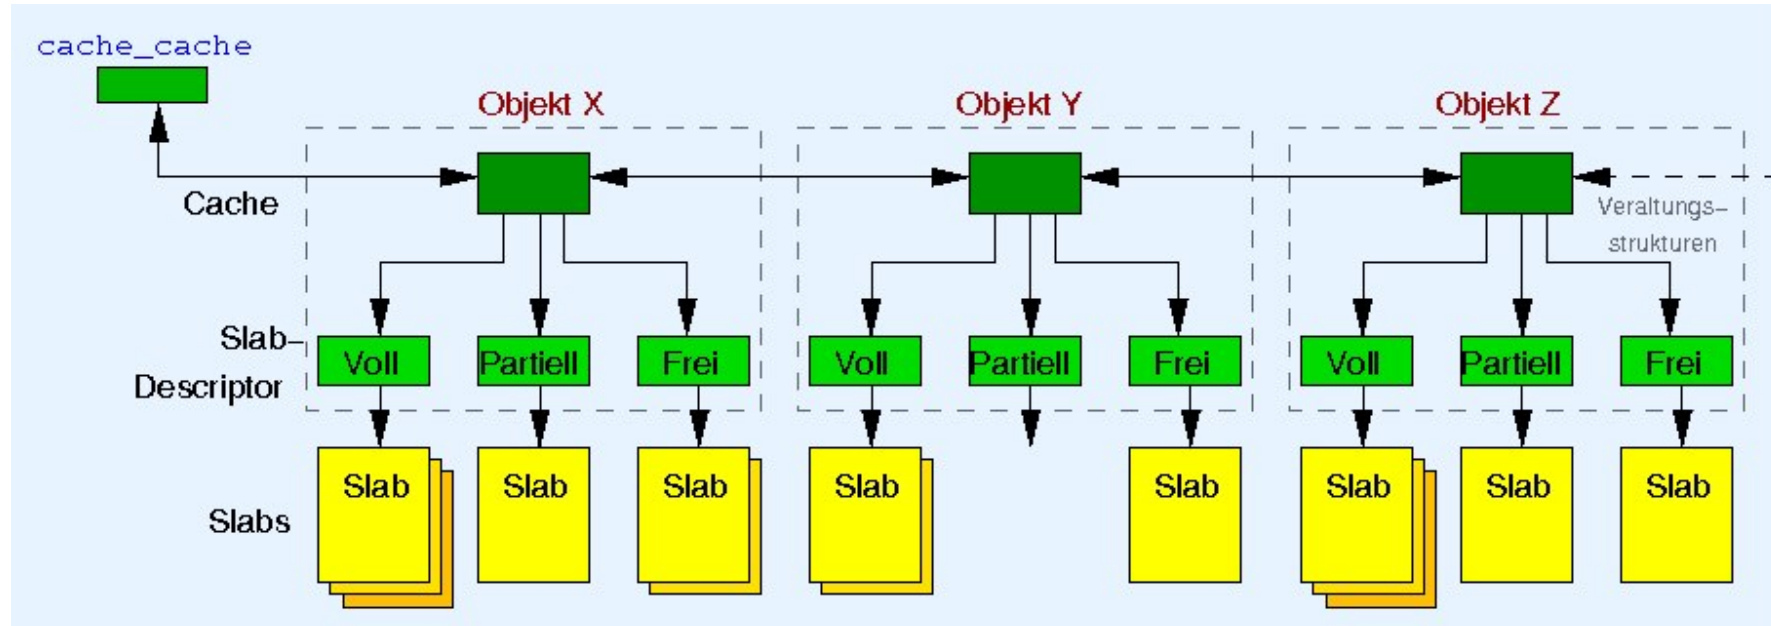


Page

Page

$$\text{Fragmentierung} = \frac{2 \cdot 4096\text{Bytes} - (3 \cdot 2052)\text{Bytes}}{3 \text{ Objekte}} = 678,7 \frac{\text{Bytes}}{\text{Objekt}}$$

SLAB: Cache Objects



A close-up photograph of a computer motherboard. The image shows various components including blue plastic connectors, yellow plastic connectors, and a large blue heat sink. The motherboard is black with visible circuitry and components like capacitors.

Thanks

Questions