

A detailed close-up photograph of a computer motherboard. The image shows various components including blue plastic connectors, yellow RAM modules, and a large blue heat sink. The circuitry of the motherboard is visible, with gold-plated pins and silver capacitors. The lighting is bright, highlighting the textures and colors of the hardware.


# Lecture Operating System

## **18. Paging Introduction**



# 18. Paging Introduction

- 1. How can we virtualize memory with pages, so as to avoid the problems of segmentation?**
- 2. What are the basic techniques?**
- 3. How do we make those techniques work well, with minimal space and time overheads?**



# Concept of Paging

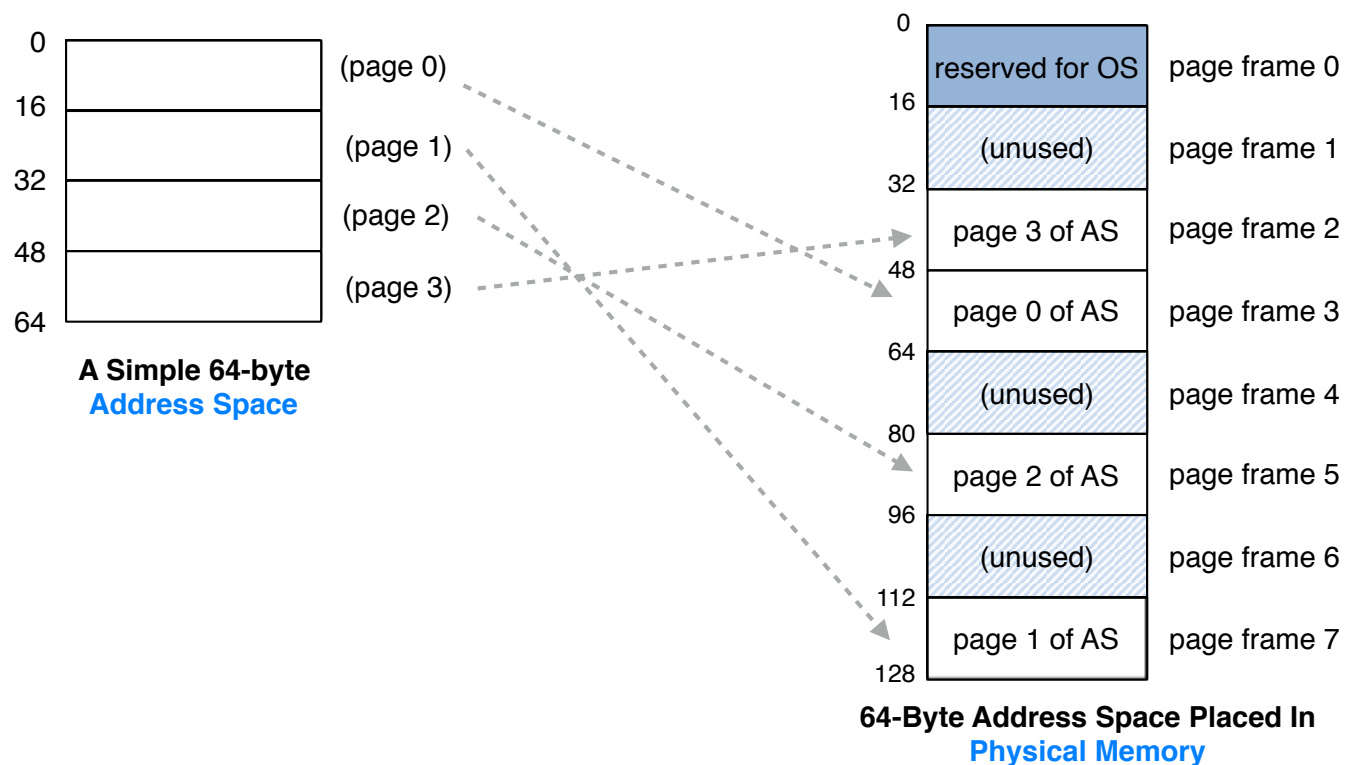
- Paging **splits up** address space into **fixed-sized** unit called a **page**.
  - Segmentation: variable size of logical segments (code, stack, heap, etc.)
- With paging, **physical memory** is also **split** into some number of pages called a **page frame**.
- **Page table** per process is needed to **translate** the virtual address to physical address.

# Advantages of Paging

- **Flexibility:** Supporting the abstraction of address space effectively
  - Don't need assumption how heap and stack grow are used.
- **Simplicity:** ease of free-space management
  - The page in address space and the page frame are the same size.
  - Easy to allocate and keep a free list



# Example: A Simple Paging

- 128-byte physical memory with 16 bytes **page frames**
- 64-byte address space with 16 bytes **pages**



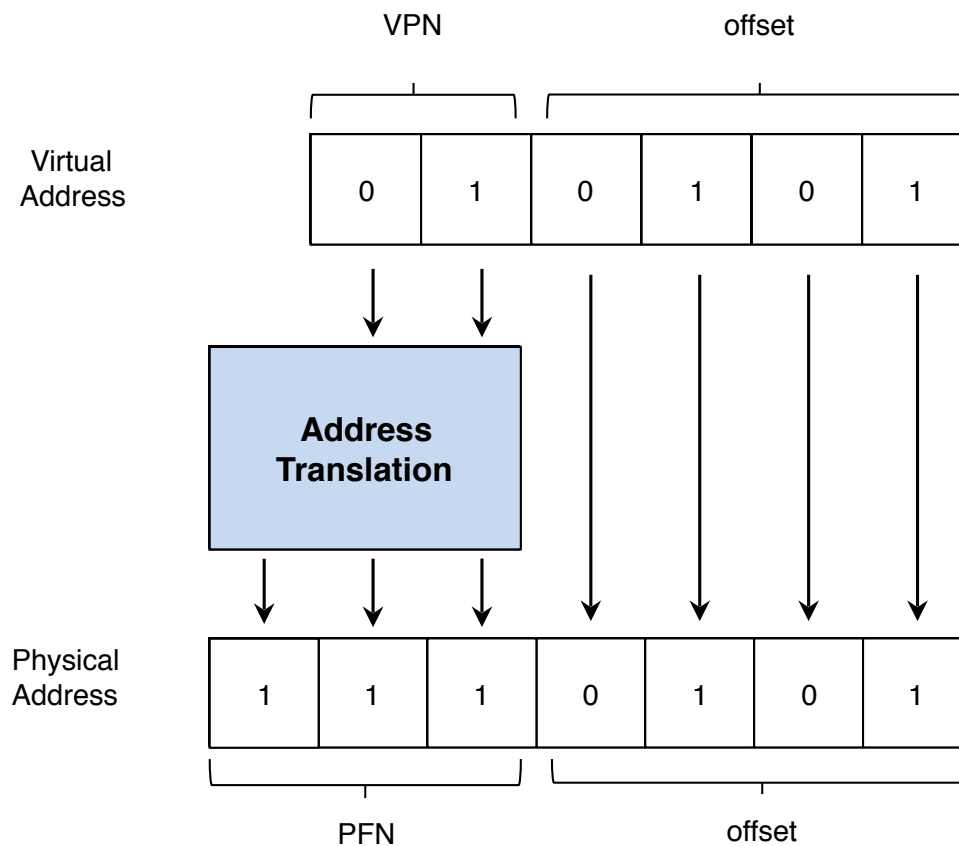
# Address Translation

- Two components in the virtual address:
  - VPN: virtual page number
  - Offset: offset within the page
- Example: 64-byte address space
  - 21 in 64-byte address space
    - 010101



Example: virtual address 21 in 64-byte address space

# Address Translation



# Page Tables


- The page table is just a data structure that is used to map the virtual address to physical address.
  - Simplest form: a linear page table, an array
- The OS indexes the array by VPN, and looks up the page-table entry.
- Page tables can get awfully **large**
  - 32-bit address space with 4-KB pages, 20 bits for VPN
    - $4MB = 2^{20} \text{ entries} * 4 \text{ Bytes per page table entry}$
- Page tables for **each** process are stored in memory.



# Common Flags of Page Table Entry

- **Valid Bit:** Indicating whether the particular translation is valid.
- **Protection Bit:** Indicating whether the page could be read from, written to, or executed from
- **Present Bit:** Indicating whether this page is in physical memory or on disk(swapped out)
- **Dirty Bit:** Indicating whether the page has been modified since it was brought into memory
- **Reference Bit (Accessed Bit):** Indicating that a page has been accessed

# Page Table Entry



## An x86 Page Table Entry(PTE)

P:	present
R/W:	read/write bit
U/S:	supervisor
A:	accessed bit
D:	dirty bit
PFN:	the page frame number

# Paging: Too Slow

- To find a location of the desired PTE, the **starting location** of the page table is needed.
- For every memory reference, paging requires the OS to perform one **extra memory reference**.

```
1 // Extract the VPN from the virtual address
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4 // Form the address of the page-table entry (PTE)
5 PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7 // Fetch the PTE
8 PTE = AccessMemory(PTEAddr)
9
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and fetch it
17     offset = VirtualAddress & OFFSET_MASK
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19     Register = AccessMemory(PhysAddr)
```

# A Memory Trace

## Example: A Simple Memory Access

```
int array[1000];  
...  
for (i = 0; i < 1000; i++)  
    array[i] = 0;
```



```
prompt> gcc -o array array.c -Wall -o  
prompt> ./array
```




```
0x1024 movl $0x0, (%edi,%eax,4)  
0x1028 incl %eax  
0x102c cmpl $0x03e8,%eax  
0x1030 jne 0x1024
```

# A Virtual(And Physical) Memory Trace

```

0x1024 movl $0x0, (%edi,%eax,4)
0x1028 incl %eax
0x102c cmpl $0x03e8,%eax
0x1030 jne 0x1024

```





# Paging: Summary

### ■ Advantages:

#### ■ Does **not** lead to **external fragmentation**

- as paging (by design) divides memory into fixed-sized units.

#### ■ Is quite **flexible**

- enabling the sparse use of virtual address spaces.

### ■ Disadvantages:

#### ■ implementing without care will lead to a **slower** machine

- with many extra memory accesses to access the page table.

#### ■ and to **memory waste**

- with memory filled with page tables instead of useful application data.

A close-up photograph of a computer motherboard. The image shows various components including blue plastic connectors, yellow plastic connectors, and a large blue heat sink. The motherboard is black with visible circuitry and components like capacitors.

# Thanks

## Questions?