

A detailed close-up photograph of a computer motherboard. The image shows various components including blue plastic connectors, yellow RAM modules, and a large blue heat sink. The motherboard itself is black with intricate circuitry and gold-plated pins visible in the connectors.

Lecture


# Operating System

## 21. Swapping: Mechanisms




## 21. Swapping: Mechanisms

- 1. To support large address spaces, the OS will need a place to stash away portions of address spaces (e.g. on discs) that currently aren't in great demand**
- 2. How can the OS make use of a larger, slower device to transparently provide the illusion of a large virtual address space?**



# Beyond Physical Memory: Mechanisms

- Require an additional level in the **memory hierarchy**.
  - OS need a place to stash away portions of address space that currently aren't in great demand.
  - In modern systems, this role is usually served by a **hard disk drive**
- Memory Hierarchie in modern systems:

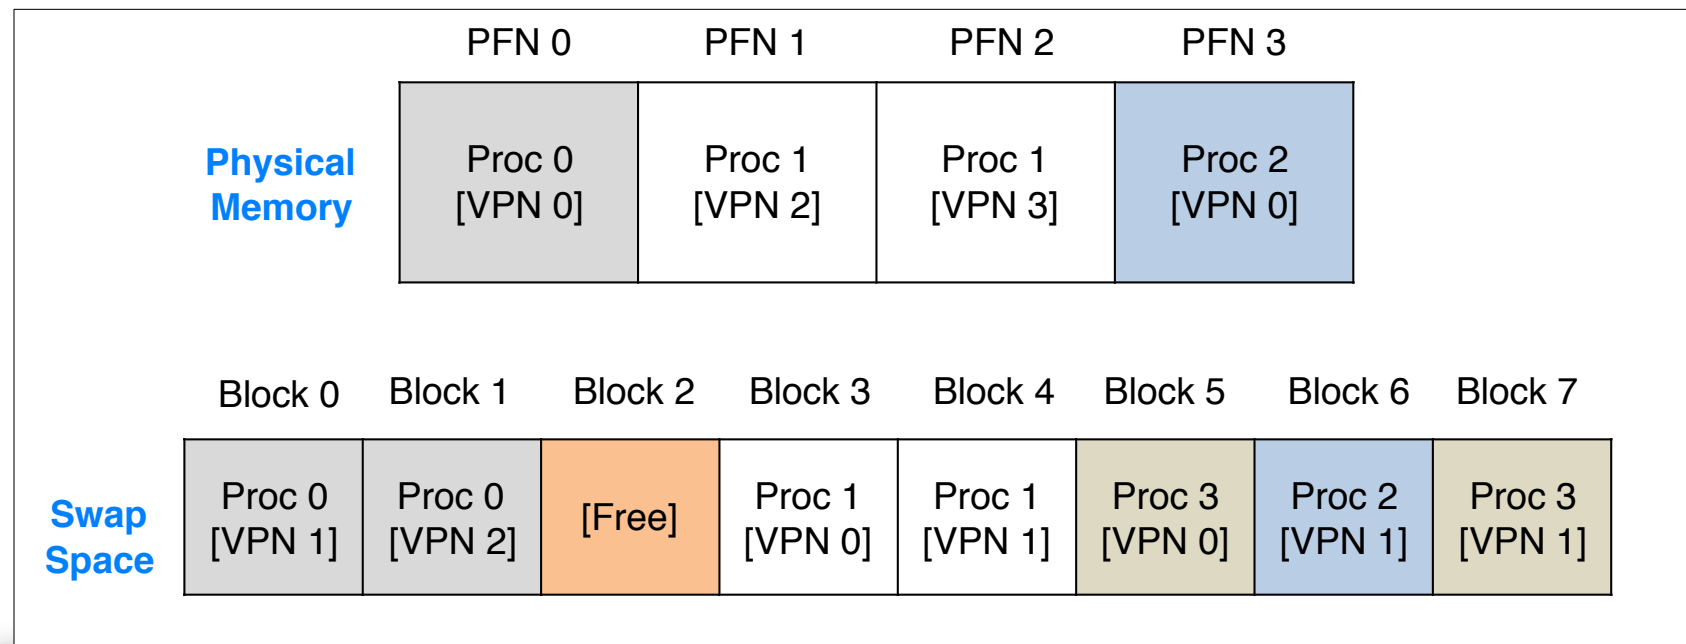


# Single large address for a process

- Always need to first arrange for the code or data to be in memory when before calling a function or accessing data.
- To Beyond just a **single process**.
  - The addition of **swap space** allows the OS to support the illusion of a large virtual memory for multiple concurrently-running process

# Swap Space

- Reserve some space on the disk for moving pages back and forth.
- OS need to remember to the swap space, **in page-sized unit**



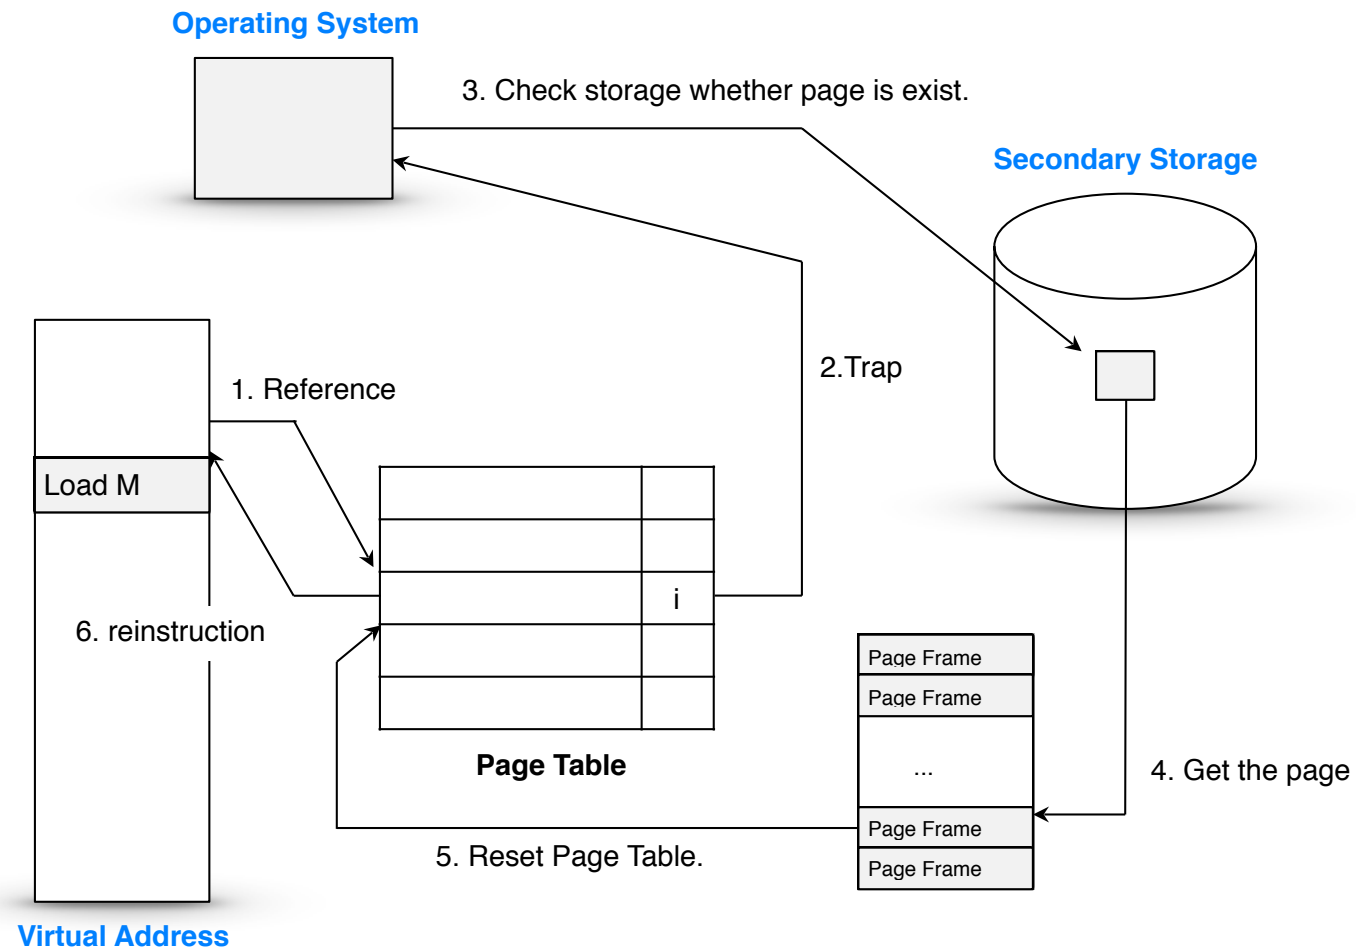
Physical Memory and Swap Space

# Present Bit

- Add some machinery higher up in the system in order to support swapping pages to and from the disk.
  - When the hardware looks in the PTE, it may find that the page is not present in physical memory.
- Accessing page that is **not in physical memory**.
  - If a page is **not present** and has been swapped disk, the OS need to swap the page into memory in order to service the **page fault**.

Value	Meaning
1	page is present in physical memory
0	The page is not in memory but rather on disk.

# Page Fault Control Flow



When the OS receives a page fault, it looks in the PTE and issues the request to disk.

# Page Fault Control Flow

## Hardware

```
1: VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2: (Success, TlbEntry) = TLB_Lookup(VPN)
3: if (Success == True) // TLB Hit
4:     if (CanAccess(TlbEntry.ProtectBits) == True)
5:         Offset = VirtualAddress & OFFSET_MASK
6:         PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:         Register = AccessMemory(PhysAddr)
8:     else RaiseException(PROTECTION_FAULT)
9: else // TLB Miss
10:    PTEAddr = PTBR + (VPN * sizeof(PTE))
11:    PTE = AccessMemory(PTEAddr)
12:    if (PTE.Valid == False)
13:        RaiseException(SEGMENTATION_FAULT)
14:    else
15:        if (CanAccess(PTE.ProtectBits) == False)
16:            RaiseException(PROTECTION_FAULT)
17:        else if (PTE.Present == True)
18:            // assuming hardware-managed TLB
19:            TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
20:            RetryInstruction()
21:        else if (PTE.Present == False)
22:            RaiseException(PAGE_FAULT)
```

page was both **present** and **valid**  
so simple grab PFN from PTE

page is **not present**  
so PageFault Handler must run



# Page Fault Control Flow

## Software

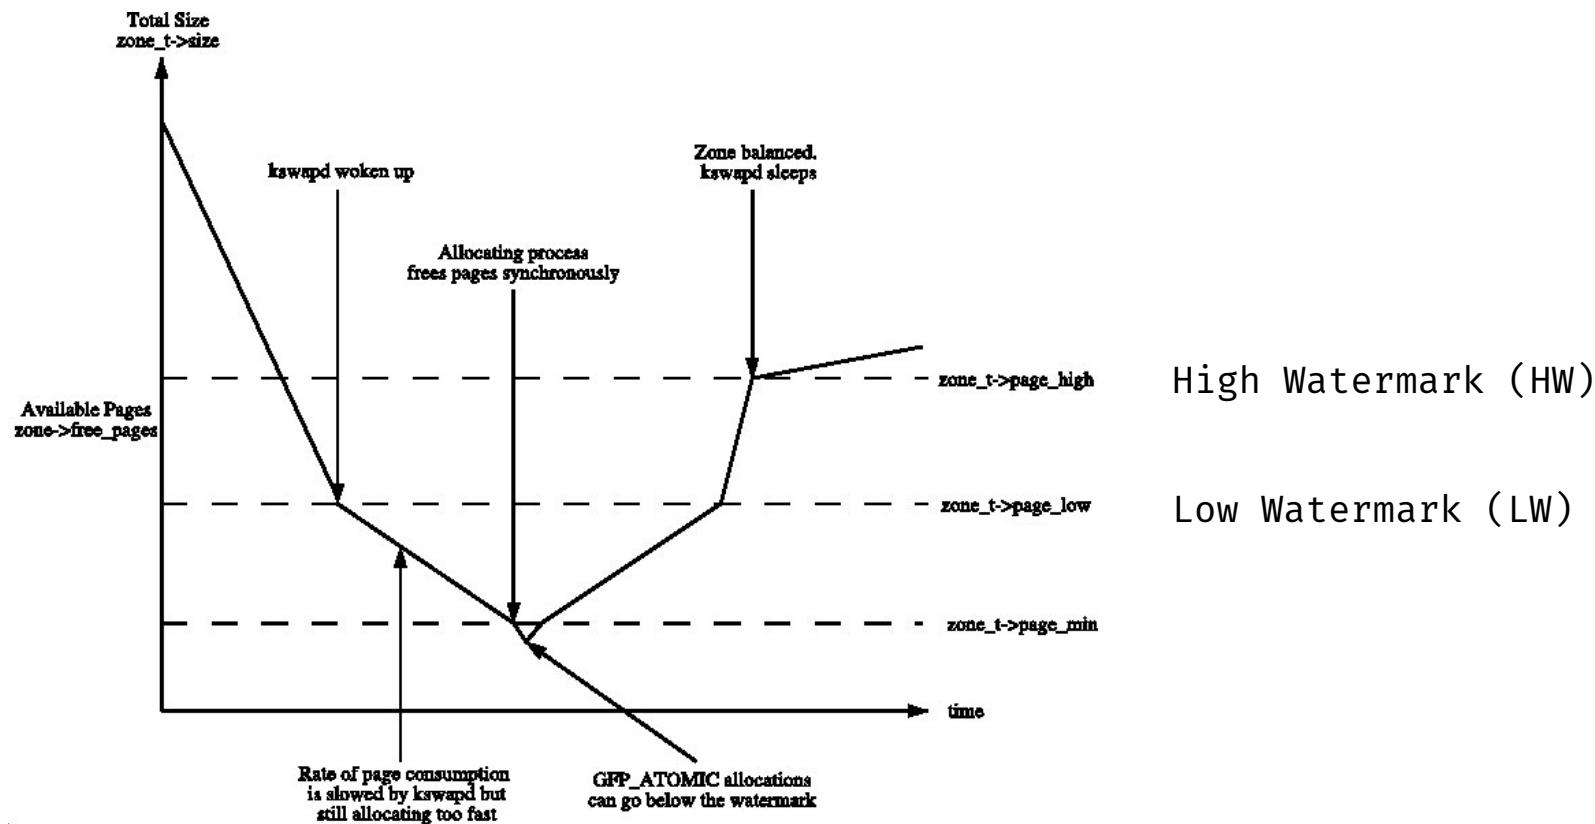
```
1:  PFN = FindFreePhysicalPage()
2:  if (PFN == -1)                // no free page found
3:      PFN = EvictPage()         // run replacement algorithm
4:  DiskRead(PTE.DiskAddr, pfn)   // sleep (waiting for I/O)
5:  PTE.present = True           // update page table with present
6:  PTE.PFN = PFN                // bit and translation (PFN)
7:  RetryInstruction()           // retry instruction
```

# Page Replacement

- The OS like to page out pages to make room for the new pages the OS is about to bring in.
  - The process of picking a page to kick out, or replace is known as **page-replacement** policy
- OS waits until memory is entirely full, and only then replaces a page to make room for some other page
  - This is a little bit unrealistic, and there are many reason for the OS to keep a small portion of memory free more proactively.

# Swap Daemon, Page Daemon

- There are fewer than **LW pages** available, a background thread that is responsible for freeing memory runs.
- The thread evicts pages until there are **HW pages** available.



A close-up photograph of a computer motherboard. The image shows several blue RAM modules installed in yellow DIMM slots. The motherboard is black with various electronic components like capacitors and circuit traces visible. The lighting is bright, highlighting the metallic and plastic surfaces.

# Thanks

## Questions