

Lecture

Operating System

29. Lock Based Data Structures

29. Lock Based Data Structures

1. Counter
2. Linked List
3. Queue
4. Hash Table



29. Lock Based Data Structures

1. Counter

2. Linked List

3. Queue

4. Hash Table



Criteria

- Adding locks to a data structure makes the structure **thread safe**.
- How locks are added determine both the **correctness** and **performance** of the data structure.

Example: Concurrent Counter

simple, but not scalable

```
typedef struct __counter_t {  
    int value;  
} counter_t;  
  
void init(counter_t *c) {  
    c->value = 0;  
}  
  
void increment(counter_t *c) {  
    c->value++;  
}  
  
void decrement(counter_t *c) {  
    c->value--;  
}  
  
int get(counter_t *c) {  
    return c->value;  
}
```

Example: Concurrent Counter

Add a **single lock**

```
void init(counter_t *c) {
    c->value = 0;
    Pthread_mutex_init(&c->lock, NULL);
}

void increment(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    c->value++;
    Pthread_mutex_unlock(&c->lock);
}


void decrement(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    c->value--;
    Pthread_mutex_unlock(&c->lock);
}

int get(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    int rc = c->value;
    Pthread_mutex_unlock(&c->lock);
    return rc;
}
```

```
typedef struct __counter_t {
    int value;
    pthread_lock_t lock;
} counter_t;
```


The performance costs

- Each thread updates a single shared counter.
 - Each thread updates the counter one million times.
 - iMac with four Intel 2.7GHz i5 CPUs.



**Performance of
Traditional vs. Sloppy Counters**
(Threshold of Sloppy, S , is set to 1024)

Perfect Scaling

- Even though more work is done, it is **done in parallel**.
- The time taken to complete the task **is not increased**.

Sloppy Counter

- The sloppy counter works by representing ...
 - A single **logical counter** via numerous local physical counters, *on per CPU core*
 - A single **global counter**
 - There are **locks**:
 - One for each local counter and one for the global counter
- Example: on a machine with four CPUs
 - Four local counters
 - One global counter

The basic idea of sloppy counting

When a thread running on a core wishes to increment the counter.

- It increment its local counter.
- Each CPU has its own local counter:
 - Threads across CPUs can update local counters *without contention*.
 - Thus counter updates are **scalable**.
- The local values are periodically transferred to the global counter.
 - Acquire the global lock
 - Increment it by the local counter's value
 - The local counter is then reset to zero.

The basic idea of sloppy counting

- **How often** the local-to-global transfer occurs is determined by a threshold, S (sloppiness).
 - The smaller S :
 - The more the counter behaves like the non-scalable counter.
 - The bigger S :
 - The more scalable the counter.
 - The further off the global value might be from the actual count.


Sloppy counter example

- Tracing the Sloppy Counters
 - The threshold S is set to 5.
 - There are threads on each of four CPUs
 - Each thread updates their local counters $L_1 \dots L_4$.

Time	L_1	L_2	L_3	L_4	G
0	0	0	0	0	0
1	0	0	1	1	0
2	1	0	2	1	0
3	2	0	3	1	0
4	3	0	3	2	0
5	4	1	3	3	0
6	5 → 0	1	3	4	5 (from L_1)
7	0	2	4	5 → 0	10 (from L_4)

Importance of the threshold value S

- Each four threads increments a counter 1 million times on four CPUs.
- Low $S \rightarrow$ Performance is **poor**, The global count is always quire **accurate**.
- High $S \rightarrow$ Performance is **excellent**, The global count **lags**.



Scaling Sloppy Counters

Sloppy Counter Implementation

```
typedef struct __counter_t {
    int global;           // global count
    pthread_mutex_t glock; // global lock
    int local[NUMCPUS];   // local count (per cpu)
    pthread_mutex_t llock[NUMCPUS]; // ... and locks
    int threshold;        // update frequency
} counter_t;

// init: record threshold, init locks, init values
//       of all local counts and global count
void init(counter_t *c, int threshold) {
    c->threshold = threshold;

    c->global = 0;
    pthread_mutex_init(&c->glock, NULL);

    int i;
    for (i = 0; i < NUMCPUS; i++) {
        c->local[i] = 0;
        pthread_mutex_init(&c->llock[i], NULL);
    }
}
```

Sloppy Counter Implementation

```
// update: usually, just grab local lock and update local amount
//           once local count has risen by 'threshold', grab global
//           lock and transfer local values to it
void update(counter_t *c, int threadID, int amt) {
    pthread_mutex_lock(&c->llock[threadID]);
    c->local[threadID] += amt;    // assumes amt > 0
    if (c->local[threadID] ≥ c->threshold) { // transfer to global
        pthread_mutex_lock(&c->glock);
        c->global += c->local[threadID];
        pthread_mutex_unlock(&c->glock);
        c->local[threadID] = 0;
    }
    pthread_mutex_unlock(&c->llock[threadID]);
}

// get: just return global amount (which may not be perfect)
int get(counter_t *c) {
    pthread_mutex_lock(&c->glock);
    int val = c->global;
    pthread_mutex_unlock(&c->glock);
    return val;    // only approximate!
}
```

29. Lock Based Data Structures

1. Counter
2. **Linked List**
3. Queue
4. Hash Table



Concurrent Linked Lists

```
// basic node structure
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;

// basic list structure (one used per list)
typedef struct __list_t {
    node_t *head;
    pthread_mutex_t lock;
} list_t;

void List_Init(list_t *L) {
    L->head = NULL;
    pthread_mutex_init(&L->lock, NULL);
}
```

Concurrent Linked Lists

```
int List_Insert(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        pthread_mutex_unlock(&L->lock);
        return -1; // fail
    }
    new->key = key;
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
    return 0; // success
}
```

```
int List_Lookup(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *curr = L->head;
    while (curr) {
        if (curr->key == key) {
            pthread_mutex_unlock(&L->lock);
            return 0; // success
        }
        curr = curr->next;
    }
    pthread_mutex_unlock(&L->lock);
    return -1; // failure
}
```

Discussion

- The code **acquires** a lock in the insert routine upon entry.
- The code **releases** the lock upon exit.
 - If `malloc()` happens to fail, the code must also **release the lock** before failing the insert.
 - This kind of exceptional control flow has been shown to be **quite error prone**.
- **Solution:** The lock and release *only surround* the actual critical section in the insert code

Concurrent Linked Lists rewritten

```
void List_Insert(list_t *L, int key) {  
    // synchronization not needed  
    node_t *new = malloc(sizeof(node_t));  
    if (new == NULL) {  
        perror("malloc");  
        return;  
    }  
    new->key = key;  
  
    // just lock critical section  
    pthread_mutex_lock(&L->lock);  
    new->next = L->head;  
    L->head = new;  
    pthread_mutex_unlock(&L->lock);  
}
```

```
int List_Lookup(list_t *L, int key) {  
    int rv = -1;  
    pthread_mutex_lock(&L->lock);  
    node_t *curr = L->head;  
    while (curr) {  
        if (curr->key == key) {  
            rv = 0;  
            break;  
        }  
        curr = curr->next;  
    }  
    pthread_mutex_unlock(&L->lock);  
    return rv; // now both success and failure  
}
```


Scaling Linked List

- How does the single lock linked list perform?
- Alternative: Hand-over-hand locking (lock coupling)
 - **Add a lock per node** of the list instead of having a single lock for the entire list.
 - When traversing the list,
 - First grabs the next node's lock.
 - And then releases the current node's lock.
 - Enable a high degree of concurrency in list operations.
 - However, in practice, **the overheads** of acquiring and releasing locks for each node of a list traversal is **prohibitive**.

29. Lock Based Data Structures

1. Counter
2. Linked List
- 3. Queue**
4. Hash Table



Concurrent Queue

- Simple Concurrent Queue
 - Add one lock
 - Scaling?
- Alternative:
 - Michael and Scott Concurrent Queues

Michael and Scott Concurrent Queues

- There are two locks.
 - One for the **head** of the queue.
 - One for the **tail**.
 - The goal of these two locks is to enable concurrency of *enqueue* and *dequeue* operations.
- Add a dummy node
 - Allocated in the queue initialization code
 - Enable the separation of head and tail operations
- Run the code in the book (OSTEP) !

29. Lock Based Data Structures

1. Counter
2. Linked List
3. Queue
- 4. Hash Table**




Concurrent Hash Table

- Focus on a simple hash table
 - The hash table does not resize.
 - Built using the concurrent lists
 - It uses a **lock per hash bucket** each of which is represented by a *list*.

Performance of Concurrent Hash Table

- From 10,000 to 50,000 concurrent updates from each of four threads.
- iMac with four Intel 2.7GHz i5 CPUs.



The simple concurrent hash table scales magnificently.

Concurrent Hash Table

```
#define BUCKETS (101)

typedef struct __hash_t {
    list_t lists[BUCKETS];
} hash_t;

void Hash_Init(hash_t *H) {
    int i;
    for (i = 0; i < BUCKETS; i++) {
        List_Init(&H->lists[i]);
    }
}

int Hash_Insert(hash_t *H, int key) {
    int bucket = key % BUCKETS;
    return List_Insert(&H->lists[bucket], key);
}


int Hash_Lookup(hash_t *H, int key) {
    int bucket = key % BUCKETS;
    return List_Lookup(&H->lists[bucket], key);
}
```

Learned important lessons

- **be careful** with acquisition and release of locks around control flow changes
- **enabling more** concurrency does not **necessarily increase** performance
- performance problems should only be remedied once they **exist**
- avoid premature optimization
 - there is no value in making **something faster** if doing so **will not improve** the overall **performance** of the **application**

“Premature optimization is the
root of all evil.”

— Donald Knuth [2]



Thanks

Questions?

