

Lecture

Operating System

31. Concurrency: Semaphore

Concurrency Objectives

- **Mutual exclusion** (e.g., A and B don't run at same time)
 - solved with locks
- **Ordering** (e.g., B runs after A does something)
 - solved with condition variables and **semaphores**

31. Semaphore

- 1. Protect Critical Sections**
- 2. Use as Condition Variable**
- 3. Solve Bounded Buffer Problem**
- 4. Reader-Writer Locks**



31. Semaphore

- 1. Protect Critical Sections**
2. Use as Condition Variable
3. Solve Bounded Buffer Problem
4. Reader-Writer Locks



Semaphore: A definition

- An object **with an integer value**
 - We can manipulate with two routines;
 - `sem_wait()` and `sem_post()`.
 - Initialization
 - Declare a semaphore `s` and initialize it to the value `1`
 - The second argument, `0`, indicates that the semaphore is shared between threads in the same process.

```
#include <semaphore.h>
sem_t s;
sem_init(&s, 0, 1); // initialize s to the value 1
```

Interact with semaphore

- `sem_wait()`
 - If the value of the semaphore was **one** or **higher** when called `sem_wait()`, **return right away**.
 - It will cause the caller to **suspend execution** waiting for a subsequent post.
 - When negative, the value of the semaphore is equal to the number of waiting threads.

```
int sem_wait(sem_t *s) {  
    decrement the value of semaphore s by one  
    wait if value of semaphore s is negative  
}
```

Interact with semaphore (Cont.)

- `sem_post()`
 - Simply **increments** the value of the semaphore.
 - If there is a thread waiting to be woken, **wakes** one of them up.

```
int sem_post(sem_t *s) {  
    increment the value of semaphore s by one  
    if there are one or more threads waiting, wake one  
}
```


Possible Implementation

```
struct semaphore {  
    int value;  
    queueType list;  
}  
  
sem_wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this task  
        to S->list;  
        block();  
    }  
}  
  
sem_post(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a task P  
        from S->list;  
        wakeup(P);  
    }  
}
```


Binary Semaphores (Locks)

- What should **X** be?
 - The initial value should be 1.


```
sem_t m;  
sem_init(&m, 0, X); // initialize semaphore to  
                    // X; what should X be?  
  
sem_wait(&m);  
//critical section here  
sem_post(&m);
```

Thread Trace: Single Thread Using A Semaphore

Value of Semaphore	Thread 0
1	
1	call sema_wait()
0	sem_wait() returns
0	(crit sect)
0	call sem_post()
1	sem_post() returns

Critical Section (buffer) unprotected

- Two periodic threads: T(0) and T(1)
 - Priority scheduling: T(1) has higher Priority than T(0)
- T(0) and T(1) share a buffer during runtime each instance



Content of buffer in undefined state



Two Threads Using A Semaphore

Thread Trace

Value	Thread 0	State	Thread 1	State
1		Running		Ready
1	call sem_wait()	Running		Ready
0	sem_wait() retruns	Running		Ready
0	(crit set: begin)	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	call sem_wait()	Running
-1		Ready	decrement sem	Running
-1		Ready	(sem < 0) → sleep	sleeping
-1		Running	<i>Switch → T0</i>	sleeping
-1	(crit sect: end)	Running		sleeping
-1	call sem_post()	Running		sleeping
0	increment sem	Running		sleeping
0	wake(T1)	Running		Ready
0	sem_post() returns	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	sem_wait() retruns	Running
0		Ready	(crit sect)	Running
0		Ready	call sem_post()	Running
1		Ready	sem_post() returns	Running

31. Semaphore

1. Protect Critical Sections
- 2. Use as Condition Variable**
3. Solve Bounded Buffer Problem
4. Reader-Writer Locks



Semaphores As Condition Variables

```
sem_t s;

void *
child(void *arg) {
    printf("child\n");
    sem_post(&s); // signal here:
                  // child is done

    return NULL;
}

int
main(int argc, char *argv[]) {
    sem_init(&s, 0, X); // what should X be?
    printf("parent: begin\n");
    pthread_t c;
    pthread_create(&c, NULL, child, NULL);
    sem_wait(&s); // wait here for child
    printf("parent: end\n");
    return 0;
}
```

- What should **X** be?
- The initial value should be 0.

What we would like to see here is:

```
parent: begin
child
parent: end
```


Parent Waiting For Child (Case 1)

Thread Trace

The parent call `sem_wait()` before the child has called `sem_post()`.

Val	Parent	State	Child	State
0	Create(Child)	Running	<i>(Child exists; is runnable)</i>	Ready
0	call <code>sem_wait()</code>	Running		Ready
-1	decrement sem	Running		Ready
-1	$(sem < 0) \rightarrow \text{sleep}$	sleeping		Ready
-1	<i>Switch</i> →Child	sleeping	child runs	Running
-1		sleeping	call <code>sem_post()</code>	Running
0		sleeping	increment sem	Running
0		Ready	wake(Parent)	Running
0		Ready	<code>sem_post()</code> returns	Running
0		Ready	<i>Interrupt; Switch</i> →Parent	Ready
0	<code>sem_wait()</code> returns	Running		Ready

Parent Waiting For Child (Case 2)

Thread Trace

The child runs to completion before the parent call `sem_wait()`.

Val	Parent	State	Child	State
0	Create(Child)	Running	(Child exists; is runnable)	Ready
0	<i>Interrupt; switch→Child</i>	Ready	child runs	Running
0		Ready	call <code>sem_post()</code>	Running
1		Ready	increment sem	Running
1		Ready	wake(nobody)	Running
1		Ready	<code>sem_post()</code> returns	Running
1	parent runs	Running	<i>Interrupt; Switch→Parent</i>	Ready
1	call <code>sem_wait()</code>	Running		Ready
0	decrement sem	Running		Ready
0	$(sem < 0) \rightarrow \text{awake}$	Running		Ready
0	<code>sem_wait()</code> returns	Running		Ready

31. Semaphore

1. Protect Critical Sections
2. Use as Condition Variable
- 3. Solve Bounded Buffer Problem**
4. Reader-Writer Locks



The Producer/Consumer (Bounded-Buffer) Problem

■ **Producer:** put() interface

- Wait for a buffer to become **empty** in order to put data into it.

■ **Consumer:** get() interface

- Wait for a buffer to become **filled** before using it.

```
int buffer[MAX];
int fill = 0;
int use = 0;

void put(int value) {
    buffer[fill] = value;    // line f1
    fill = (fill + 1) % MAX; // line f2
}

int get() {
    int tmp = buffer[use];   // line g1
    use = (use + 1) % MAX;   // line g2
    return tmp;
}
```

A Solution: Adding empty/full (Bounded-Buffer) Problem

```
sem_t empty;
sem_t full;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&empty);    // line P1
        put(i);              // line P2
        sem_post(&full);     // line P3
    }
}

void *consumer(void *arg) {
    int i, tmp = 0;
    while (tmp != -1) {
        sem_wait(&full);     // line C1
        tmp = get();         // line C2
        sem_post(&empty);    // line C3
        printf("%d\n", tmp);
    }
}
```

Result of First Attempt?

- Imagine that MAX is greater than 1 .
 - If there are **multiple** producers, **race condition** can happen at line **f1**.
 - It means that the old data there is overwritten.
- We've forgotten here is **mutual exclusion**.
 - The filling of a buffer and incrementing of the index into the buffer is a **critical section**.

```
int buffer[MAX];
int fill = 0;
int use = 0;

void put(int value) {
    buffer[fill] = value;    // line f1
    fill = (fill + 1) % MAX; // line f2
}
```


A Solution: Adding Mutual Exclusion (Bounded-Buffer) Problem

```
sem_t empty;
sem_t full;
sem_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex);    // line p0 (NEW LINE)
        sem_wait(&empty);    // line p1
        put(i);              // line p2
        sem_post(&full);     // line p3
        sem_post(&mutex);    // line p4 (NEW LINE)
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex);    // line c0 (NEW LINE)
        sem_wait(&full);     // line c1
        int tmp = get();     // line c2
        sem_post(&empty);    // line c3
        sem_post(&mutex);    // line c4 (NEW LINE)
        printf("%d\n", tmp);
    }
}
```



Imagine two thread:
one producer and
one consumer....

No Solution: It's a deadlock

- Imagine two thread: one producer and one consumer.
 - The consumer **acquire** the mutex (line `c0`).
 - The consumer **calls** `sem_wait()` on the full semaphore (line `c1`).
 - The consumer is **blocked** and **yield** the CPU.
 - The consumer still holds the mutex!
 - The producer **calls** `sem_wait()` on the binary mutex semaphore (line `p0`).
 - The producer is now **stuck** waiting too.
 - **a classic deadlock!**

A Solution: Adding Mutual Exclusion (Bounded-Buffer) Problem

Correctly!

```
sem_t empty;
sem_t full;
sem_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&empty); // line p1
        sem_wait(&mutex); (MOVED MUTEX HERE...)
        put(i);           // line p2
        sem_post(&mutex); (MOVED MUTEX HERE...)
        sem_post(&full);  // line p3
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&full); // line c1
        sem_wait(&mutex); (MOVED MUTEX HERE...)
        int tmp = get(); // line c2
        sem_post(&mutex); (MOVED MUTEX HERE...)
        sem_post(&empty); // line c3
    }
}
```

31. Semaphore

1. Protect Critical Sections
2. Use as Condition Variable
3. Solve Bounded Buffer Problem
- 4. Reader-Writer Locks**



Reader-Writer Locks

- Imagine a number of concurrent list operations, including **inserts** and simple **lookups**.
 - **insert:**
 - Change the state of the list
 - A traditional **critical section** makes sense.
 - **lookup:**
 - Simply read the data structure.
 - As long as we can guarantee that no insert is on-going, we can allow many lookups to proceed **concurrently**.
- This special type of lock is known as a **reader-writer lock**.

Reader-Writer Locks

- Only **a single writer** can acquire the lock.
- Once a reader has acquired **a read lock**,
 - **More readers** will be allowed to acquire the read lock too.
 - A writer will have to wait until all readers are finished.

```
typedef struct _rwlock_t {
    sem_t lock;           // binary semaphore (basic lock)
    sem_t writelock;      // used to allow ONE writer or MANY readers
    int readers;          // count of readers reading in critical section
} rwlock_t;

void rwlock_init(rwlock_t *rw) {
    rw->readers = 0;
    sem_init(&rw->lock, 0, 1);
    sem_init(&rw->writelock, 0, 1);
}
```

Acquire/Release Reader-Writer-Lock

```
void rwlock_acquire_readlock(rwlock_t *rw) {
    sem_wait(&rw->lock);
    rw->readers++;
    if (rw->readers == 1)
        sem_wait(&rw->writelock); // first reader acquires writelock
    sem_post(&rw->lock);
}
```

```
void rwlock_release_readlock(rwlock_t *rw) {
    sem_wait(&rw->lock);
    rw->readers--;
    if (rw->readers == 0)
        sem_post(&rw->writelock); // last reader releases writelock
    sem_post(&rw->lock);
}
```

```
void rwlock_acquire_writelock(rwlock_t *rw) {
    sem_wait(&rw->writelock);
}
```

```
void rwlock_release_writelock(rwlock_t *rw) {
    sem_post(&rw->writelock);
}
```

Summary Reader-Writer-Lock

- The reader-writer locks have fairness problem.
 - It would be relatively easy for reader to starve writer.
 - How to prevent more readers from entering the lock once a writer is waiting?

Thanks

Questions?

