

Betriebssysteme und Systemnahe Programmierung

Kapitel 6 • MINIX Internals

Winter 2016/17

Marcel Waldvogel

The Internal Structure of MINIX

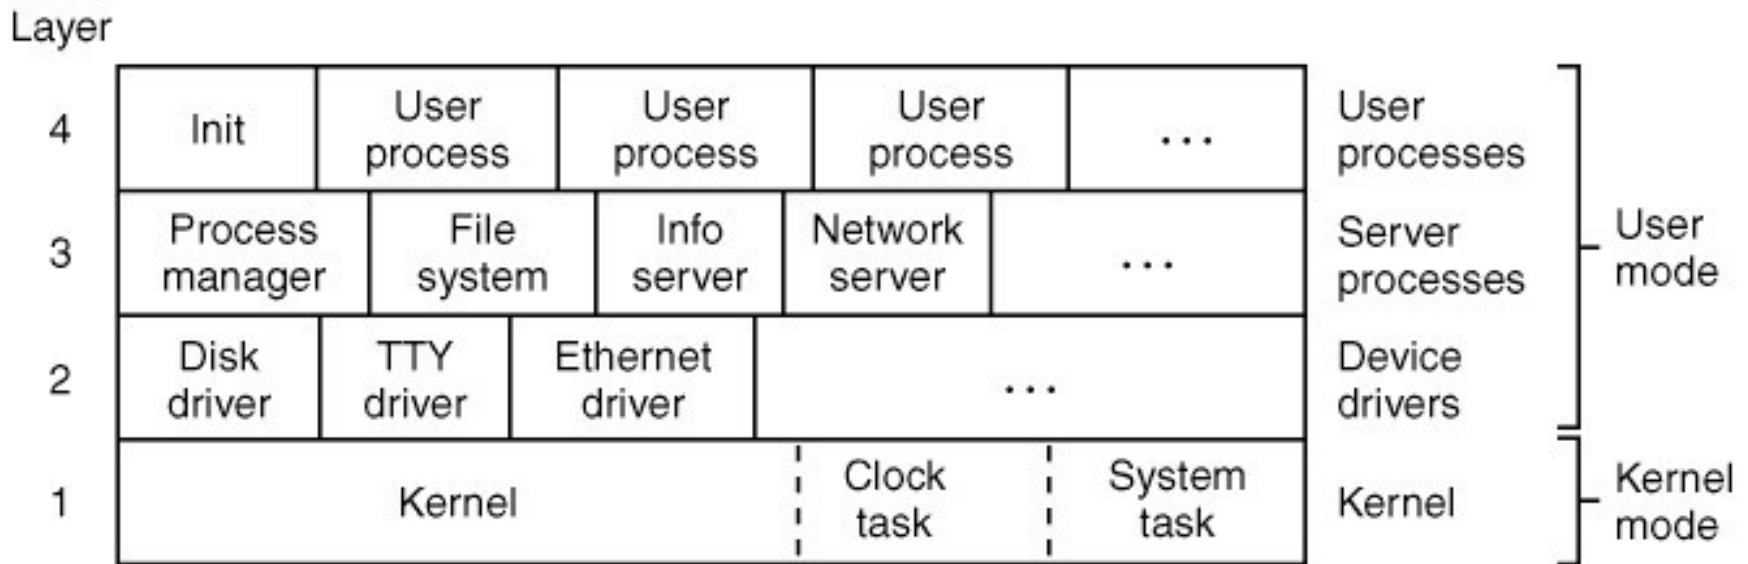


Figure 2-29. MINIX 3 is structured in four layers.
Only processes in the bottom layer may use privileged (kernel mode) instructions.

MINIX 3 Startup

| Component | Description | Loaded by |
|-----------|--|-----------------|
| kernel | Kernel + clock and system tasks | (in boot image) |
| pm | Process manager | (in boot image) |
| fs | File system | (in boot image) |
| rs | (Re)starts servers and drivers | (in boot image) |
| memory | RAM disk driver | (in boot image) |
| log | Buffers log output | (in boot image) |
| tty | Console and keyboard driver | (in boot image) |
| driver | Disk (at, bios, or floppy) driver | (in boot image) |
| init | parent of all user processes | (in boot image) |
| floppy | Floppy driver (if booted from hard disk) | /etc/rc |
| is | Information server (for debug dumps) | /etc/rc |
| cmos | Reads CMOS clock to set time | /etc/rc |
| random | Random number generator | /etc/rc |
| printer | Printer driver | /etc/rc |

Figure 2-30. Some important MINIX 3 system components. Others such as an Ethernet driver and the inet server may also be present.

MINIX Memory (1)

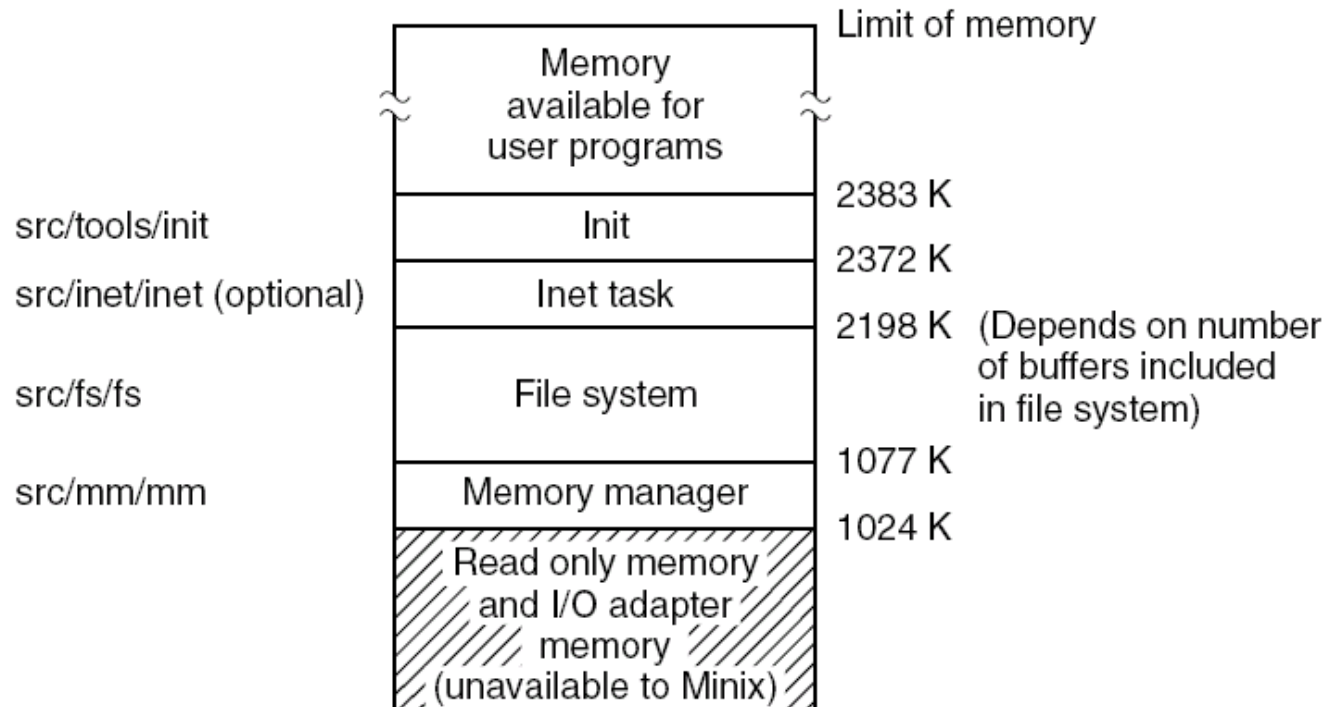


Figure 2-31. Memory layout after MINIX 3 has been loaded from the disk into memory. The kernel, servers, and drivers are independently compiled and linked programs, listed on the left. Sizes are approximate and not to scale.

MINIX Memory (2)




Figure 2-31. Memory layout after MINIX 3 has been loaded from the disk into memory. The kernel, servers, and drivers are independently compiled and linked programs, listed on the left. Sizes are approximate and not to scale.

MINIX Header File

```
#include <minix/config.h>           /* MUST be first */
#include <ansi.h>                   /* MUST be second */
#include <sys/types.h>
#include <minix/const.h>
#include <minix/type.h>
#include <limits.h>
#include <errno.h>
#include <minix/syslib.h>
#include "const.h"
```

Figure 2-32. Part of a master header which ensures inclusion of header files needed by all C source files. Note that two *const.h* files, one from the *include/* tree and one from the local directory, are referenced.

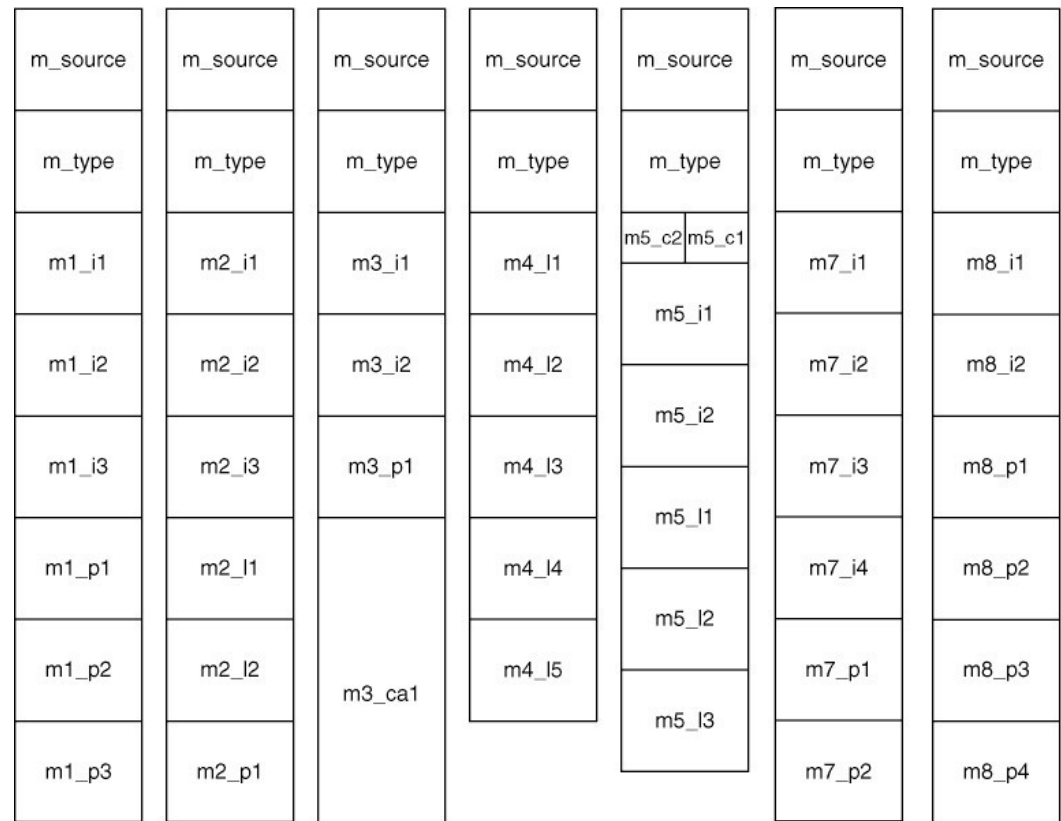
Sizes of Types in MINIX

| Type | 16-Bit MINIX | 32-Bit MINIX |
|-------|--------------|--------------|
| gid_t | 8 | 8 |
| dev_t | 16 | 16 |
| pid_t | 16 | 32 |
| ino_t | 16 | 32 |

Figure 2-33. The size, in bits, of some types on 16-bit and 32-bit systems.

MINIX Message Types

Figure 2-34. The seven message types used in MINIX 3. The sizes of message elements will vary, depending upon the architecture of the machine; this diagram illustrates sizes on CPUs with 32-bit pointers, such as those of Pentium family members.



Debug Dump

| --nr- | -id- | -name- | -flags- | -traps- | -ipc_to mask----- |
|-------|------|--------|---------|---------|-------------------|
| (-4) | (01) | IDLE | P-BS- | ----- | 00000000 00001111 |
| [-3] | (02) | CLOCK | ---S- | --R-- | 00000000 00001111 |
| [-2] | (03) | SYSTEM | ---S- | --R-- | 00000000 00001111 |
| [-1] | (04) | KERNEL | ---S- | ----- | 00000000 00001111 |
| 0 | (05) | pm | P--S- | ESRBN | 11111111 11111111 |
| 1 | (06) | fs | P--S- | ESRBN | 11111111 11111111 |
| 2 | (07) | rs | P--S- | ESRBN | 11111111 11111111 |
| 3 | (09) | memory | P--S- | ESRBN | 00110111 01101111 |
| 4 | (10) | log | P--S- | ESRBN | 11111111 11111111 |
| 5 | (08) | tty | P--S- | ESRBN | 11111111 11111111 |
| 6 | (11) | driver | P--S- | ESRBN | 11111111 11111111 |
| 7 | (00) | init | P-B-- | E--B- | 00000111 00000000 |

Figure 2-35. Part of a debug dump of the privilege table. The clock task, file server, tty, and init processes privileges are typical of tasks, servers, device drivers, and user processes, respectively. The bitmap is truncated to 16 bits.

Bootstrapping MINIX (1)





Figure 2-36. Disk structures used for bootstrapping.
(a) Unpartitioned disk. The first sector is the bootblock.

Bootstrapping MINIX (2)



(b)

Figure 2-36. Disk structures used for bootstrapping.

(b) Partitioned disk. The first sector is the master boot record, also called masterboot.

Boot Time in MINIX

```
rootdev=256  
ramimagedev=916  
ramsize=4096  
processor=586  
bus=at  
video=vga  
chrome=color  
memory=800:92880,100000:2F00000  
c0=at  
image=/minix/2.0.3r5
```

Figure 2-37. Boot parameters passed to the kernel at boot time in a typical MINIX 3 system.

System Initialization in MINIX

```
#include <minix/config.h>
#if _WORD_SIZE == 2
#include "mpx88.s"
#else
#include "mpx386.s"
#endif
```

Figure 2-38. How alternative assembly language source files are selected.

Interrupt Handling in MINIX (1)

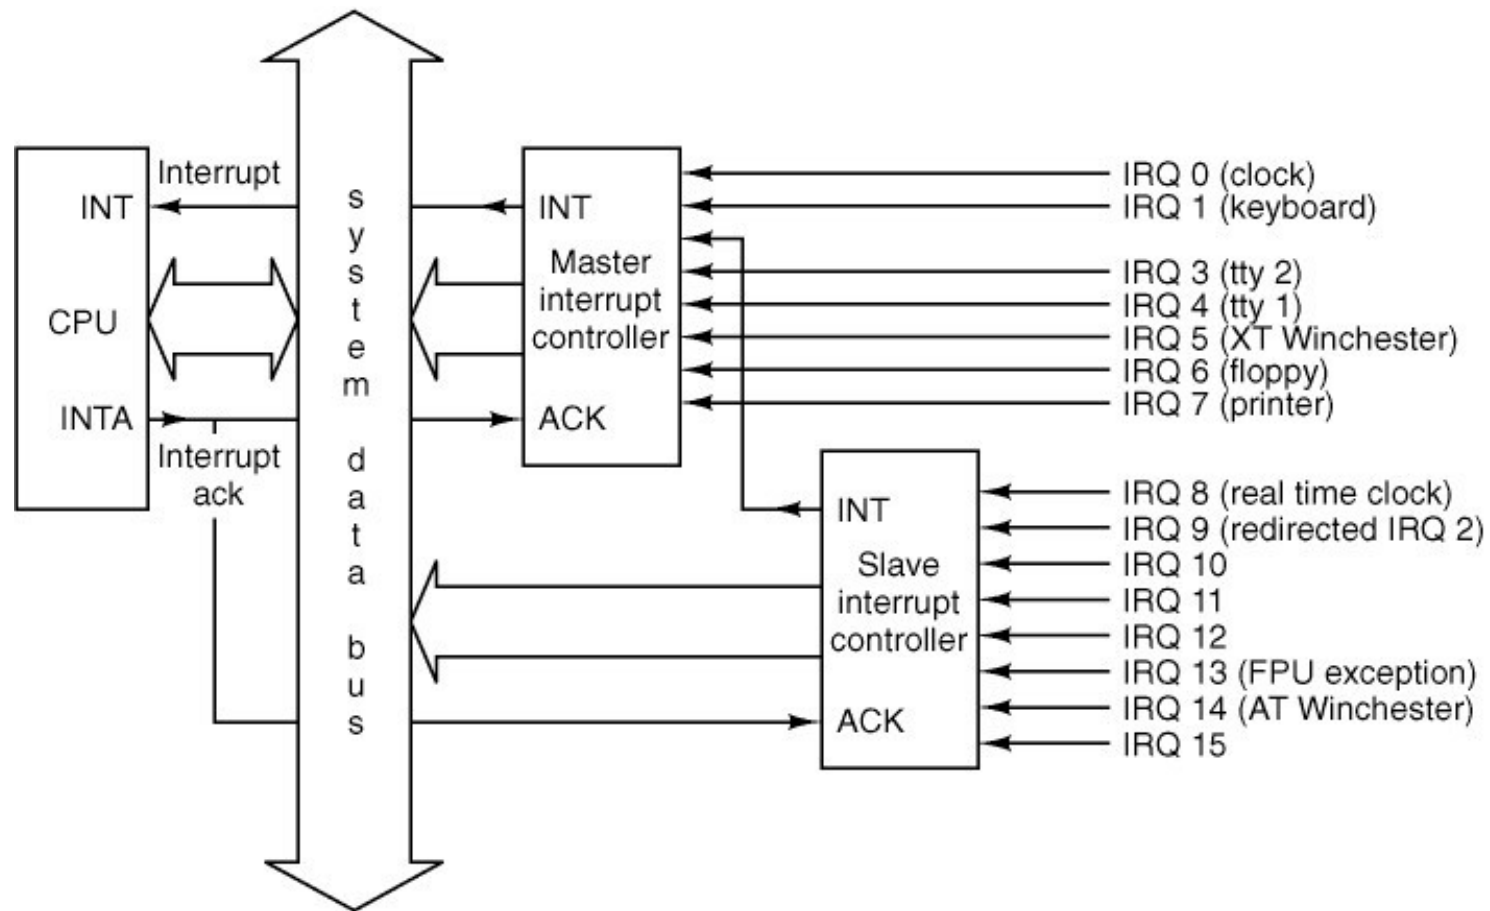
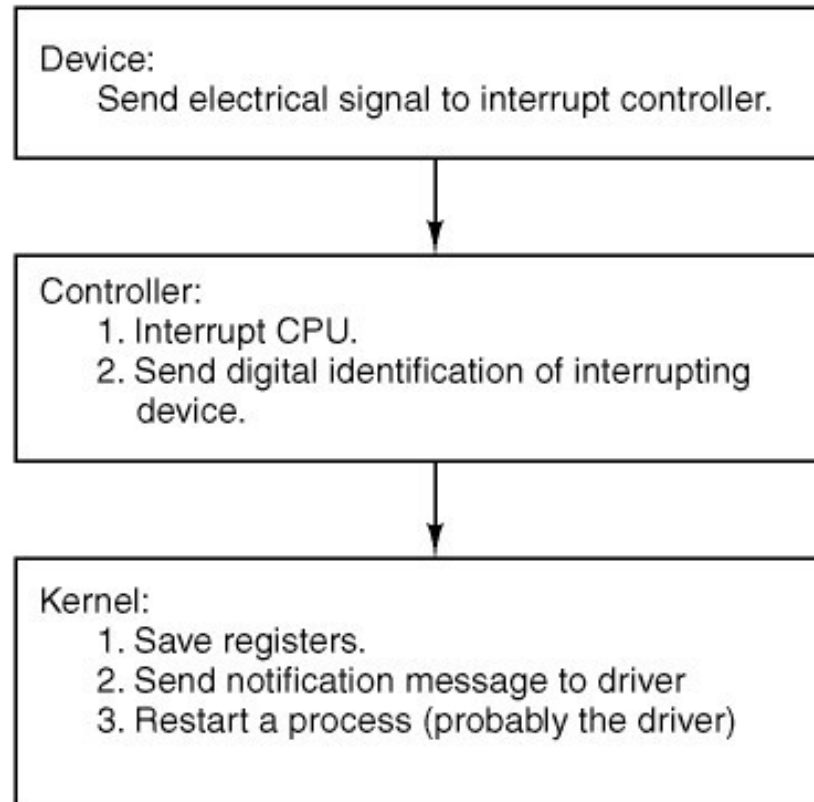


Figure 2-39. Interrupt processing hardware on a 32-bit Intel PC.

Interrupt Handling in MINIX (2)



(a)

Figure 2-40. (a) How a hardware interrupt is processed.

Interrupt Handling in MINIX (3)

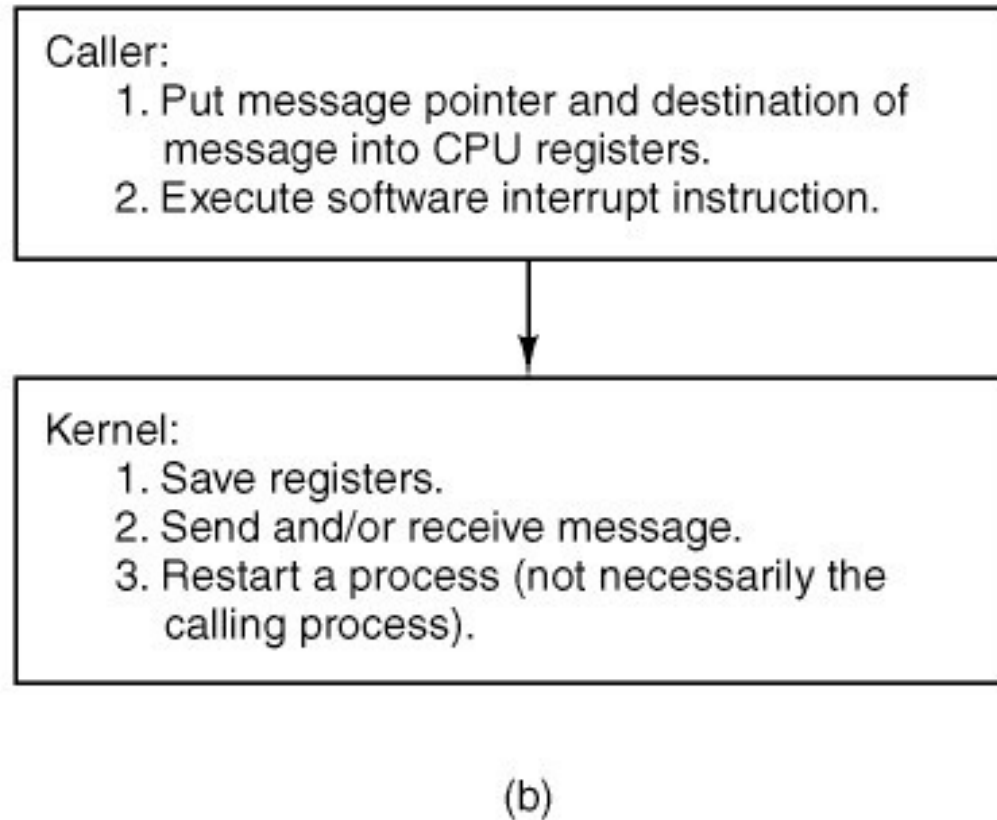


Figure 2-40. (b) How a system call is made.

Restart

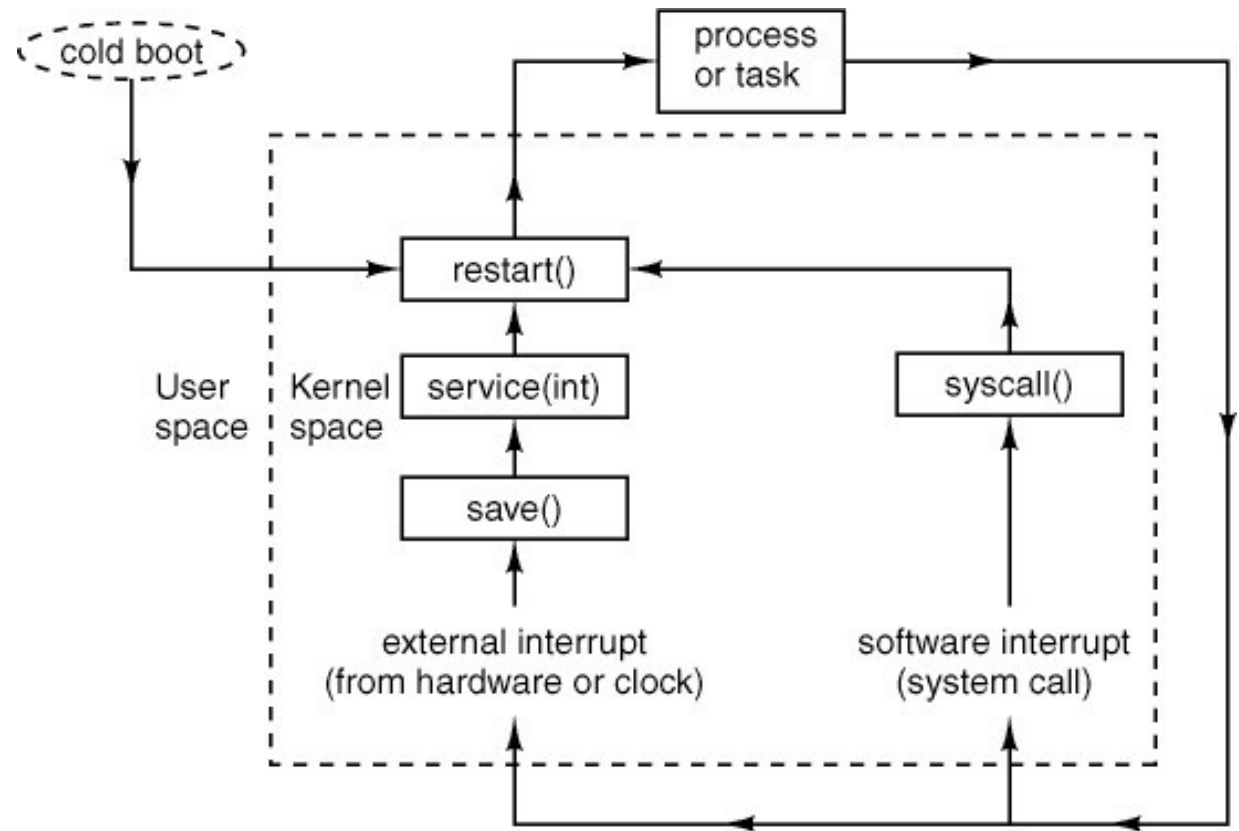




Figure 2-41. Restart is the common point reached after system startup, interrupts, or system calls. The most deserving process (which may be and often is a different process from the last one interrupted) runs next. Not shown in this diagram are interrupts that occur while the kernel itself is running.

Queueing



(a)



(b)

Figure 2-42. Queueing of processes trying to send to process 0.

Scheduling in MINIX

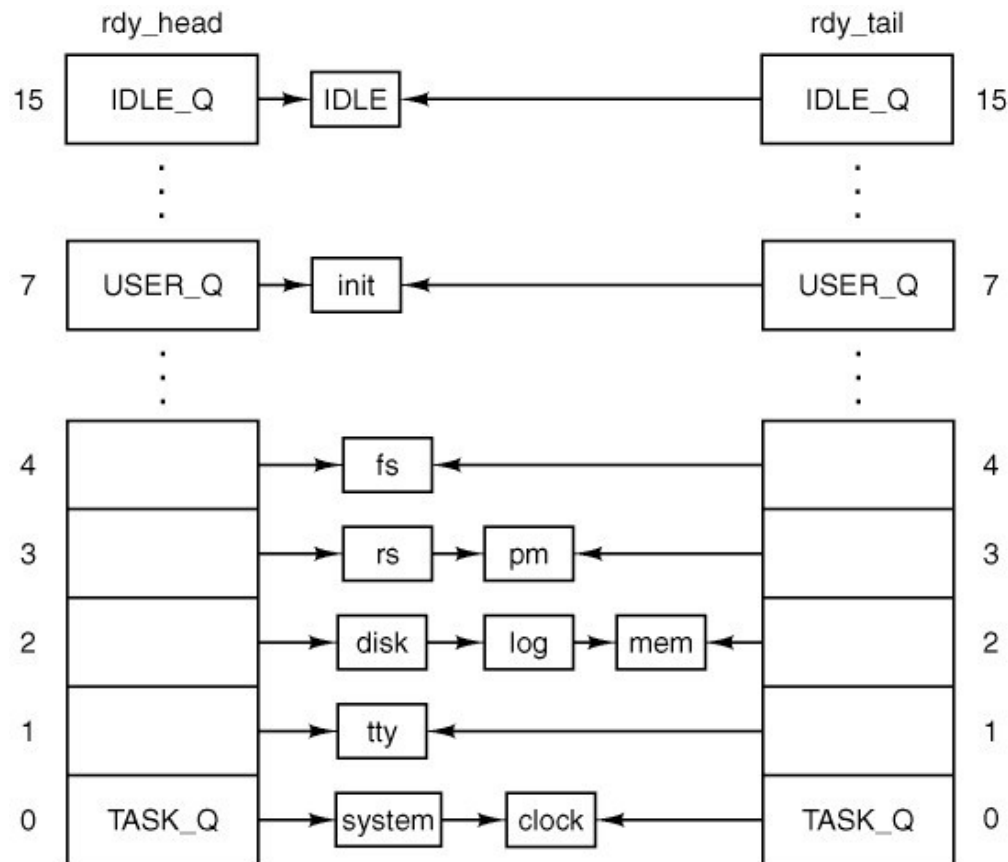


Figure 2-43. The scheduler maintains sixteen queues, one per priority level. Shown here is the initial queuing process as MINIX 3 starts up.

Hardware-Dependent Kernel Support

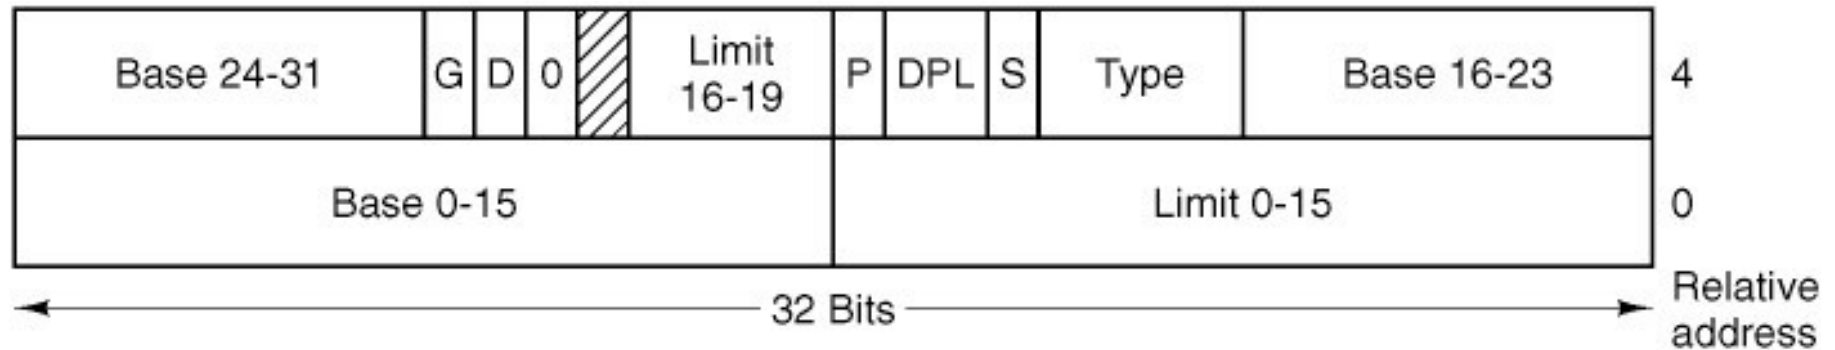


Figure 2-44. The format of an Intel segment descriptor.

Overview of System Task (1)

| Message type | From | Meaning |
|---------------|------------|---|
| sys_fork | PM | A process has forked |
| sys_exec | PM | Set stack pointer after EXEC call |
| sys_exit | PM | A process has exited |
| sys_nice | PM | Set scheduling priority |
| sys_privctl | RS | Set or change privileges |
| sys_trace | PM | Carry out an operation of the PTRACE call |
| sys_kill | PM,FS, TTY | Send signal to a process after KILL call |
| sys_getksig | PM | PM is checking for pending signals |
| sys_endksig | PM | PM has finished processing signal |
| sys_sigsend | PM | Send a signal to a process |
| sys_sigreturn | PM | Cleanup after completion of a signal |
| sys_irqctl | Drivers | Enable, disable, or configure interrupt |

Figure 2-45. The message types accepted by the system task.

“Any” means any system process; user processes cannot call the system task directly

Overview of System Task (2)

| Message type | From | Meaning |
|---------------|-----------------|--|
| sys_devio | Drivers | Read from or write to an I/O port |
| sys_sdevio | Drivers | Read or write string from/to I/O port |
| sys_vdevio | Drivers | Carry out a vector of I/O requests |
| sys_int86 | Drivers | Do a real-mode BIOS call |
| sys_newmap | PM | Set up a process memory map |
| sys_segctl | Drivers | Add segment and get selector (far data access) |
| sys_memset | PM | Write char to memory area |
| sys_umap | Drivers | Convert virtual address to physical address |
| sys_vircopy | FS, Drivers | Copy using pure virtual addressing |
| sys_physcopy | Drivers | Copy using physical addressing |
| sys_virvcopy | Any | Vector of VCOPY requests |
| sys_physvcopy | Any | Vector of PHYSCOPY requests |
| sys_times | PM | Get uptime and process times |
| sys_setalarm | PM, FS, Drivers | Schedule a synchronous alarm |
| sys_abort | PM, TTY | Panic: MINIX is unable to continue |
| sys_getinfo | Any | Request system information |

Figure 2-45. The message types accepted by the system task.

“Any” means any system process; user processes
cannot call the system task directly

Clock Hardware




Figure 4-47. A programmable clock.

Clock Software (1)

Typical duties of a clock driver.

1. Maintain time of day
2. Prevent processes from running longer than allowed
3. Accounting for CPU usage
4. Handling alarm system call by user processes
5. Providing watchdog timers for parts of system itself
6. Doing profiling, monitoring, and statistics gathering

Clock Software (2)

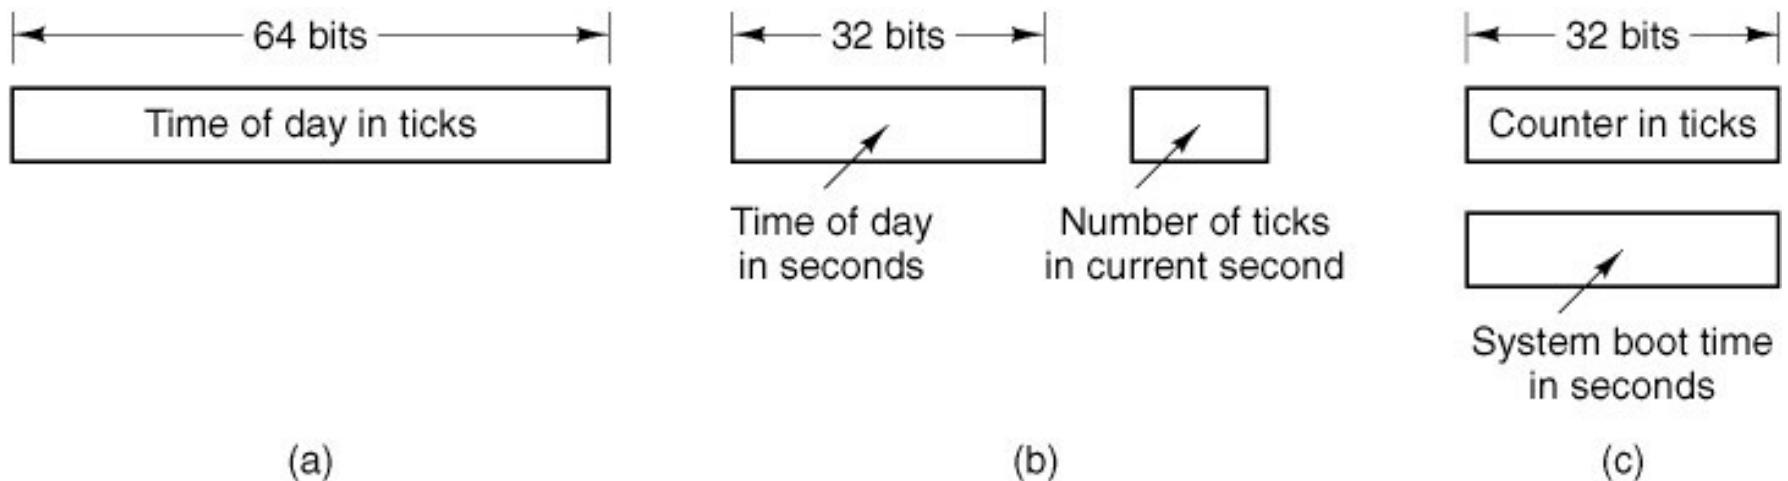


Figure 2-48. Three ways to maintain the time of day.

Clock Software (3)

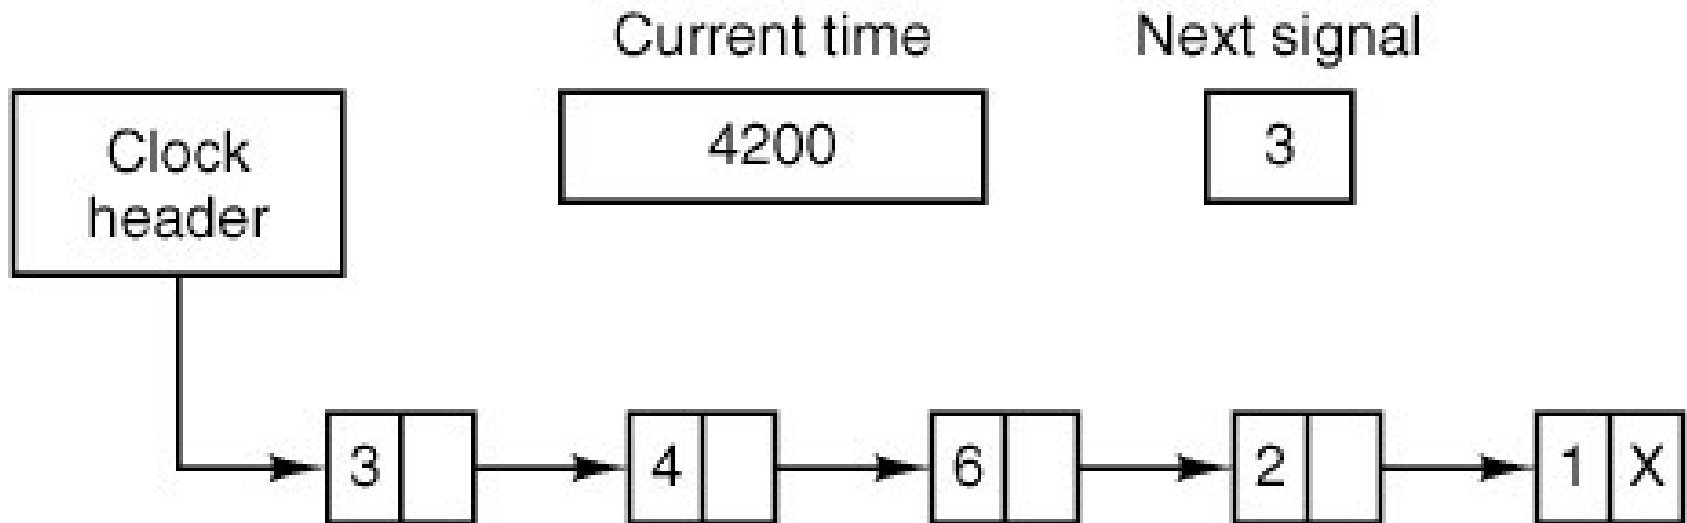


Figure 2-49. Simulating multiple timers with a single clock.

Summary of Clock Services

| Service | Access | Response | Clients |
|-------------------|---------------|--------------|-----------------------------------|
| get_uptime | Function call | Ticks | Kernel or system task |
| set_timer | Function call | None | Kernel or system task |
| reset_timer | Function call | None | Kernel or system task |
| read_clock | Function call | Count | Kernel or system task |
| clock_stop | Function call | None | Kernel or system task |
| Synchronous alarm | System call | Notification | Server or driver, via system task |
| POSIX alarm | System call | Signal | User process, via PM |
| Time | System call | Message | Any process, via PM |

Figure 2-50. The time-related services supported by the clock driver.