

10.3 Run a program

```
1 #include <unistd.h>
2 int execve(const char *file, char *const argv[], char *const env[]);
```

- ▶ The `execve(2)` function **executes** the program stored in `file` (no `$PATH` search), passing arguments `argv` and environment `env`.
 - **Note** that arrays `argv` and `env` must be `NULL`-terminated!
 - By convention `argv[0]` shall be the **filename** of the program.
 - ▶ On success, `execve` **does not return**.
 - ▶ The **calling program is replaced** by the program loaded from `file`, replacing code, data, and stack of the caller.
 - ▶ Most process attributes are preserved, most notably:
 - The **process ID** stays the same.
 - By default, **file descriptors** remain open.
- See `execve(2)` for a list of exceptions!

The `exec` family

- ▶ Setting up `argv[]` and `env[]` sometimes is cumbersome. The `exec(3)` family of functions facilitates this task:

Example `execlp`, One of the seven `exec` functions:

```
1 int execlp(const char *file, const char *arg0, ..., NULL);
```

- ▶ Consumes a variable list of **arguments** passed as `argv` to `main`. The **last one** must be `NULL`!
- ▶ If there's no slash `/` in the string `file`, then a `$PATH` search is performed, just like the shell would do.
Warning: Insecure defaults if `$PATH` is unset!
- ▶ `execl` (without `p`) is similar, but does not perform a `$PATH` search.

```
1 int main(void)
2 {
3     int status;
4     pid_t pid = fork();
5
6     if (pid == 0) {
7         execl("/usr/bin/sleep", "sleep", "100", NULL);
8         err(1, "execl");
9     }
10
11     if (pid < 0) err(1, "fork");
12     printf("Child %u\n", pid);
13
14     if (wait(&status) < 0) err(1, "wait");
15     if (WIFEXITED(status))
16         printf("Child exit status was: %d\n", WEXITSTATUS(status));
17     else
18         printf("Child terminated without exit status.\n");
19     return 0;
20 }
```

```
1 $ ./a.out
2 # ...
3 -rwx----- 1 marcel users 8624 Jan  7 13:25 a.out
4 -rw----- 1 marcel users  432 Jan  7 13:24 exec.c
5 # ...
6 Child exit status was: 0
```

10.4 Duplicating file descriptors

```
1 #include <unistd.h>
2 int dup(int oldfd);
3 int dup2(int oldfd, int newfd);
```

- ▶ Both functions return -1 on error, or a **new** file descriptor referring to the **same open file description** as `oldfd`.
 - They share offset and other flags.
- ▶ `dup(2)` uses the **lowest unused** number for the new file descriptor.
 - Old and new descriptor can be used **interchangably**.
 - Closing one does **not close** the other.
- ▶ `dup2(2)` uses the **given newfd** for the new file descriptor.
 - If `newfd` is an open file descriptor, it is closed first, ignoring errors!
 - Careful programming: Check `close(newfd)` for errors before using `dup2`.
- ▶ Both functions can be replaced with `fcntl(2)` and additional work.

Note Only **closing the last file descriptor** referring to the underlying open file description frees the associated resources, cf. `close(2)`.

Redirection of *stdout* to a file

```
1  int main(void)
2  {
3      int fd = open("data.txt", O_WRONLY|O_CREAT|O_EXCL|O_TRUNC, 0666);
4      if (fd < 0) err(1, "open");
5      pid_t pid = fork();    if (pid < 0) err(1, "fork");
6
7      if (pid == 0) { /* the child */
8          printf("C: Writing to data.txt\n");
9          close(STDOUT_FILENO);
10         if (dup2(fd, STDOUT_FILENO) < 0) err(1, "dup2");
11         close(fd); /* not needed any more */
12         execlp("date", "date", "+%Y-%m-%d %H:%M:%S", NULL); /* see date(1) */
13         err(1, "execlp");
14     }
15
16     /* the parent */
17     close(fd); /* not needed any more */
18     printf("P: Waiting for child %d\n", pid);
19     int status;
20     wait(&status);
21     if (WIFEXITED(status))
22         printf("P: Child exited with %d\n", WEXITSTATUS(status));
23     return 0;
24 }
```

Running the previous program:

```
1 $ ./a.out
2 Parent: Waiting for child 4333
3 Child: Writing to foo
4 Parent: Child exited with 0
5 $ cat foo
6 2014-01-05 17:18:19
```

11

Inter-Process Communication

aka. IPC

11.1 Pipes & FIFOs

- ▶ Pipes are the **oldest** form of UNIX System IPC.
- ▶ An **overview** discussion of pipes & FIFOs is in `pipe(7)`.

Pipes (detailed discussion on next slides)

- ▶ A pipe has two **file descriptors**, one for **reading** and one for **writing**.
I.e., access is done via `read(2)` and `write(2)`.
- ▶ Pipes can be used only between processes that have a **common ancestor**.

FIFOs Like pipes, but with a **name** in the file system.

- ▶ Created with C function `mkfifo(3)`, or command line tool `mkfifo(1)`.
- ▶ The FS is just a convenient place to put the FIFO's name, no data is stored.
- ▶ Any process may `open(2)` a FIFO through its **file system name** (subject to permissions).
- ▶ Open with `O_RDONLY` to get the **reading end** of the FIFO, and `O_WRONLY` for the **writing end**.

Using pipes

```
1 #include <unistd.h>
2 int pipe(int pipefd[2]);
```

- ▶ `pipe(2)` creates a pipe, **returning** the file descriptors in `pipefd`.
- ▶ `pipefd[0]` is the read end, `pipefd[1]` is the write end.
- ▶ Returns `0` on success. On error `-1` is returned, and `errno` is set.

```
1 #define MAXLINE 100
2
3 int main(void)
4 {
5     int fd[2];
6     if (pipe(fd) < 0)
7         err(1, "pipe");
8
9     pid_t pid = fork();
10    if (pid < 0)
11        err(1, "fork");
```

```
12         if (pid > 0) {                               /* parent */
13             close(fd[0]);
14             write(fd[1], "hello world", 11);
15             close(fd[1]);
16
17         } else {                                       /* child */
18             close(fd[1]);
19             char line[MAXLINE];
20             ssize_t n = read(fd[0], line, MAXLINE-1);
21             line[n] = '\0';
22             printf("From parent: %s\n", line);
23             close(fd[0]);
24         }
25         return 0;
26     }
```

Further remarks

- ▶ Data written to a pipe can be **read only once**.
 - Multiple readers may see only parts of the information.
- ▶ Writing big data chunks may be **non-atomic**:
 - Multiple **reads** may be necessary to get **all data**.
 - Data written by different processes may be **interleaved**.
- ▶ Pipes have **limited capacity**:
 - 64kiB since Linux 2.6.11.
 - **Writing blocks** if the pipe is full, **reading blocks** if the pipe is empty.
 - If **O_NONBLOCK** is set for the according file descriptors, these operations fail instead (cf. **fcntl(2)**).
- ▶ Properties of pipe descriptors:
 - **fstat(2)** returns a file type of FIFO for both ends of a pipe.
 - We can test for a pipe with the **S_ISFIFO** macro, cf. **sys_stat.h(0)**.
- ▶ What happens if the **other end** has been closed?
 - **Reading** returns 0 after all the data has been read.
 - **Writing** generates **signal SIGPIPE**. If we either ignore the signal or catch it and return from the signal handler, **write** returns -1 with **errno** set to **EPIPE**.

11.2 Signals

Definition A signal is a message notifying a process of an event.

- ▶ A signal is identified by an **integer**, and traditional signals carry **no extra data**.
 - Newer implementations allow to pass an additional `int`, and provide more information to the recipient.
- ▶ Linux knows 32 “standard” signals (plus 32 more).
- ▶ An overview of Linux’ signal interface is in `signal(7)`.

How to react to a signal

- ▶ **Ignore** the signal — not recommended in general.
- ▶ The predefined **default action** — termination in most cases.
- ▶ Catch a signal by installing a **handler** function.
- ▶ **Block** the signal — it becomes pending, until unblocked.

Some signals

Name	Description	Default Action
SIGINT	terminal interrupt (e.g., pressed C-c)	terminate
SIGALRM	timer expired	terminate
SIGSEGV	invalid memory reference	terminate
SIGCHLD	change in status of a child	ignore
SIGUSR1	user-defined	terminate
SIGUSR2	user-defined	terminate
SIGTERM	termination	terminate
SIGSTOP	stop	stop process
SIGCONT	continue stopped process	continue/ignore
SIGKILL	termination	terminate

A complete list is in [signal\(7\)](#).

Note The default disposition of [SIGKILL](#) and [SIGSTOP](#) cannot be changed — they are not handled by the receiving program, but by the OS.

Sending signals

```
1 #include <signal.h>
2
3 int kill(pid_t pid, int sig);
4 int raise(int sig);
```

```
1 #include <unistd.h>
2
3 unsigned int alarm(unsigned int n);
```

- ▶ `kill(2)` sends signal `sig` to the process identified by `pid`.
 - `pid` may be zero, or even negative, to send to a group of processes.
- ▶ `raise(3)` sends signal `sig` to the calling process.
- ▶ Both return `0` on success, or `-1` on error, setting `errno`.
- ▶ `alarm(2)` schedules `SIGALRM` to be sent to the caller after `n` seconds.
 - Any previous alarm is **canceled**. If `n` is `0`, no alarm is set.
 - Returns number of **remaining seconds** before delivery of a previously set alarm, or `0` if none was set.

From the shell

`kill [-signal] pid...` Send `signal`, or `SIGTERM` to the processes identified by `pids`.

- ▶ This is a shell builtin.
- ▶ `kill(1)` describes the program `/bin/kill` instead.

Install a signal handler

```
1 #define _POSIX_C_SOURCE 199309L
2 #include <signal.h>
3
4 int sigaction(int sig, const struct sigaction *act,
5               struct sigaction *old);
6
7 struct sigaction {
8     void (*sa_handler)(int);
9     /* incomplete, we will extend this on the following slides */
10 };
```

- ▶ `sigaction(2)` sets/queries the action to be taken for signal `sig`.
- ▶ If not `NULL`, `act` describes the new action.
- ▶ The old setting is returned via `old`, if that is not `NULL`.
- ▶ Struct member `sa_handler` points to the function to be called, or is `SIG_DFL` to restore the default action, or `SIG_IGN` to ignore the signal.
- ▶ Returns `0` on success, or `-1` on error, setting `errno`.

Question What is the type of the handler?

First simple example

```
1 void handler(int s) /* this is called with the signal number as argument */
2 {
3     printf("caught signal %d\n", s);
4     sleep(1);
5 }
6
7 int main(void)
8 {
9     struct sigaction sa;
10
11     memset(&sa, 0, sizeof(sa)); /* clear memory, cf. memset(3) */
12     sa.sa_handler = handler;
13
14     sigaction(SIGINT, &sa, NULL); /* install handler for SIGINT */
15     sigaction(SIGUSR1, &sa, NULL); /* same for SIGUSR1 */
16
17     printf("pid is %d\n", getpid());
18     while (1)
19         pause(); /* sleep until signal arrives, cf. pause(2) */
20
21     return 0;
22 }
```

Running the example

- ▶ Use two terminals, one to run the program, and one to send signals to it.

```
1 $ pk-cc -o signal0 signal0.c
2 $ ./signal0
3 pid is 10226
4 ^Ccaught signal 2      # caused by pressing C-c
5 caught signal 2       # other terminal, line 1
6 caught signal 10      # ... line 2
7 Terminated          # ... line 3
8 $
```

```
1 $ kill -INT 10226
2 $ kill -USR1 10226
3 $ kill -TERM 10226
```

- ▶ Pressing C-c causes the terminal to send **SIGINT** to the process.
- ▶ The program does not handle **SIGTERM**, so the default action is taken.

Notes

- ▶ Signals **SIGKILL**, and **SIGSTOP** **cannot be caught**, blocked, or ignored.
- ▶ Again, there is more than one **interface** to signals. We will not discuss the older **signal(2)**, due to several flaws.