

12.2 Symbol resolution

- ▶ For each **local symbol**, the compiler guarantees exactly one definition. The name is modified to be unique (e.g. `count` above).
- ▶ If *the compiler finds no definition*, it expects it to come from another module, and leaves it to the linker, (e.g. `buf` above).
- ▶ When **the linker** resolves *global* symbols, several conditions can occur:
 - **No definition** is found in the symbol table of any input object file.
 - **Multiple definitions** are found in different object files, and one must be chosen.

Example No `main` function, and `buf` undefined.

```
1 $ pk-cc swap.o #without -c, try to build an executable
2 ../lib/crt1.o: In function '_start':
3 (.text+0x20): undefined reference to 'main'
4 swap.o: In function 'swap':
5 ../swap.c:12: undefined reference to 'buf'
6 swap.o(.data+0x0): undefined reference to 'buf'
7 collect2: error: ld returned 1 exit status
```

- ▶ The linker tries to link with `crt1.o`, which refers to the `main` function.

Choosing one among multiple definitions

- ▶ The linker distinguishes weak and strong symbols:
 - Functions and initialized global variables are **strong symbols**.
 - Uninitialized global variables are **weak symbols**.
- ▶ If a conflict arises, the strategy is as follows:
 - Multiple strong symbols → raise error.
 - One strong, and multiple weak symbols → choose the strong one.
 - Multiple weak symbols → choose an arbitrary one.

Example Two strong symbols

```
1 $ pk-cc -c foo1.c bar1.c          # compilation is fine
2 $ pk-cc foo1.o bar1.o
3 bar1.o: In function 'main':
4 .../bar1.c:2: multiple definition of 'main'
5 foo1.o:../foo1.c:2: first defined here
6 collect2: error: ld returned 1 exit status
```

foo1.c

```
1 int main(void)
2 {
3     return 0;
4 }
```

bar1.c

```
1 int main(void)
2 {
3     return 0;
4 }
```

Question What will happen here?

foo2.c

```
1 #include <stdio.h>
2
3 void f(void);
4 int x = 12345;
5
6 int main(void)
7 {
8     f();
9     printf("x = %d\n", x);
10
11     return 0;
12 }
```

bar2.c

```
1 int x = 54321;
2
3 void f(void)
4 {
5     x++;
6 }
```

Example One strong, and one weak symbol.

foo3.c

```
1 #include <stdio.h>
2
3 void f(void);
4
5 int x = 12345; /* strong */
6
7 int main(void)
8 {
9     f();
10    printf("x = %d\n", x);
11
12    return 0;
13 }
```

bar3.c

```
1 int x; /* weak */
2
3 void f(void)
4 {
5     x = 54321;
6 }
```

The result is probably expected:

```
1 $ pk-cc -c foo3.c bar3.c
2 $ pk-cc foo3.o bar3.o
3 $ ./a.out
4 x = 54321
```

► Maybe check out the symbol tables? `readelf -s {foo,bar}3.o`

Example Two weak symbols.

foo4.c

```
1 #include <stdio.h>
2
3 void f(void);
4
5 int x; /* weak */
6
7 int main(void)
8 {
9     x = 12345;
10    f();
11    printf("x = %d\n", x);
12
13    return 0;
14 }
```

bar4.c

```
1 int x; /* weak */
2
3 void f(void)
4 {
5     x = 54321;
6 }
```

Again, no surprise:

```
1 $ pk-cc -c foo4.c bar4.c
2 $ pk-cc foo4.o bar4.o
3 $ ./a.out
4 x = 54321
```

Note The linker has unified both variables `x`, and makes references to both symbols address the same space in memory.

Question What about this one?

foo5.c

```
1 #include <stdio.h>
2
3 void f(void);
4
5 int x = 12345; /* strong */
6 int y = 54321; /* strong */
7
8 int main(void)
9 {
10     f();
11     printf("x = %d, y = %d\n", x, y);
12
13     return 0;
14 }
```

bar5.c

```
1 double x; /* weak */
2
3 void f(void)
4 {
5     x = 1;
6 }
```

Explain this:

```
1 $ pk-cc -c foo5.c bar5.c
2 $ pk-cc foo5.o bar5.o
3 /usr/bin/ld: Warning: alignment 4
4 of symbol 'x' in foo5.o is smaller
5 than 8 in bar5.o
6 $ ./a.out
7 x = 0, y = 1072693248
```

Notes

- ▶ Not long ago, the linker would give **no warning** at all.
- ▶ It is extremely difficult to **debug** such code.

What else?

- ▶ After resolving symbols, the linker knows which definition belongs to each symbol.

Recall

- ▶ Machine code does not use variable names any more.
- ▶ The compiler produced code that accesses variables and functions only by their **memory addresses**.

⇒ How does this go together with separate compilation and symbol resolution?

12.3 The program in memory

How does a program start?

- ▶ When a program is run, it is **copied into memory** by the **loader**.
 - Copy **text segment**, *i.e.*, the actual machine code,
 - copy **initialized data**,
 - **initialize** uninitialized data,
 - *etc.*
- ▶ We want to minimize the amount of data to be copied!
 - Only load parts that are **actually required**,
 - and only load them **when** they are needed.
- ▶ We want to save memory!
 - Do not load the same code into memory multiple times.
 - **Share** already loaded code between processes.
- ▶ Avoid expensive **transformations**
 - Store program on disk in a **format** that allows fast setup of the process image.

Virtual memory

- ▶ VM is a mapping from the process' **virtual address space** into the machine's physical address space (organized in **pages**).
- ▶ The VM system may flag pages as, e.g., **read only**, **executable**, or **private**, cf. `mmap(2)`.
- ▶ A physical page may reside on disk, until **loaded on demand**.
 - So we compile the memory layout into the executable file,
 - the loader just **maps the file** into the process' virtual address space, and
 - the VM system gets the pages into memory when actually referenced.
- ▶ Multiple running instances of a program share their text (machine code) through a *read only* mapping to **the same physical** address space.

Note To achieve all this, the structure of the program file depends on the process' memory layout!

Example

- ▶ The mapping of virtual memory can be observed in `/proc/$PID/maps`.
- ▶ I have several `xterms` running, one with PID 23172.

```

1 $ cat /proc/23172/maps                                     # this is on a 64bit machine
2 # address          perms offset  dev   inode  pathname
3 00400000-00472000    r-xp 00000000 00:10 987773 /usr/bin/xterm
4 00672000-00673000    r--p 00072000 00:10 987773 /usr/bin/xterm
5 00673000-0067c000    rw-p 00073000 00:10 987773 /usr/bin/xterm
6 0067c000-0067e000    rw-p 00000000 00:00 0
7 006f7000-00718000    rw-p 00000000 00:00 0      [heap]
8 # ...
9 7f05d802c000-7f05d81cc000 r-xp 00000000 00:10 953240 /usr/lib/libc-2.18.so
10 # ...
11 ffffffff600000-ffffffff601000 r-xp 00000000 00:00 0 [vsyscall]

```


- ▶ The executable pages are readonly,
- ▶ the writable area is private (*i.e.*, copy-on-write).
- ▶ Other instances of `xterm` have the same mapping for the binary `xterm`,
- ▶ but may have others for the shared libraries.

(*cf.* `proc(5)` for more information)

Process memory layout

When running, a process has the following **virtual memory layout**.

(This is for 32bit Linux)



- ▶ Kernel memory (1GiB) is not accessible by the process.
 - **Shared** among all processes.
- ▶ The stack maintains local variables and function calls.
- ▶ Shared libraries may even be added at runtime.
- ▶ The heap contains allocated memory.
- ▶ `.data` and `.bss` store global variables.
- ▶ `.text` and `.rodata` are marked `ro`, so can be shared with other processes.

12.4 Object file layout

- ▶ There are various formats to store binary programs.
- ▶ Linux uses ELF, the **Executable and Linking Format**.
- ▶ COFF and `a.out` are others, the latter coined the name used by `gcc` for default binaries (in ELF on Linux!).
- ▶ All formats have the concept of **sections** in common.
- ▶ A section is the unit of organization in a binary.

Some section names (but there are many more)

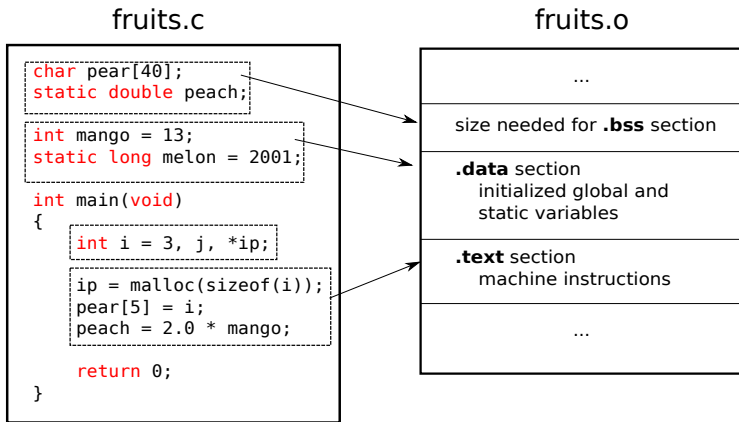
`.text` The program code, *i.e.*, processor instructions.

`.rodata` Read-only data, *e.g.*, string literals.

`.data` *Initialized global* variables.

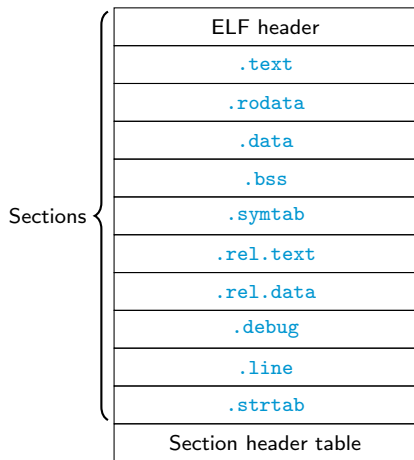
`.bss` *Uninitialized global* variables.

`.symtab` The symbol table, displayed with `readelf -s`.



A typical *relocatable* object file

This is what the compiler produces out of the **individual C files**.

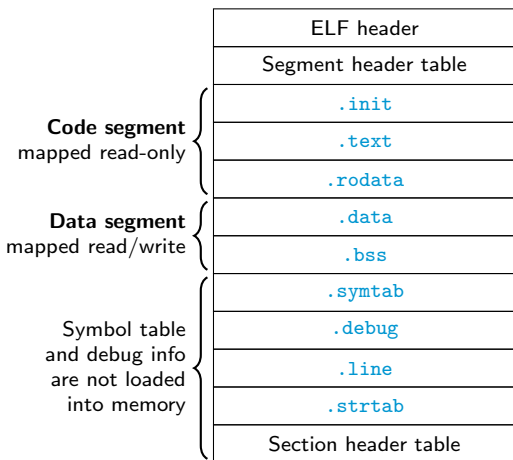


- ▶ The **ELF header** describes word size, endian, object file type, machine type, offset and format of the section header table, and other information.
- ▶ The **section header table** describes the locations of the various sections.
- ▶ Try

```
1 $ pk-cc -c -m32 swap.c
2 $ readelf -S swap.o
```

A typical *executable* object file

That is what we want to have in the **final binary program**.




- ▶ The **segment header table** describes the mapping of contiguous file sections into memory.

- ▶ Try this

```
1 $ pk-cc -m32 -static swap.c \  
2 > main.c  
3 $ readelf -l a.out
```

12.5 Relocation

- So after resolving the symbols, the linker needs to put all the code from the individual object files' sections into the final program's sections:



- This process is called **Relocation**.

Relocation involves **two tasks**:

1. Relocating **sections** and symbol **definitions**.
 - Merge sections of the same type.
 - Assign **run-time memory addresses** to the new aggregate sections, the input sections, and each symbol defined in the input.⇒ After this step, every global symbol has a known run-time memory address.
2. Relocating symbol **references** everywhere in the code.
 - Modification of each reference in `.text` and `.data`, so that they point to correct location.

Relocation entries

- ▶ The assembler does not know where data and code will be stored ultimately,
 - ▶ nor does it know addresses of the external objects.
- ⇒ In such situations a
- relocation entry*
- is generated by the assembler.

Relocation entries

```
1 typedef struct {  
2     int offset;      /* Offset of the reference to relocate. */  
3     int symbol: 24,  /* Symbol the reference should point to. */  
4         type:8;      /* Relocation type. */  
5 } Elf32_Rel;
```

- ▶ ELF defines 11 different relocation types
- ▶ We will look at [R_386_32](#) only.
This is used to relocate 32bit absolute addresses.

Example: Relocation at work

- ▶ Function `f` simply assigns `0xbeef` to the global variable `x`.
- ▶ The final memory location of `x` is not yet known.
- ▶ Relocation will fix the runtime address of `x`.

`foo6.c`

```
1 int x = 0xdead;
2
3 void f(void)
4 {
5     x = 0xbeef;
6 }
```

- ▶ First we **compile** `foo.c` into a *relocatable object file*:

```
1 $ pk-cc -m32 -c foo6.c
```

- ▶ Have a look at the `.data` section:

```
1 $ objdump -d -j.data foo6.o
2 Disassembly of section .data:
3 00000000 <x>:
4      0:  ad de 00 00                                ....
```

- Variable `x` appears at address `0x0` in the `.data` section.
- The value `0xdead` is stored there.

► Have a look at the `.text` section:

```

1 $ objdump -d -j.text foo6.o
2 Disassembly of section .text:
3 00000000 <f>:
4   0:   55                push    %ebp
5   1:   89 e5             mov     %esp,%ebp
6   3:  c7 05 00 00 00 00 ef  movl    $0xbeef,0x0
7   a:  be 00 00
8   d:   5d                pop     %ebp
9   e:   c3                ret

```

- In line 6, the value `0xbeef` is copied to address `0x0`.
- This address `0x0` appears at **offset 5** in the `.text` section.

► These are the relocation entries:

```

1 $ objdump -r -j.text foo6.o
2 RELOCATION RECORDS FOR [.text]:
3 OFFSET      TYPE          VALUE
4 00000005 R_386_32          x

```

- So on relocation of symbol `x`, the absolute 32bit address at **offset 5** in the `.text` section must be updated.

- Then we **link** `foo.o` with something that uses `f`.

```
1 $ pk-cc -m32 foo6.o bar6.c /* main in bar6.c simply calls f in foo6.c */
```

- Have a look at the `.data` section **after relocation**:

```
1 $ objdump -d -j.data a.out
2 Disassembly of section .data:
3 08049698 <x>:
4 8049698:      ad de 00 00      ....
```

- Variable `x` has been moved to address `0x8049698` in the `.data` section.
- The value `0xdead` is stored there.

- Have a look at the `.text` section **after relocation**:

```
1 $ objdump -d -j.text a.out | grep -C3 beef
2 080483cd <f>:
3 80483cd:      55                push    %ebp
4 80483ce:      89 e5             mov     %esp,%ebp
5 80483d0:      c7 05 98 96 04 08 ef movl    $0xbeef,0x8049698
6 80483d7:      be 00 00
7 80483da:      5d                pop     %ebp
8 80483db:      c3                ret
```

- The reference to variable `x` has been **updated to address** `0x8049698`.

Example: Relocation in the .data section

- ▶ Sometimes, updating references in the `.text` section is not enough.
- ▶ Here we have a global variable `xp` initialized with an address!
- ▶ The compiler **cannot even fix a value** for `xp`!

foo7.c

```
1 int x = 0xdead;
2 int *xp = &x;
3
4 void f(void)
5 {
6     *xp = 0xbeef;
7 }
```

- ▶ Again, we compile and link our object files:

```
1 $ pk-cc -m32 -c foo7.c
2 $ pk-cc -m32 foo7.o bar6.c
```

► Relocation in the `.data` section.

```

1  $ objdump -d -j.data foo7.o
2  Disassembly of section .data:
3  00000000 <x>:
4      0:    ad de 00 00                                ....
5  00000004 <xp>:
6      4:    00 00 00 00    # This value has to be updated on relocation of x!
7  $ objdump -r -j.data foo7.o    # Note: in .data this time!
8  RELOCATION RECORDS FOR [.data]:
9  OFFSET      TYPE              VALUE
10 00000004 R_386_32          x
11 $ objdump -d -j.data a.out
12 Disassembly of section .data:
13 08049698 <x>:
14 8049698:    ad de 00 00                                ....
15 0804969c <xp>:
16 804969c:    98 96 04 08                                ....

```

► Relocation in the `.text` section:

```

1  $ objdump -d -j.text foo7.o
2  Disassembly of section .text:
3  00000000 <f>:
4      0:  55                push    %ebp
5      1:  89 e5             mov     %esp,%ebp
6      3:  a1 00 00 00 00    mov     0x0,%eax
7      8:  c7 00 ef be 00 00 movl    $0xbeef, (%eax)
8      e:  5d                pop     %ebp
9      f:  c3                ret
10 $ objdump -r -j.text foo7.o
11 RELOCATION RECORDS FOR [.text]:
12 OFFSET      TYPE          VALUE
13 00000004 R_386_32          xp
14 $ objdump -d -j.text a.out | grep -C3 beef
15 80483cd:      55                push    %ebp
16 80483ce:      89 e5             mov     %esp,%ebp
17 80483d0:      a1 9c 96 04 08    mov     0x804969c,%eax
18 80483d5:      c7 00 ef be 00 00 movl    $0xbeef, (%eax)
19 80483db:      5d                pop     %ebp
20 80483dc:      c3                ret

```

► Recall: `0x804969c` is the address of the variable `xp`!